

# Writing good C++14

Bjarne Stroustrup

Morgan Stanley, Columbia University

[www.stroustrup.com](http://www.stroustrup.com)



# The big question

- “**What is good modern C++?**”
  - *Many* people want to write “Modern C++”
- Guidelines project
  - Produce a useful answer
  - Enable *many* people to use that answer
    - For most programmers, not just language experts
  - Please help!



# The problem and the opportunity

- We have a great modern language
  - C++11 (good)
    - -> C++14 (better)
      - -> C++17 (much better still, I hope)
  - Technical specifications
  - Shipping
    - in wide-spread production work
      - and more facilities well in the works
  - C++1\*
    - is easier to write and maintain
    - runs faster
    - can express more than older C++
      - with less code



# The problem and the opportunity

- Many people
  - Use C++ in archaic or foreign styles
  - Get lost in details
  - Are obsessed with language-technical details


Doctor, doctor, it hurts when I do X!!!  
So don't do X



- “Within C++ is a smaller, simpler, safer language struggling to get out”
  - Code can be simpler
  - as efficient as ever
  - as expressive as ever

# A smaller, simpler C++

- Let's get it out
  - **Now!**
  - Without inventing a new language
  - 100% compatibility – compile with current compilers
- Coding guidelines
  - Supported by a “guidelines support library” (GSL)
  - Supported by analysis tools
- Don't sacrifice
  - Generality
  - Performance
  - Simplicity
  - Portability across platforms



# A smaller, simpler C++

- I think **we** can do it
  - I can't do it alone
  - No individual can
  - No single company can
- Please help!



# Initial work (still incomplete)

- I describe significant initial work
  - Microsoft (Herb Sutter and friends)
  - Morgan Stanley (Bjarne Stroustrup and friends)
  - CERN (Axel Naumann and friends)
- Available
  - Core guidelines (now)
  - Guidelines support library (now; Microsoft, GCC, Clang; Windows, Linux, Mac)
  - Analysis tool (Microsoft in October; ports later (November?))
  - MIT License
- Related CppCon talks
  - Herb Sutter: ***Writing Good C++14 By Default*** (Tuesday)
  - Gabriel Dos Reis: ***Modules*** (Tuesday)
  - Gabriel Dos Reis: ***Contracts*** (Wednesday)
  - Neil MacIntosh: ***Static analysis*** (Wednesday)
  - Neil MacIntosh: ***array\_view, string\_view, etc.*** (Wednesday)

# We all hate coding rules<sup>\*†</sup>

- Rules are (usually)
  - Written to prevent misuse by poor programmers
    - “don’t do this and don’t do that”
  - Written by people with weak experience with C++
    - At the start of an organization’s use of C++
- Rules (usually) focus on
  - “layout and naming”
  - Restrictions on language feature use
  - Not on programming principles
- Rules (usually) are full of bad advice
  - Write “pseudo-Java” (as some people thought was cool in 1990s)
  - Write “C with Classes” (as we did in 1986)
  - Write C (as we did in 1978)
  - ...

\*Usual caveats

†and thanks



# Coding rules<sup>\*</sup>

- Are outdated
  - Become a drag of their users
- Are specialized
  - but used outside their intended domain
- Are not understood by their users
  - Enforced by dictate: Do this or else!
  - Require detailed language-lawyer knowledge to follow
- Are not well supported by tools
  - Platform dependencies
  - Compiler dependencies
  - Expensive
- Do not provide guidance
  - Telling what not to do is not enough

<sup>\*</sup>Usual caveats



# Coding guidelines

- Let's build a **good** set!
  - Comprehensive
  - Browsable
  - Supported by tools (from many sources)
  - Suitable for gradual adoption
- For modern C++
  - Compatibility and legacy code be damned! (initially)
- Prescriptive
  - Not punitive
- Teachable 
  - Rationales and examples
- Flexible
  - Adaptable to **many** communities and tasks
- Non-proprietary
  - But assembled with taste and responsiveness
- We aim to offer guidance
  - What is good modern C++?
  - Confused, backwards-looking teaching is a big problem

# High-level rules

- Provide a conceptual framework
  - Primarily for humans
- Many can't be checked completely or consistently
  - *P.1: Express ideas directly in code*
  - *P.2: Write in ISO Standard C++*
  - *P.3: Express intent*
  - *P.4: Ideally, a program should be statically type safe*
  - *P.5: Prefer compile-time checking to run-time checking*
  - *P.6: What cannot be checked at compile time should be checkable at run time*
  - *P.7: Catch run-time errors early*
  - *P.8: Don't leak any resource*
  - *P.9: Don't waste time or space*



# Lower-level rules

- Provide enforcement
  - Some complete
  - Some heuristics
  - Many rely on static analysis
  - Some beyond our current tools
  - Often easy to check “mechanically”
- Primarily for tools
  - To allow specific feedback to programmer
- Help to unify style
- Not minimal or orthogonal
  - *F.16: Use **T\*** or **owner<T\*>** to designate a single object*
  - *C.49: Prefer initialization to assignment in constructors*
  - *ES.20: Always initialize an object*

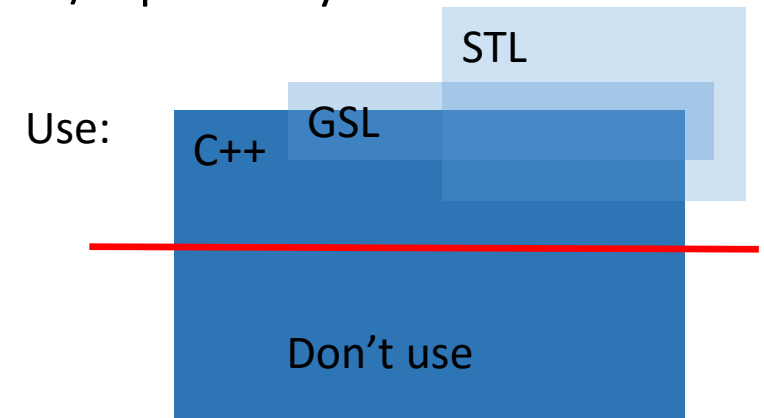


# The structure of a rule

- The rule itself - e.g., ***no naked `new`***
- **Reference number** - e.g., **C.7** (the 7th rule related to classes).
- **Reason** (rationale) - because programmers find it hard to follow rules they don't understand
- **Example** - because rules are hard to understand in the abstract; can be positive or negative
- **Alternative** - for "don't do this" rules
- **Exception** - we prefer simple general rules. However, many rules apply widely, but not universally
- **Enforcement** - ideas about how the rule might be checked "mechanically"
- **See also** - references to related rules and/or further discussion (in this document or elsewhere)
- **Note** (comments) - something that needs saying that doesn't fit the other classifications
- **Discussion** - references to more extensive rationale and/or examples placed outside the main lists of rules


# Subset of superset

- Simple sub-setting doesn't work
  - We need the low-level/tricky/close-to-the-hardware/error-prone/expert-only features
    - For implementing higher-level facilities efficiently
    - Many low-level features can be used well
  - We need the standard library
- Extend language with a few abstractions
  - **Use** the STL
  - **Add** a small library (the GSL)
    - **No** new language features
    - Messy/dangerous/low-level features can be used to implement the GSL
  - **Then** subset
- What we want is “**C++ on steroids**”
  - Simple, safe, flexible, and fast
  - Not a neutered subset



# Some rules rely on libraries

- The ISO C++ standard library
  - E.g., **vector<T>** and **unique\_ptr<T>**
- The Guideline Support Library
  - E.g., **array\_view<T>** and **not\_null<T>**
- Some rules using the GSL
  - *l.11: Never transfer ownership by a raw pointer ( $T^*$ )*
    - Use an ownership pointer (e.g. **unique\_ptr<T>**) or **owner<T\*>**
  - *l.12: Declare a pointer that may not be the **nullptr** as **not\_null***
    - E.g., **not\_null<int\*>**
  - *l.13 Do not pass an array as a single pointer*
    - Use a handle type, e.g., **vector<T>** or **array\_view<T>**





# Double our productivity

- “Imitate experienced programmers”
  - Most programmer don’t know what “everybody knows”
- Eliminate whole classes of errors
  - Fewer crashes and security violations
- Simplify
  - Simplicity aids maintenance
  - Consistent style speeds up learning
  - Guide people away from obscure corners and exotic technique
  - Emphasis on avoiding waste improves performance
  - Separate rules for exceptional needs
- Do not compromise performance






# Have you gone mad? (no)

- We attack the most common and the most serious sources of errors
  - I hate debugging
- We eliminate whole classes of errors
  - Eliminate resource leaks
    - Without loss of performance
  - Eliminate dangling pointers
    - Without loss of performance
  - Eliminate out-of-range access
    - With minimal cost
- Tool support is essential
  - Static analysis
  - Support library (tiny)
  - Reinforce the type system

# Core Rules

- Some people will not be able to apply all rules
  - At least initially
  - Gradual adoption will be very common
- Many people will need additional rules
  - For specific needs
- We initially focus on the core rules
  - The ones we hope that everyone eventually could benefit from
- The core of the core
  - No leaks
  - No dangling pointers
  - No type violations through pointers



# No resource leaks

- We know how
  - Root every object in a scope
    - `vector<T>`
    - `string`
    - `ifstream`
    - `unique_ptr<T>`
    - `shared_ptr<T>`
  - RAII
    - “No naked **new**”
    - “No naked **delete**”



# Dangling pointers – the problem

- One nasty variant of the problem

```
void f(X* p)
{
    // ...
    delete p;           // looks innocent enough
}

void g()
{
    X* q = new X;       // looks innocent enough
    f(q);
    // ... do a lot of work here ...
    q->use();            // Ouch! Read/scramble random memory
}
```

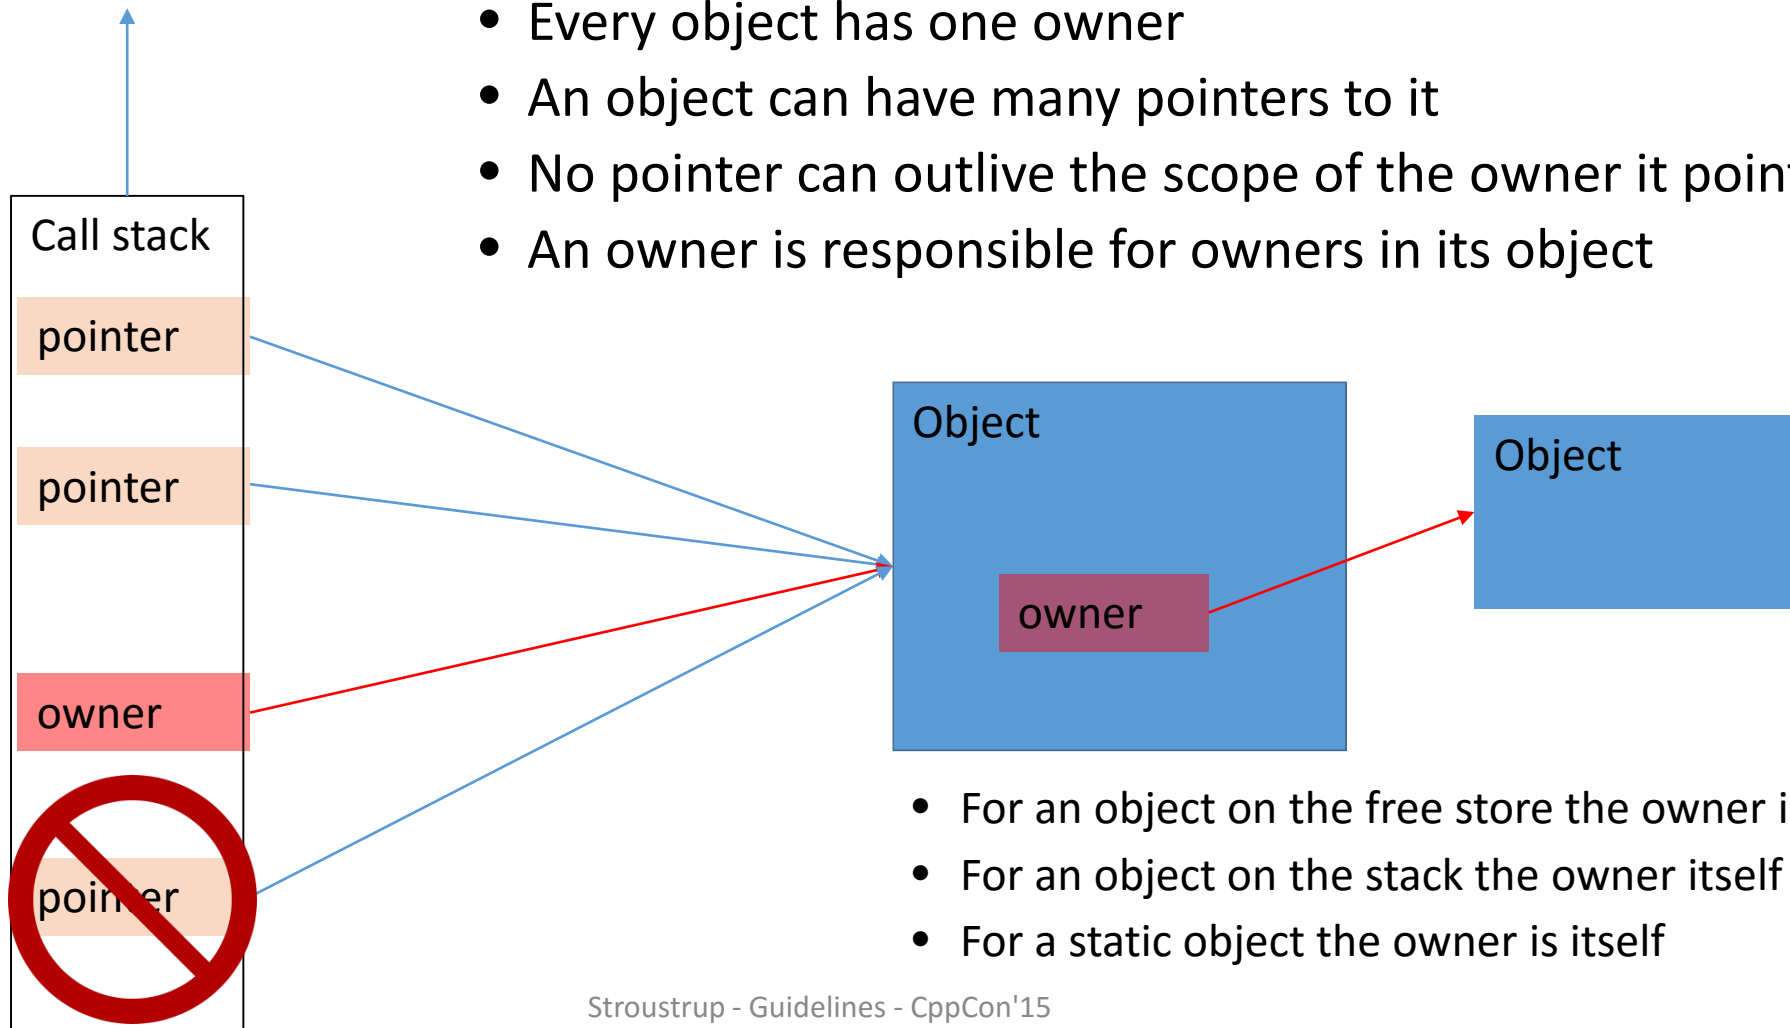


# Dangling pointers

- We ***must*** eliminate dangling pointers
  - Or type safety is compromised
  - Or memory safety is compromised
  - Or resource safety is compromised
- Eliminated by a combination of rules
  - Distinguish owners from non-owners
  - Assume raw pointers to be non-owners
  - Catch all attempts for a pointer to “escape” into a scope enclosing its owner’s scope
    - **return**, **throw**, out-parameters, long-lived containers, ...
  - Something that holds an owner is an owner
    - E.g. `vector<owner<int*>>, owner<int*>[], ...`



# Owners and pointers





# Dangling pointers

- Ensure that no pointer outlives the object it points to

```
void f(X* p)
{
    // ...
    delete p;           // bad: delete non-owner
}

void g()
{
    X* q = new X;       // bad: assign object to non-owner
    f(q);
    // ... do a lot of work here ...
    q->use();            // Make sure we never get here
}
```



# How do we represent ownership?

- High-level: Use an ownership abstraction
- Low-level: mark owning pointers **owner**
  - An **owner** must be **deleted** or passed to another **owner**
  - A non-**owner** may not be **deleted**
- Note
  - I talk about pointers
  - What I say applies to anything that refers to an object
    - References
    - Containers of pointers
    - Smart pointers
    - ..



# How do we represent ownership

- Mark an owning **T\***: **owner<T\*>**
  - Initial idea
    - **owner<T\*>** would hold a **T\*** and an “owner bit”
    - Costly: bit manipulation
    - Not ABI compatible
    - Not C compatible
  - So our **GSL owner** is
    - A handle for static analysis
    - Documentation
    - Not a type with it's own operations
    - Cost free: No run-time cost (time or space)
    - ABI compatible
    - **template<typename T> owner = T;**

# GSL: owner<T>

- How do we implement ownership abstractions?

```
template<SemiRegular T>
```

```
class vector {
```

```
    owner<T*> elem;           // the anchors the allocated memory
```

```
    T* space;                 // just a position indicator
```

```
    T* end;                   // just a position indicator
```

```
    // ...
```

```
};
```

- **owner<T\*>** is just an alias for **T\***

# GSL: owner<T>

- How about code we cannot change?

```
void foo(owner<int*>);           // foo requires an owner
```

```
void f(owner<int*> p, int* q, owner<int*> p2, int* q2)
{
    foo(p);                     // OK: transfer ownership
    foo(q);                     // bad: q is not an owner
    delete p2;                 // necessary
    delete q2;                 // bad: not an owner
}
```

- A static analysis tool can tell us where our code mishandles ownership

# owner is a low-level mechanism

- Use proper ownership abstractions
  - E.g., **unique\_ptr** and **vector**
    - Implemented using **owner**
- **owner** is intended to simplify static analysis
  - **owners** in application code is a sign of a problem
    - Usually, C-style interfaces

# How to avoid/catch dangling pointers

- Rules (giving pointer safety):
  - Don't transfer to pointer to a local to where it could be accessed by a caller
  - A pointer passed as an argument can be passed back as a result
  - A pointer obtained from new can be passed back as a result as an owner

```
int* f(int* p)
{
    int x = 4;
    return &x;           // No! would point to destroyed stack frame
    return new int{7};    // OK (sort of: doesn't dangle, but returns an owner as an int*)
    return p;            // OK: came from caller
}
```

# How to avoid/catch dangling pointers

- It's not just pointers
  - All ways of “escaping”
    - **return**, **throw**, place in long-lived container, ...
  - Same for containers of pointers
    - E.g. **vector<int\*>**, **unique\_ptr<int>**, iterators, built-in arrays, ...
  - Same for references
- Never let a “pointer” point to an out-of-scope object

# How to avoid/catch dangling pointers

- Classify pointers according to ownership

```
vector<int*> f(int* p)  
{  
    int x = 4;  
    int* q = new int{7};  
    vector<int*> res = {p, &x, q};    // Bad: { unknown, pointer to local, owner }  
    return res;  
}
```

- Don't mix different ownerships in an array
- Don't let different return statements of a function mix ownership

# How to avoid/catch dangling pointers

- Try to be explicit about ownership

```
vector<int*> f(int* p)
{
    int x = 4;
    owner<int*> q = new int{7};
    vector<int*> res = {p, &x, q};           // Bad: { unknown, pointer to local, owner }
    vector<owner<int*>> r2 = {p, &x, q};      // Bad: { unknown, pointer to local, owner }
    return res;
}
```

- Some convoluted code cannot be represented in a statically type-safe manner
  - Avoid such code
  - If you really need it, encapsulate it in an expression that include run-time representation of ownership (pointer, ownership bit)



# Other problems

- Other ways of misusing pointers
  - Range errors: `array_view<T>`
  - `nullptr` dereferencing: `not_null<T>`
- Wasteful ways of addressing pointer problems
  - Misuse of smart pointers
- Other ways of breaking the type system (beyond the scope of this talk)
  - Unions
  - Casts
- “Just test everywhere at run time” is ***not*** an acceptable answer
  - Hygiene rules
  - Static analysis
  - Run-time checks

# GSL - array\_view<T>

- Common style

```
void f(int* p, int n)           // what is n? (How would a tool know?)
{
    p[7] = 9;                   // OK?
    for (int i=0; i<n; ++i) p[i] = 7; // OK?
}
```

- Better

```
void f(array_view<int> a)
{
    a[7] = 9;                   // OK? Checkable against a.size()
    for (int& x : a) x = 7;     // OK
}
```

# GSL - array\_view<T>

- Common style

```
void f(int* p, int n);  
int a[100];  
// ...  
f(a,100);  
f(a,1000);    // likely disaster
```

- “Make simple things simple”

- Simpler than “old style”
- Shorter
- At least as fast
- Sometimes using the GSL
- Sometimes using the STL

- Better

```
void f(array_view<int> a)  
int a[100];  
// ...  
f(array_view<int>{a});  
f(a);  
f({a,1000}); // easily checkable
```

# nullptr problems

- Mixing **nullptr** and pointers to objects
  - Causes confusion
  - Requires (systematic) checking

- Caller

```
void f(char*);
```

```
f(nullptr);           // OK?
```

- Implementer

```
void f(char* p)
```

```
{
```

```
    if (p==nullptr)    // necessary?
```

```
    // ...
```

```
}
```

- Can you trust the documentation?
- Compilers don't read manuals, or comments
- Complexity, errors, and/or run-time cost

# GSL - not\_null<T>

- Caller

```
void f(not_null<char*>);
```

```
f(nullptr);    // Obvious error: caught by static analysis
```

```
char* p = nullptr;
```

```
f(p);          // Constructor for not_null can catch the error
```

- Implementer

```
void f(not_null<char*> p)
```

```
{
```

```
    // if (p==nullptr) // not necessary
```

```
    // ...
```

```
}
```

# GSL - `not_null<T>`

- **`not_null<T>`**
  - A simple, small class
  - **`not_null<T*>`** is **`T*`** except that it cannot hold **`nullptr`**
  - Can be used as input to analyzers
    - Minimize run-time checking
  - Checking can be “debug only”
  - For any **`T`** that can be compared to **`nullptr`**
    - E.g. **`not_null<array_view<T>>`**

# To summarize

- Type and resource safety:
  - RAI (scoped objects with constructors and destructors)
  - No dangling pointers
  - No leaks (track ownership pointers)
  - Eliminate range errors
  - Eliminate nullptr dereference
- That done we attack other sources of problems
  - Logic errors
  - Performance bugs
  - Maintenance hazards
  - Verbosity
  - ...



# (Mis)uses of smart pointers

- “Smart pointers” are popular
  - To represent ownership
  - To avoid dangling pointers
- “Smart pointers” are overused
  - Can be expensive
    - E.g., **shared\_ptr**
  - Can mess up interfaces for otherwise simple functions
    - E.g. **unique\_ptr** and **shared\_ptr**
  - Often, we don’t need a pointer
    - Scoped objects
    - We need pointers
      - For OO interfaces
      - When we need to change the object referred to

But ordinary pointers don’t dangle any more





# (Mis)uses of smart pointers

- Consider
  - `void f(T*);` *// use; no ownership transfer or sharing*
  - `void f(unique_ptr<T>);` *// transfer unique ownership and use*
  - `void f(shared_ptr<T*>);` *// share ownership and use*
- Taking a raw pointer (`T*`)
  - Is familiar
  - Is simple, general, and common
  - Is cheaper than passing a smart pointer (usually)
  - Doesn't lead to dangling pointers
  - Doesn't lead to replicated versions of a function for different shared pointers
- In terms of tradeoffs with smart pointers, other simple “object designators” are equivalent to `T*`
  - iterators, references, **`array_view`**, etc.

# (Mis)uses of smart pointers


- Don't use ownership pointers unless you change ownership

```
void f(X*);                // just uses X; no ownership transfer or sharing – good
void g(shared_ptr<X>);      // just uses X – bad
unique_ptr<X> h(unique_ptr<X>); // just uses X – bad (give pointer back to prevent destruction)
```

```
void use()
{
    auto p = make_shared<X>{};
    f(p.get());                // extract raw pointer (note: pointers do not dangle)
    g(p);                      // mess with use count (probably a mistake)
    auto q = h(make_unique<X>(p.get())); // transfer ownership to just use (a mistake)
                                // extract raw pointer, then wrap it and copy
    q.release();               // prevent destruction
}
```

# Rules, standards, and libraries

- Could the rules be enforced by the compiler?
  - Some could, but we want to use the rules **now**
    - Some compiler support would be very nice; let's talk
  - Many could not
  - Rules will change over time
  - Compilers have to be more careful about false positives
  - Compilers cannot ban legal code
- Could the GSL be part of the standard?
  - Maybe, but we want to use it **now**
  - The GSL is tiny and written in portable C++11
  - The GSL does not depend on other libraries
  - The GSL is similar to, but not identical to **boost::** and **experimental::** components
    - So they may become standard
- We rely on the standard library



# Too many rules

- For
  - Novices, experts, infrastructure, ordinary large applications, low-latency, high-reliability, security targets, hard-real time
- You can't remember all of those rules!
- You don't need all of those rules
- You couldn't learn all of those rules before writing code
- You'd hate to even look through all of those rules
- The rule set must be extensible
  - you'll never know them all
- The tools know the rules
  - And will point you to the relevant ones

# Rule classification

- P: Philosophy
- I: Interfaces
- F: Functions
- C: Classes and class hierarchies
- Enum: Enumerations
- ES: Expressions and statements
- E: Error handling
- R: Resource management
- T: Templates and generic programming
- CP: Concurrency
- The Standard library
- SF: Source files
- CPL: C-style programming
- GSL: Guideline support library

## Supporting sections

- NL: Naming and layout
- PER: Performance
- N: Non-Rules and myths
- RF: References
- Appendix A: Libraries
- Appendix B: Modernizing code
- Appendix C: Discussion
- To-do: Unclassified proto-rules

# We are not unambitious

- Type and resource safety
  - No leaks
  - No dangling pointers
    - No bad accesses
  - No range errors
  - No use of uninitialized objects
  - No misuse of
    - Casts
    - Unions
- We think we can do it
  - At scale
    - 4+ million C++ Programmers, N billion lines of code
  - Zero-overhead principle



# We aim to change the way we write code

- That means ***you***
- What would you like your code to look like in 5 years?
  - Once we know, we can aim to achieve that
  - Modernizing a large code base is not easy
  - The answer is not “just like my code today”
  - Think “gradual adoption” (except for brand-new code)
- Not everybody will agree what the code should look like
  - Not all code should look the same
  - We think there can be a common core
  - We need discussion, feedback, and a variety of tools
- Help wanted!
  - Rules, tools, reviews, comments
  - Editors





# Current status

- Available
  - About 350 Rules (<https://github.com/isocpp/CppCoreGuidelines>)
  - GSL for Clang, GCC, and Microsoft (<https://github.com/microsoft/gsl>)
  - First tools: October for Microsoft; ports later (November?)
  - MIT License
- We need help
  - Review of rules
    - More examples and refinements for existing rules
  - Specialized rule sets
    - For particular application areas, projects, ...
    - For concurrency
    - For libraries
    - ...
- Continuous development
  - “forever”





# The basic C++ model is now complete

- C++ (using the guidelines) is type safe and resource safe
  - Which other language can claim that?
  - Eliminate dangling pointers
  - Eliminate resource leaks
  - Check for range errors (optionally and cheaply)
  - Check for **nullptr** (optionally and cheaply)
  - Have concepts
- Why not a new C++-like language?
  - Competing with C++ is hard
    - Most attempts fail, C++ constantly improves
  - It would take 10 years (at least)
    - And we would still have lots of C and C++
  - A new C++-like language might damage the C++ community
    - Dilute support, divert resources, distract



# Questions

- P: Philosophy
- I: Interfaces
- F: Functions
- C: Classes and class hierarchies
- Enum: Enumerations
- ES: Expressions and statements
- E: Error handling
- R: Resource management
- T: Templates and generic programming
- CP: Concurrency
- The Standard library
- SF: Source files
- CPL: C-style programming
- GSL: Guideline support library

## Supporting sections

- NL: Naming and layout
- PER: Performance
- N: Non-Rules and myths
- RF: References
- Appendix A: Libraries
- Appendix B: Modernizing code
- Appendix C: Discussion
- To-do: Unclassified proto-rules





# Coding guidelines

- Boost Library Requirements and Guidelines
- Bloomberg: BDE C++ Coding
- Facebook: ???
- GCC Coding Conventions
- Google C++ Style Guide
- JSF++: JOINT STRIKE FIGHTER AIR VEHICLE C++ CODING STANDARDS
- Mozilla Portability Guide.
- Geosoft.no: C++ Programming Style Guidelines
- Possibility.com: C++ Coding Standard
- SEI CERT: Secure C++ Coding Standard
- High Integrity C++ Coding Standard
- [llvm.org/docs/CodingStandards.html](http://llvm.org/docs/CodingStandards.html)

# Non-aims

- Create “the one true C++ subset”
  - There can be no such marvel
  - Core guidelines + guidelines for specific needs
- Making a totally flexible set of rules to please everybody
  - Our rules are **not** value neutral
    - Total freedom is chaos
  - We want “modern C++”
    - **not** “everything anyone ever thought was cool and/or necessary”
- Turning C++ into Java, Haskell, C, or whatever
  - “If you want Smalltalk you know where to find it”
- What we want is “**C++ on steroids**”
  - Simple, safe, flexible, and fast
  - Not a neutered subset

# Philosophy

- Attack hard problems
  - Resources, interfaces, bounds, ...
- Be prescriptive
  - “don’t do that” is not very helpful
- Give rationale
  - “because I say so” is not very helpful
- Offer machine-checkable rules
  - Machines are systematic, fast, and don’t get bored
- Don’t limit generality
  - For most of us most of the time
- Don’t compromise performance
  - Of course
- Subset of superset
  - Don’t fiddle with subtle language rules