

Being Smart About Pointers

Michael VanLoon

Adapted from CPPcon 2014
m.vanloon@F5.com



A Quick Diversion

- Let's talk about RAI...

RAII:

Resource Acquisition is Initialization

- Scope-Based Resource Management
- An idiom for managing resources based on container object lifetime
- Most often used in scoped variable scenarios
- Automatic and low maintenance
- Reduces the opportunity for negligence
- How? Deterministic, Encapsulation, Locality, Exception Safety

RAII

- Deterministic:
 - Deallocation of resources happens immediately, during container object destruction
 - Standard scoping and object lifetime rules apply
 - Standard ordering rules apply
 - No delayed garbage collection

RAII

- Encapsulation:
 - Cleanup code is kept with other logically-related object management code
 - Cleanup code is not sprinkled throughout your program wherever an object might exit scope

RAII

- Locality:
 - Constructor, Destructor, other object management code, and container methods are all logically and physically grouped together

RAII

- Exception Safety:
 - Stack-based variables will exit scope as stack is unwound following exception
 - SBRM-based variables will automatically clean up resources as they exit scope

(lack of) RAI example

old and busted >> new hotness

```
class C
{
public:
    C(int i): x(i) { }
    ~C() { }
private:
    int x;
};
```

```
int
main(int argc, const char* argv[])
{
    for (int i = 0; i < 100; ++i)
    {
        C* c = new C(i);
        /* do something with c */
        delete c;
    }
    return 0;
}
```


RAII example

old and busted >> new hotness

```
class C;

class SBRM
{
public:
    SBRM(C* c) : pc(c)
        { }
    ~SBRM()
        { delete pc; }
    operator C*()
        { return pc; }
private:
    C* pc;
};
```

```
int
main(int argc, const char* argv[])
{
    for (int i = 0; i < 100; ++i)
    {
        SBRM cc(new C(i));
        C* c = cc;
        /* do something with c */
    }
    return 0;
}
```

Smart pointers, the toolbox

- `Boost::scoped_ptr` and `std::unique_ptr` (C++11)
- `shared_ptr`, with `make_shared` and `enable_shared_from_this`
- `weak_ptr`

Smart pointers, what they can do

- Features
 - RAI, with all its benefits
 - Custom deleters – not just new/delete!
 - Reference counted in some cases
 - Guarantee of deterministic delete
- Benefits
 - No more memory leaks or multiple frees
 - Code easier to read, easier to maintain, with less bugs
 - Simpler code logic; no messy ownership tracking

Pointer Best Practices

- Avoid explicit new/delete whenever possible
 - e.g. make_shared, make_unique, or custom factory
- Avoid delete except for exceptional circumstances
 - It's a sign you need to re-think pointer handling code
 - It's easy to forget
 - It's easy to accidentally maintain around
- Avoid handling raw pointers whenever reasonable
 - Put them into a container as soon as you get them
 - Even better: use make_shared or make_unique
- How?
 - Boost::scoped_ptr or std::unique_ptr
 - shared_ptr
 - Custom container of your own design

`boost::scoped_ptr` and `std::unique_ptr`

- Both adhere to RAII principles
- Both appropriate for managing local, scope-based resources
- Both actively prevent unintentional sharing of resource
- Both are faster and simpler than `shared_ptr`
- Prefer `unique_ptr` if you have it. If not, `scoped_ptr` and `shared_ptr` solve many of the same problems.

boost::scoped_ptr

- Available in boost
- scoped_ptr has neither Move nor Copy semantics; it is truly local scope only
- Separate scoped_array for new [] / delete []
- Does not support custom deleter

std::unique_ptr

- C++11 improvement on scoped_ptr
- unique_ptr has Move semantics but not Copy semantics (i.e. can return out of function and be embedded in standard containers)
- unique_ptr natively supports delete []
- unique_ptr allows custom deleter
- In C++14, has make_unique which is more exception safe and supports avoidance of explicit new and pointer handling

boost::scoped_ptr example

```
class C {  
public:  
    C() { cout << "C!" << endl; }  
    ~C() { cout << "~C" << endl; }  
};  
  
C* Cfactory() {  
    return new C;  
}  
  
void spFunc() {  
    boost::scoped_ptr<C> c(Cfactory());  
    // Do something with c...  
}
```

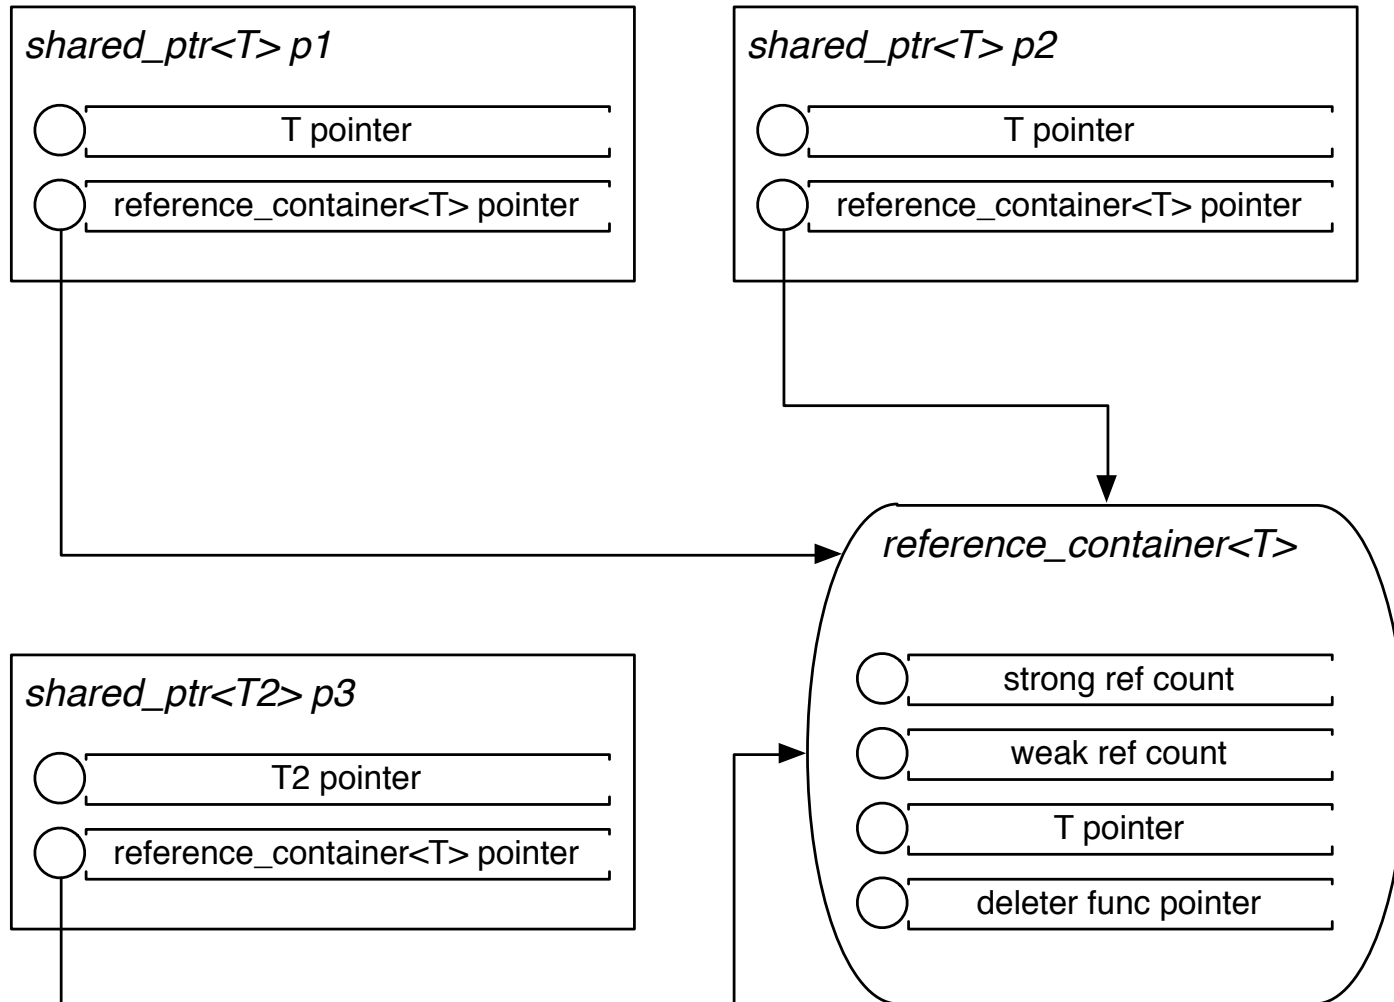

std::unique_ptr example

```
unique_ptr<C> returnUP() {  
    unique_ptr<C> c(new C);  
    // Do some magic with c, but we're not final owner  
    return c;  
}  
  
void upFunc() {  
    unique_ptr<C> c(returnUP());  
    // Do something with c; we are the final owner  
}  
  
void upVecFunc() {  
    vector<unique_ptr<C>> cVec;  
    for (int i = 0; i < 5; ++i)  
        cVec.emplace_back(new C);  
    // cVec goes out of scope and cleans up nicely  
}
```

shared_ptr, make_shared and enable_shared_from_this

- shared_ptr adheres to RAII principles
- Reference-counted container
- Copyable
- Custom deleter supported for cleaning up resource allocations other than new
- make_shared allocates object more efficiently, is exception safe, and avoids explicit new
- enable_shared_from_this allows object to return shared_ptr containing itself
- Prefer unique/scoped_ptr over shared_ptr

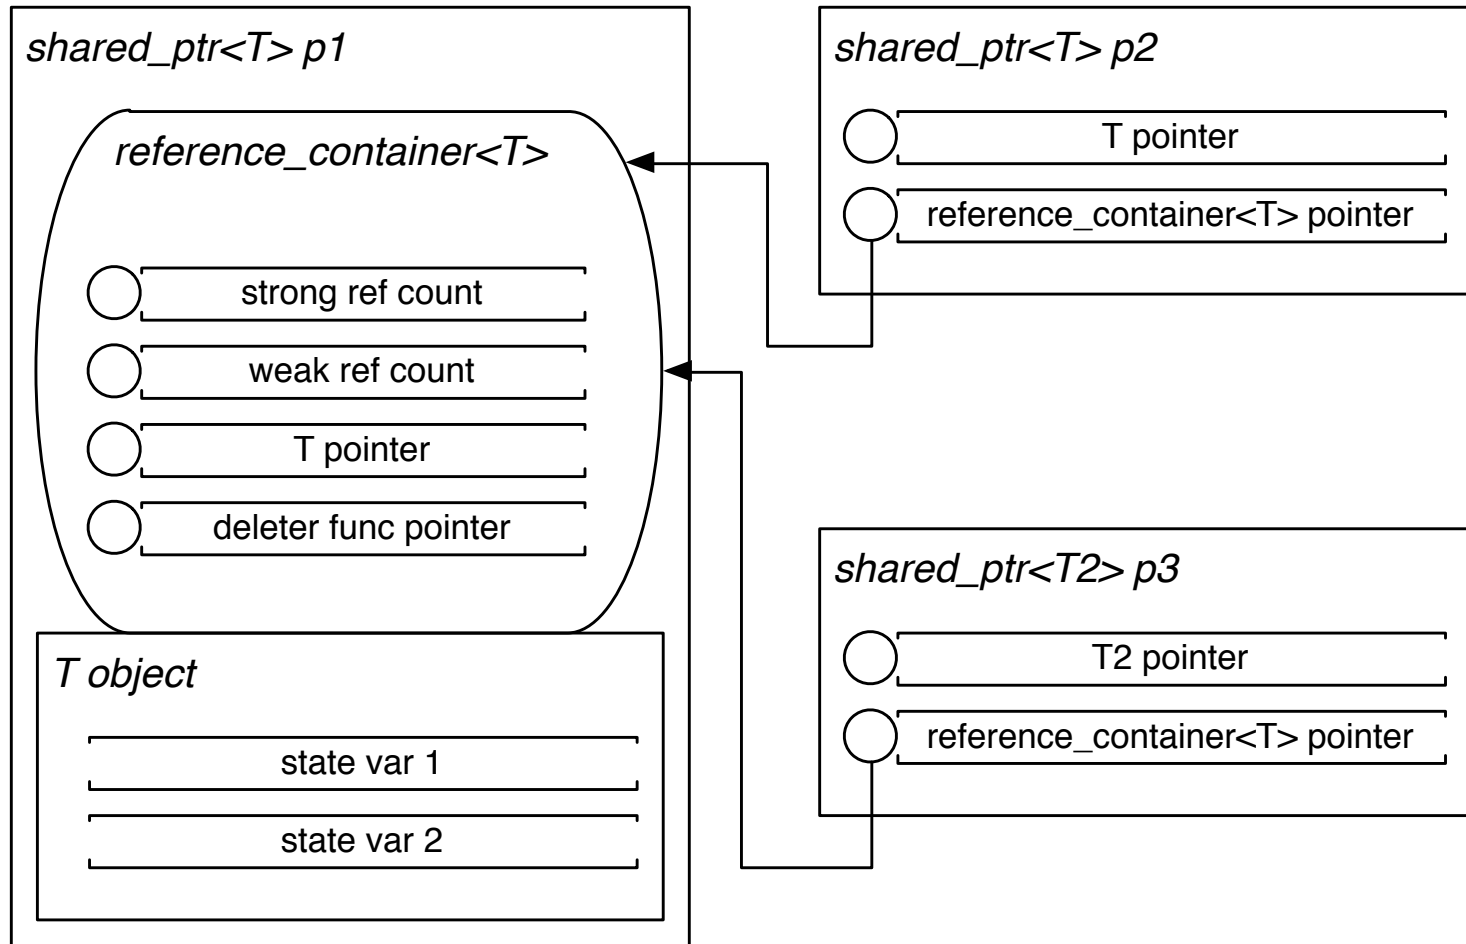
Example shared_ptr implementation



shared_ptr construction/assignment and make_shared

```
class C {  
public:  
    C() { cout << "C!" << endl; }  
    ~C() { cout << "~C" << endl; }  
};  
  
void spAssign() {  
    shared_ptr<C> p1 = make_shared<C>();  
    assert(p1.use_count() == 1);  
    shared_ptr<C> p2 = p1;  
    assert(p2.use_count() == 2);  
    shared_ptr<C> p3(new C);  
    assert(p3.use_count() == 1);  
    p2 = p3;  
    assert(p1.use_count() == 1 && p3.use_count() == 2);  
}
```

Example make_shared optimization



shared_ptr copy/move and embedding in containers

```
shared_ptr<C> Cfactory() {  
    return make_shared<C>();  
}
```

```
void spVec() {  
    vector< shared_ptr<C> > cVec;  
    cVec.emplace_back(Cfactory()); // first C  
    cVec.emplace_back(Cfactory()); // second C  
    // Do something with cVec...  
    // everything cleans up nicely  
}
```

shared_ptr share-aware class and enable_shared_from_this

```
class BAD {  
public:  
    shared_ptr<BAD> getSP()  
    { return shared_ptr<BAD>(this); }  
  
    static shared_ptr<BAD> create()  
    { return make_shared<BAD>(); }  
};  
  
shared_ptr<BAD> beBAD() {  
    shared_ptr<BAD> p1 = BAD::create();  
    BAD* raw = p1.get();  
    shared_ptr<BAD> p2(raw->getSP());  
    return p1;  
}
```

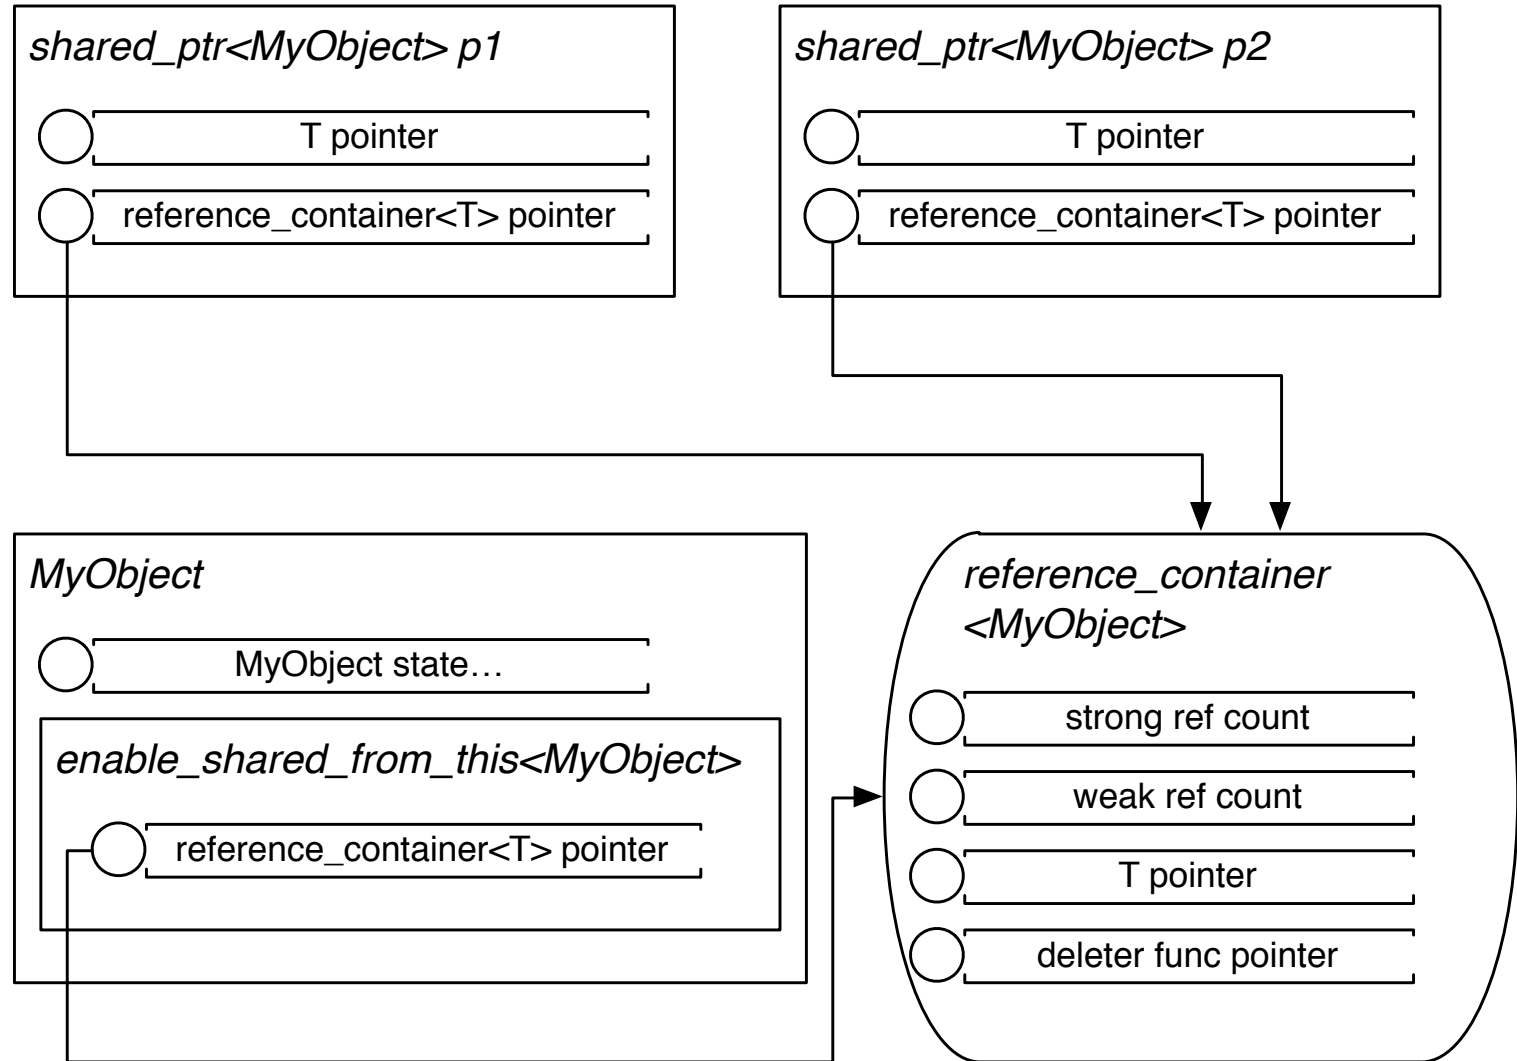
shared_ptr share-aware class and enable_shared_from_this

```
class Good: public enable_shared_from_this<Good> {
public:
    shared_ptr<Good> getSP()
    { return shared_from_this(); }

    static shared_ptr<Good> create()
    { return make_shared<Good>(); }
};

shared_ptr<Good> doGood() {
    shared_ptr<Good> p1 = Good::create();
    Good* raw = p1.get();
    shared_ptr<Good> p2 (raw->getSP());
    return p1;
}
```


Example enable_shared_from_this



shared_ptr custom deleter

```
shared_res* allocShared() {  
    // Do some magic to allocate shared resource  
    return malloc(100);  
}  
  
void releaseShared(shared_res* p) {  
    // Do some magic to release shared resource  
    free(p);  
}  
  
void testSharedResource() {  
    shared_ptr<shared_res> sr(allocShared(), releaseShared);  
    // Do something with sr...  
    // safe and automatic cleanup  
}
```

weak_ptr

- Helps break circular references
- Holds a non-owning “weak” reference
- Contained resource can only be used by requesting a shared_ptr from the weak_ptr
- A weak_ptr will not prevent resource from being destroyed (but does delay destruction of shared ref count block)

weak_ptr

```
void wpTwiddler() {  
    boost::weak_ptr<C> w1;  
    {  
        boost::shared_ptr<C> p1 = boost::make_shared<C>();  
        assert(p1.use_count() == 1);  
        w1 = p1;  
        assert(p1.use_count() == 1); // weak adds no ref  
        assert(!w1.expired()); // w1 is usable?  
        boost::shared_ptr<C> p2 = w1.lock();  
        assert(p2); // not empty  
        assert(p2.use_count() == 2);  
    }  
    assert(w1.expired() && w1.use_count() == 0);  
}
```

What is a smart pointer?

- It's a container for a pointer
 - It holds your pointer and gives it back to you
 - It manages the lifetime of the held resource
 - It facilitates RAI
 - Constructs with reasonable defaults
 - Cleans up pointer at appropriate time
- (Ideally) properly conveys interface semantics
 - Shared, unique, etc.

What is a smart pointer?

- It's a proxy for a pointer
 - It stands in for the pointer in most cases where you might use the raw pointer
 - In many cases it can transparently act as if it is the original pointer type

Various kinds of smart pointers

- “Generic” smart pointers:
 - `auto_ptr`
 - `scoped_ptr`
 - `unique_ptr`
 - `shared_ptr`
 - `weak_ptr`
 - `_com_ptr_t`
 - `CComPtr`

Various kinds of smart pointers

- Domain-specific smart pointers:
 - `std::string`
 - `_bstr_t`
 - `_variant_t`

The different kinds of reference counting

- Object-based reference counting
 - The count is stored in the object
 - Microsoft COM model
 - Sometimes referred to as “intrusive” (`boost::intrusive_ptr`)
 - Example: Microsoft COM
- Container-based reference counting
 - The count is stored in the smart pointer container
 - Example: `shared_ptr`

Object-based reference counting

- Advantages
 - Very simple smart pointer logic
 - Can easily mix with raw pointer usage
 - Object has greater visibility into its lifetime and cleanup
 - Lighter weight overall
- Disadvantages
 - Outside the COM world, less familiar
 - Outside the COM world, no ready-made solution
 - Object itself becomes slightly bigger and more complex

Container-based reference counting

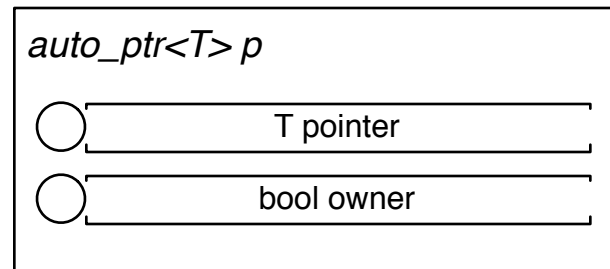
- Advantages
 - Can contain pointer to any object; no added object knowledge or complexity
 - Very standard and widely used
- Disadvantages
 - Difficult to mix with raw pointer usage
 - Heavier weight overall because of complexity of good `shared_ptr` implementation

How it works: object-based reference counting

- Sorry this talk is too short, so here's what you get:
 - Stick an atomic count in an object
 - Provide methods to add and release references
 - Object deletes itself when reference goes to zero

How it works: auto_ptr

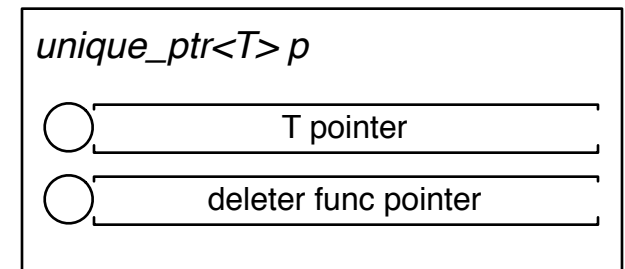
- Contains raw pointer and an ownership flag
- Ownership is transferred to last auto_ptr that takes ownership of pointer
- Problematic and error-prone ownership semantics, so it's deprecated!



How it works:

scoped_ptr and unique_ptr

- Contains raw pointer
- Does now allow transfer of ownership (except via move in unique_ptr), which simplifies logic
- unique_ptr optionally contains pointer to custom deleter func, which will be called instead of delete if desired



How it works: shared_ptr

- High level overview:
- Contains a pointer and a reference count that is shared among all instances
- Contains optional custom deleter
- Last releaser decrements reference to zero, and object is deleted

Interface Semantics

or How to Pass Your Stuff Around

- Functions that receive resources controlled by `shared_ptr` can be written with several different interfaces:
 - `foo(shared_ptr<Widget>)`
 - `foo(const shared_ptr<Widget>&)`
 - `foo(shared_ptr<Widget>&&) [C++11]`
 - `foo(Widget&)`
 - `foo(const Widget&)`
 - `foo(Widget*)`
 - `foo(const Widget*)`
- The interface you choose says a lot about what you intend for both the caller and the callee

Interface Semantics

or How to Pass Your Stuff Around

```
foo(shared_ptr<Widget>)
```

- Caller contract:
 - Caller is required to wrap object in `shared_ptr`, which implies that `Widget` object cannot be stack-based
- Callee contract:
 - Callee is allowed to take an ownership stake in object (i.e. copy `shared_ptr`)
- Other implications:
 - `shared_ptr` is copied during call, resulting in increment of atomic variable

Interface Semantics

or How to Pass Your Stuff Around

```
foo(const shared_ptr<Widget>&)
```

- Caller contract:
 - Caller is required to wrap object in `shared_ptr`, which implies that `Widget` object cannot be stack-based
- Callee contract:
 - Callee is may take an ownership stake in object (i.e. copy `shared_ptr`), but is not required to do so
- Other implications:
 - `shared_ptr` reference is not copied during function call, resulting in better performance

Interface Semantics

or How to Pass Your Stuff Around

```
foo(shared_ptr<Widget>&&) [C++11]
```

- Caller contract:
 - Caller is required to wrap object in `shared_ptr`, which implies that `Widget` object cannot be stack-based
 - Caller will receive object that has been emptied out by move operation
- Callee contract:
 - Callee must take ownership stake of object (i.e. move `shared_ptr`)
- Other implications:
 - `shared_ptr` reference is moved during function call from caller's copy to callee's copy

Interface Semantics

or How to Pass Your Stuff Around

`foo(Widget&) and foo(const Widget&)`

- Caller contract:
 - Caller is not required to wrap object in `shared_ptr`, which allows object to be allocated either on stack or heap
 - Caller retains full ownership of object
 - Object is always valid (reference)
- Callee contract:
 - Callee cannot take ownership stake of object, because `shared_ptr` container is absent
- Other implications:
 - No churn of atomic reference count

Interface Semantics

or How to Pass Your Stuff Around

`foo (Widget*)` and `foo (const Widget*)`

- Caller contract:
 - Caller is not required to wrap object in `shared_ptr`, which allows object to be allocated either on stack or heap
 - Caller retains full ownership of object
 - Object may be optional (i.e. NULL pointer)
- Callee contract:
 - Callee cannot take ownership stake of object, because `shared_ptr` container is absent
- Other implications:
 - No churn of atomic reference count
 - Prefer reference over pointer unless object must be optional

Interface Semantics

Key Points

- Choose to pass by `Widget&` or `Widget*`, unless the callee is expected to take an ownership stake in the object
 - Passing by `Widget` type does not impose container or allocation constraints on caller
 - Choose to pass by `Widget&` over `Widget*`, unless the `Widget` object is optional (may be NULL)
- Choose to pass by `const shared_ptr<Widget>&` if the callee may take ownership in object
 - `shared_ptr` reference may be copied if needed, but will not make an unnecessary copy on function entry

Summary

- Use a smart pointer. It avoids the explicit, problematic manual delete
- Prefer `unique_ptr` (or `boost::scoped_ptr` if you don't have C++11 yet) because of its efficiency and simplicity if possible
- Use `shared_ptr` for any variables that need shared semantics
- Use `make_shared` when possible; it may be more efficient, and avoids explicit `new`
- Consider object-based reference counting for cases where it fits a paradigm better