

C++ devirtualization in clang

Piotr Padlewski, University of Warsaw,
intern at Google

Why devirtualization matters?

- Because we want to inline virtual functions!
- even if we will not inline, optimizer will take advantage of direct call
- we can make binaries smaller

Everything will lead to performance boost.

How virtual calls works

```
struct A {  
    virtual void foo();  
};  
void g(A * a) {  
    a->foo();  
}
```

```
%vtable = load void (%struct.A*)**, void (%struct.A*)*** %a, align 8
```

```
%1 = load void (%struct.A*)*, void (%struct.A**) %vtable, align 8
```

```
call void %1(%struct.A* %a)
```

Old devirtualization

```
struct A {  
    virtual void foo();  
};
```

```
void f() {  
    A a;  
    a.foo();  
    a.foo();  
}
```

Both will be direct call

Old devirtualization

```
struct A {  
    virtual void foo();  
};  
void A::foo() { }  
void g(A& a) {  
    a.foo();  
}
```

```
void f() {  
    A a;  
    g(a);  
}
```

Works because g will be inlined

... and also because foo is defined in this TU.

Old devirtualization

```
struct A {  
    A();  
    virtual void foo();  
};  
void A::foo() { }  
void g(A& a) {  
    a.foo();  
}  
void f() {  
    A a;  
    g(a);  
}
```

Constructor is now outline.
Won't devirtualize because optimizer doesn't know what is stored to vptr.

Old devirtualization

```
struct A {  
    virtual void bar();  
    virtual void foo();  
};  
void A::bar() { }  
void g(A& a) {  
    a.foo();  
    a.foo();  
}  
void f() {  
    A a;  
    g(a);  
}
```

**Won't devirtualize because
optimizer doesn't know that foo
didn't change vptr.**

Old devirtualization

```
struct A {  
    virtual void foo();  
    virtual ~A() = default;  
};  
void g(A& a) {  
    a.foo();  
}  
void f() {  
    A a;  
    g(a);  
}
```

**Won't devirtualize because
vtable is now external.**

@vtable for A = external unnamed_addr constant [5 x i8*]

Calling the same virtual function

```
struct A {  
    virtual void foo();  
};  
void g(A * a) {  
    a->foo();  
    a->foo();  
}
```

```
%vtable = load (...) %a
```

```
%1 = load void (...) %vtable
```

```
call void %1(%struct.A* %a)
```

```
%vtable2 = load (...) %a
```

```
%2 = load void (...) %vtable2
```

```
call void %2(%struct.A* %a)
```

Placement new is evil

```
struct A {  
    virtual void foo();  
};  
struct B : A {  
    virtual void foo();  
};  
void A::foo() { new(this) B; } //in .cc  
void g() {  
    A *a = new A;  
    a->foo();  
    a->foo(); // Undefined behaviour  
}
```

```
struct A {  
    virtual A* foo();  
};  
struct B : A {  
    virtual A* foo();  
};  
A* A::foo() { return new(this) B; }  
void g() {  
    A *a = new A;  
    A *a2 = a->foo();  
    a2->foo(); // this is fine  
}
```

assuming type after ctor call

After calling the constructor, it would be very nice to know what is the vptr value.

```
call void @_ZN1AC1Ev(%struct.A* %a)
```

```
%vtable = load i64*, i64** %a, align 8
```

```
%cmp = icmp eq i64* %vtable, @vtable
```

```
tail call void @llvm.assume(i1 %cmp)
```

```
%vtable1 = load void (%struct.A*)**, void (%struct.A*)*** %a, align 8
```

available_externally vtables

@vtable for A = external unnamed_addr constant [6 x i8*]

@vtable for a = available_externally unnamed_addr constant [6 x i8*] (

definition)

Unfortunately it is not safe to generate available_externally vtables in cases like:

- class has inline virtual functions
- class has hidden virtual functions - `__attribute__((visibility("hidden")))`

propagating “constness” of vptr

```
struct A {  
    virtual void foo();  
};  
void g(A * a) {  
    a->foo();  
    a->foo();  
}
```

```
%vtable = load (...) %a
```

```
%1 = load void (...) %vtable
```

```
call void %1(%struct.A* %a)
```

```
%vtable2 = load (...) %a
```

```
%2 = load void (...) %vtable2
```

```
call void %2(%struct.A* %a)
```

propagating “constness” of vptr

```
%vtable = load (...) %a
```

```
%1 = load void (...) %vtable
```

```
call void %1(%struct.A* %a)
```

```
%vtable2 = load (...) %a
```

```
%2 = load void (...) %vtable2
```

```
call void %2(%struct.A* %a)
```

```
%vtable = load (...) %a
```

```
%1 = load void (...) %vtable
```

```
call void %1(%struct.A* %a)
```

```
%2 = load void (...) %vtable
```

```
call void %2(%struct.A* %a)
```

introducing !invariant.group

```
%vtable = load (...) %a, !invariant.group !0
```

```
%1 = load void (...) %vtable
```

```
call void %1(%struct.A* %a)
```

```
%vtable2 = load (...) %a, !invariant.group !0
```

```
%2 = load void (...) %vtable2
```

```
call void %2(%struct.A* %a)
```

Dealing with placement new

```
int main() {  
    MyClass c;  
    c.foo();  
    auto c2 = new (&c) MyOtherClass();  
    c2->foo();  
    c2->~MyOtherClass();  
    new (&c) MyClass();  
}
```


introducing invariant.group.barrier

```
auto c2 = new (&c) MyOtherClass();
```

```
call void @_ZN7MyClassD1Ev(%struct.MyClass* nonnull %c) #1
```

```
%2 = bitcast %struct.MyClass* %c to %struct.MyOtherClass*
```

```
%3 = call @llvm.invariant.group.barrier(%2)
```

introducing -fstrict-vtable-pointers

For reasons like:

- LTO between code with **!invariant.group** and without is unsafe
- Optimizers don't know how to deal with **invariant.group.barrier**
- **UBSan** doesn't check for UBs from changing dynamic type
- Whole thing is still experimental.

Corner cases

```
void g() {  
    A *a = new A;  
    a->foo();  
    A *b = new(a) B;  
    assert(a == b);  
    b->foo();  
}
```

After assert llvm knows that `a == b` and can replace one with another.

Results

Unknown...

We know that devirtualization happen.

**But the project is not yet done, there are so many things that we can still do, and the largest thing will come with !invariant.
group.**

Sources

Devirtualization paper:

<https://goo.gl/r5dOxR>

How devirtualization works in gcc (*Jan Hubička*)

<http://hubicka.blogspot.com/2014/01/devirtualization-in-c-part-1.html>

<https://www.youtube.com/watch?v=oN15LjHX9xY>

Thank you