

`completion<T>`

Improving the `future<T>` with monads

Travis Gockel

C++ Pirate
SolidFire

22 September 2015

Outline

1 Introduction

- SolidFire
- Why a different `future<T>?`

2 `completion<T>`

- `expected<T, E>`
- `completion<T>`
- `completion_promise<T>`
- Edge Cases

3 Extra

- Aggregation functions
- Workers/Executors
- Implementation

Outline

1 Introduction

- SolidFire
- Why a different `future<T>`?

2 `completion<T>`

- `expected<T, E>`
- `completion<T>`
- `completion_promise<T>`
- Edge Cases

3 Extra


- Aggregation functions
- Workers/Executors
- Implementation

Who is SolidFire?

- Distributed all-flash block storage vendor
 - Data loss or unavailability is unacceptable
 - ...even in the face of hardware failure
 - Deployed as an appliance – software must maintain itself
 - ...in environments we do not control
- Small engineering team (\approx 40 C++ programmers)
 - Team efficiency is critical
 - Growing – code can't look completely foreign

The Problem Space

- Request for data comes into the transport service over iSCSI
- Look up the internal IDs for the request
- Sometimes, the data is in the local cache
- For those that are not, we fetch associated chunks of data from other services
→ User-visible latency is slowest-responding service



Outline

1 Introduction

- SolidFire
- Why a different `future<T>?`

2 `completion<T>`

- `expected<T, E>`
- `completion<T>`
- `completion_promise<T>`
- Edge Cases

3 Extra

- Aggregation functions
- Workers/Executors
- Implementation

What is the goal of `future<T>`?

- A mechanism for asynchronous communication between a provider (`promise<T>`) and a receiver (`future<T>`)
- Make concurrent programming easier

std::future<T>

```
template <typename T> class future {  
public:  
    T get();  
    bool valid() const noexcept;  
    void wait() const;  
    template <typename TRep, typename TPeriod>  
    future_status wait_for(const chrono::duration<TRep,  
        TPeriod>& rel_time) const;  
    template <typename TClock, typename TDuration>  
    future_status wait_until(const  
        chrono::time_point<TClock, TDuration>& abs_time)  
        const;  
    shared_future<T> share();  
};
```


How do I use this?

```
buffer read_multi(location x, location y) {  
    future<buffer> fut_z1 = async(&read, x);  
    buffer z2 = read(y);  
    return fut_z1.get() + z2;  
}
```

How do I use this?

```
buffer read_multi(location x, location y) {  
    future<buffer> fut_z1 = async(&read, x);  
    buffer z2 = read(y);  
    return fut_z1.get() + z2;  
}
```

Only sort of asynchronous.

What is wrong with `future<T>`?

get **and** wait

What is wrong with `future<T>`?

get **and** wait

Should never be used.

What about `then<F>`?

```
template <typename F>
auto then(F&& continuation)
    -> future<decltype(continuation(*this))>;
```

Usage...

```
future<int>    x = async(&thing);
future<double> y = x.then([] (future<int> f) { return
    double(f.get()); });
```

Awkward for Aggregation

```
std::vector<std::future<int>> all_tasks = foo();  
std::future<std::vector<std::future<int>>> all_done  
    = when_all(make_move_iterator(begin(all_tasks)),  
               make_move_iterator(end(all_tasks))  
               );
```

- Which `future<T>::get` will block?
- Slow – `future<T>` is multithreaded, so it *must* deal with locking
→ Adding the values in `std::vector<int>` is much faster than adding `std::vector<std::future<int>>`

```
try {  
    try {  
        try {  
            foo();  
        } catch (...) {  
            throw;  
        }  
        bar();  
    } catch (...) {  
        throw;  
    }  
    baz();  
} catch (const  
    std::exception& ex) {  
    report(ex);  
}
```

```
try {  
    foo();  
    bar();  
    baz();  
} catch (const  
    std::exception& ex) {  
    report(ex);  
}
```

```
try {  
    try {  
        try {  
            foo();  
        } catch (...) {  
            throw;  
        }  
        bar();  
    } catch (...) {  
        throw;  
    }  
    baz();  
} catch (const  
    std::exception& ex) {  
    report(ex);  
}
```

```
foo_async()  
    .then([] (std::future<void> x)  
        {  
            x.get(); return bar();  
        })  
    .then([] (std::future<void> y)  
        {  
            y.get(); return baz();  
        })  
    .then([] (std::future<void> z)  
        {  
            try {  
                z.get();  
            } catch (const  
                std::exception& ex) {  
                report(ex);  
            }  
        })  
    );
```


Outline

- 1 Introduction
 - SolidFire
 - Why a different `future<T>`?

- 2 `completion<T>`
 - `expected<T, E>`
 - `completion<T>`
 - `completion_promise<T>`
 - Edge Cases

- 3 Extra
 - Aggregation functions
 - Workers/Executors
 - Implementation

Error Reporting

- `future<T>` combines two ideas: asynchronous communication *and* error reporting.
- Can we split this?

expected<T, E>

“N4015: A proposal to add a utility class to represent expected monad.”

```
template <typename T, typename E> class expected {
public:
    auto status() const -> bool;
    auto get() -> T;
    template <typename F> auto map(F f) ->
        expected<decltype(f(std::declval<T>()))>;
    template <typename F> auto catch_error(F f)
        -> expected<std::common_type_t<T,
            decltype(f(std::declval<std::exception_ptr>()))>>;
    // ...
};
```

expected<T, E>

```
expected<int>    x = foo();  
expected<double> y = x.map([] (int a) { return  
    double(a); });  
expected<double> z = y.catch_error([] (auto) { return  
    0.0; });  
expected<string> w = z.map([] (double a) { return  
    to_string(a); });  
cout << w.get() << endl;
```

Outline

- 1 Introduction
 - SolidFire
 - Why a different `future<T>`?

- 2 `completion<T>`
 - `expected<T, E>`
 - `completion<T>`
 - `completion_promise<T>`
 - Edge Cases

- 3 Extra
 - Aggregation functions
 - Workers/Executors
 - Implementation

completion<T>

```
template <typename T> class completion {  
public:  
    auto status() const -> completion_status;  
    auto on_complete(F f) -> void;  
    template <typename F> auto then(F f) ->  
        completion<decltype(f(std::declval<expected<T>>()))>;  
    template <typename F> auto map(F f) ->  
        completion<decltype(f(std::declval<T>()))>;  
    template <typename F> auto catch_error(F f)  
        -> completion<std::common_type_t<T,  
            decltype(f(std::declval<std::exception_ptr>()))>>;  
    auto disable() -> void;  
};
```

Missing features?

get and wait Why bother?


But we can for compatibility with `std::future<T>`

Overloads for `T&` and `void`? Not needed – let `expected<T>` deal with that

completion_data<T>

```
template <typename T> struct completion_data {  
    mutable spin_mutex          protect;  
    completion_status           status;  
    optional<expected<T>>      value;  
    function<void (expected<T>&&)> callback;  
  
    reference_tracking_data<completion_data> tracker;  
};
```


completion_status Transitions



completion<T>::then(Func&&)

```
template <typename T>
template <typename Func>
completion<T>::then(Func&& func)
->
    completion<decltype(func(std::declval<expected<T>>()))>
```

Purpose: Call func with an expected<T>

- If status is has_value, call func immediately
- If status is no_value, store func for later use
- All other cases → error

Arg func: U (*) (expected<T>&&)

completion<T>::on_complete(Func&&)

```
template <typename T>
template <typename Func>
void completion<T>::on_complete(Func&& func);
```

Purpose: Chaining the last step in the process

Arg func: void (*) (expected<T>&&)

completion<T>::map(Func&&)

```
template <typename T>
template <typename Func>
auto completion<T>::map(Func&& func);
    -> completion<decltype(func(std::declval<T>()))>
```

Purpose: Similar to `then(F&&)`, *but* checks the status of `expected<T>`

- If completed in success, call the function
- If completed in failure, forward the `exception_ptr` to the next step
→ Not throwing is significantly faster!

Arg func: `U (*) (T&&)`

Note: `completion<T>::catch_error(Func&& func)` has the opposite rules

completion<T>::get()

```
template <typename T>
T
completion<T>::get() {
    std::promise<expected<T>> promise;
    on_complete([&promise] (expected<T>&& x) {
        promise.set_value(std::move(x));
    });
    return promise.get_future().get().get();
}
```

Only incur the slow `mutex` if somebody asks for the slow call.

→ The `wait` family requires internal tricks, but is still implementable.

Outline

- 1 Introduction
 - SolidFire
 - Why a different `future<T>`?

- 2 `completion<T>`
 - `expected<T, E>`
 - `completion<T>`
 - `completion_promise<T>`
 - Edge Cases

- 3 Extra
 - Aggregation functions
 - Workers/Executors
 - Implementation

completion_promise<T>

- Same role as `std::promise<T>`
- Allows for the delivery of values to the `completion`
- Shares the `completion_data`

completion_promise<T>::complete(expected<T> value);

```
template <typename T>
void completion_promise<T>::complete(expected<T> value);
```

Purpose: Deliver a value to the associated completion<T>

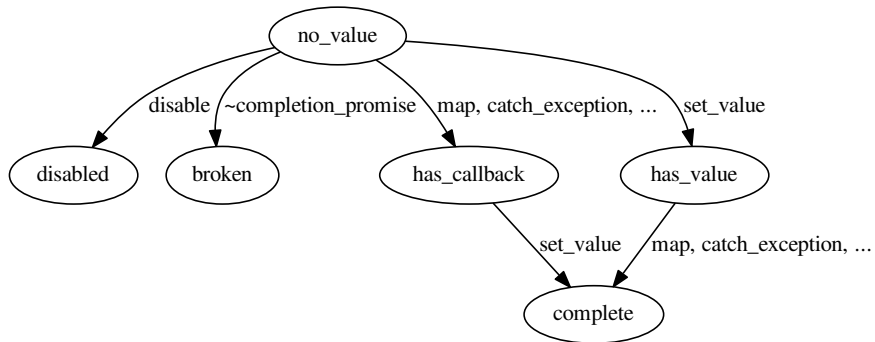
- If status is `has_callback`, call it immediately
- If status is `no_value`, store value for later retrieval
- If status is `disabled`, do nothing
- All other cases → error

std::promise<T> Compatibility Functions

```
template <typename... U>
void set_value(U&&... args) {
    complete(make_expected(std::forward<U>(args) ...));
}

void set_exception(std::exception_ptr ex) {
    complete(make_unexpected(std::move(ex)));
}
```

completion_status Transitions



Outline

- 1 Introduction
 - SolidFire
 - Why a different `future<T>`?

- 2 `completion<T>`
 - `expected<T, E>`
 - `completion<T>`
 - `completion_promise<T>`
 - **Edge Cases**

- 3 Extra
 - Aggregation functions
 - Workers/Executors
 - Implementation

Dealing with Broken Promises

When the `completion_promise` is destroyed without having set the value, what should happen?


Do nothing? Potentially causes deadlock elsewhere.

Complete with a `broken_promise`? This can trigger a chain of actions, thus blocking in the destructor.

throw? Throwing in a destructor is bad.

`assert(false)`? Panic – something terrible has happened.

Disabling



- Disable *H* – should *F* be disabled?
 - Okay if *F* has no side-effects
 - But how can I know?
→ Allow opt-in
- Disable *C* – what about *D*, *E* and *F*?
 - Only allow disabling leaves (do nothing)
- Is this useful?
 - Perhaps “disabling” should be handled by a different construct

Declarative “Pureness”

```
template <typename T>
template <typename Func>
auto completion<T>::then(Func&& func) -> ...
{
    bool is_pure = continuation_traits<Func>::is_pure;
    // ...
}
```

```
template <typename F> struct pure_func_wrapper { ... };
template <typename F> pure_func_wrapper<F>
    pure_func(F&&);
```

Better: Pick up `__attribute__((pure_function))` automatically?

Outline

- 1 Introduction
 - SolidFire
 - Why a different `future<T>?`

- 2 `completion<T>`
 - `expected<T, E>`
 - `completion<T>`
 - `completion_promise<T>`
 - Edge Cases

- 3 Extra
 - Aggregation functions
 - Workers/Executors
 - Implementation

Introduction

- Algorithms are range-based where possible

concat

```
template <typename... T>
completion_concat_result_t<completion<T>...>
concat(completion<T>&&...);
```

Purpose: Combines heterogeneous completions and notifies you when all of them are completed.

Complexity: $O(n)$

Return Type: Similar rules to `tuple_cat`

concat usage

```
template <typename... T>
using expected_tuple = tuple<expected<T>...>;

completion<int> a = foo();
completion<void> b = bar();
completion<tuple<double, string>> c = baz();

completion<expected_tuple<int, void, double, string>> d =
    concat(move(a), move(b), move(c));
// Do things with d
```

for_each

```
template <typename TInputRange, typename FUnitFunc>
completion<void>
for_each(TInputRange&& input,
        const FUnitFunc& func
        );
```

Purpose: Call a callback for some completions and notifies you when all callbacks have completed running.

Complexity: $O(n)$

for_each usage

```
std::vector<completion<int>> input = foo();  
completion<void> all_done = for_each(std::move(input),  
    bar);
```

collect

```
template <template <typename...> class TOutput,  
          template <typename...> class TInput,  
          typename  
          >  
completion<TOutput<expected<T>>>  
collect (TInput<completion<T>>& input);  
  
template <typename  
          T,  
          template <typename...> class TInput  
          >  
completion<TInput<expected<T>>>  
collect (TInput<completion<T>>& input);
```

Purpose: Transform a `collection<completion<T>>` into a `completion<collection<T>>`

Complexity: $O(n)$

collect

```
template <typename T,  
         template <typename...> class TTInput  
         >  
completion<TTInput<expected<T>>>  
collect (TTInput<completion<T>>& input) {  
    return collect<TTInput>(input);  
}
```

collect_n

```
template <template <typename...> class TOutput,  
          template <typename...> class TInput,  
          typename  
          T  
        >  
completion<pair<TOutput<std::pair<size_t, expected<T>>>>  
collect_n(TInput<completion<T>>& input, size_t min);  
  
template <typename  
          T,  
          template <typename...> class TInput  
        >  
completion<pair<TOutput<std::pair<size_t, expected<T>>>>  
collect_n(TInput<completion<T>>& input, size_t min);
```

Purpose: Transform a `collection<completion<T>>` into a `completion<collection<T>>`

Complexity: $O(n)$

Arg min: Minimum amount of items to wait for

sequence

```
template <template <typename...> class TOutput,  
          template <typename...> class TInput,  
          typename  
          >  
completion<TOutput<expected<T>>>  
sequence (TInput<completion<T>>& input);  
  
template <typename  
          T,  
          template <typename...> class TInput  
          >  
completion<TInput<expected<T>>>  
sequence (TInput<completion<T>>& input);
```

Purpose: Transform a `collection<completion<T>>` into a `completion<collection<T>>` *in the same order*.

Complexity: $O(n \log n)$ (data structure dependent)

Outline

- 1 Introduction
 - SolidFire
 - Why a different `future<T>?`

- 2 `completion<T>`
 - `expected<T, E>`
 - `completion<T>`
 - `completion_promise<T>`
 - Edge Cases

- 3 **Extra**
 - Aggregation functions
 - **Workers/Executors**
 - Implementation

Who performs the next step?

- When binding a continuation, who should perform the action?
- It depends on `completion::status...`
 - `has_value`: The binder invokes immediately
 - `no_value`: The deliverer will invoke (later)
- This can be problematic (sometimes)
 - If the deliverer is on a “never block” thread (example: RPC message delivery), we can quickly lose control
 - If the receiver is client-facing, immediately invoking is “unfair” – the result is usually strange-looking latency in unrelated activities
- Enter the `executor`!

What is an Executor?

- Decides when and where to perform some action
- Launch policy describes if we should execute now or later
 - `immediate`: Always execute the next step inline
 - `async`: Always post the task to work later (potentially on a different thread)
 - `sharing`: If the executor has queued work, post to work later; otherwise, execute immediately
- Executors use various methods of determining what order to run queued tasks
 - `fifo`: Run tasks in almost the same order they were queued
 - `fifo_strict`: Run tasks in exactly the same order they were queued
 - `qos`: Use information about the task to prioritize (requires task introspection)
 - `random`: Run in a random order (useful for testing)

Informing the next step

```
template <typename Func>
auto then_with(execute_token exec, Func&& continuation)
    -> ...;
template <typename Func>
auto map_with(execute_token exec, Func&& continuation)
    -> ...;
// etc.
```

What is this `execute_token`?

- **executor-defined data structure** for informing it how to run
→ Often just a `shared_ptr<executor>`
- **`execute_tokens` are created** through the base `executor::get_token()` method or through implementation-specific named functions
Example: `priority_worker::get_priority_token(int)`

QoS-based Execution

```
execute_token token = worker->create_job_token(volume);  
read_metadata(volume, lba, length)  
    .map_with(token, check_cache, _1)  
    .map_with(token, remote_data_fetch, m_rpc, _1)  
    ...
```

- Create a new token at the beginning of each user request (really: draw them from a pool)
- Prioritization happens at the beginning of each “next step” – if work units are small enough, we do not need to support reprioritization
→ Reprioritization isn't beneficial in our use, anyway

Outline

- 1 Introduction
 - SolidFire
 - Why a different `future<T>`?

- 2 `completion<T>`
 - `expected<T, E>`
 - `completion<T>`
 - `completion_promise<T>`
 - Edge Cases

- 3 **Extra**
 - Aggregation functions
 - Workers/Executors
 - **Implementation**

Note

- The remainder of this presentation
- These are meant only as *examples* to demonstrate the idea – they are not “production-ready”

completion<T>::on_complete(F&& func)

```
template <typename F> void on_complete(F&& func) {
    unique_lock lock(impl->protect);
    if (impl->status == completion_status::no_value) {
        impl->callback = std::forward<F>(func);
        impl->status    = completion_status::has_callback;
    } else if (impl->status ==
        completion_status::has_value) {
        auto completer = on_scope_exit([this] { impl->status
            = completion_status::complete; });
        std::forward<F>(func)(std::move(impl->value));
    } else {
        throw std::logic_error("invalid_status");
    }
}
```


completion<T>::then(F&& func)

```
template <typename F>
then_result_t<F> then(F&& func) {
    unique_lock lock(impl->protect);
    if (impl->status == completion_status::no_value) {
        result_promise promise;
        impl->callback = [promise, func = std::forward<F>(func)] (expected<T>&& x) {
            promise.complete(std::move(func), std::move(x));
        };
        impl->status = completion_status::has_callback;
        return promise.get_completion();
    } else if (impl->status == completion_status::has_value) {
        auto completer = on_scope_exit([this] { impl->status = completion_status::complete; });
        return result_promise::create_completed(try_to(std::move(func), std::move(impl->value)))
            .get_completion();
    } else {
        throw std::logic_error("invalid_status");
    }
}
```

completion<T>::map(F&& func)

```
template <typename F>
completion_map_result_t<completion, F>
map(F&& func) {
    return then([func = std::forward<F>(func)]
        (expected<T> x) {
            return std::move(x).map(func).get();
        });
}
```

Yuck! Should set the result directly instead of relying on `get`.

for_each

```
template <typename TInputRange, typename FUnitFunc>
completion<void>
for_each(TInputRange&& input, const FUnitFunc& func) {
    completion_promise<void> promise;
    auto remaining =
        make_shared<atomic<size_t>>(distance(input));
    for (auto&& step : input)
        step.then(func)
            .on_complete([promise, remaining]
                (expected<void>&&) mutable {
                    if (remaining->fetch_sub(1,
                        std::memory_order_relaxed) ==
                        1)
                        promise.set_result();
                });
    return promise.get_completion();
}
```

for_each

```

template <typename TInputRange, typename FUnitFunc>
completion<void>
for_each(TInputRange&& input, const FUnitFunc& func) {
    completion_promise<void> promise;
    auto remaining =
        make_shared<atomic<size_t>>(distance(input));
    for (auto&& step : input)
        step.then(func)
            .on_complete([promise, remaining]
                (expected<void>&&) mutable {
                    if (remaining->fetch_sub(1,
                        std::memory_order_relaxed) ==
                        1)
                        promise.set_result();
                });
}

```

Why `std::memory_order_relaxed`?

The user-defined `func` is called under a different completion action, which happens under a lock, so the action is sequentially-consistent.

collect

```

template <template <typename...> class TOutput,
          template <typename...> class TInput,
          typename T
        >
completion<TOutput<expected<T>>>
collect (TInput<completion<T>>& input)
{
    auto collected = make_shared<pair<spin_mutex, TOutput<expected<T>>> >();
    auto done = for_each(input,
                          [collected] (expected<T>&& x) {
                              spin_lock lock(collected->first);
                              collected->second.insert(end(collected->second), std::move(x))
                          })
        .then([collected] (expected<void>) { return std::move(collected->second); });

    #if COMPLETION_COLLECT_SHUFFLE
    done = done.map([] (auto x) {
        random_shuffle_if_possible(x);
        return x;
    });
    #endif
    return done;
}

```

completion_promise<T>::complete(expected<T>)

```
template <typename U> void complete(expected<U> value) {
    if (impl->status == completion_status::no_value) {
        impl->value = std::move(value);
        impl->status = completion_status::has_value;
    } else if (impl->status ==
        completion_status::has_callback) {
        auto completer = on_scope_exit([this] { impl->status
            = completion_status::complete; });
        impl->callback(std::move(value));
    } else if (impl->status ==
        completion_status::disabled) {
        // do nothing?
    } else {
        throw std::logic_error("invalid_status");
    }
}
```

concat

```
template <size_t KInputCount, typename... TResultArgs>
struct concat_completer {
    using result_type = tuple<TResultArgs...>;

    atomic<size_t>                remaining{KInputCount};
    result_type                    values;
    completion_promise<result_type> promise;

    template <size_t KStartIdx, typename... TValues>
    void complete_part(integral_constant<size_t,
        KStartIdx>, TValues&&... values);

    template <size_t... DestIdxs, typename... TValues>
    void set_data(index_tuple<DestIdxs...>, TValues&&...
        values);
};
```

concat

```
template <size_t... DestIdxs, typename... TValues>
void set_data(index_tuple<DestIdxs...>, TValues&&...
    values) {
    std::tie(std::get<DestIdxs>(values)...) =
        std::make_tuple<TValues...>(values...);
}
```


concat

```
template <size_t KStartIdx, typename... TValues>
void complete_part(integral_constant<size_t, KStartIdx>,
    TValues&&... values) {
    set_data(make_index_tuple_range<KStartIdx,
        sizeof...(TValues)>(),
        std::forward<TValues>(values)...);

    if (remaining.fetch_sub(1, std::memory_order_seq_cst)
        == 1)
        deliver_promise();
}
```

concat

```
template <size_t KStartIdx, typename... TValues>
void complete_part(integral_constant<size_t, KStartIdx>,
    TValues&&... values) {
    set_data(make_index_tuple_range<KStartIdx,
        sizeof...(TValues)>(),
        std::forward<TValues>(values)...);

    if (remaining.fetch_sub(1, std::memory_order_seq_cst)
        == 1)
        deliver_promise();
}
```

Why `std::memory_order_seq_cst`?

Before delivering the promise, we must make sure that other threads which are *not* calling `deliver_promise` have their memory modifications from `set_data` visible to all other threads.