

# 3D Face Tracking and Reconstruction using Modern C++

Patrik Huber

Centre for Vision, Speech and Signal Processing  
University of Surrey, UK

[patrikhuber@gmail.com](mailto:patrikhuber@gmail.com) | [www.patrikhuber.ch](http://www.patrikhuber.ch)




# The Centre for Vision, Speech and Signal Processing



- University of Surrey, Guildford, UK
- One of the largest UK computer vision, audio-visual signal processing and multimedia communication groups
- Founded 1986
- 120 people (70 PhDs)
- Grant portfolio £12M



# The Centre for Vision, Speech and Signal Processing



Abnormalities due to Alzheimer's



# The Centre for Vision, Speech and Signal Processing

Live Event




S3A



Listeners at home



# The Centre for Vision, Speech and Signal Processing



<http://www.bbc.co.uk/taster/projects/dancer>



# A few words about myself

- PhD student in computer vision
- Working on faces, 3D face models from 2D images
- Very interested in modern C++
- Open source software:
  - infrequent bug reports, recently started to contribute to OpenCV
  - recently released quite a few bits of our research on GitHub




# Face tracking and 3D face reconstruction


2D image or video



Find the face and  
landmark points



3D face representation



Applications

Frontalisation  
Recognition  
Expression Analysis  
HCI  
Animation

- Landmark detection & generic library
  - Algorithm, C++ library, interesting C++ bits


- 3D face reconstruction



Main goal for us:  
Make it easier to work with 3D models



# Regression-based landmark detection



# Regression-based landmark detection


- Learn «shape-update» to ground-truth location

$$\delta s = \mathbf{A}_n \mathbf{f}(\mathbf{I}, s)$$

Shape update (x, y)

Learned regressor

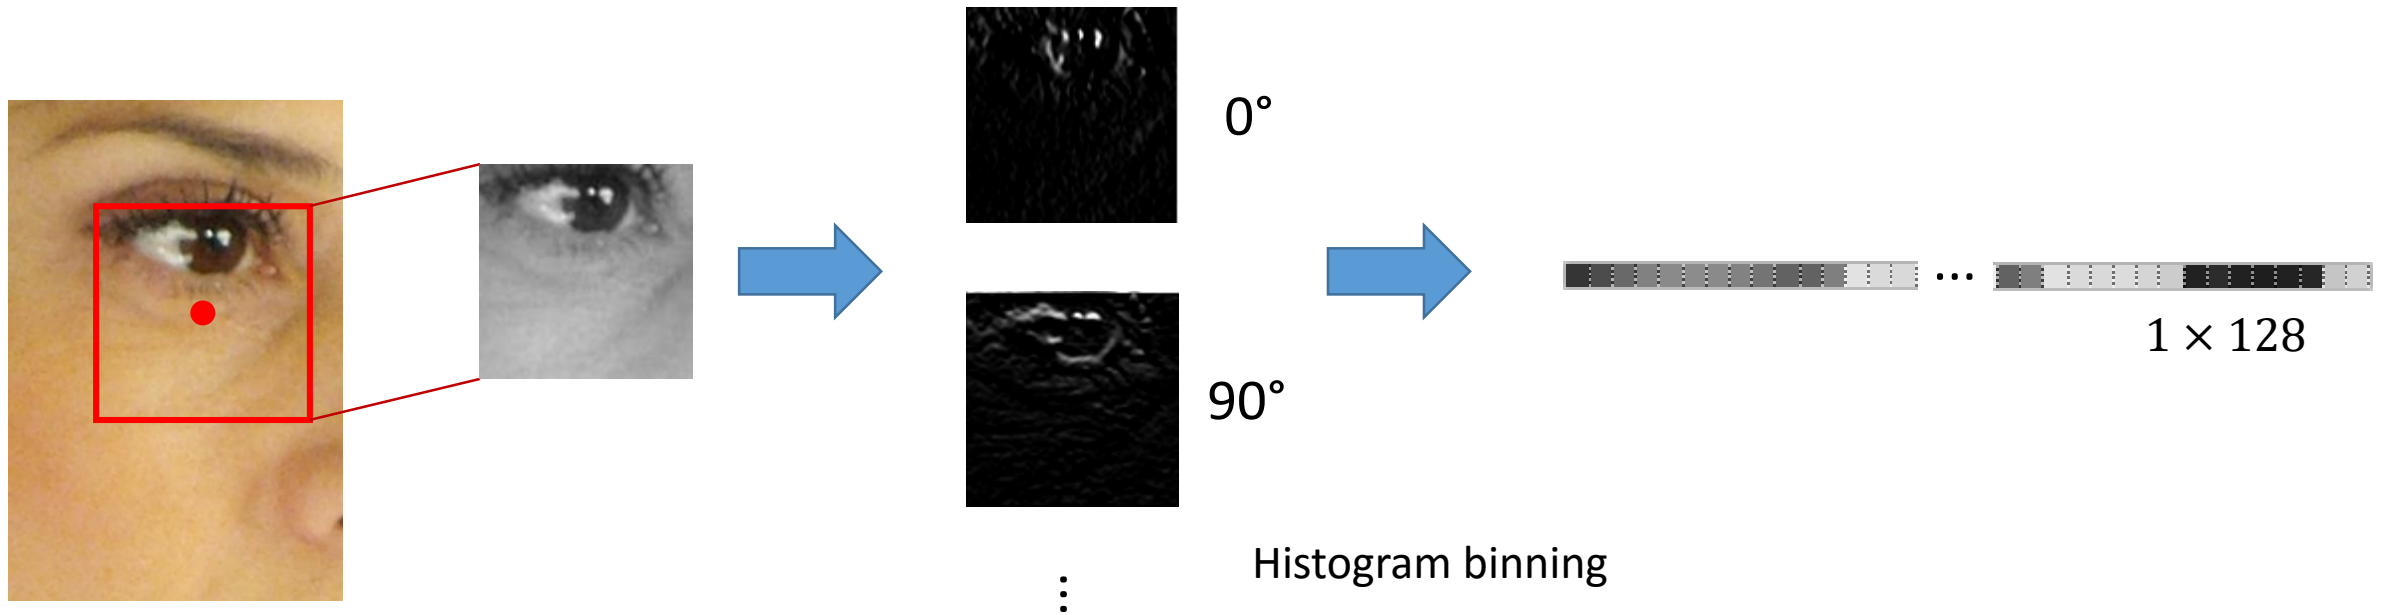
Feature extraction



*Supervised Descent Method and Its Applications to Face Alignment*, X. Xiong and F. De la Torre, CVPR 2013

# A better representation: Local image features

- Feature extraction: Histogram of Oriented Gradients (HOG)



(simplified diagram)

# Learning the model

- Learn using a bunch of training data  
(given images and landmarks)

$$\begin{bmatrix} \delta s \end{bmatrix} = \begin{bmatrix} \mathbf{A}_n \end{bmatrix} \begin{bmatrix} \mathbf{f}(\mathbf{I}, s) \end{bmatrix}$$

# Generic supervised descent

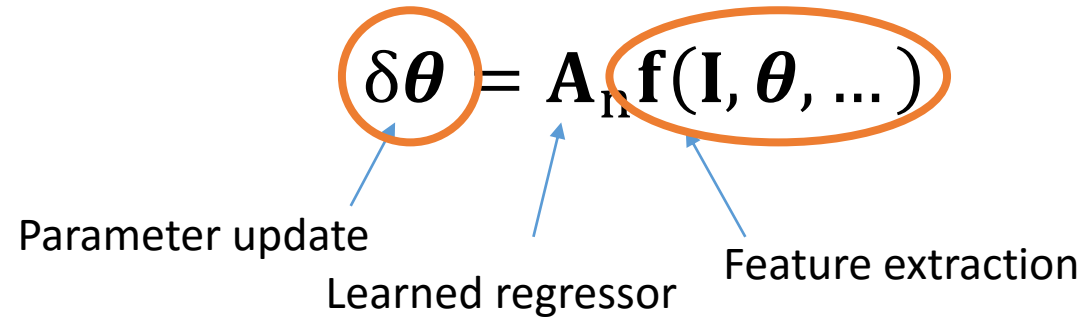
- Generic formulation

$$\delta\theta = A_n f(I, \theta, \dots)$$

Parameter update

Learned regressor

Feature extraction

The diagram shows the equation  $\delta\theta = A_n f(I, \theta, \dots)$ . Three orange ovals highlight the terms  $\delta\theta$ ,  $A_n$ , and  $f(I, \theta, \dots)$ . Blue arrows point from these ovals to the labels 'Parameter update', 'Learned regressor', and 'Feature extraction' respectively.

We might also estimate  
3D head pose parameters  
 $\theta = [R_x, R_y, R_z]$

- In previous case:  $\theta = [x_1, \dots, x_n, y_1, \dots, y_n]$

# Modeling $f(\dots)$

- Model projection  $f(\dots)$  as a function in the code
- Can be a lambda in trivial case
- Or function object with state or additional data in more complex scenarios

# Hello world example

- $\mathbf{f}(\mathbf{I}, \mathbf{x})$  for 2D landmark detection:


```
class HogTransform {
public:
    HogTransform(vector<Mat> images, ...HOG parameters...) { ... };

    Mat operator()(Mat parameters, size_t regressor_level, int training_index = 0) {
        // shortened, to get the idea across:
        Mat hog_descriptors = extract_hog_features(images[training_index], parameters);
        return hog_descriptors;
    }
private:
    vector<Mat> images;
};
```



# Interlude – OpenCV matrix class

- `cv::Mat`
  - Reference-counted
  - «Header» contains: rows, cols, flags
  - Images = matrices



CV\_8UC4  
CV\_32SC1  
CV\_64FC3

```
cv::Mat m = cv::Mat::ones(5, 5, CV_32FC1);  
cv::Mat sub = m.rowRange(0, 2); // no data copy  
  
float element = m.at<float>(1, 0);
```

# *Hello world* example

- $\mathbf{f}(\mathbf{I}, \mathbf{x})$  for 2D landmark detection:

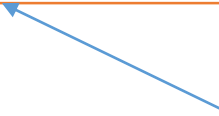
```
class HogTransform {
public:
    HogTransform(vector<Mat> images, ...HOG parameters...) { ... };

    Mat operator()(Mat parameters, size_t regressor_level, int training_index = 0) {
        // shortened, to get the idea across:
        Mat hog_descriptors = extract_hog_features(images[training_index], parameters);
        return hog_descriptors;
    }
private:
    vector<Mat> images;
};
```

# Building blocks overview

```
template<class Solver = PartialPivLUSolver>  
class LinearRegressor
```

Solves " $y=mx+b$ "



```
template<class RegressorType>  
class SupervisedDescentOptimiser
```

```
-template<class ProjectionFunction>  
  void train(cv::Mat parameters, cv::Mat initialisations,  
             ProjectionFunction projection)
```

# *Hello world* example

- Training the model:

```
vector<cv::Mat> training_images = ...;
cv::Mat training_landmarks = ...; // each row is the landmarks for one image, i.e. [x_0, x_1, ...]
cv::Mat initialisations = ...; // generated initialisations

HogTransform hog(training_images, ...HoG parameters...);

vector<LinearRegressor<>> regressors(5);
supervised_descent_optimiser<LinearRegressor<>> model(regressors);


auto print_residual = [&training_landmarks](const Mat& current_predictions) {
    cout << cv::norm(current_predictions, training_landmarks, cv::NORM_L2)
        / cv::norm(training_landmarks, cv::NORM_L2) << endl;
};

model.train(training_landmarks, initialisations, hog, print_residual);

// store the model.
```

# Using the model

```
landmark_model<LinearRegressor<>> model;  
{  
    std::ifstream f("model.bin", std::ios::binary);  
    cereal::BinaryInputArchive archive(f);  
    archive(model);  
}  
  
cv::Mat image = cv::imread(imagefile);  
auto landmarks = model.detect(image, hog, facebox);
```




# From *hello world* to the real world

- Perturbing the initialisations to get more diverse training data
- Learn the x/y update in normalised coordinates, not in pixels
- Change the size of the extraction window with each iteration



# Using the generic code

- Estimate 3D face model parameters:
  - $\theta = [r_x, r_y, r_z, t_x, t_y, t_z, \alpha_0, \alpha_1]$
  - $\mathbf{f}(\mathbf{I}, \theta)$  performs a projection from 3D to 2D using the current  $\theta$



# C++ challenges

# Solving the linear system of equations

- Matlab: `x = A \ b;`
  - Fast
  - Parallelised
  - Automatically chooses the best suitable algorithm (can be pseudo-inverse)
  - Warns if the system is not well-conditioned
- Our «A» can be 15000 x 8000 or larger

# Solving the linear system of equations

- OpenCV `inv()`?
- Eigen?
  - PartialPivLU?
  - FullPivLU?
  - ColPivHouseholderQR?
  - LDLT?
  - `jacobiSvd`?
  - ...
- ...

# Importing models from Matlab

- Not fun!
- Matlab SDK (ugly, hard to use C API)
- Text files and manual parsing

# Storing & distributing the models in C++

- Simple task: Train a model on a local PC, upload it to repo, everybody should be able to load it
- Don't do by hand, use a serialisation library
- XML and text format out of the question

# Storing & distributing the models in C++

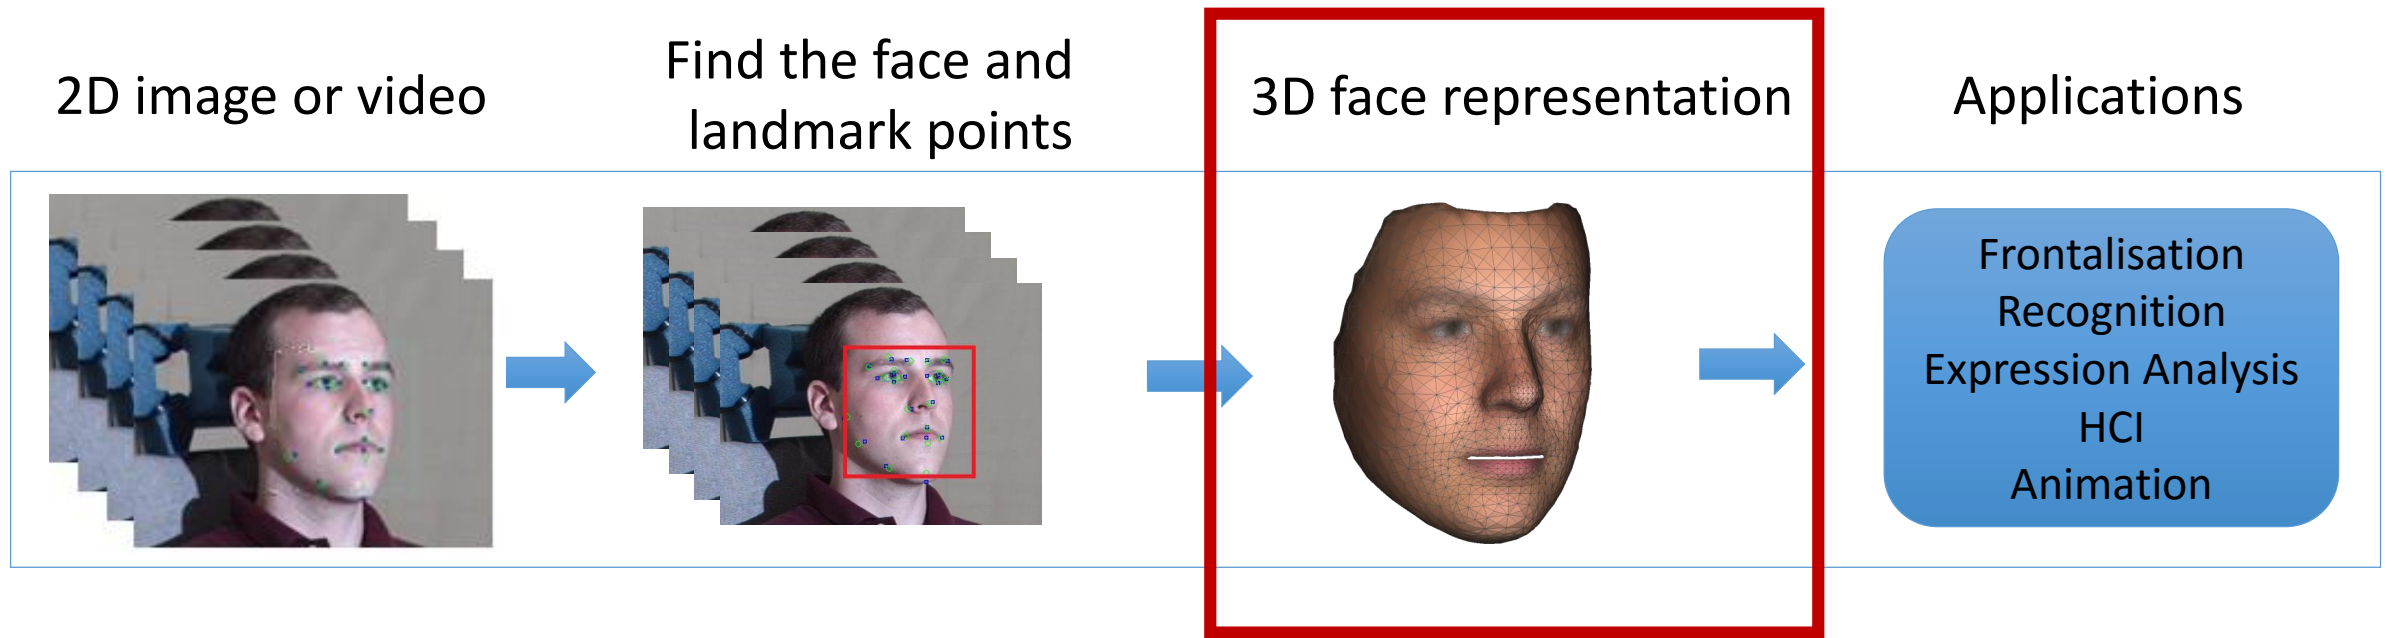
- Boost serialization:
  - It's awesome!
  - But: Files stored with version x can't be loaded with version <x
  - Older versions don't compile on VS2015
- cereal:
  - C++11 library for serialization, header-only  
(<https://github.com/USCiLab/cereal>)
  - Can embed headers into our own code repo
  - Not so nice: No upgrade guarantee – new versions can break old files



# The superviseddescent library





- C++11, fully cross-platform
- Header-only
- CMake (for tests/examples)
- $\geq$ VS2013,  $\geq$ gcc-4.8.2,  $\geq$ clang-3.5
- License: Apache
- Repo: <https://github.com/patrikhuber/superviseddescent>
- Dependencies:
  - OpenCV core
  - Eigen
  - cereal

# Face tracking and 3D face reconstruction



# The face model

- A 3D Morphable Model
  - Shape and colour model («PCA» models), learned from 3D scans
  - Camera (rendering) parameters



# Using the 3D model

```
MorphableModel morphable_model = morphablemodel::load(filename); // loaded using cereal
vector<Vec2f> image_points;
vector<Vec4f> model_points;

Mat affine_cam = estimate_affine_camera(image_points, model_points);

vector<float> shape_coeffs = fit_shape_to_landmarks_linear(morphable_model, affine_cam,
                                                         image_points, vertex_indices);

Mesh mesh = morphable_model.draw_sample(shape_coeffs, vector<float>());
write_obj(mesh, "out.obj");
```

\*All namespaces omitted. Complete example:

<https://github.com/patrikhuber/eos/blob/master/examples/fit-model.cpp>

# 3D model fitting library

- C++11/14, fully cross-platform
- Header-only
- Includes low-resolution 3D shape model
- CMake (for examples)
- >=VS2015, >=gcc-4.8.2, >=clang-3.5
- License: Apache
- Repo: <https://github.com/patrikhuber/eos>
- Dependencies:
  - OpenCV core
  - Eigen
  - cereal

# Demo

# Summing up...

- Other alternatives available for 2D-stuff
  - Without source code (e.g. Intraface) or only either Linux **or** Windows
  - Most source code is C-style code (e.g. clandmark)
  - A few notable exceptions: dlib, (OpenCV)
    - (and others, non-C++, e.g. menpo)



# Summing up...

- Probably coming soon-ish: Better dealing with facial expressions

# Summing up...

- Header-only libraries are great!

# Summing up...

- OpenCV or not...
  - Try to integrate parts into OpenCV  
or...
  - Get rid of it, purely header-only easier for mobile & web

# Team

- Zhenhua Feng (Uni Surrey)
- Guosheng Hu (previously Uni Surrey)
- Philipp Kopp (Reutlingen Uni)
- Rafael Tena (previously Uni Surrey)
- Pouria Mortazavian (previously Uni Surrey)
- Willem Koppen (Uni Surrey)
- Michael Grupp (Reutlingen Uni)
- Dr. Matthias Rätsch (Reutlingen Uni)
- Dr. William Christmas (Uni Surrey)
- Prof. Josef Kittler (Uni Surrey)



# References & links

- Own & related publications:
  - **A Multiresolution 3D Morphable Face Model and Fitting Framework**, P. Huber, G. Hu, R. Tena, P. Mortazavian, W. Koppen, W. Christmas, M. Rätsch, J. Kittler, *in peer review (2015)*
  - **Fitting 3D Morphable Models using Local Features**, P. Huber, Z. Feng, W. Christmas, J. Kittler, M. Rätsch, *ICIP 2015*
  - **Random Cascaded-Regression Copse for Robust Facial Landmark Detection**, Z. Feng, P. Huber, J. Kittler, W. Christmas, X.J. Wu, *IEEE Signal Processing Letters*, 2015
  - **Supervised Descent Method and Its Applications to Face Alignment**, X. Xiong and F. De la Torre, *CVPR 2013*
  - **A Morphable Model for the Synthesis of 3D Faces**, V. Blanz and T. Vetter , *SIGGRAPH 1999*
- Links:
  - [www.opencv.org](http://www.opencv.org)

# Thank you!

Questions?