

# **std::allocator Is to Allocation what std::vector Is to Vexation**

**Prepared for CppCon  
Bellevue, WA, Sep 20-25, 2015**

Andrei Alexandrescu, Ph.D.  
andrei@erdani.com

## **Working Titles**

**std::allocator Is to  
Allocation what  
Alligator Is to  
Allegation**

# Designing Allocators That Don't Not Work

## A Cyclical Business

- 1983: “malloc is awesome. Use it.”
- 1993: “malloc is awful. Replace it.”
- 2003: “malloc is great. Use it.”
- 2013: “malloc is wanting. Supplement it.”

# History

## It was a Dark and Stormy Night

- It all started with C:

```
void* malloc(size_t size);
```

```
void free(void* p);
```

- Caller knows the size to allocate
- Caller knows the size to deallocate
- Caller doesn't tell **free** the size to deallocate

⇒ Allocator must manage size

Size management adds undue difficulties and inefficiencies to any allocator design

## Turning Time Back

```
struct Blk { void* ptr; size_t length; };
```

```
Blk malloc(size_t size);
```

```
void free(Blk block);
```

# operator new: The Next Generation

## operator new

- Had to work with `malloc`
  - Added type specificity
  - Several syntactic oddities
  - No communication with the constructor
  - "The Great Array Distraction"
- 
- Generally unrecommended for any serious use

## To Wit: N3536

“[...] omission [of size information in global `operator delete`] has unfortunate performance consequences. Modern memory allocators often allocate in size categories, and, for space efficiency reasons, do not store the size of the object near the object. Deallocation then requires searching for the size category store that contains the object. This search can be expensive, particularly as the search data structures are often not in memory caches.”

Along came  
`std::allocator`

## std::allocator

- Failure?
  - But it's still used!
  - “Making Allocators Work”, Alisdair Meredith, CppCon 2014
- 
- What's good?
  - What could be done better?

## Development

- Meant as a fix for near/far pointers
  - No intent to expand scope beyond that
  - Just distraction in the midst of a large design
- 
- `std::allocator` isn't a good allocator because it wasn't designed to be a good allocator

## Obvious Mistakes

- Type as parameter
  - Misunderstanding of charter
  - Allocator, not factory!
  - Yet still traffics in `void*`
  - Allocators should trade in `Blk`
  - See also: `rebind<U>::other`, indexed under “awful crutch” in Merriam-Webster

## More Obvious Mistakes

- Stateless assumption
  - *All* allocators have state!
  - Some is monostate
  - Assume state present
  - A global allocator object optional



## Subtle Issue: Composition

- Most actual allocators are compositions
  - Battery specialized by size classes
  - Arrays/lists/trees of other allocators
  - Hooks on top of others (freelists, debug, stats)
- 
- Getting composition right is getting allocators right

**New Design**

# A Leg Up

- Efficiency
  - Leaving size to the client
  - Scoped allocation patterns
  - Thread-local allocation patterns
- Features
  - Better configurability (debug/stats)
  - Specialization/adaptation
  - No legacy, no nonsense

## Simplest Composite Allocator

Let's define FallbackAllocator

```
template <class Primary, class Fallback>
class FallbackAllocator
    : private Primary
    , private Fallback {
public:
    Blk allocate(size_t);
    void deallocate(Blk);
};
```

## FallbackAllocator::allocate

```
template <class P, class F>
Blk FallbackAllocator<P, F>::allocate(
    size_t n) {
    Blk r = P::allocate(n);
    if (!r.ptr) r = F::allocate(n);
    return r;
}
```

## FallbackAllocator::deallocate

```
template <class P, class F>
void FallbackAllocator<P, F>::deallocate(
    Blk b) {
    if (P::owns(b)) P::deallocate(b);
    else F::deallocate(b);
}
```

- New optional method: `bool owns(Blk)`
- Primary must define it

## Good Citizenry

```
template <class P, class F>
bool FallbackAllocator::owns(Blk b) {
    return P::owns(b) || F::owns(b);
}
```

- Relies on MDFINAE

## Suddenly!

```
template <size_t s> class StackAllocator {
    char d_[s];
    char* p_;
    ...
};
class Mallocator { ... };
...
using Localloc = FallbackAllocator<
    StackAllocator<16384>,
    Mallocator>;
...
```

## Use

```
void fun(size_t n) {
    Localloc a;
    auto b = a.allocate(n * sizeof(int));
    SCOPE_EXIT { a.deallocate(b); };
    int* p = static_cast<int*>(b.ptr);
    ... use p[0] through p[n - 1] ...
}
```

## StackAllocator

```
template <size_t s> class StackAllocator {
    ...
    StackAllocator() : p_(d_) {}
    Blk allocate(size_t n) {
        auto n1 = roundToAligned(n);
        if (n1 > (d_ + s) - p_) {
            return { nullptr, 0 };
        }
        Blk result = { p_, n };
        p_ += n1;
        return result;
    }
    ...
};
```

# StackAllocator

```
...
void deallocate(Blk b) {
    if (b.ptr + roundToAligned(n) == p_) {
        p_ = b.ptr;
    }
}
bool owns(Blk b) {
    return b.ptr >= d_ && b.ptr < d_ + s;
}
// NEW: deallocate everything in O(1)
void deallocateAll() {
    p_ = d_;
}
...
```

# Freelist

- Keeps list of previous allocations of a given size

```
template <class A, size_t s> class Freelist {
    A parent_;
    struct Node { Node* next; } root_;
public:
    Blk allocate(size_t n) {
        if (n == s && root_) {
            Blk b = { root_, n };
            root = root_.next;
            return b;
        }
        return parent_.allocate(n);
    }
    ...
};
```

# Freelist

```
...  
bool owns(Blk b) {  
    return b.length == s || parent_.owns(b);  
}  
void deallocate(Blk b) {  
    if (b.length != s) return parent_.deallocate(b);  
    auto p = (Node*)b.ptr;  
    p.next = root_;  
    root_ = p;  
}
```

## Freelist improvements

- Add tolerance: min...max
- Allocate in batches
- Add upper bound: no more than top elems

## Freelist improvements

```
Freelist<A, // parent allocator
    17, // Use list for object sizes 17...
    32, // .. through 32
    8, // Allocate 8 at a time
    1024 // No more than 1024 remembered
> alloc8r;
```

## Adding Affixes

```
template <class A, class Prefix,
          class Suffix = void>
class AffixAllocator;
```

- Add optional prefix and suffix
- Construct/destroy appropriately
- Uses: debug, stats, info, ...



## “How are we doing?”

```
template <class A, ulong flags>
class AllocatorWithStats;
```

- Collect info about any other allocator
- Per-allocation and global
  - primitive calls, failures, byte counts, high tide
- Per-allocation
  - caller file/line/function/time

## Bitmapped Block

```
template <class A, size_t blockSize>
class BitmappedBlock;
```

- Organized by constant-size blocks
- Tremendously simpler than malloc's heap
- Faster, too
- Layered on top of any memory hunk
- 1 bit/block sole overhead (!)
- Multithreading tenuous (needs full interlocking for >1 block)

## When one fills up...

```
template <class Creator>
class CascadingAllocator;
...
auto a = cascadingAllocator({
    return Heap<...>();
});
```

- Keep a list of allocators, grown lazily
- Coarse granularity
- Linear search upon allocation
- May return memory back

## Segregator

```
template <size_t threshold,
class SmallAllocator,
class LargeAllocator>
struct Segregator;
```

- Sizes  $\leq$  threshold goes to SmallAllocator
- All else goes to LargeAllocator

## Segregator: self-composable!

```
typedef Segregator<4096,  
    Segregator<128,  
        Freelist<Mallocator, 0, 128>,  
        MediumAllocator>,  
    Mallocator>  
Allocator;
```

- Could implement any “search” strategy
- Works only for a handful of size classes

## Bucketizer: Size Classes

```
template <class Allocator,  
    size_t min,  
    size_t max,  
    size_t step>  
struct Bucketizer;
```

- [min, min + step),  
 [min + step, min + 2 \* step)...
- Within a bucket allocates the maximum size

**We're done!**

Most major allocation  
tropes covered

## **Approach to copying**

- Allocator-dependent
- Some noncopyable & immovable: stack/in-situ allocator
- Some noncopyable, movable: non-stack regions
- Many refcounted, movable
- Perhaps other models, too

# Granularity

- Where are types, factories, ...?
- Design surprise: these are extricable
- Focus on **Blk**

## Realistic Example

```
using FList = Freelist<Mallocator, 0, -1>;
using A = Segregator<
    8, Freelist<Mallocator, 0, 8>,
    128, Bucketizer<FList, 1, 128, 16>,
    256, Bucketizer<FList, 129, 256, 32>,
    512, Bucketizer<FList, 257, 512, 64>,
    1024, Bucketizer<FList, 513, 1024, 128>,
    2048, Bucketizer<FList, 1025, 2048, 256>,
    3584, Bucketizer<FList, 2049, 3584, 512>,
    4072 * 1024, CascadingAllocator<
        decltype(newHeapBlock)>,
    Mallocator
>;
```

## Complete API

```
static constexpr unsigned alignment;  
static constexpr goodSize(size_t);  
Blk allocate(size_t);  
Blk allocateAll();  
bool expand(Blk&, size_t delta);  
void reallocate(Blk&, size_t);  
bool owns(Blk);  
void deallocate(Blk);  
void deallocateAll();
```

## Complete API: aligned versions

```
Blk alignedAllocate(size_t, unsigned);  
Blk alignedReallocate(size_t, unsigned);
```

- AlignedMallocator
- posix\_memalign on Posix
- \_aligned\_malloc on Windows
- Regions support this as well!

# Summary

## Summary

- Fresh approach from first principles
- Understanding history
  - Otherwise: "...doomed to repeat it"
- Composability is key