

Beyond Sanitizers


Fuzzing and Hardening your C++ apps
for Security and Reliability

CppCon, September 21, 2015

Kostya Serebryany kcc@google.com

Agenda

- What are the Sanitizers?
 - ASan, TSan, MSan, UBSan
 - Talk at CppCon 2014
- Fuzz testing
- Code hardening



Sanitizers

Dynamic Testing Tools
based on compile-time instrumentation

ASan (AddressSanitizer)

Bugs related to addressing memory

ASan report example: global-buffer-overflow

```
int global_array[100] = {-1};  
int main(int argc, char **argv) {  
    return global_array[argc + 100];    // BOOM  
}
```

```
% clang++ -O1 -fsanitize=address a.cc ; ./a.out
```

```
==10538== ERROR: AddressSanitizer global-buffer-overflow  
READ of size 4 at 0x000000415354 thread T0
```

```
#0 0x402481 in main a.cc:3
```

```
#1 0x7f0a1c295c4d in __libc_start_main ??:0
```

```
#2 0x402379 in _start ??:0
```

```
0x000000415354 is located 4 bytes to the right of global  
variable 'global_array' (0x4151c0) of size 400
```

ASan report example: use-after-free

```
int main(int argc, char **argv) {
```

```
    int *array = new int[100];
```

```
    delete [] array;
```

```
    return array[argc]; } // BOOM
```

```
% clang++ -O1 -fsanitize=address a.cc && ./a.out
```

```
==30226== ERROR: AddressSanitizer heap-use-after-free
```

```
READ of size 4 at 0x7faa07fce084 thread T0
```

```
#0 0x40433c in main a.cc:4
```

```
0x7faa07fce084 is located 4 bytes inside of 400-byte region
```

```
freed by thread T0 here:
```

```
#0 0x4058fd in operator delete[](void*) _asan_rtl_
```

```
#1 0x404303 in main a.cc:3
```

```
previously allocated by thread T0 here:
```

```
#0 0x405579 in operator new[](unsigned long) _asan_rtl_
```

```
#1 0x4042f3 in main a.cc:2
```

ASan report example: stack-use-after-return

```
int *g;                                int main() {
void LeakLocal() {                      LeakLocal();
    int local;                          return *g;
    g = &local;                          }
}
```

```
% clang -g -fsanitize=address a.cc
```

```
% ASAN_OPTIONS=detect_stack_use_after_return=1 ./a.out
```

```
==19177==ERROR: AddressSanitizer: stack-use-after-return
READ of size 4 at 0x7f473d0000a0 thread T0
#0 0x461ccf in main a.cc:8
```

Address is located in stack of thread T0 at offset 32 in frame

```
#0 0x461a5f in LeakLocal() a.cc:2
```

This frame has 1 object(s):

```
[32, 36) 'local' <== Memory access at offset 32
```

TSan (ThreadSanitizer)

Bugs related to concurrency

TSan report example: data race

```
int X;  
std::thread t([&]{x = 42;});  
x = 43;  
t.join();
```

```
% clang -fsanitize=thread -g race.cc && ./a.out
```

```
WARNING: ThreadSanitizer: data race (pid=25493)
```

```
Write of size 4 at 0x7fff7f10e338 by thread T1:
```

```
#0 main::_0::operator()() const race.cc:4 ...
```

```
Previous write of size 4 at 0x7...8 by main thread:
```

```
#0 main race.cc:5
```

MSan (MemorySanitizer)

Bugs related to contents of memory

MSan report example

```
int main(int argc, char **argv) {  
    int x[10];  
    x[0] = 1;  
    return x[argc]; }  

```

```
% clang -fsanitize=memory a.c -g; ./a.out
```

WARNING: Use of uninitialized value

#0 0x7f1c31f16d10 in main a.cc:4

Uninitialized value was created by an
allocation of 'x' in the stack frame of
function 'main'

UBSan

(UndefinedBehaviorSanitizer)

Many other kinds of undefined behavior

UBSan report example: int overflow

```
int main(int argc, char **argv) {  
    int t = argc << 16;  
    return t * t;  
}
```

```
% clang -fsanitize=undefined a.cc -g; ./a.out
```

```
a.cc:3:12: runtime error:
```

```
signed integer overflow: 65536 * 65536
```

```
cannot be represented in type 'int'
```

Sanitizers have found thousands of bugs everywhere

Proof links: [\[1\]](#), [\[2\]](#), [\[3\]](#)

© Sanitizers' marketing department

But Sanitizers are not enough

- ASan, TSan, MSan, UBSan are “best-effort tools”:
 - Only as good as the tests are
 - Do not prove correctness
- Beyond Sanitizers:
 - Improve test quality (aka test coverage) by fuzzing
 - Protect from security-sensitive bugs in production (hardening)



What's "Fuzzing"?

https://en.wikipedia.org/wiki/Fuzz_testing

Fuzz testing or fuzzing is a software **testing** technique, often automated or semi-automated, that involves providing invalid, unexpected, or random data to the inputs of a computer program.

Generation-based Fuzzing

- Generate millions of inputs, feed them to the target app
 - Can (and should) be used with Sanitizers
 - May generate invalid inputs (stresses the parser)
 - Or may produce valid inputs by design (e.g. “csmith” -- C fuzzer)
 - Actively used by Chromium security team, found thousands of bugs
- Extremely effective, yet often barely scratches the surface

Mutation-based Fuzzing

- Acquire a test corpus (e.g. crawl the web)
 - Minimize the corpus according to, e.g. `code_coverage/execution_time`
- Mutate tests from the corpus and run them
- Often better results compared to generation-based fuzzing
 - But harder for highly structured inputs, e.g. C++

Control-flow-guided (coverage-guided) fuzzing

- Same as mutation-based fuzzing, but also
 - Run the mutations with code coverage instrumentation
 - Add the mutations to the corpus if new coverage is discovered
- 1-3 orders of magnitude faster than plain mutation-based fuzzing

AFL-fuzz

AFL-fuzz, a control-flow guided fuzzer

- Instrument the binary at compile-time
 - Regular mode: instrument assembly
 - Recent addition: LLVM compiler instrumentation mode
- Provide 64K counters representing all edges in the app
 - 8 bits per edge (# of executions: 1, 2, 3, 4-7, 8-15, 16-31, 32-127, 128+)
 - Imprecise (edges may collide) but very efficient
- AFL-fuzz is the driver process, the target app runs as separate process(es)

AFL-fuzz is not a toy!

IJG jpeg [1](#) libjpeg-turbo [1](#) [2](#) libpng [1](#) libtiff [1](#) [2](#) [3](#) [4](#) [5](#) mozjpeg [1](#) PHP [1](#) [2](#) [3](#) [4](#) Mozilla Firefox [1](#) [2](#) [3](#)
[4](#) Internet Explorer [1](#) [2](#) [3](#) [4](#) Apple Safari ⁽¹⁾ ⁽²⁾ Adobe Flash / PCRE [1](#) [2](#) sqlite [1](#) [2](#) [3](#) [4](#)... OpenSSL [1](#)
[2](#) [3](#) [4](#) LibreOffice [1](#) [2](#) [3](#) [4](#) poppler [1](#) freetype [1](#) [2](#) GnuTLS [1](#) GnuPG [1](#) [2](#) [3](#) [4](#) OpenSSH [1](#) [2](#) [3](#) bash
(post-Shellshock) [1](#) [2](#) tcpdump [1](#) [2](#) [3](#) [4](#) [5](#) [6](#) [7](#) [8](#) JavaScriptCore [1](#) [2](#) [3](#) [4](#) pdftium [1](#) [2](#) ffmpeg [1](#) [2](#) [3](#)
[4](#) libmatroska [1](#) libarchive [1](#) [2](#) [3](#) [4](#) [5](#) [6](#) ... wireshark [1](#) [2](#) ImageMagick [1](#) [2](#) [3](#) [4](#) [5](#) [6](#) [7](#) [8](#) ... lcms [1](#) libbpg ⁽¹⁾
lame [1](#) FLAC audio library [1](#) [2](#) libsndfile [1](#) [2](#) [3](#) less / lesspipe [1](#) [2](#) [3](#) strings (+ related tools) [1](#) [2](#) [3](#)
[4](#) [5](#) [6](#) [7](#) file [1](#) [2](#) [3](#) [4](#) dpkg [1](#) rcs [1](#) systemd-resolved [1](#) [2](#) libyaml [1](#) Info-Zip unzip [1](#) [2](#) libtasn1 [1](#)
[2](#) OpenBSD pfctl [1](#) NetBSD bpf [1](#) man & mandoc [1](#) [2](#) [3](#) [4](#) [5](#) ... IDA Pro ^[reported by authors] clamav [1](#) [2](#)
[3](#) [4](#) [5](#) libxml2 [1](#) glibc [1](#) clang / llvm [1](#) [2](#) [3](#) [4](#) [5](#) [6](#) [7](#) [8](#) ... nasm [1](#) [2](#) ctags [1](#) mutt [1](#) procmail [1](#) fontconfig
[1](#) pdksh [1](#) [2](#) Qt [1](#) wavpack [1](#) redis / lua-cmsgpack [1](#) taglib [1](#) [2](#) [3](#) privoxy [1](#) perl [1](#) [2](#) [3](#) [4](#) [5](#) [6](#) [7](#)...
libxmpradare2 [1](#) [2](#) SleuthKit [1](#) fwknop ^[reported by author] X.Org [1](#) exifprobe [1](#) jhead ^[?] capnproto
[1](#) Xerces-C [1](#) metacam [1](#) djvulibre [1](#) exiv [1](#) Linux btrfs [1](#) [2](#) [3](#) [4](#) Knot DNS [1](#) curl [1](#) [2](#) wpa_supplicant
[1](#) libde265 [□] dnsmasq [1](#) imlib2 [1](#) libraw [1](#) libwmf [1](#) uudecode [1](#) MuPDF [1](#) libbson [1](#) libsass [1](#)

LLVM libFuzzer

Sanitizer Coverage instrumentation

- **Clang/LLVM flag** `-fsanitize-coverage=`
 - `func/bb/edge`: records if a function, basic block or edge was executed
 - `indirect-calls`: records unique indirect caller-callee pairs
 - `8bit-counters`: similar to AFL, provides 8-state counter for edges
 - (1, 2, 3, 4-7, 8-15, 16-31, 32-127, 128+)
- Provides the status in-process and dumps on disk at exit
 - i.e. supports in-process and out-of-process clients
- Should be combined with ASan, MSan, or UBSan
- Typical slowdown: within 10%
 - 8bit counters may be unfriendly to multi-threaded apps

LLVM libFuzzer

- Lightweight in-process control-flow guided fuzzer

- Provide your own target function

- `void LLVMFuzzerTestOneInput (const uint8_t *Data, size_t Size);`

- `-fsanitize-coverage=edge[,indirect-calls][,8bit-counters]`
- `-fsanitize={address,memory,undefined,leak}`
- Link with libFuzzer

- Younger than AFL-fuzz and is not as algorithmically sophisticated. Yet quite capable!
- Targeted at libraries/APIs, not at large applications

libFuzzer usage

- Acquire a test corpus, put it into a directory `CORPUS`
 - empty corpus is OK
- Run `./my-fuzzer CORPUS`
 - `-jobs=N`: `N` parallel jobs, all working on the same corpus
 - `-max_len=N`: limit the input size (default: 64)
 - `-help`: more knobs
- Newly discovered test inputs are written to `CORPUS`
- Bug/timeout will stop the process & dump input on disk
- Optional: feed the produced corpus to AFL-fuzz

Example: FreeType (font rendering library) fuzzer

```
void TestOneInput(const uint8_t *data, size_t size) {  
    FT_Face face;  
    if (size < 1) return;  
    if (!FT_New_Memory_Face(library, data, size, 0, &face)) {  
        FT_Done_Face(face);  
    }  
}
```

Results with FreeType (ASan+UBsan)

#45999 **left shift of negative value -4592**

#45989 **leak** in t42_parse_charstrings

#45987 512 byte input **consumes 1.7Gb** / 2 sec to process

#45986 **leak** in ps_parser_load_field

#45985 **signed integer overflow**: -35475362522895417 * -8256 cannot be represented in t

#45984 **signed integer overflow**: 2 * 1279919630 cannot be represented in type 'int'

#45983 runtime error: **left shift of negative value -9616**

#45966 **leaks** in parse_encoding, parse_blend_design_map, t42_parse_encoding

#45965 left shift of 184 by 24 places cannot be represented in type 'int'

#45964 **signed integer overflow**: 6764195537992704 * 7200 cannot be represented in type

#45961 FT_New_Memory_Face consumes 6Gb+

#45955 **buffer overflow** in T1_Get_Private_Dict/strncmp

#45938 **shift exponent 2816 is too large** for 64-bit type 'FT_ULONG'

#45937 **memory leak** in FT_New_Memory_Face/FT_Stream_OpenGzip

#45923 **buffer overflow** in T1_Get_Private_Dict while doing FT_New_Memory_Face

#45922 **buffer overflow** in skip_comment while doing FT_New_Memory_Face

#45920 FT_New_Memory_Face **takes infinite time** (in PS_Conv_Strtol)

#45919 FT_New_Memory_Face **consumes 17Gb** on a small input

Example: OpenSSL

```
SSL_CTX *sctx;
int Init() { ... }
extern "C" void LLVMFuzzerTestOneInput(unsigned char * Data, size_t Size) {
    static int unused = Init();
    SSL *server = SSL_new(sctx);
    BIO *sinbio = BIO_new(BIO_s_mem());
    BIO *soutbio = BIO_new(BIO_s_mem());
    SSL_set_bio(server, sinbio, soutbio);
    SSL_set_accept_state(server);
    BIO_write(sinbio, Data, Size);
    SSL_do_handshake(server);
    SSL_free(server);
}
```

Demo: OpenSSL, the “HeartBleed” bug

Exact commands: lvm.org/docs/LibFuzzer.html




Control-flow-guided fuzzing
is not the end

Concolic execution (rocket science)

- Concolic: concrete and symbolic
 - Execute with instrumentation
 - Figure out which branches are never taken
 - Feed the data to SMT solver, get new test inputs that cover more branches
- Good in theory and often in practice too, but very heavyweight

Data-flow-guided fuzzing

- Intercept the data flow, analyze the inputs of comparisons
- Modify the test inputs, observe the effect on comparisons
- Prototype in LLVM libFuzzer (and [go-fuzz](#))
 - Already have trophies (DEMO)
 - May use taint analysis, e.g. DFSan, to make smarter mutations



Code hardening

Threat #1

Buffer-overflow/use-after-free
overwrites a vptr or a function pointer
by an attacker-controlled value

Hijacked VPTR in Chromium: Pwn2Own 2013 (CVE-2013-0912)

Solution:

Control Flow Integrity (CFI)

```
clang++ -fsanitize=cfi-vcall -flto
```

CFI in Clang/LLVM

- Every disjoint class hierarchy is handled separately
 - Assumes the class hierarchy is a closed system; ok for Chrome
- Layout vtables for every hierarchy as a contiguous array
 - Align every vtable by the same power-of-2
- For every virtual function call site
 - Compile-time: compute the strict set of allowed functions
 - Run-time: perform a **range check**, **alignment check**, and a **bitset lookup**

Bitset lookup optimizations

- A bitset of ≤ 64 bits requires no memory loads
- No check if the bitset contains all ones
- Optimize the vtable layouts to minimize the bitset sizes

CFI: generated x86_64 assembler

All ones

```
mov    $0x4008f0,%ecx
mov    %rax,%rdx
sub    %rcx,%rdx
rol    $0x3b,%rdx
cmp    $0x2,%rdx
jae    CRASH
mov    %rbx,%rdi
callq  *(%rax)
...
CRASH: ud2
```

<= 64 bits

```
mov    $0x400e20,%edx
mov    %rax,%rcx
sub    %rdx,%rcx
rol    $0x3b,%rcx
cmp    $0xe,%rcx
ja     CRASH
mov    $0x4007,%edx
bt     %ecx,%edx
jae    CRASH
mov    %rbx,%rdi
callq  *(%rax)
...
CRASH: ud2
```

Full check

```
mov    $0x401810,%edx
mov    %rax,%rcx
sub    %rdx,%rcx
rol    $0x3b,%rcx
cmp    $0x40,%rcx
ja     400936 CRASH
testb  $0x1,0x402140(%rcx)
je     400936 CRASH
mov    %rbx,%rdi
callq  *(%rax)
...
CRASH: ud2
```


More CFI

- Other calls

- non-virtual member calls: `-fsanitize=cfi-nvcall`
- C-style indirect calls: `-fsanitize=cfi-icall`

- Casts for polymorphic types

- Base class => derived class: `-fsanitize=cfi-derived-cast`
- `void *` => pointer to a class: `-fsanitize=cfi-unrelated-cast`

CFI & Chromium

- Builds and runs on Linux & Android
 - ...=`cfi-vcall`, `cfi-derived-cast`, `cfi-unrelated-cast`
 - OSX and Windows are close to working too
- < 1% CPU overhead
- ~7% code size increase
- Significant cleanup was required (real bugs)

Better/different CFI

- Do not require LTO?
 - Requiring LTO is not necessary bad thing!
- Allow class hierarchies to cross the DSO boundaries
 - VS2015 Control Flow Guard (/d2guard4 + /Guard:cf)
 - Maybe not a great idea?

Threat #2

Stack-buffer-overflow
overwrites return address
by an attacker-controlled value

Solution:

SafeStack

```
clang++ -fsanitize=safe-stack
```

SafeStack

- Place local variables on a separately mmaped region
- stack-buffer-overflow can't touch the return addresses
- VPTRs and function pointers can still be affected
 - Combine with CFI
- Chromium: < 1% CPU

SafeStack: code example

```
push    %r14
push    %rbx
push    %rax
mov     0x207d0d(%rip),%r14
mov     %fs:(%r14),%rbx    # Get unsafe_stack_ptr
lea     -0x10(%rbx),%rax   # Update unsafe_stack_ptr
mov     %rax,%fs:(%r14)    # Store unsafe_stack_ptr
lea     -0x4(%rbx),%rdi
movl    $0x123456,-0x4(%rbx)
callq   40f2c0 <_Z3barPi>
mov     %rbx,%fs:(%r14)    # Restore unsafe_stack_ptr
xor     %eax,%eax
add     $0x8,%rsp
pop     %rbx
pop     %r14
retq
```

```
int main() {
    int local_var = 0x123456;
    bar(&local_var);
}
```

Summary

- ~~Rely on traditional testing to get a false sense of security~~
- Test with Sanitizers to achieve basic code sanity
 - ASan, TSan, MSan, UBSan
- Use guided fuzzing for stronger security & reliability
 - LLVM libFuzzer and AFL-fuzz make it super easy
- Harden your code for even better security
 - CFI for virtual calls, non-virtual member calls, casts, indirect calls
 - SafeStack for return addresses

Q&A



[llvm.org/docs/**LibFuzzer**.html](http://llvm.org/docs/LibFuzzer.html)

[clang.llvm.org/docs/**ControlFlowIntegrity**.html](http://clang.llvm.org/docs/ControlFlowIntegrity.html)

[clang.llvm.org/docs/**SafeStack**.html](http://clang.llvm.org/docs/SafeStack.html)