

C++ in the audio industry

An abstract graphic representing an audio waveform. It features several overlapping, wavy lines in white, red, and blue, set against a dark background. The lines are smooth and flowing, suggesting a continuous audio signal.

Timur Doumler

CppCon, 21 September 2015

real-time
programming

fast & efficient
DSP

lock-free thread
synchronisation

embedded
systems in music
hardware

SIMD and memory
alignment

cross-platform
challenges

real-time
programming

fast & efficient
DSP

lock-free thread
synchronisation

embedded
systems in music
hardware

SIMD and memory
alignment

cross-platform
challenges

- introduction
- representing audio in C++
- audio = “hard real-time” programming
 - consequences
 - dos and don'ts
- lock-free thread synchronisation
 - sharing objects between threads
 - exchanging data between threads
 - memory management / object lifetime

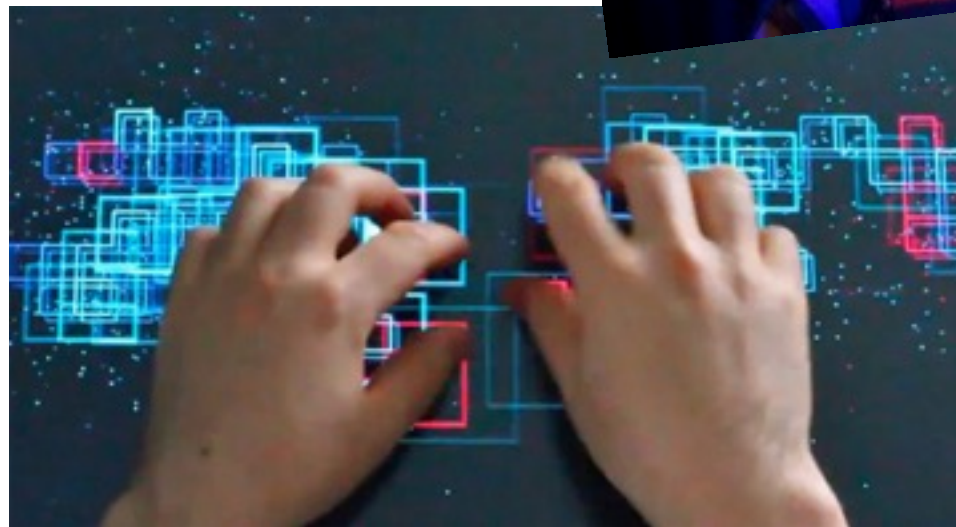


music producers
audio engineers
composers
sound designers



audio apps
games
multimedia

live performers
djs
touring musicians



science, art, creative coding



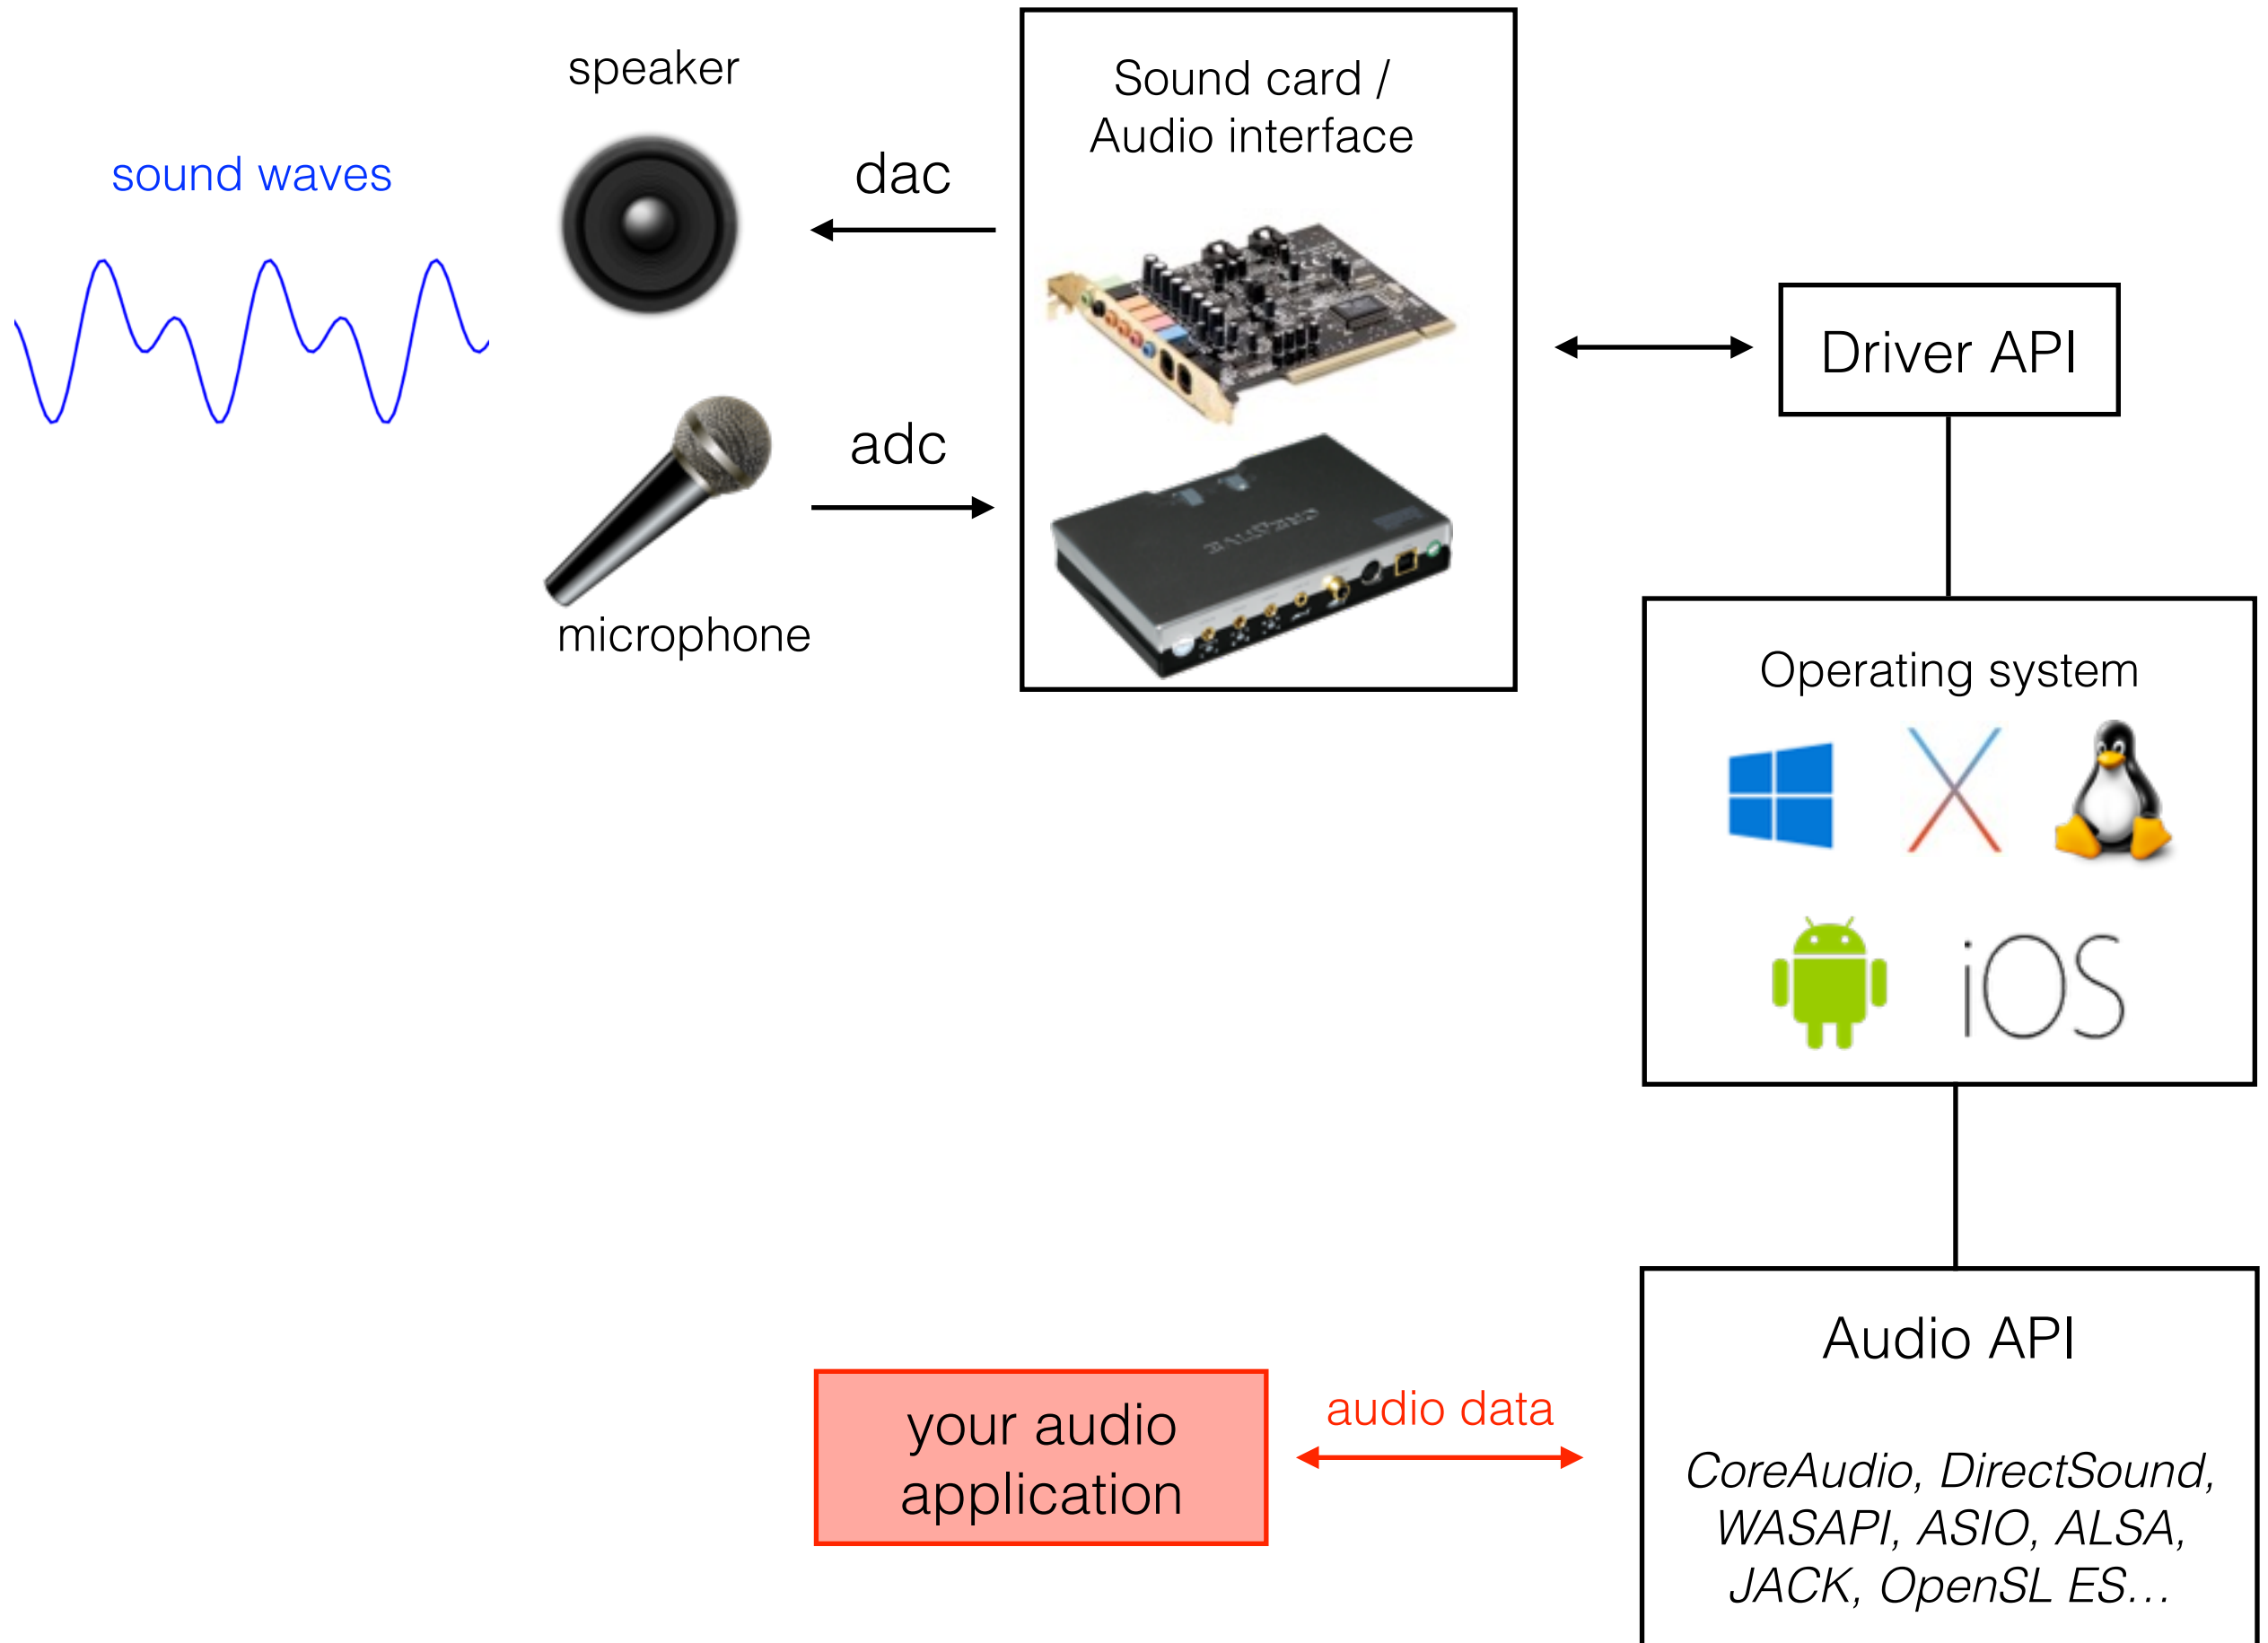
developing musical instruments,
synthesisers, effects

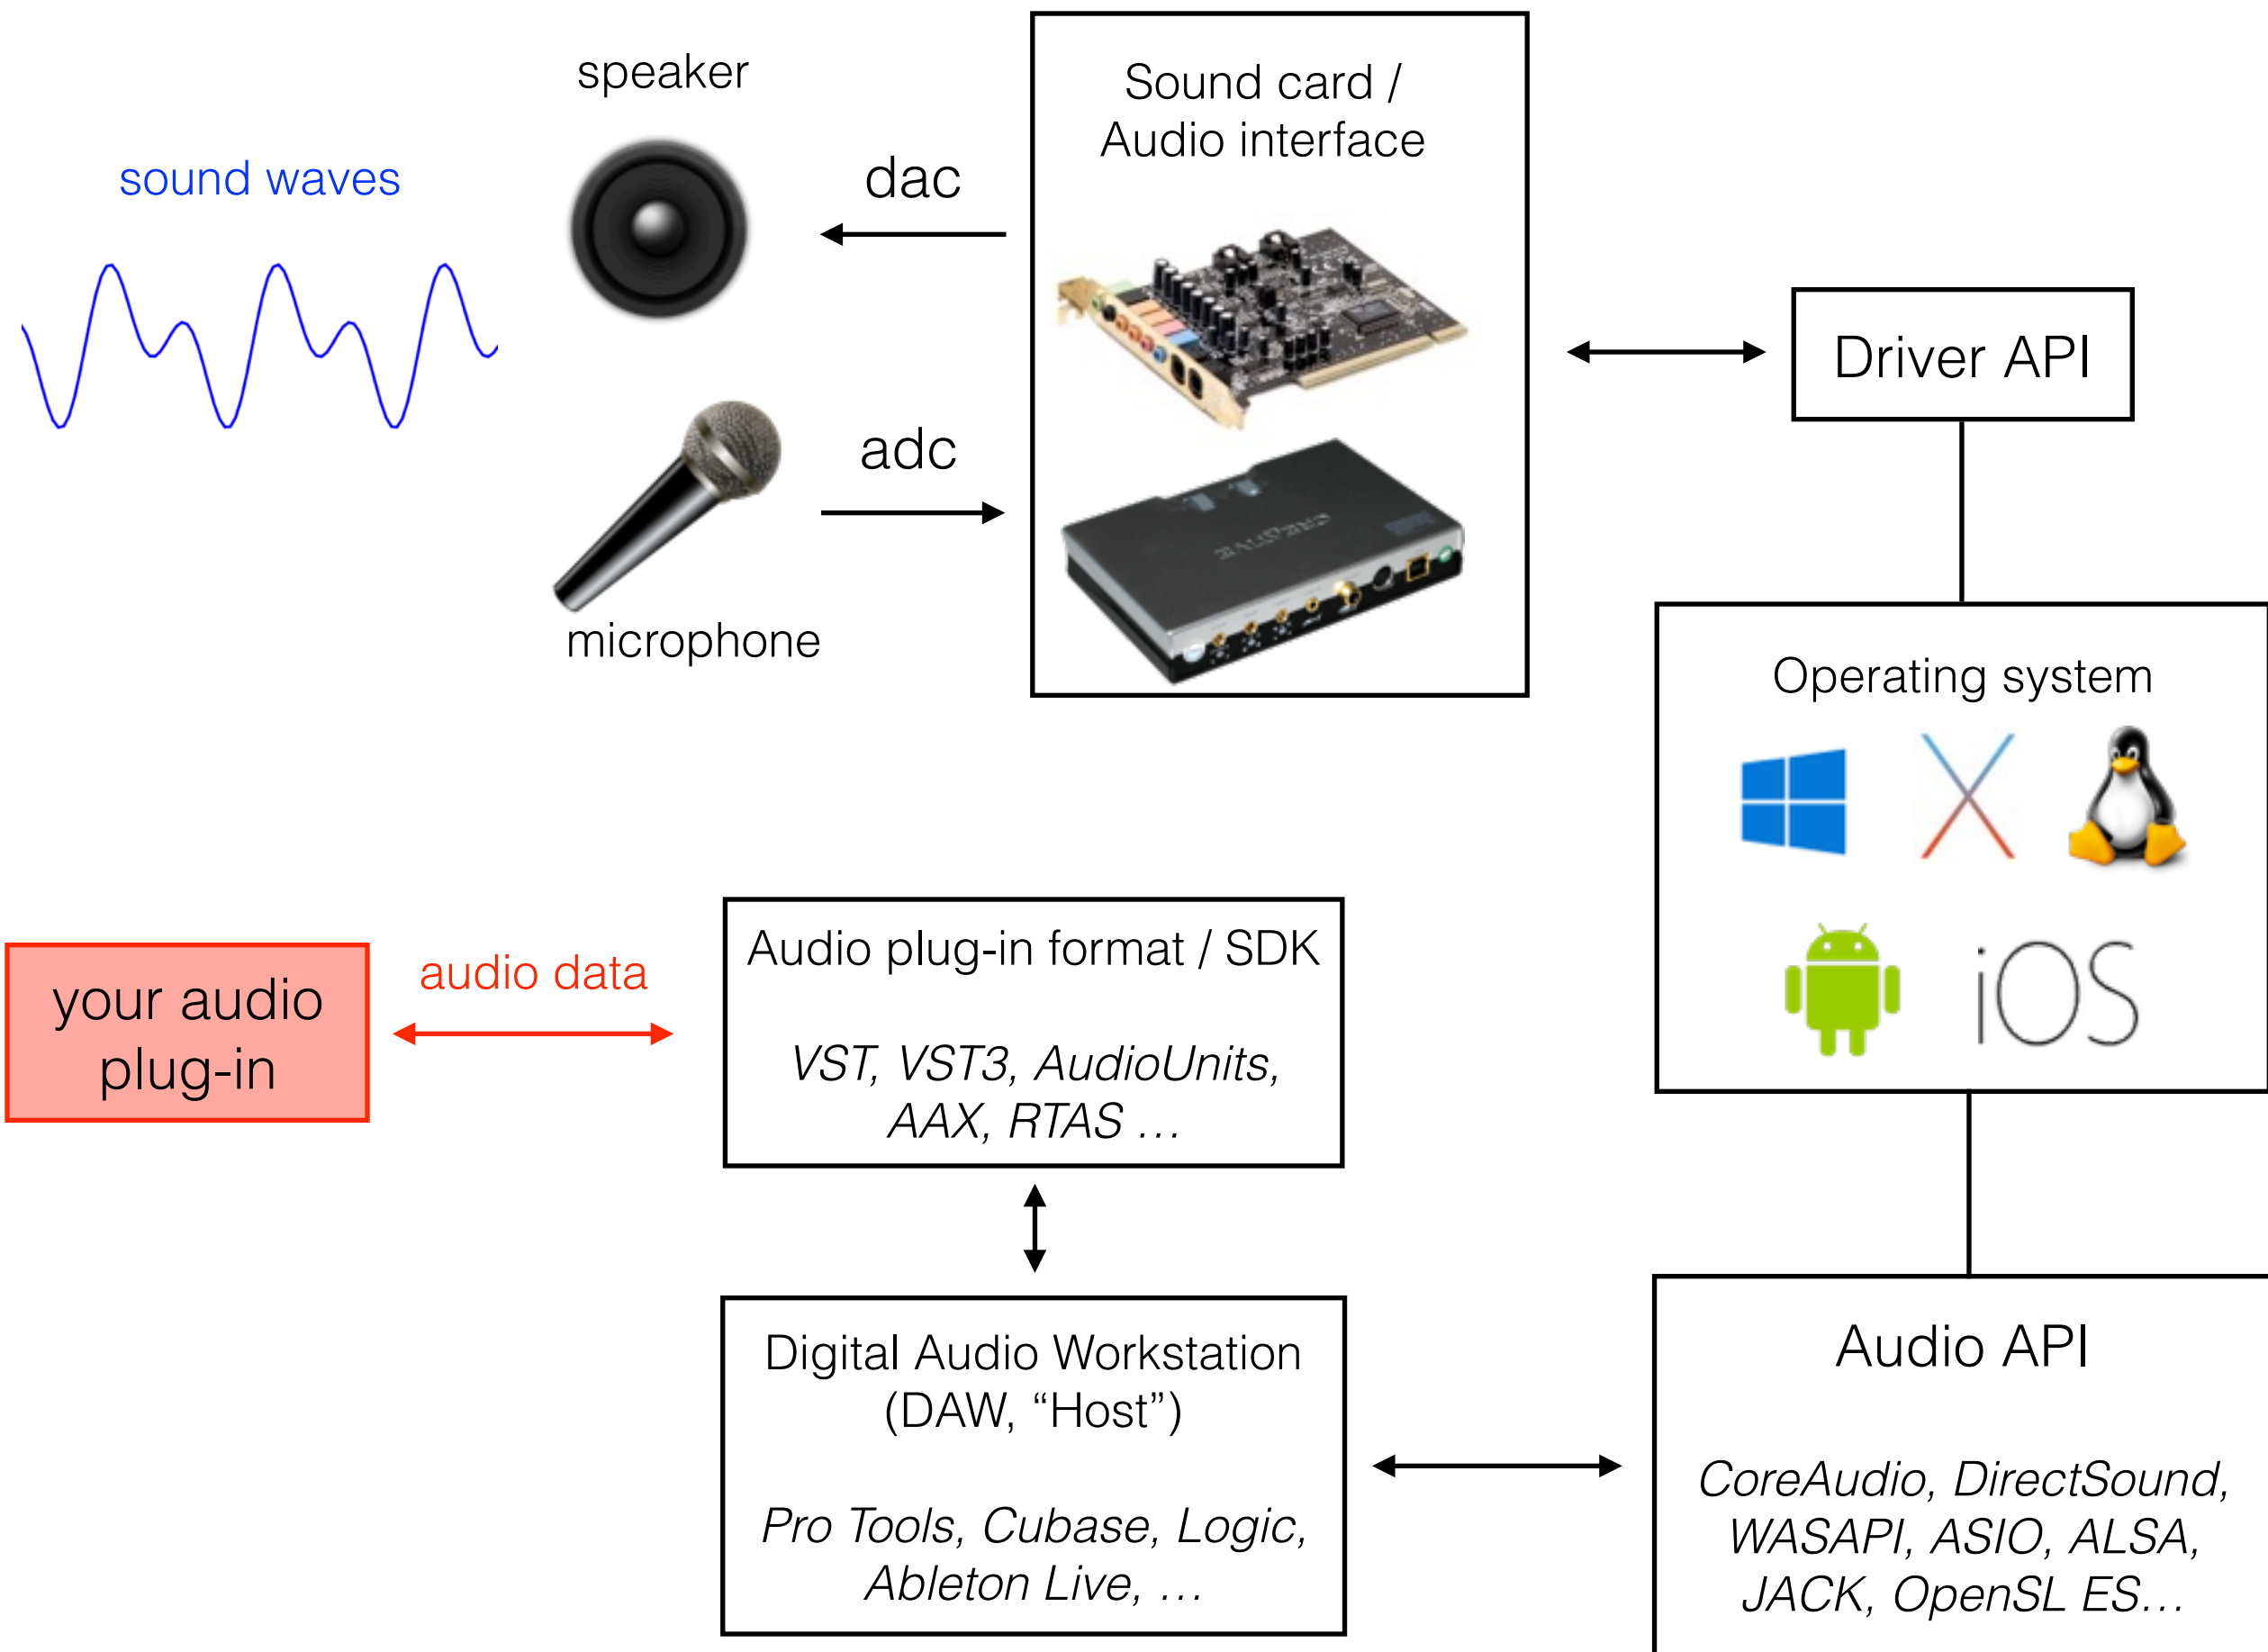


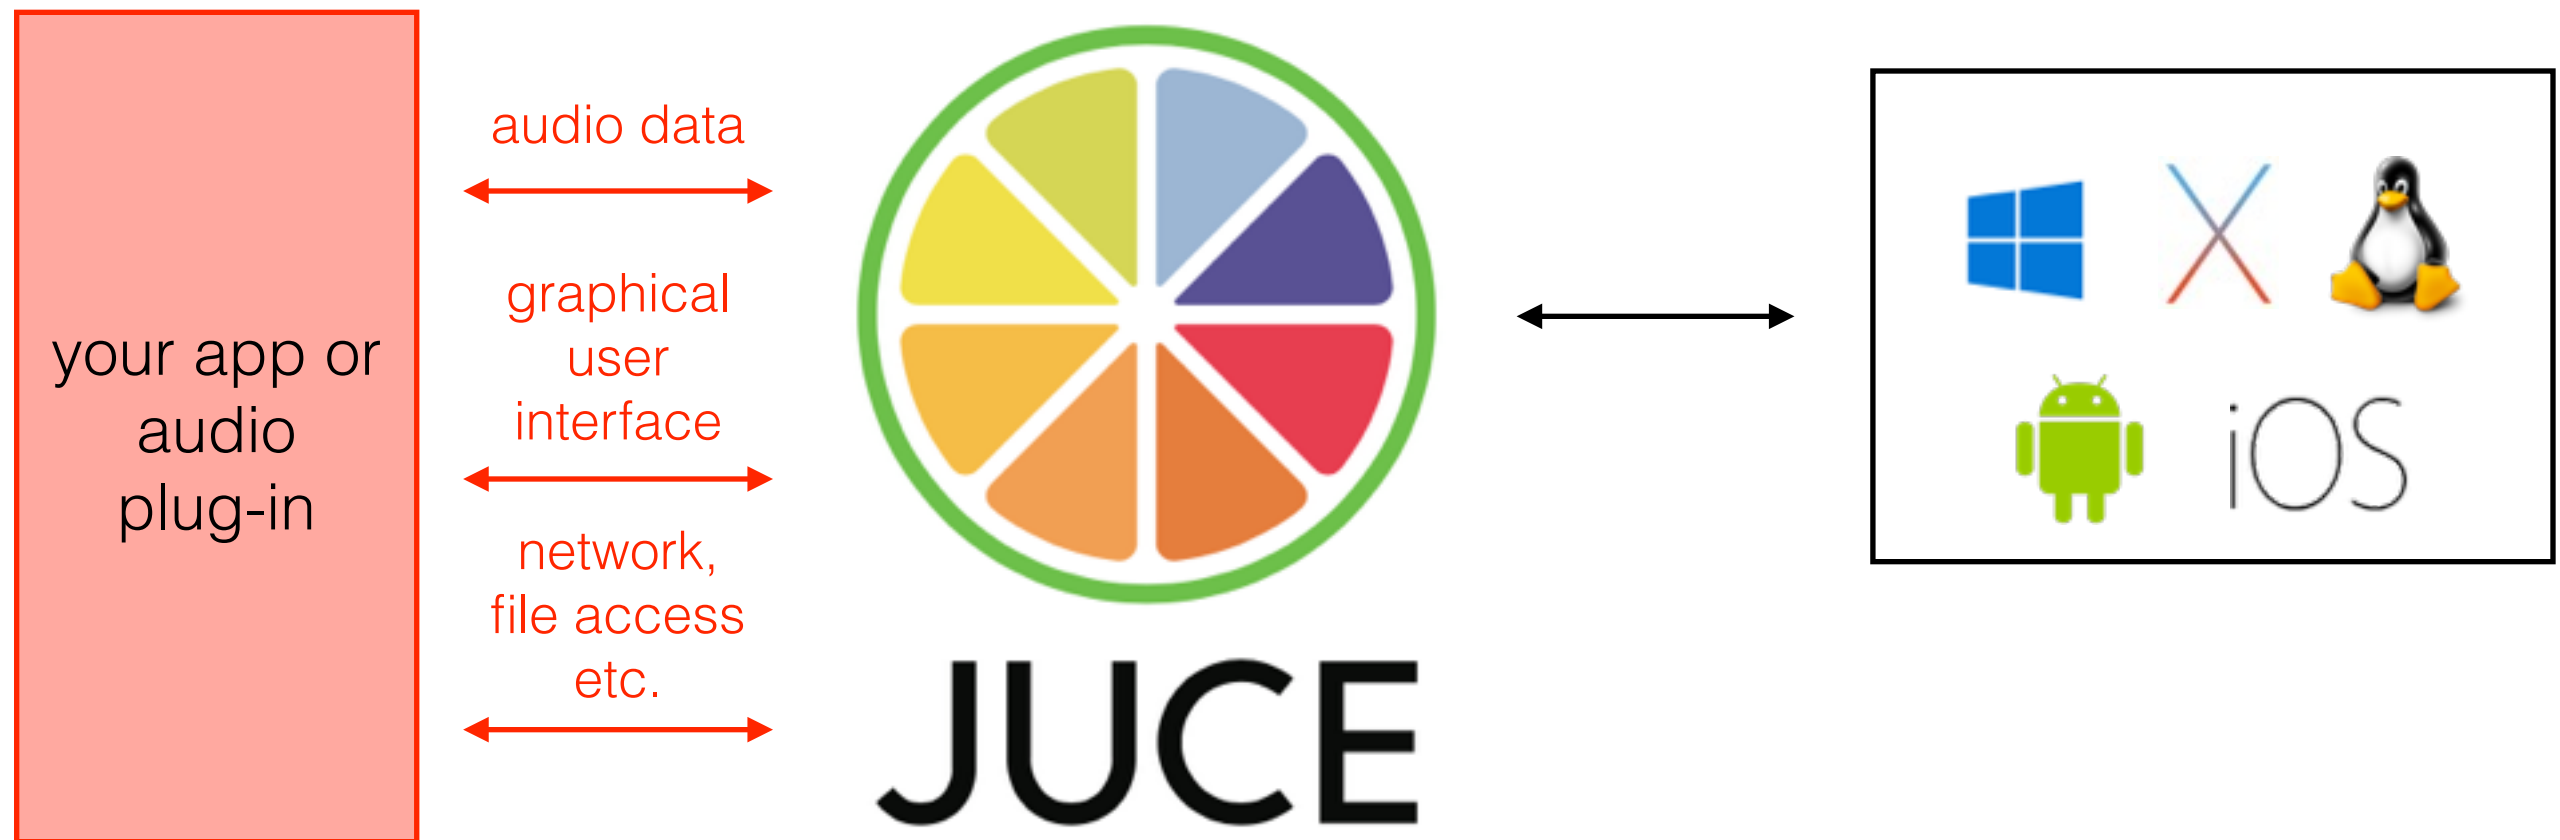
audio = data representation of sound waves

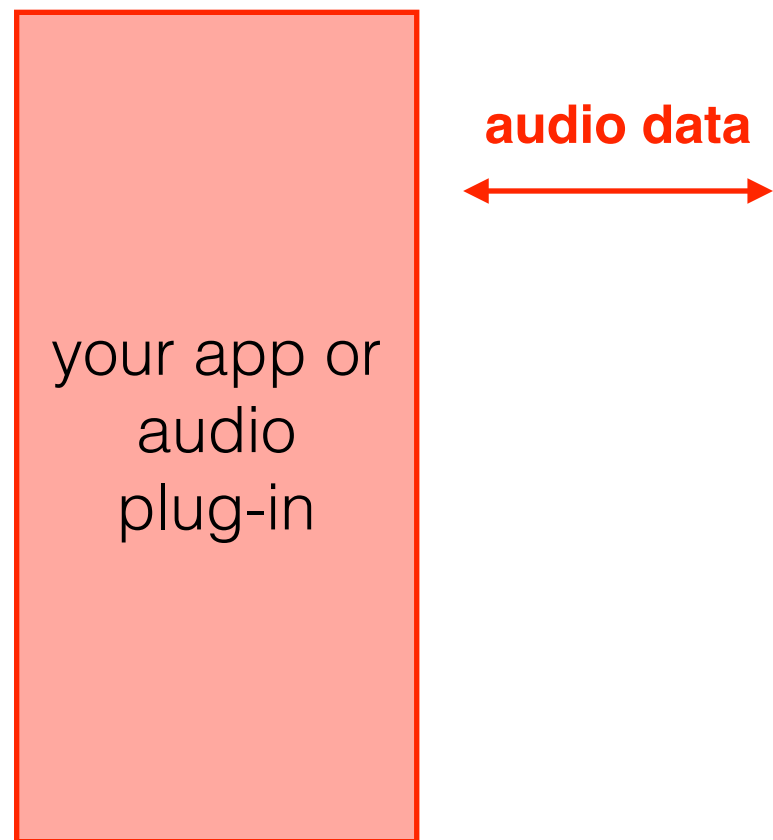


- Audio data can be represented and manipulated digitally
- The C++ language has no concept of audio
- But there are conventions, 3rd-party libraries, APIs

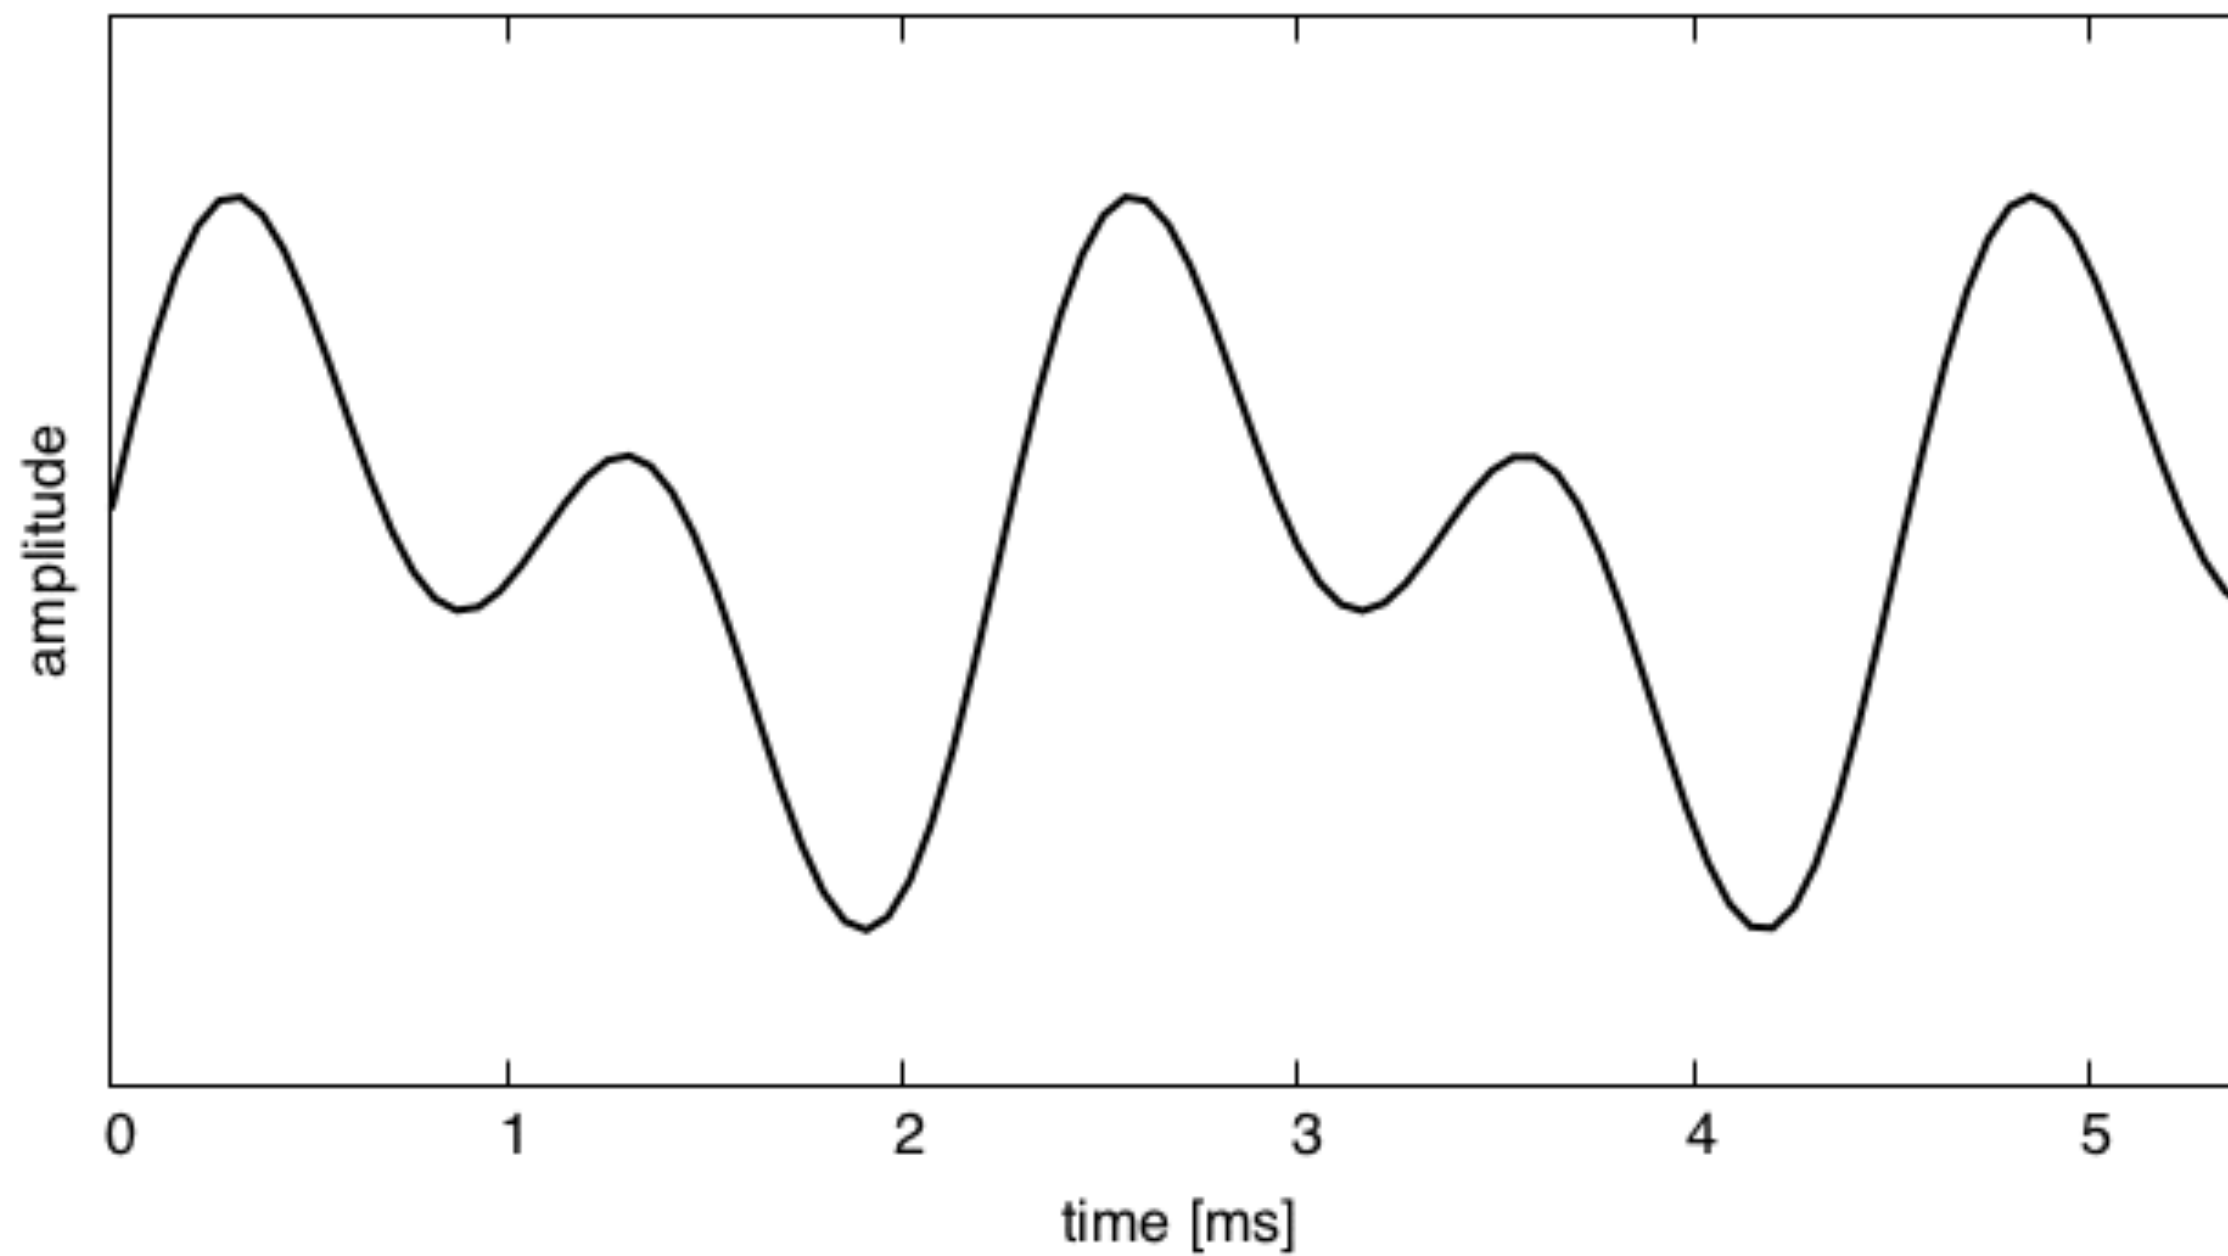








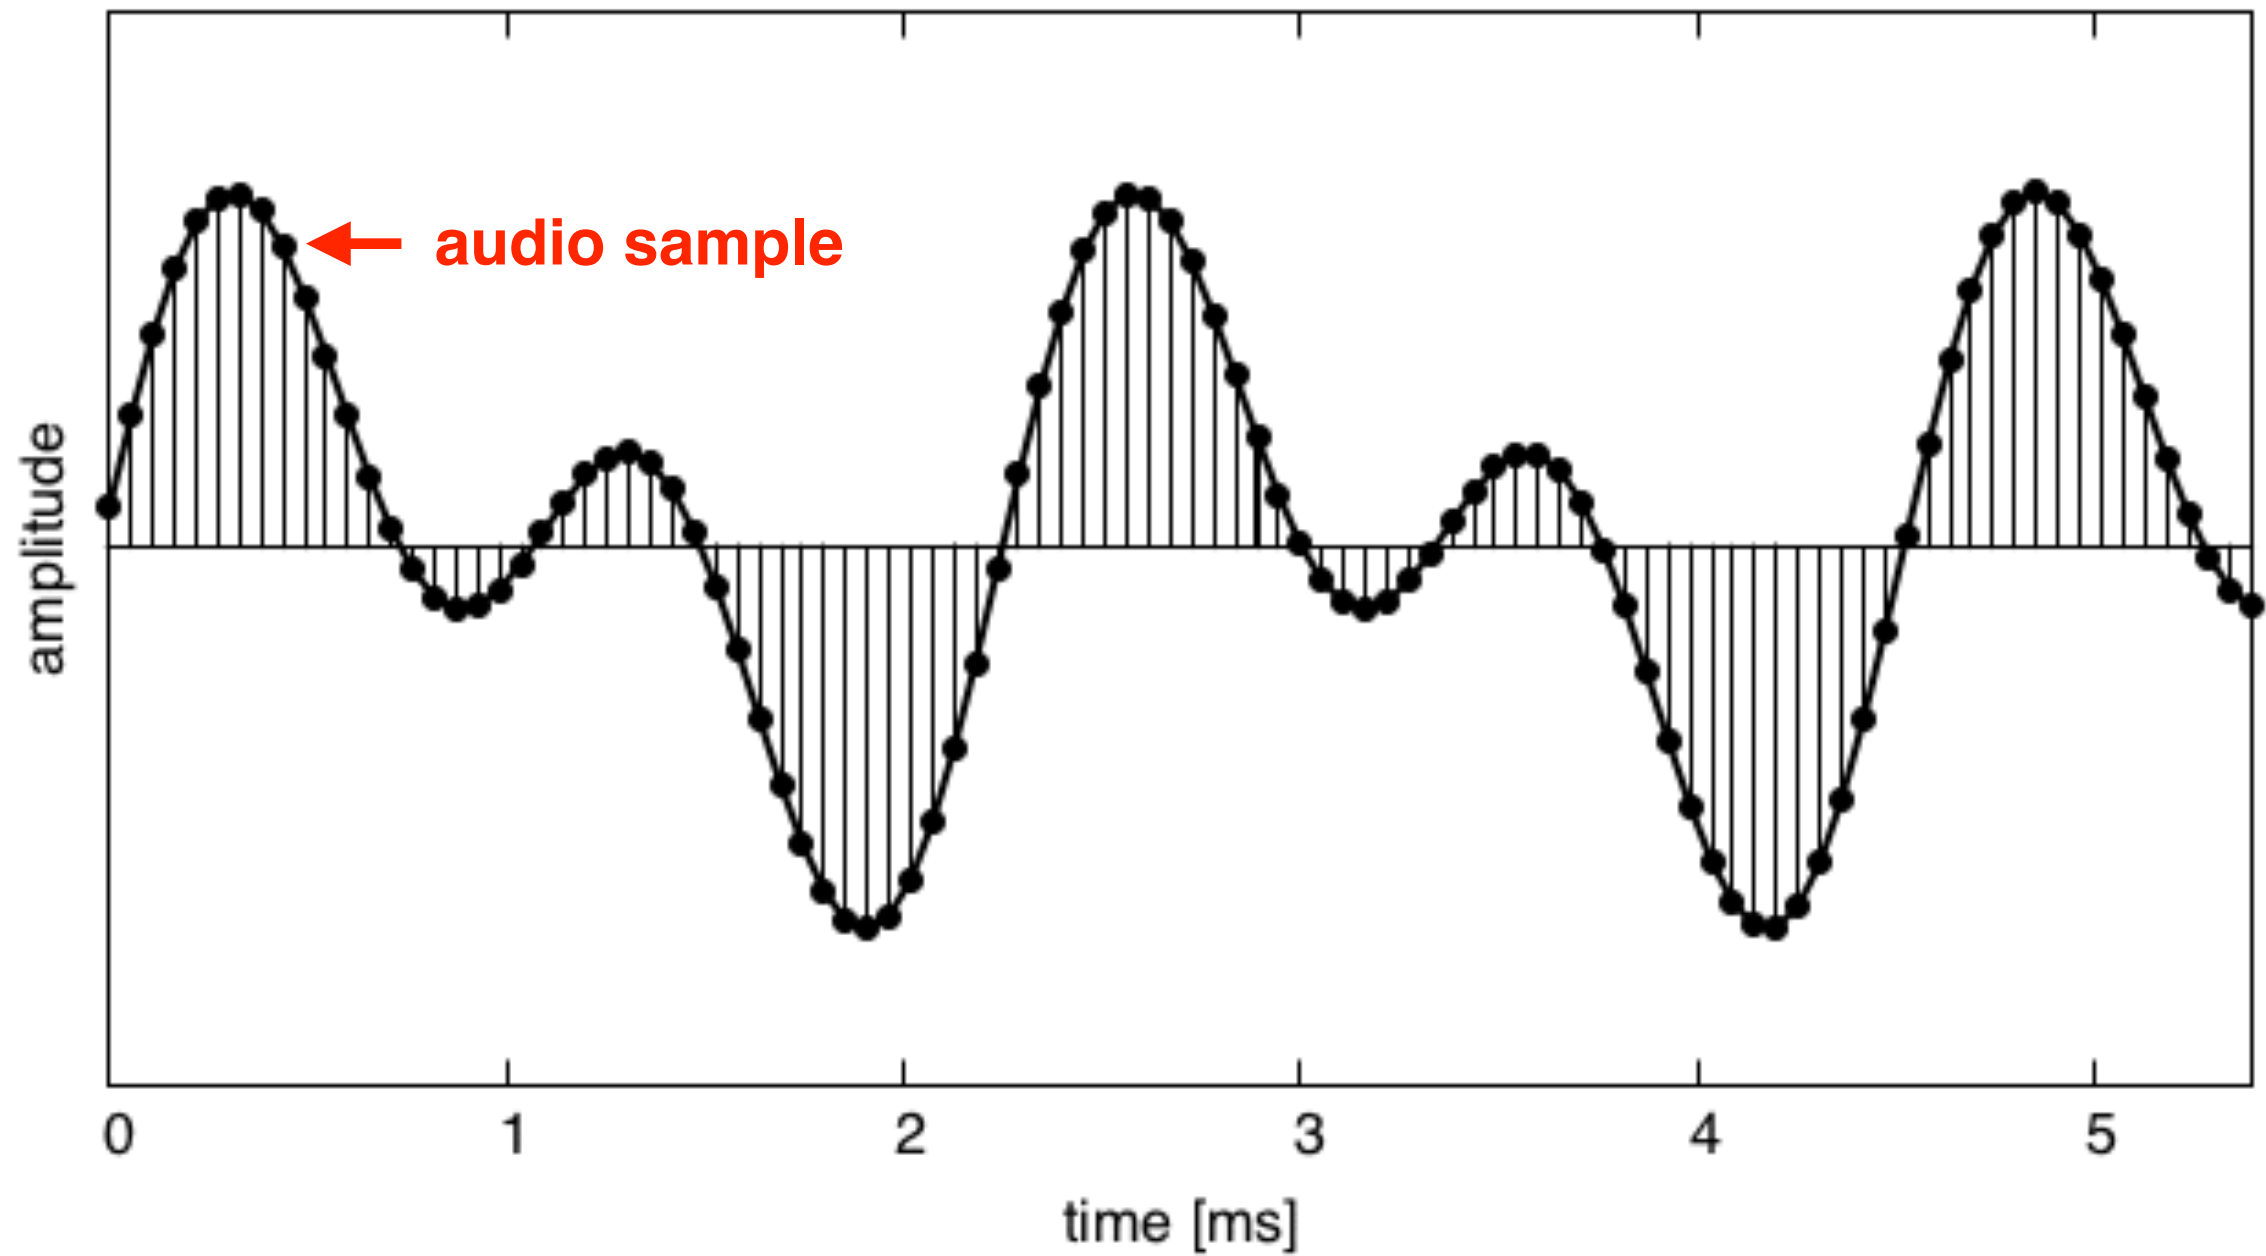
Representation of audio data in C++



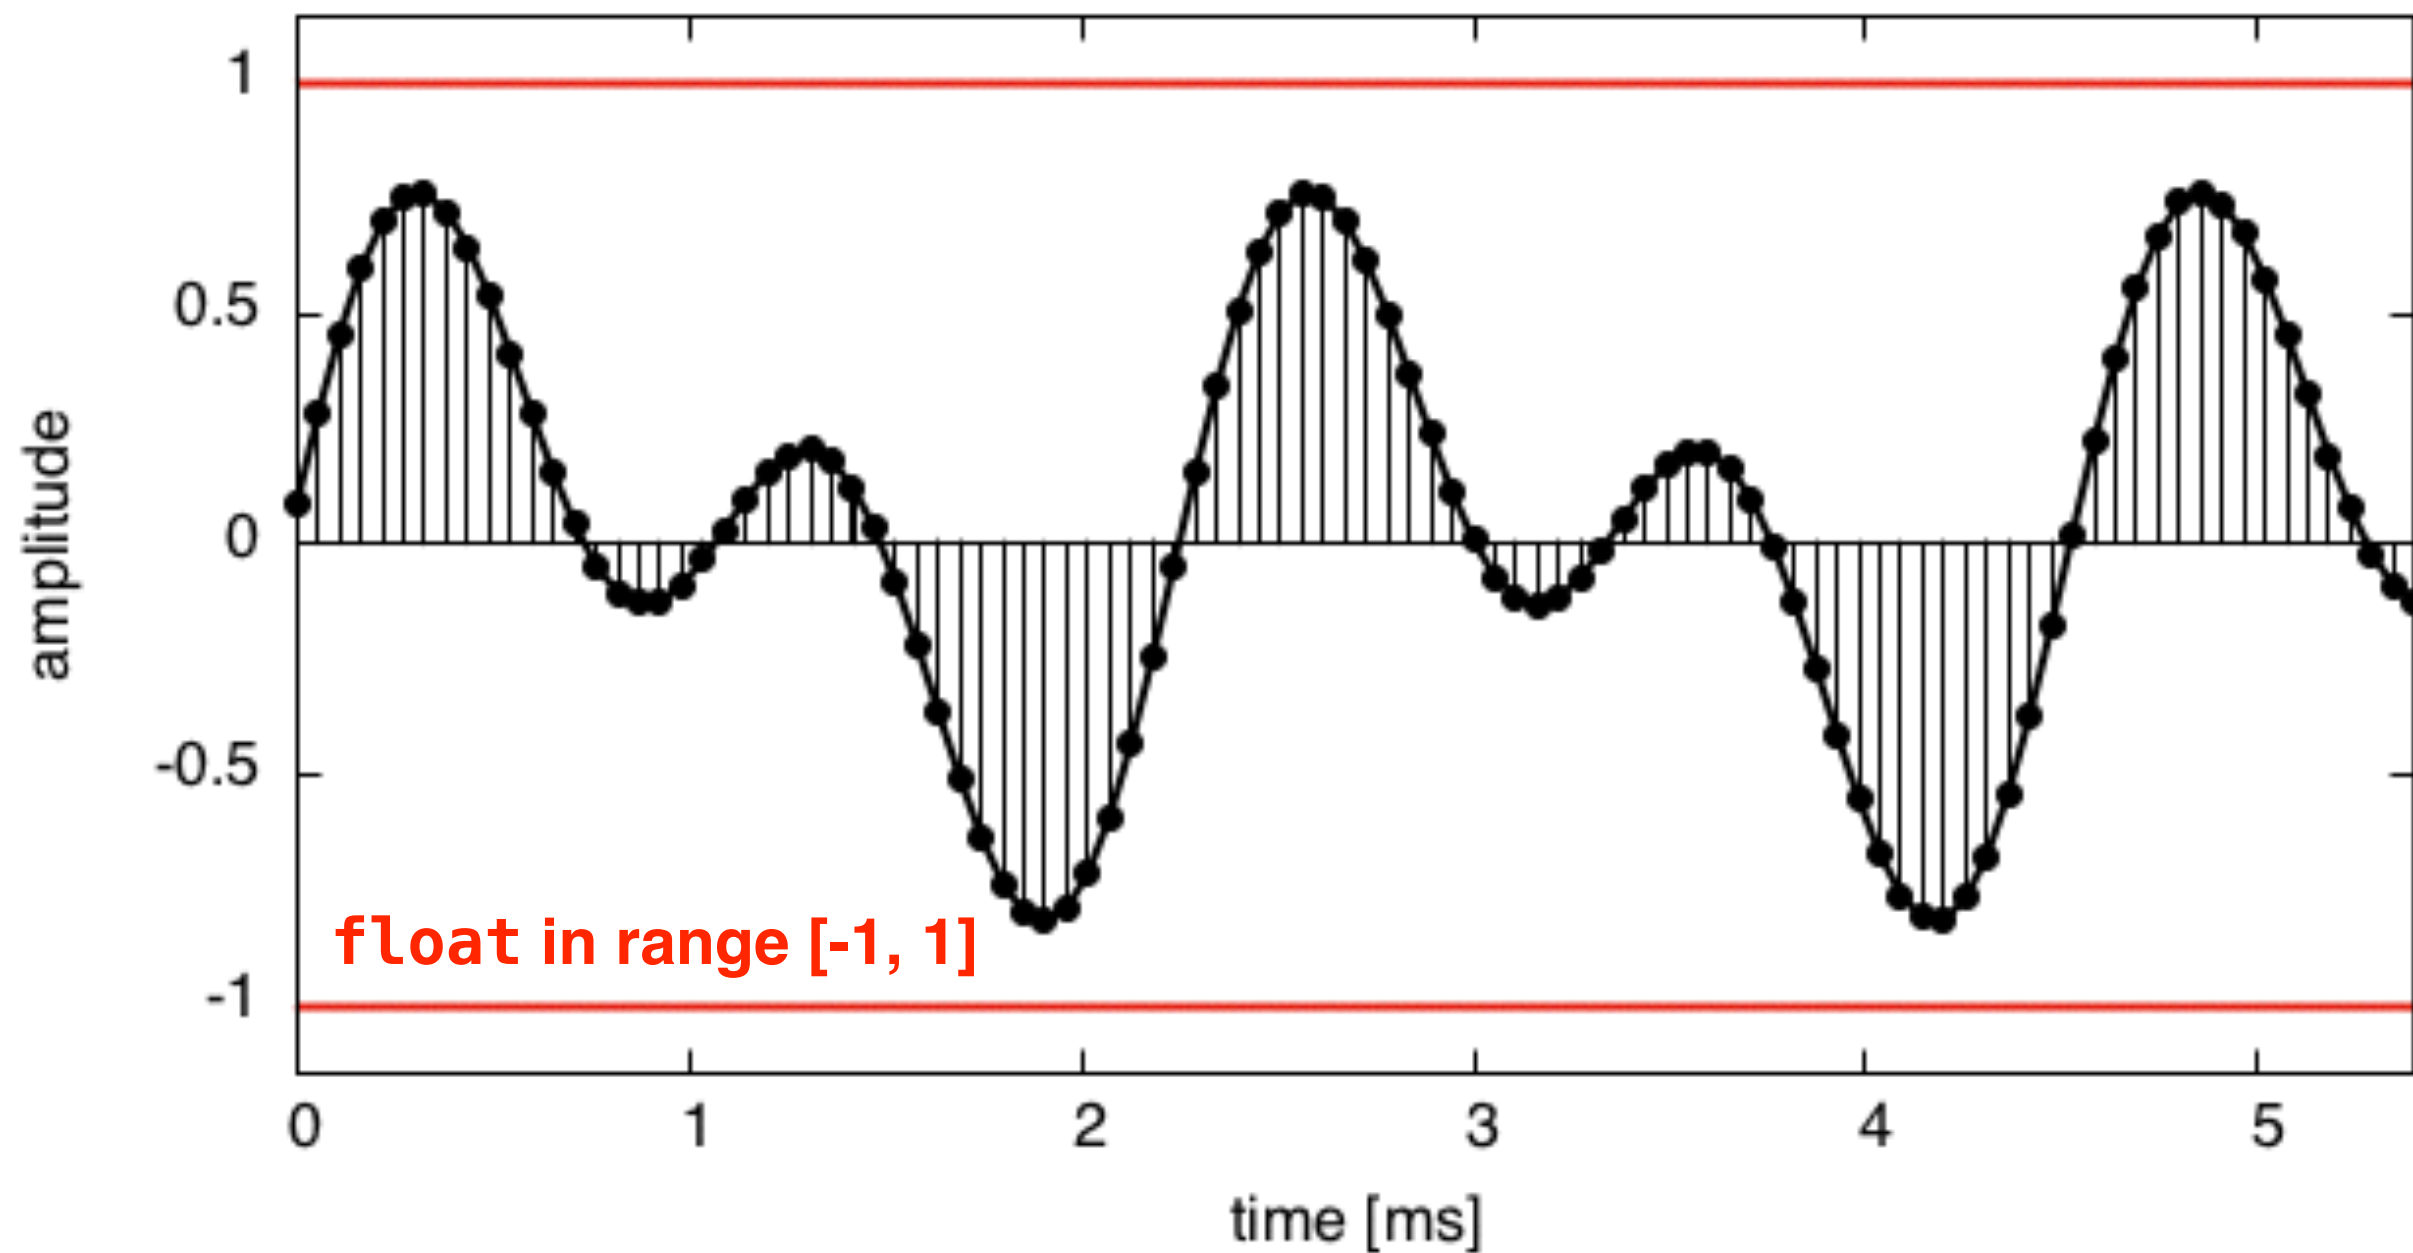
Representation of audio data in C++

sample rate

44.1 kHz, 48 kHz, 96 kHz, ...

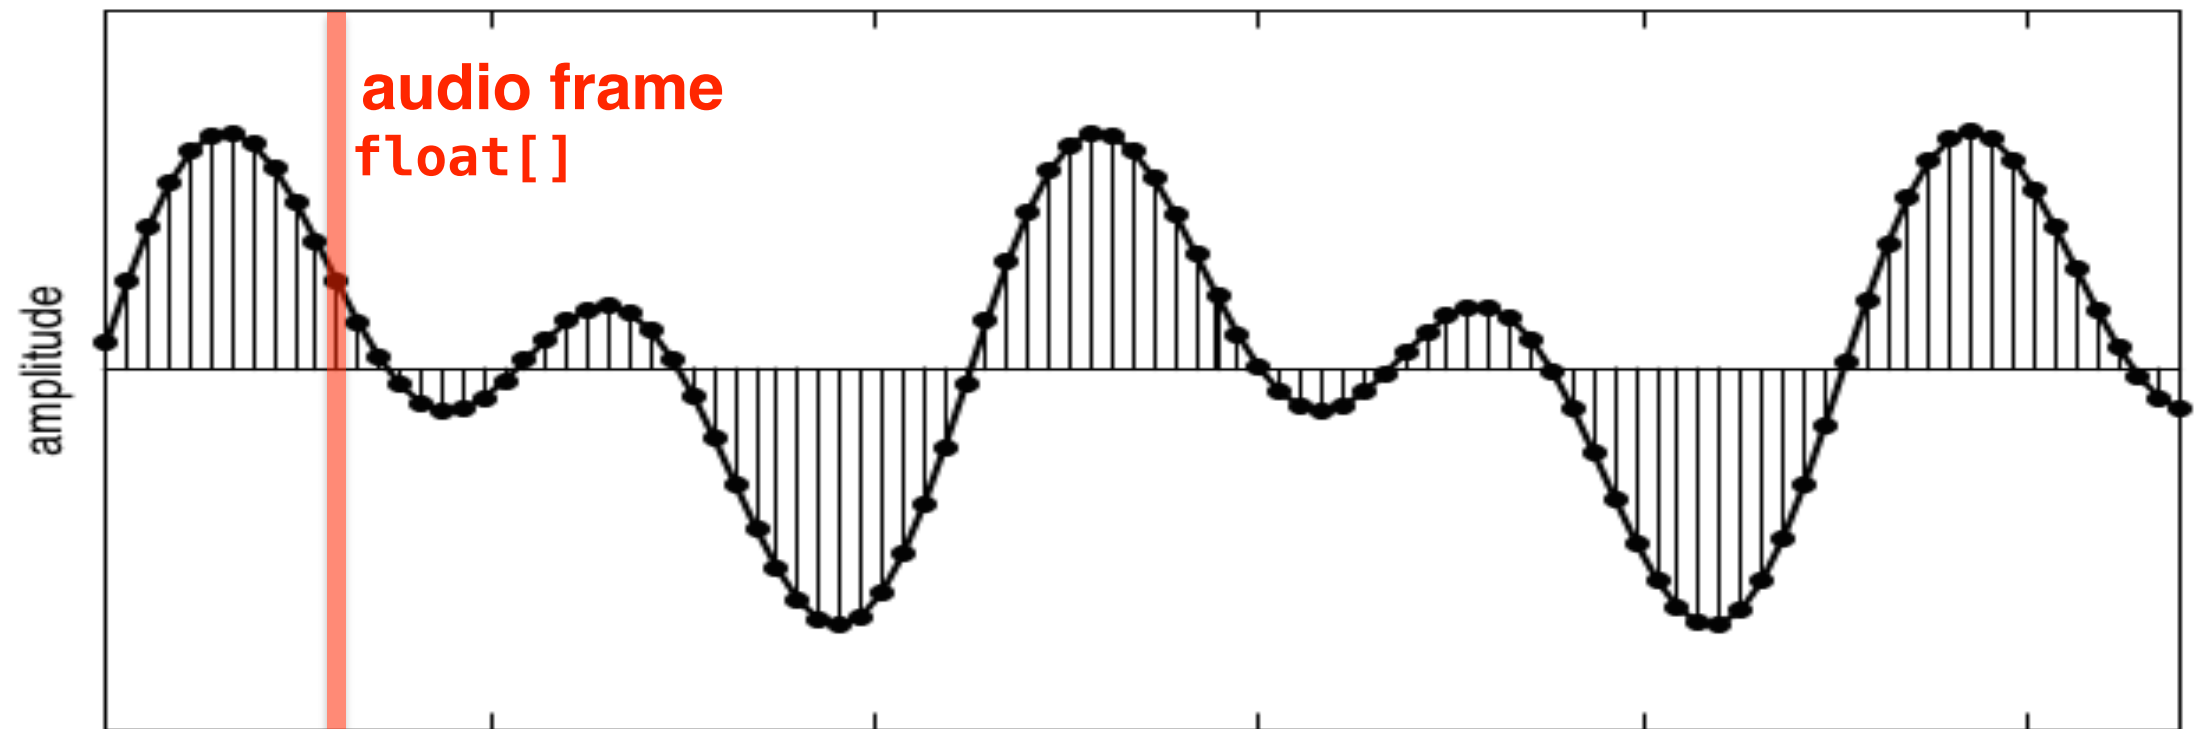


Representation of audio data in C++

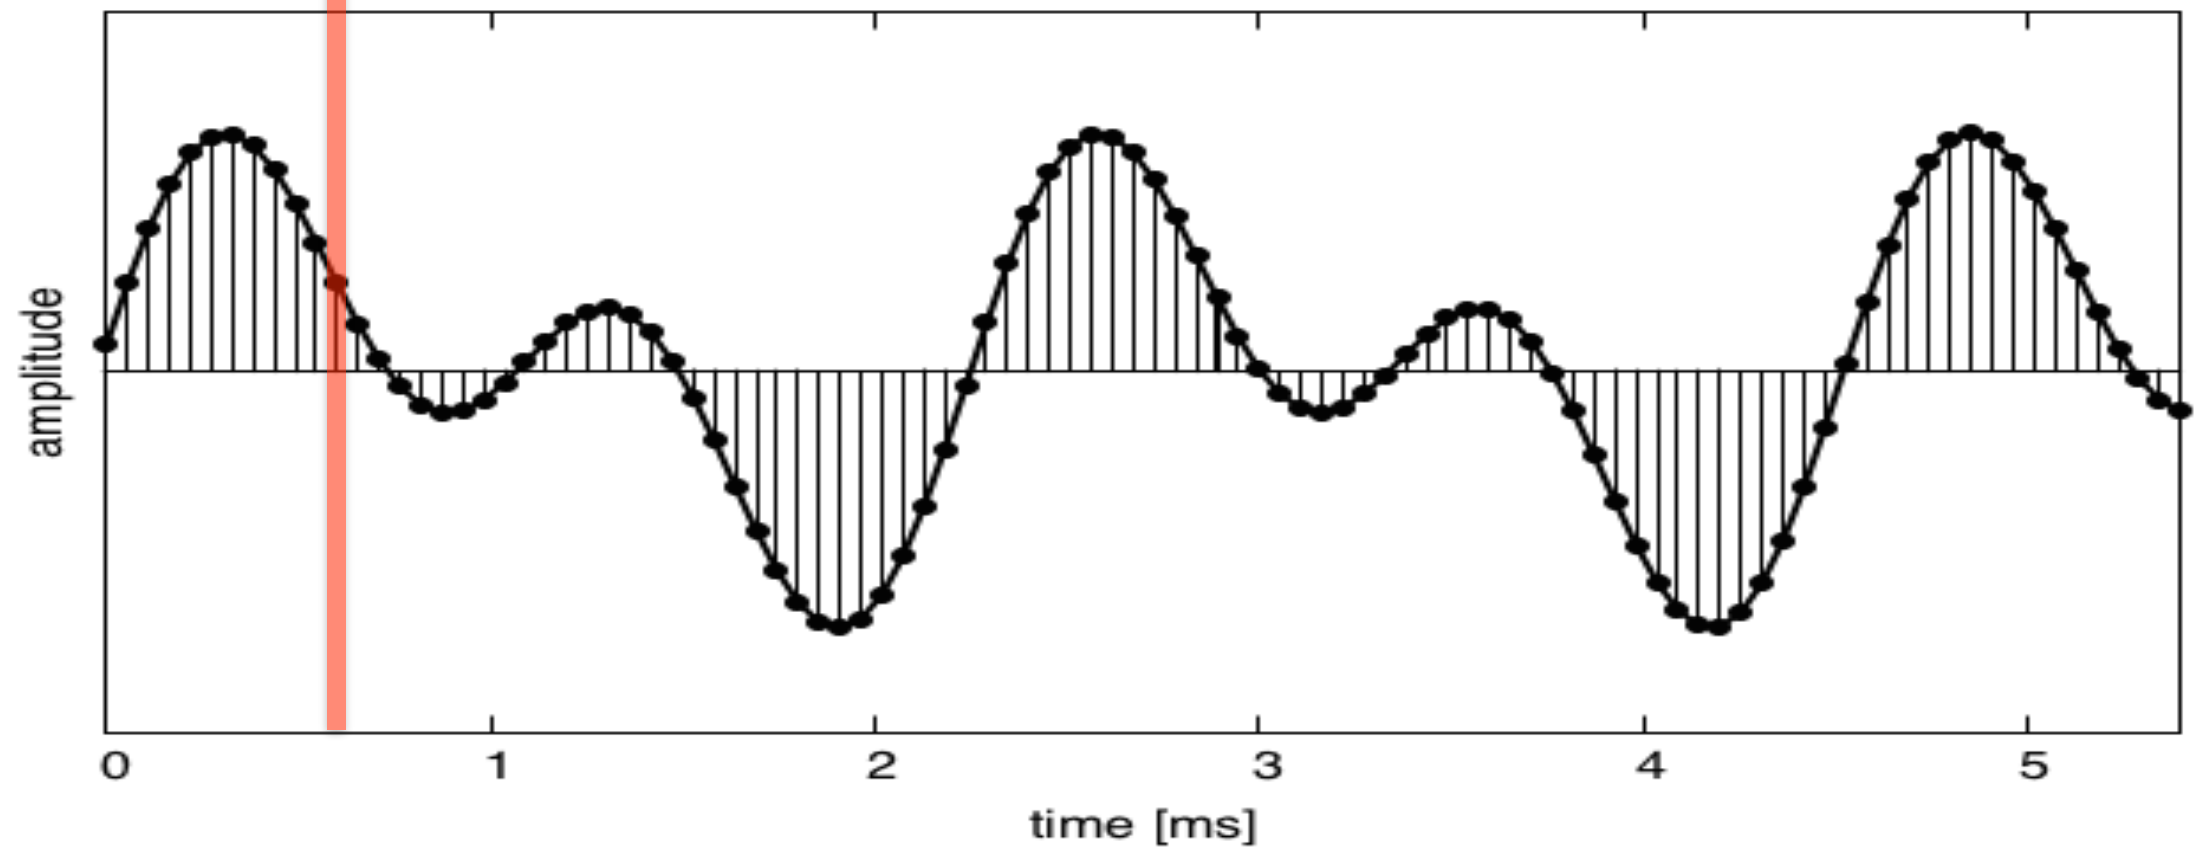


Representation of audio data in C++

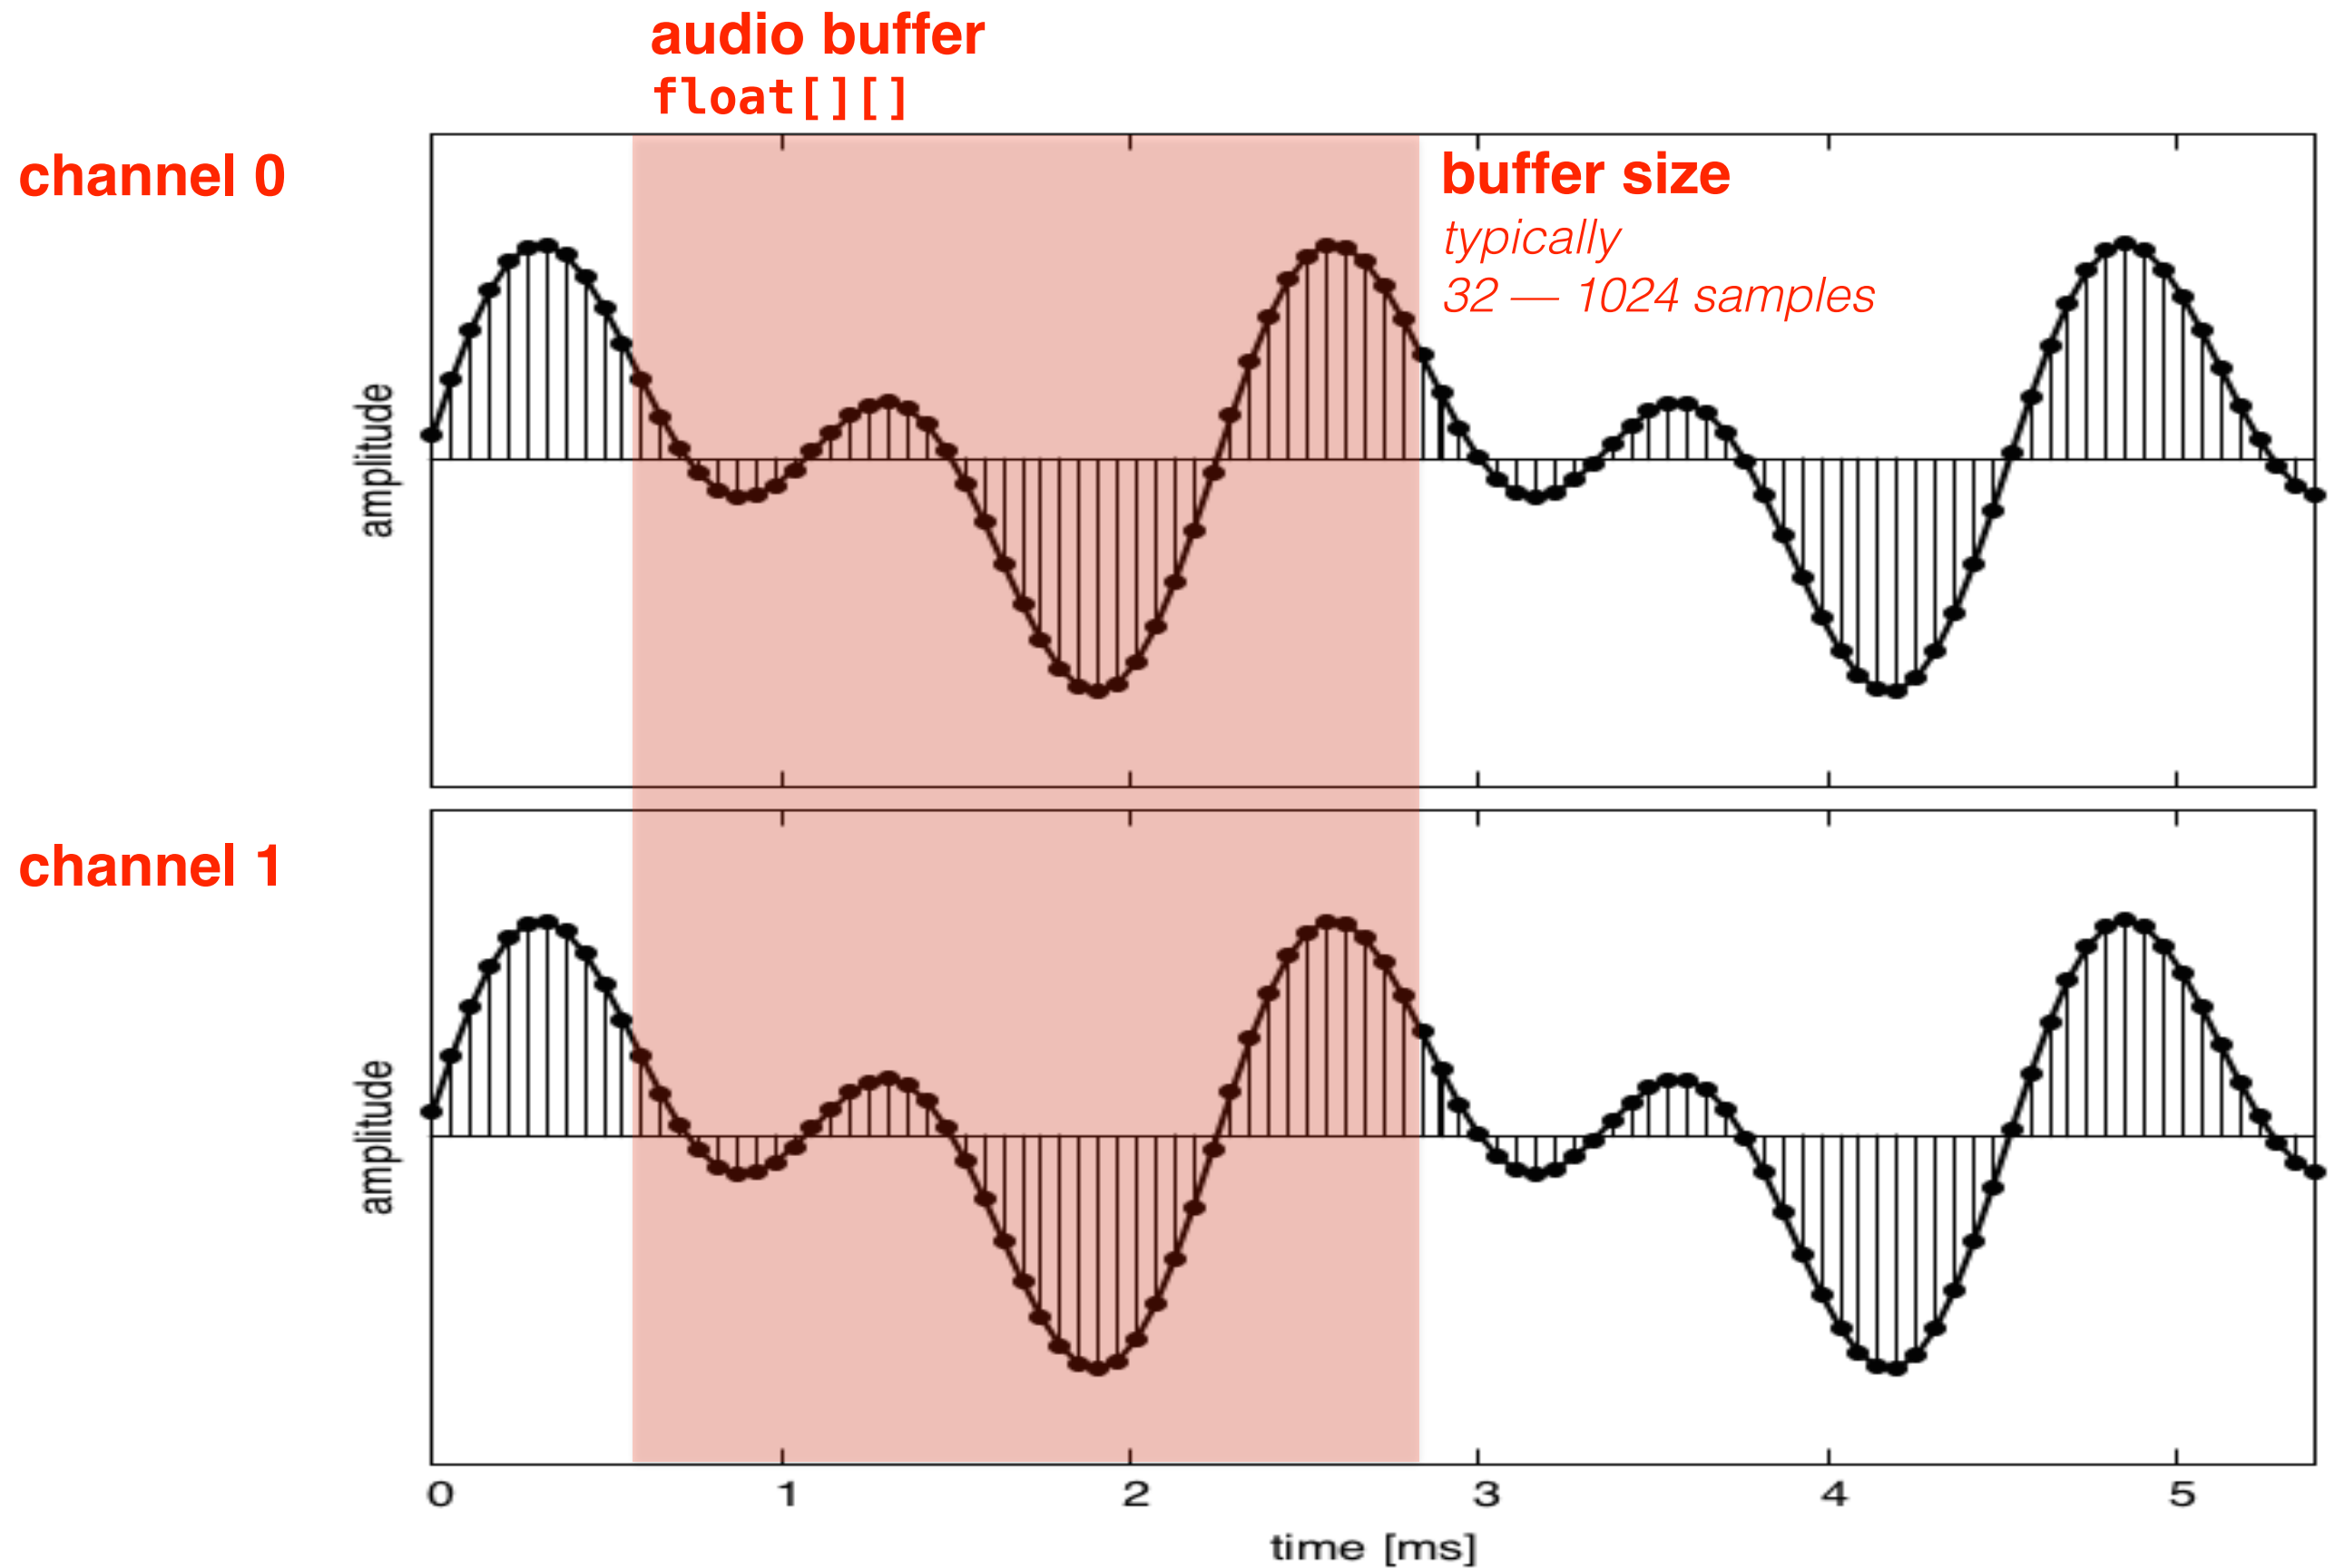
channel 0



channel 1



Representation of audio data in C++



The audio callback

```
void audioCallback (float** channelData,  
                    int numChannels,  
                    int numSamples)  
{  
    for (int channel = 0; channel < numChannels; ++channel)  
        for (int sample = 0; sample < numSamples; ++sample)  
            channelData[channel][sample] = 0.0f;  
}
```

The audio callback

```
void audioCallback (float** inputData,  
                    int numInputChannels,  
                    float** outputChannelData,  
                    int numOutputChannels,  
                    int numSamples)  
{  
    // ...  
}
```

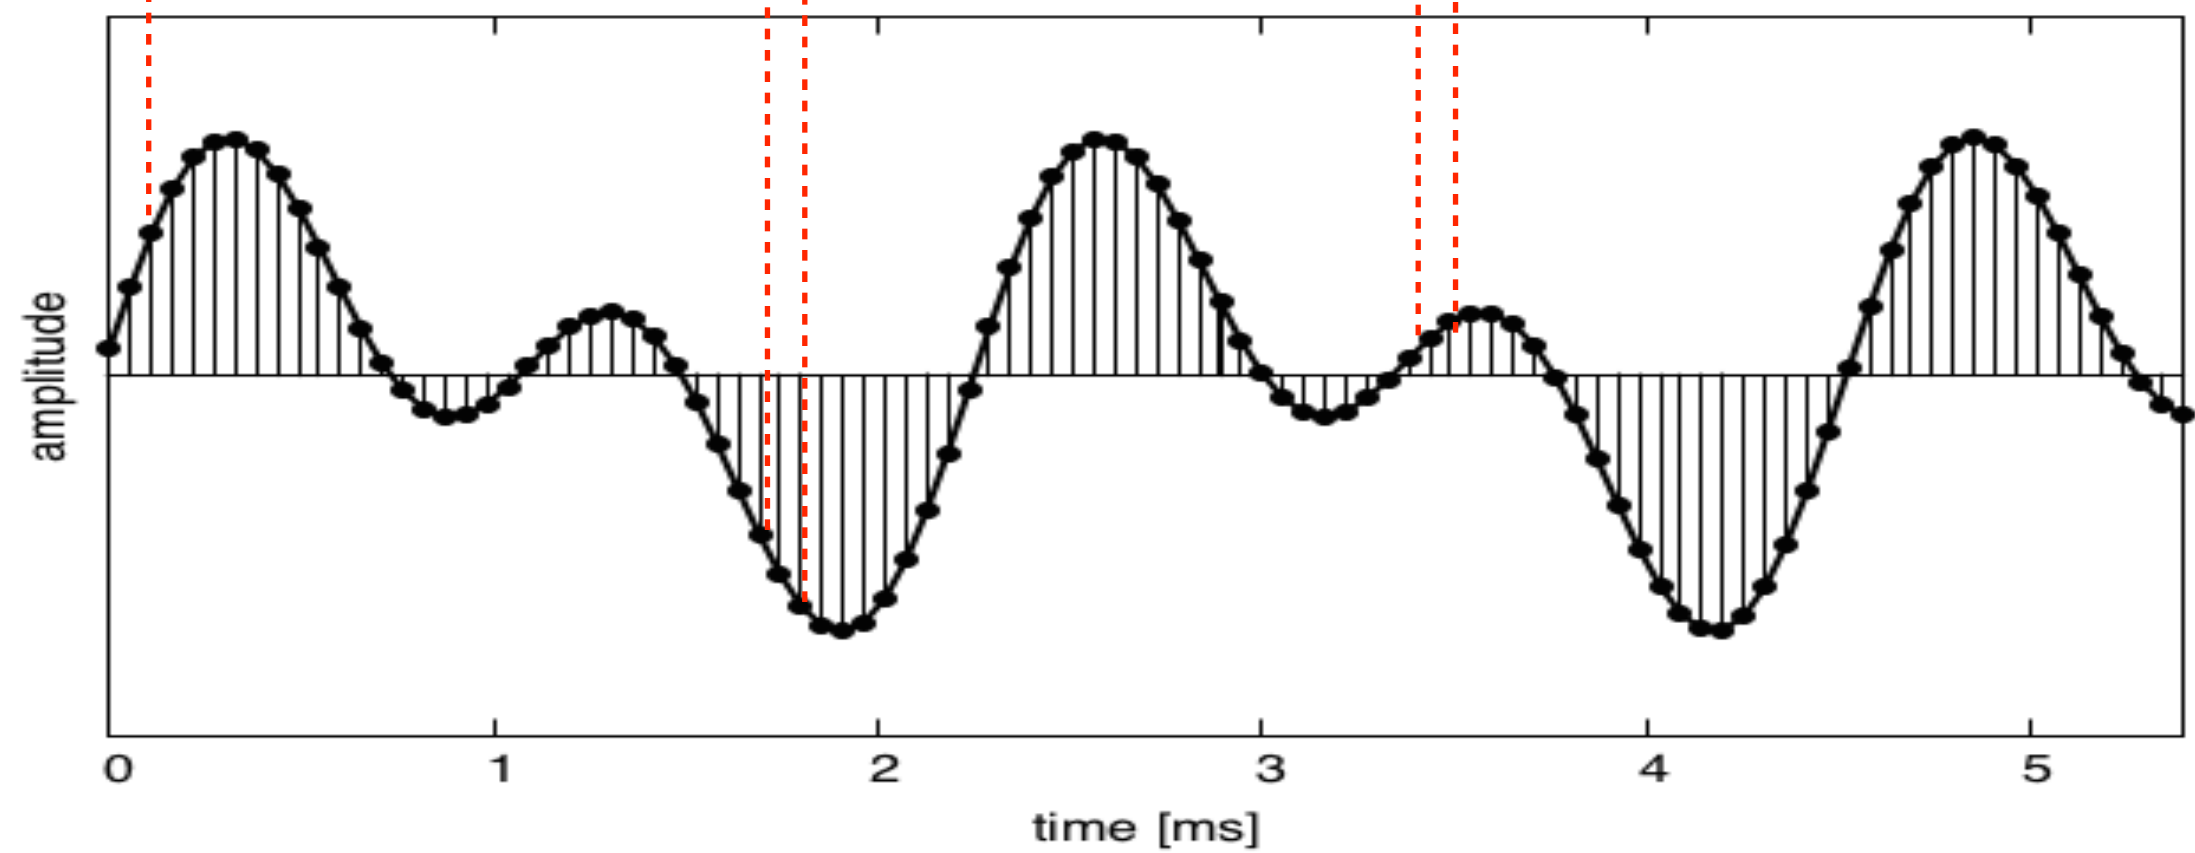


audio programming is
“hard real-time” programming

one audio buffer

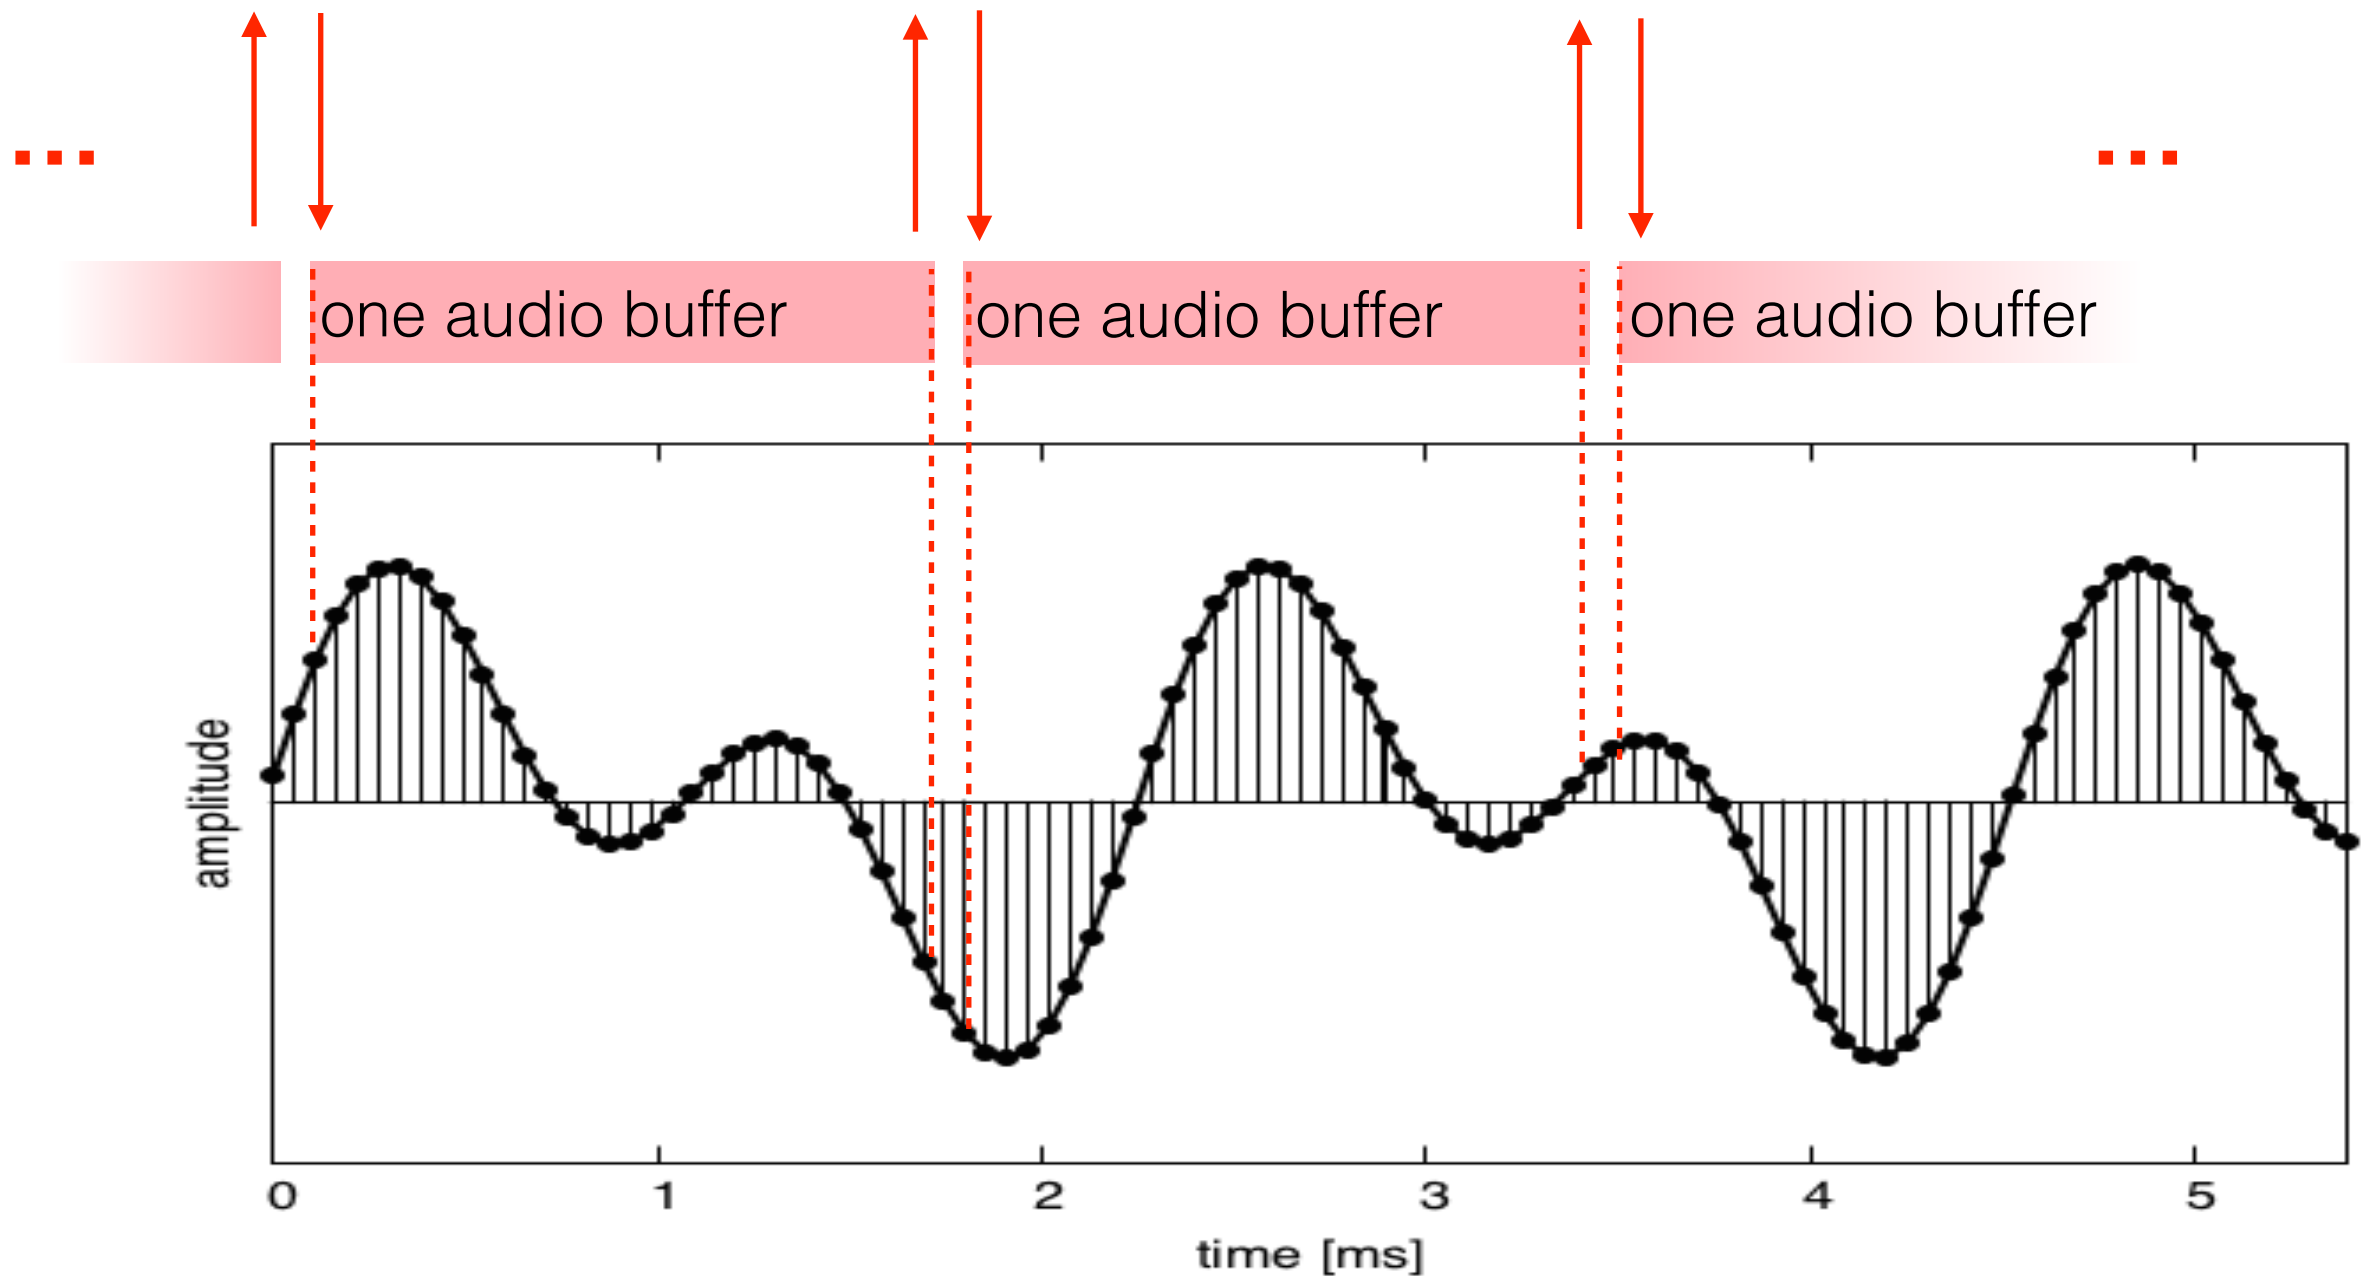
one audio buffer

one audio buffer





hard real-time audio callback





#0 rule of audio code:
the audio callback
waits for nothing.

Buffer size	Buffer length	Buffer length
	@ sample rate 44.1 kHz	@ sample rate 96 kHz
32 samples	0.73 ms	0.33 ms
64 samples	1.45 ms	0.66 ms
128 samples	2.90 ms	1.33 ms
...
1024 samples	23.2 ms	10.7 ms

Total system audio latency:

< 10 ms: *good for interactive audio performance*

10-30 ms: *noticeable*

> 30 ms: *OK if response doesn't have to feel real-time*


```
void audioCallback (float** channelData,  
                    int numChannels,  
                    int numSamples)  
{  
    for (int channel = 0; channel < numChannels; ++channel)  
        for (int sample = 0; sample < numSamples; ++sample)  
            channelData[channel][sample] = ...;  
}
```



...complex DSP calculations here...

we have to guarantee that:

- this function will return in time < buffer length
- will finish processing the buffer
- the output buffer will contain valid audio data afterwards
- there will be no error / exception / ...

```
void audioCallback (float** channelData,  
                    int numChannels,  
                    int numSamples) noexcept  
{  
    for (int channel = 0; channel < numChannels; ++channel)  
        for (int sample = 0; sample < numSamples; ++sample)  
            channelData[channel][sample] = ...;  
}
```



...complex DSP calculations here...

we have to guarantee that:

- this function will return in time < buffer length
- will finish processing the buffer
- the output buffer will contain valid audio data afterwards
- there will be no error / exception / ...

```
void audioCallback (float** channelData,  
                    int numChannels,  
                    int numSamples) noexcept  
{  
    for (int channel = 0; channel < numChannels; ++channel)  
        for (int sample = 0; sample < numSamples; ++sample)  
            channelData[channel][sample] = ...;  
}
```



...complex DSP calculations here...

we have to guarantee that:

- this function will return in time < buffer length
- will finish processing the buffer
- the output buffer will contain valid audio data afterwards
- there will be no error / exception / ...

otherwise we will cause an **audio dropout**.

#1 rule of audio code:
never cause audio dropouts.



#1 rule of audio code:
never cause audio dropouts.



#1 rule of audio code:
never cause audio dropouts.

Projects Settings SubWays T4

Double-click this scrollbar to zoom to fit 100%

Bar 1 Bar 5 Bar 9 Bar 13 Bar 17 Bar 21 Bar 25 Bar 29 Bar 33 Bar 37 Bar 41 Bar 45 Bar 49 Bar 53 Bar 57 Bar 61 Bar 65

Aux Return
Ski Rize Studios Demo Track For Tracktion 3

303
303

drums1
Drums1

samp 1
Samp one

303 again
303 again

plano samp
plano samp

sax
sax

synth stab
synth stab

Organ (not Organ)
Organ

plano hi
plano hi

plano lo
plano lo

rhodes
Rhodes

bass sonar
bass sonar

keys
keys

hi hats
hi hats

hi hats2
hi hats2

Aux Return #1 Delay Reverb

A Sampler Delay Aux Send #1

A Sampler Compressor

A Aux Send #1

A Sampler's K3 Aux Send #1

A Sampler Aux Send #1

A Aux Send #1 -5.8 dB

A Aux Send #1

A Sampler Phaser Aux Send #1

A -11.1 dB Aux Send #1

A -4.0 dB Aux Send #1

A -5.3 dB Aux Send #1

A -1.5 dB

A -4.8 dB Aux Send #1

A Sampler Aux Send #1

A -5.9 dB

Undo Redo

Save Import Timecode Snapping Options Movies Help

Clipboard Export Click Track Tracks Automation About

Audio Clip - "plano lo"

Name plano lo Gain +0.00 dB

Start 33 | 1 | 000 Pan +0.00

Length 7 | 2 | 480 Active Channels Left Right

End 40 | 3 | 480 Fade In 0 | 1 | 205

Offset 8 | 0 | 000 Fade Out 0 | 2 | 835

Colour

Disable Looping Apply X-Fade Drag X-Fade

Apply Edge Fade Copy Fade to Automation

Loop Properties

Stretch No Time-Stretching

Speed 1.000

Change Speed

Pitch 0.000

Change Pitch

Remap on Tempo Change

Reverse

View Source Info

Select Clips Auto-Tempo Split Clips Copy Marked Section Move Clip Render Clip Delete

100.00 bpm 4/4 C 32 | 1 | 912

Compressor

Loop Click

Auto Lock Punch

Snap Scroll

MIDI Learn MTC

#1 rule of audio code:
never cause audio dropouts.

- C++ in itself is not a real-time safe language.
- You have to fulfil the real-time guarantees by how you write code.
- There are certain dos and don'ts.

#2 rule of audio code:
don't do anything
where you can't predict
how much time it will take.



Don't block on the audio thread.

- Don't call anything that blocks / waits
 - `std::mutex::lock()`
 - `while (! std::mutex::try_lock()) { ... }`
 - `std::thread::join()`
- Don't call anything that blocks / waits internally

Do:

- lock-free programming
—> *examples follow...*

Why you shouldn't block / wait in the audio thread.

- reason #1: **you don't know how long it takes**
 - in most cases you are waiting for code that is not real-time safe
 - > *you get audio drop-outs*
- reason #2: **priority inversion**
 - audio thread is a high-priority system thread
 - when blocked, it is waiting for a lower-priority thread
 - the lower-priority thread may be interrupted by other threads
 - > *you get audio drop-outs*

#3 rule of audio code:

in the audio callback, only run lock-free code.



Don't (de)allocate any memory on the audio thread.

- don't call `new` or `delete`
- don't call anything that uses them internally:
 - constructing RAI objects or let them go out of scope
 - many, many other functions
 - > for example, `std::vector::push_back()` may reallocate!

Do:

- put your data on the stack
- preallocate your data
- be careful with object lifetime and ownership
- use containers from Boost.Intrusive and Boost.Lockfree
- use custom real-time safe containers (lock-free queues, stacks, lists...)
- *for larger systems:* manage your own preallocated, deterministic, lock-free heap

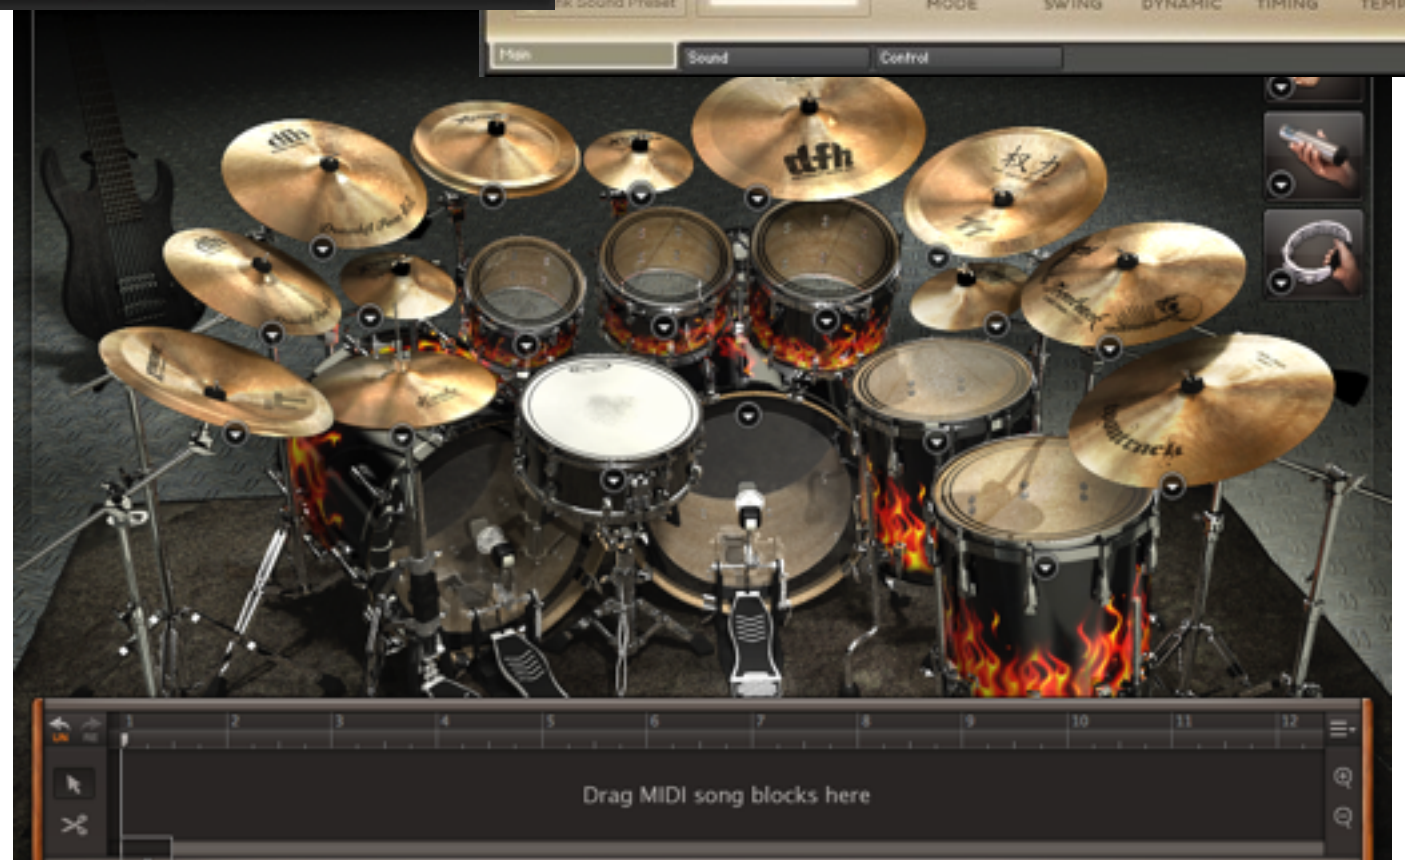
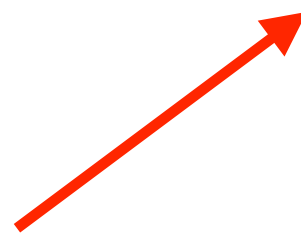
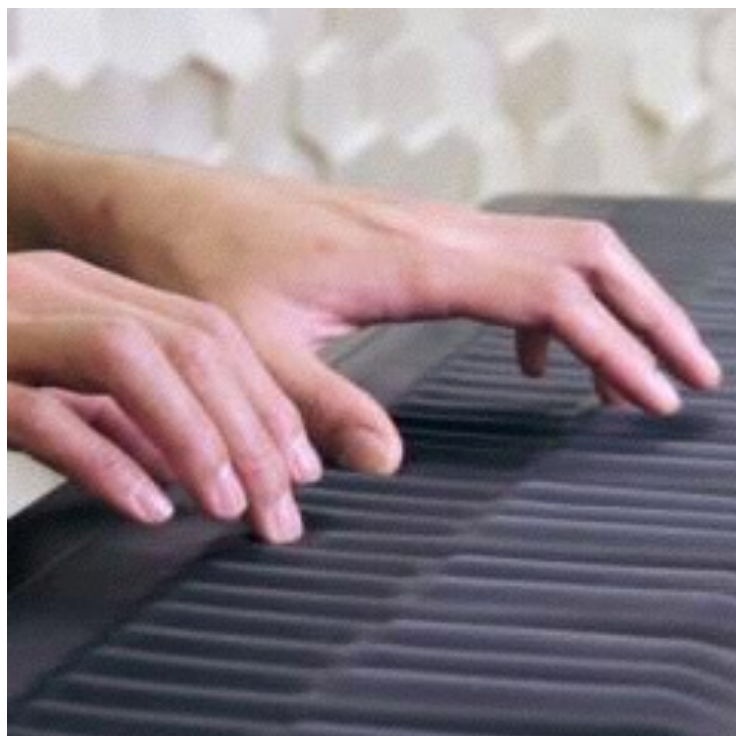
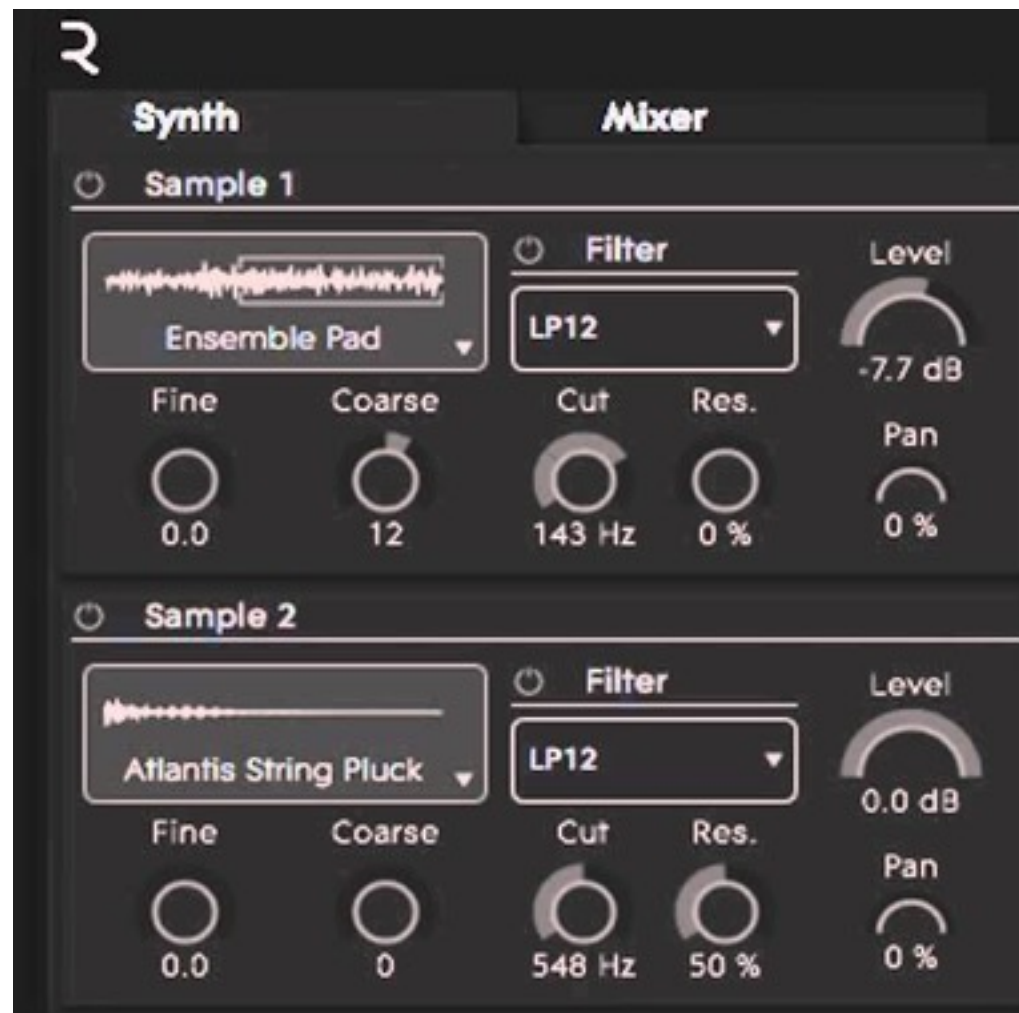
Don't do any I/O on the audio thread.

- don't call `std::cout`, `printf` etc.
- don't do any interprocess communication
- don't access files on the disk, the network, ...
- don't do any graphics work

Do:

- use other threads for all of the above
- synchronise them lock-free

- Don't call any 3rd party code.
(if you don't know what's going on inside and how long it will take)
- Don't call algorithms with unsafe worst-case behaviour.
basically: $O(1)$ is fine, amortised $O(1)$ is not
—> for example, `std::unordered_set::insert()` may rehash
- Don't run into page faults.



Don't run into page faults.

- solution #1:

use a low-priority thread to regularly “poke”
the memory you will use on the audio thread

- solution #2:

lock your real-time data into memory to prevent page-outs

- Modern OSes offer APIs for that:

`mlock()/munlock()` (POSIX systems),

`VirtualLock()/VirtualUnlock()` (Windows)

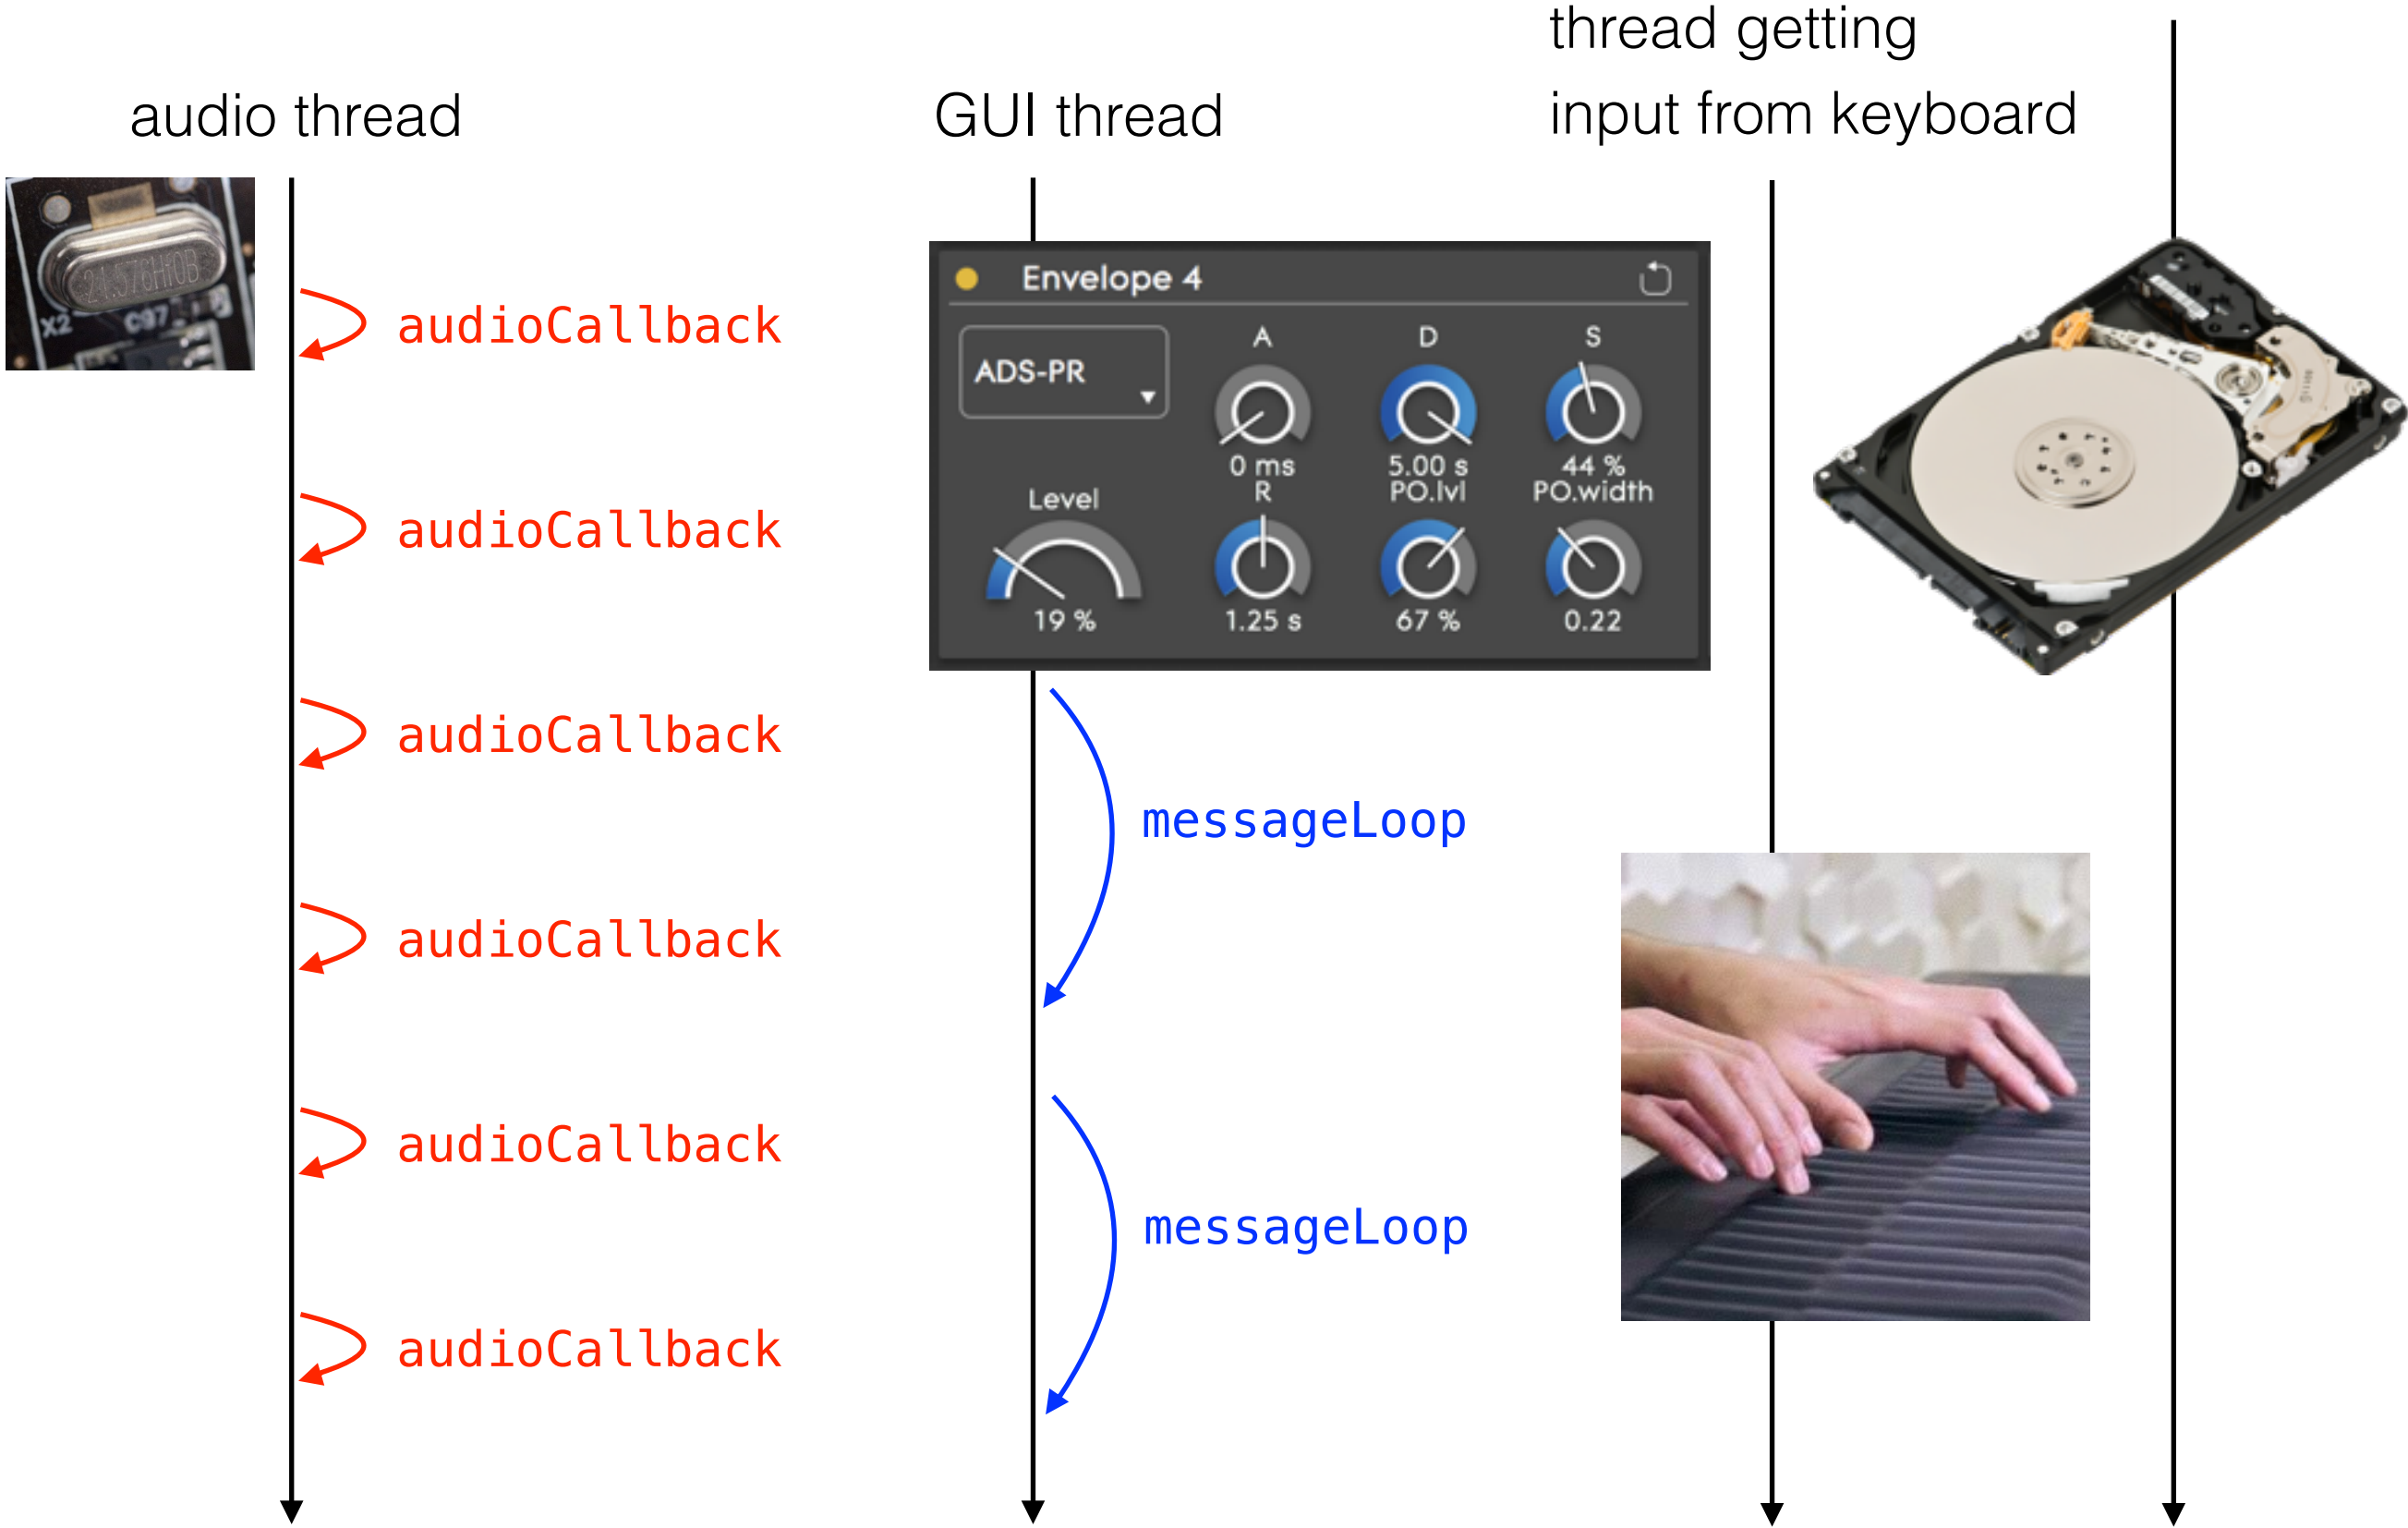
(nothing in the C++ standard or Boost 😞)

- wrap calls into a smart pointer class
- wrap calls into an STL allocator
- pre-allocate a locked memory pool and use `boost::pool_allocator`
or hand-written memory management

so why C++?

- most audio APIs use C or C++
- fast
- close to the metal
- allows you to use custom memory management
- supports concurrency, atomic types and lock-free programming
- allows you to write real-time safe code

Safe & lock-free thread synchronisation.



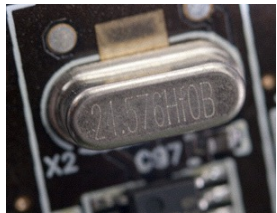
Safe & lock-free thread synchronisation.

thread reading
data from disk

thread getting
input from keyboard

audio thread

GUI thread



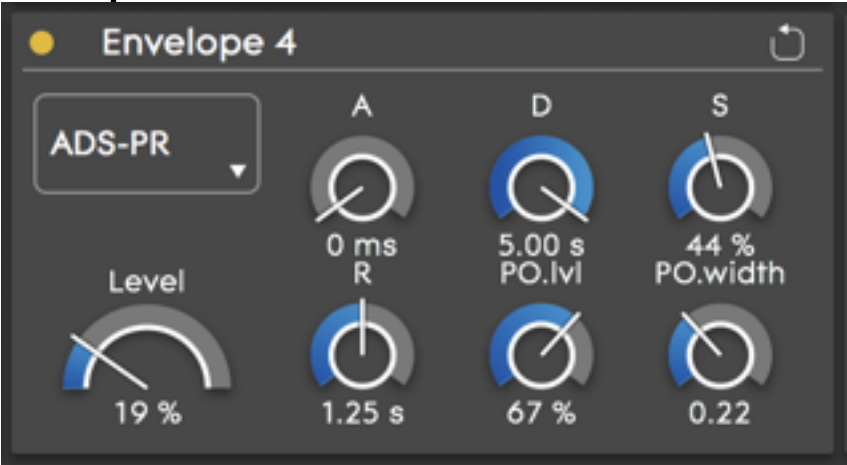
audioCallback

audioCallback

audioCallback

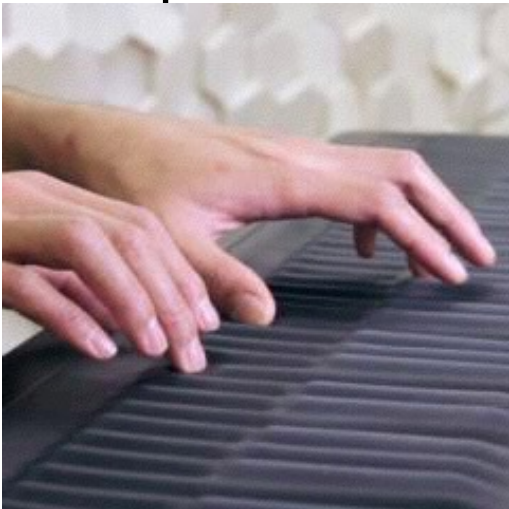


audioCallback

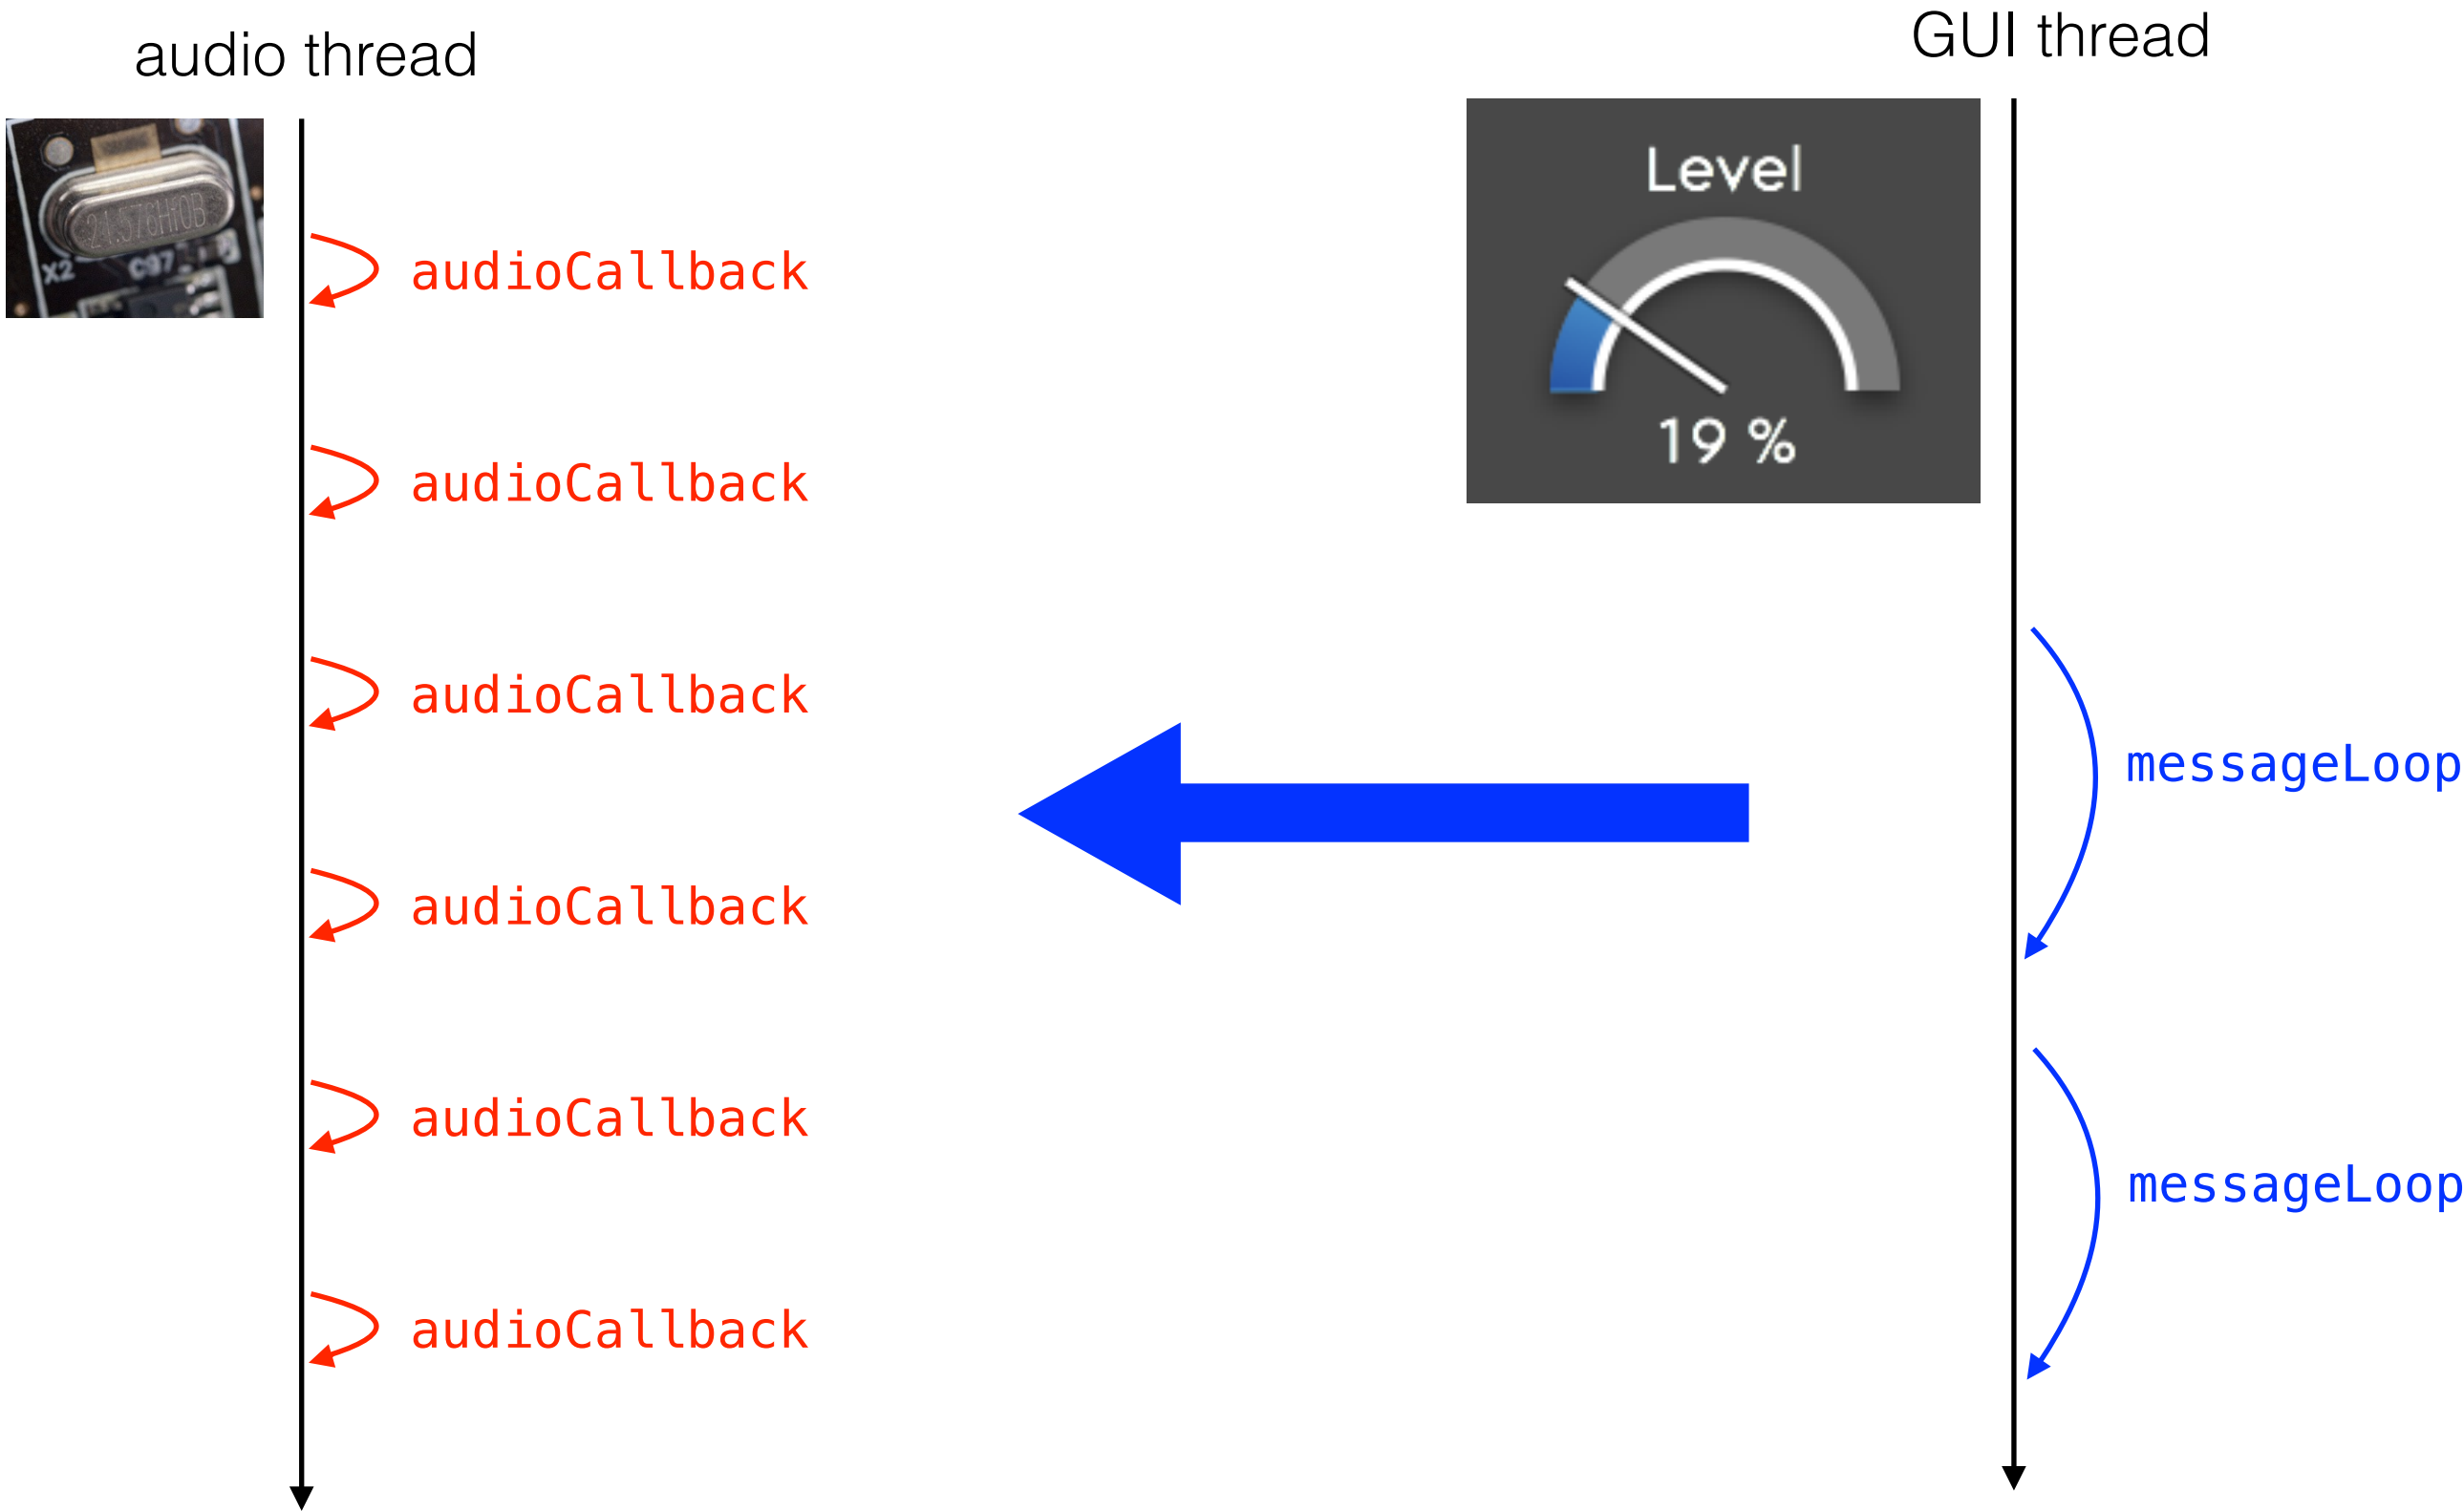


messageLoop

messageLoop



Safe & lock-free thread synchronisation.



```
class Synthesiser
{
public:

    Synthesiser() : level (1.0f) {}

    // GUI thread:
    void levelChanged (float newValue)
    {
        level = newValue;
    }

private:

    // Audio thread:
    void audioCallback (float* buffer,
                        int numSamples) noexcept
    {
        for (int i = 0; i < numSamples; ++i)
            buffer[i] = level * getNextAudioSample();
    }

    float level;
};
```

```
class Synthesiser
{
public:

    Synthesiser() : level (1.0f) {}

    // GUI thread:
    void levelChanged (float newValue)
    {
        level = newValue;
    }

private:

    // Audio thread:
    void audioCallback (float* buffer,
                        int numSamples) noexcept
    {
        for (int i = 0; i < numSamples; ++i)
            buffer[i] = level * getNextAudioSample();
    }

    float level;
};
```

```
class Synthesiser
{
public:
```

```
    Synthesiser() : level (1.0f) {}
```

```
    // GUI thread:
```

```
    void levelChanged (float newValue)
    {
        level = newValue;
    }
```

```
private:
```

```
    // Audio thread:
```

```
    void audioCallback (float* buffer,
                        int numSamples) noexcept
    {
        for (int i = 0; i < numSamples; ++i)
            buffer[i] = level * getNextAudioSample();
    }
```

```
    float level;
```

```
};
```

**data race = undefined behaviour
since C++11!**



- Undefined behaviour
- Access is not atomic - torn reads & writes
- Compiler optimisations can break program logic
- Concurrency is not expressed in the code

Torn reads and writes

```
class NotAtomic
{
public:

    void setParameter()
    {
        parameter = 0x100000002;
    }

    std::uint64_t getParameter() const
    {
        return parameter;
    }

    std::uint64_t parameter;
};
```


Torn reads and writes

```
class NotAtomic
{
public:
```

```
    void setParameter()
    {
        parameter = 0x100000002;
    }
```



```
    ...
    mov     DWORD PTR parameter, 2
    mov     DWORD PTR parameter+4, 1
    ret
    ...
```

```
std::uint64_t getParameter() const
{
    return parameter;
}
```



```
    ...
    mov     eax, DWORD PTR parameter
    mov     edx, DWORD PTR parameter+4
    ret
    ...
```

```
std::uint64_t parameter;
```

```
};
```

compiled with gcc for 32-bit x86:

Breaking optimisations

```
class Foo
{
    void bar()
    {
        flag = false;

        while (flag == false)
        {
            i++;
        }
    }

    bool flag;
    int i;
};
```

Breaking optimisations

```
class Foo
{
    void bar()
    {
        flag = false;

        while (flag == false)
        {
            i++;
        }
    }

    bool flag;
    int i;
};
```

**optimising
compiler**



```
class Foo
{
    void bar()
    {
        flag = false;

        while (true)
        {
            i++;
        }
    }

    bool flag;
    int i;
};
```

Breaking optimisations

```
class Foo
{
    void bar()
    {
        flag = false;

        while (flag == false)
        {
            i++;
        }
    }

    volatile bool flag;
    int i;
};
```

Breaking optimisations

```
class Foo
{
    void bar()
    {
        flag = false;

        while (flag == false)
        {
            i++;
        }
    }

    volatile bool flag;
    int i;
};
```

- fixes this particular problem
- but does not fix any of the other problems
- also, prevents “good” compiler optimisations
—> **wrong!!**

```
class Synthesiser
{
public:

    Synthesiser() : level (1.0f) {}

    // GUI thread:
    void levelChanged (float newValue)
    {

        level = newValue;

    }

private:

    // Audio thread:
    void audioCallback (float* buffer,
                        int numSamples) noexcept
    {

        for (int i = 0; i < numSamples; ++i)
            buffer[i] = level * getNextAudioSample();

    }

    float level;

};
```



```
class Synthesiser
{
public:

    Synthesiser() : level (1.0f) {}

    // GUI thread:
    void levelChanged (float newValue)
    {
        std::lock_guard<std::mutex> (m);
        level = newValue;
    }

private:

    // Audio thread:
    void audioCallback (float* buffer,
                        int numSamples) noexcept
    {
        std::lock_guard<std::mutex> (m);
        for (int i = 0; i < numSamples; ++i)
            buffer[i] = level * getNextAudioSample();
    }

    float level;
    std::mutex m;
};
```



```
class Synthesiser
{
public:

    Synthesiser() : level (1.0f) {}

    // GUI thread:
    void levelChanged (float newValue)
    {

        level.store (newValue);
    }

private:

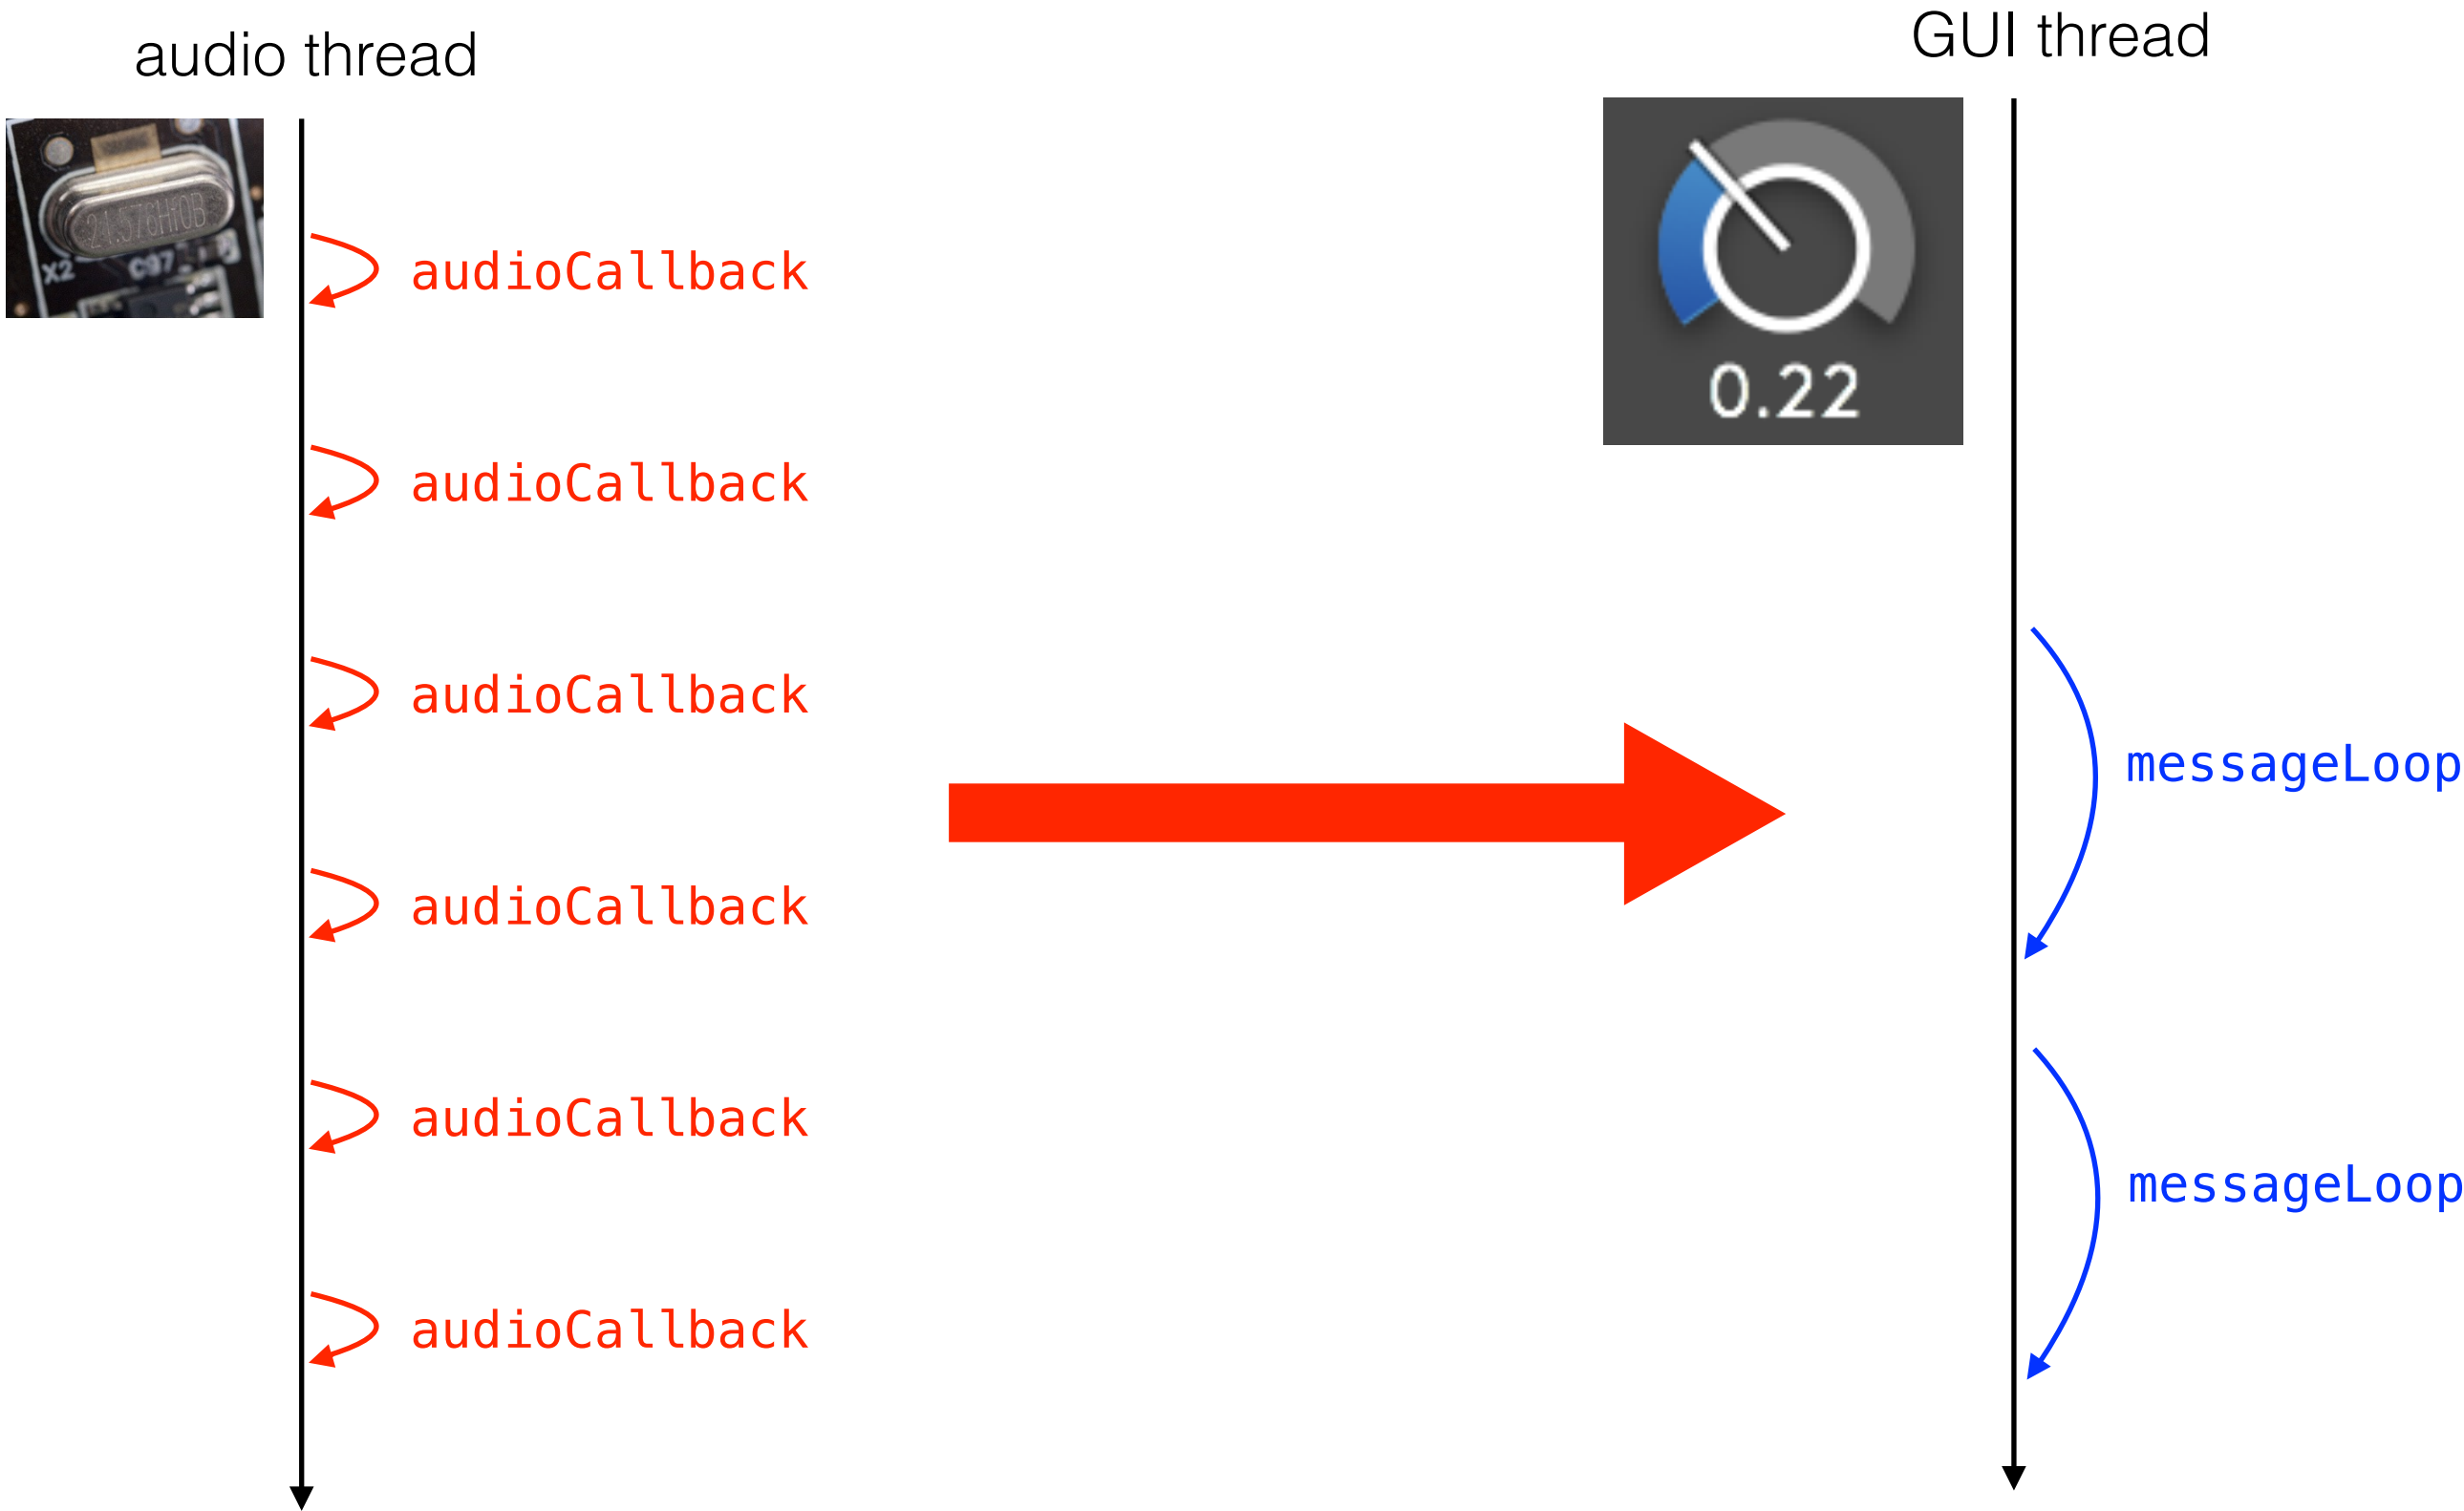
    // Audio thread:
    void audioCallback (float* buffer,
                        int numSamples) noexcept
    {

        for (int i = 0; i < numSamples; ++i)
            buffer[i] = level.load() * getNextAudioSample();
    }

    std::atomic<float> level;

};
```

Safe & lock-free thread synchronisation.



```
class Synthesiser

void audioCallback (float* buffer, int numSamples) noexcept
{
    // ... some DSP ...

    parameter.store (newValue);

    updateGui();
}

std::atomic<float> parameter;
};
```

 **never do this**

```
class Synthesiser

void audioCallback (float* buffer, int numSamples) noexcept
{
    // ... some DSP ...

    parameter.store (newValue);

    messageLoop.post (ParametersChangedNotification (parameter));
}

std::atomic<float> parameter;
};
```



- probably not lock-free
- congesting message loop with too many messages

```
class Synthesiser
```

```
void audioCallback (float* buffer, int numSamples) noexcept  
{
```

```
    // ... some DSP ...
```

```
    parameter.store (newValue);
```

```
    guiUpToDate.store (false);
```

```
}
```

```
std::atomic<float> parameter;
```

```
std::atomic<bool> guiUpToDate;
```

```
void timerCallback() // called 30x/second on a low priority thread  
{
```

```
    if (guiUpToDate.compare_exchange_strong (false, true))
```

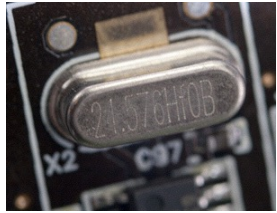
```
        updateGui (parameter.load());
```

```
}
```

```
};
```

Thread synchronisation

audio thread



audioCallback

audioCallback

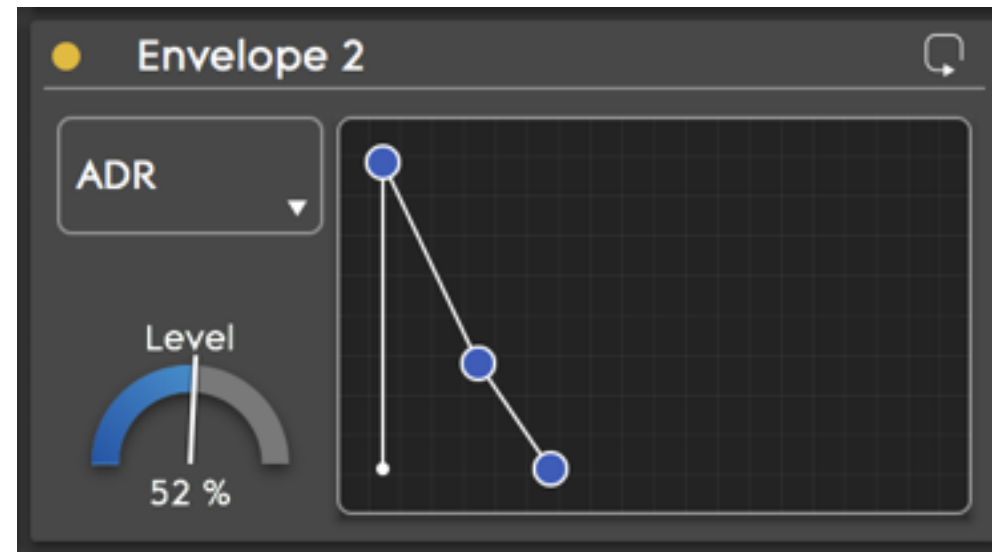
audioCallback

audioCallback

audioCallback

audioCallback

GUI thread



messageLoop

messageLoop

How to do this with an object?

- `std::atomic<Widget>` compiles if
`std::is_trivially_copyable<Widget>::value == true...`
- ...but it won't do what we want:
 - locks will be inserted by compiler
 - you can check with `std::atomic<Widget>::is_lock_free()`
- `std::atomic<Widget*>` is OK

Juggling with `std::atomic<Widget*> w` ...

audio thread

GUI thread

```
void audioCallback (...)
{
    Widget* widgetToUse;
    widgetToUse = w.load();

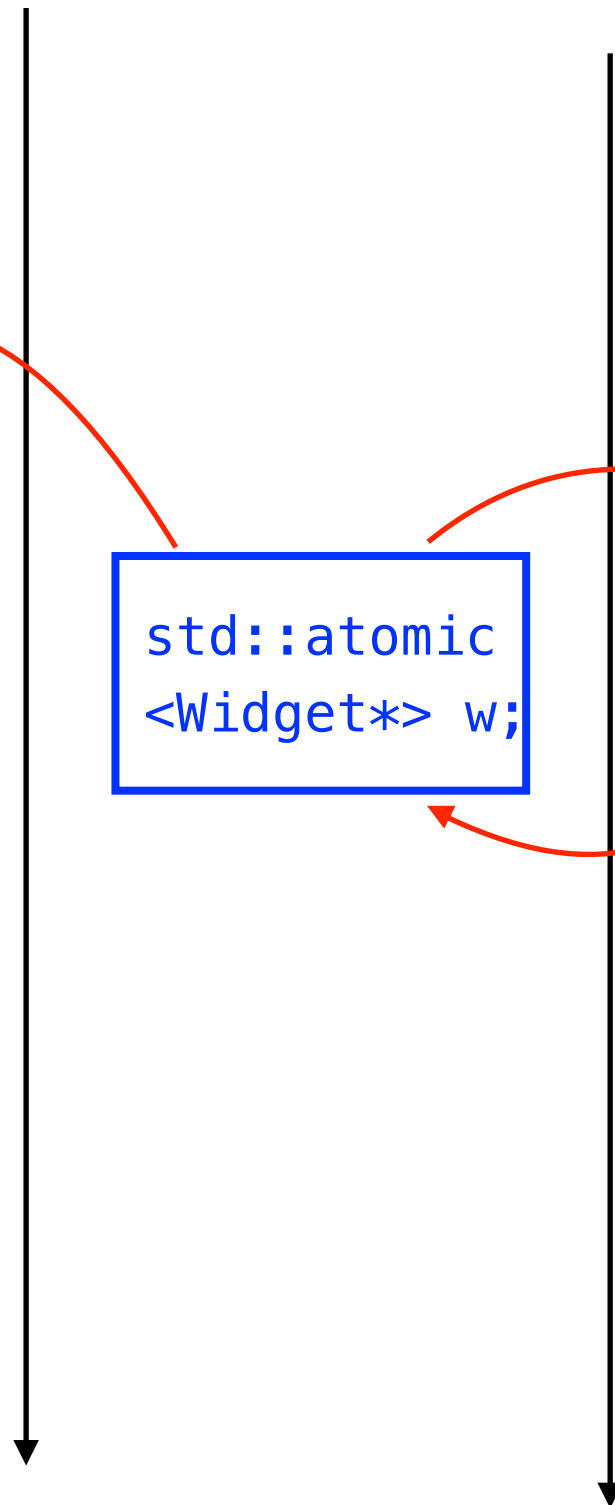
    /*
     * do something
     * with Widget...
     */
}
```

`std::atomic
<Widget*> w;`

```
Widget* newValue =
    new Widget ( /* ... */ );
```

```
if (! w.compare_exchange_weak
    (oldValue, newValue))
    /* ??? */
```

```
delete oldValue;
```



~~std::atomic<std::shared_ptr<T>>~~

std::atomic_... (std::shared_ptr<T>)

```
class Synthesiser
{
public:

    void audioCallback (float* buffer, int bufferSize)
    {
        std::shared_ptr<Widget> widgetToUse = std::atomic_load (&currentWidget);
        // do something with widgetToUse...
    }

    void updateWidget ( /* args */ )
    {
        std::shared_ptr<Widget> newWidget = std::make_shared<Widget> ( /* args */ );
        std::atomic_store (&currentWidget, newWidget);
    }


    std::shared_ptr<Widget> currentWidget;
};
```

```
class Synthesiser
{
public:

    void audioCallback (float* buffer, int bufferSize)
    {
        std::shared_ptr<Widget> widgetToUse = std::atomic_load (&currentWidget);
        // do something with widgetToUse...
    }

    void updateWidget ( /* args */ )
    {
        std::shared_ptr<Widget> newWidget = std::make_shared<Widget> ( /* args */ );
        std::atomic_store (&currentWidget, newWidget);
    }

    std::shared_ptr<Widget> currentWidget;
};
```




```
class Synthesiser
{
public:

    void audioCallback (float* buffer, int bufferSize)
    {
        std::shared_ptr<Widget> widgetToUse = std::atomic_load (&currentWidget);
        // do something with widgetToUse...
    }

    void updateWidget ( /* args */ )
    {
        std::shared_ptr<Widget> newWidget = std::make_shared<Widget> ( /* args */ );
        releasePool.add (newWidget);
        std::atomic_store (&currentWidget, newWidget);
    }

    std::shared_ptr<Widget> currentWidget;
    ReleasePool releasePool;
};
```

```

class ReleasePool : private Timer
{
public:
    ReleasePool() { startTimer (1000); }

    template<typename T> void add (const std::shared_ptr<T>& object) {
        if (object.empty())
            return;

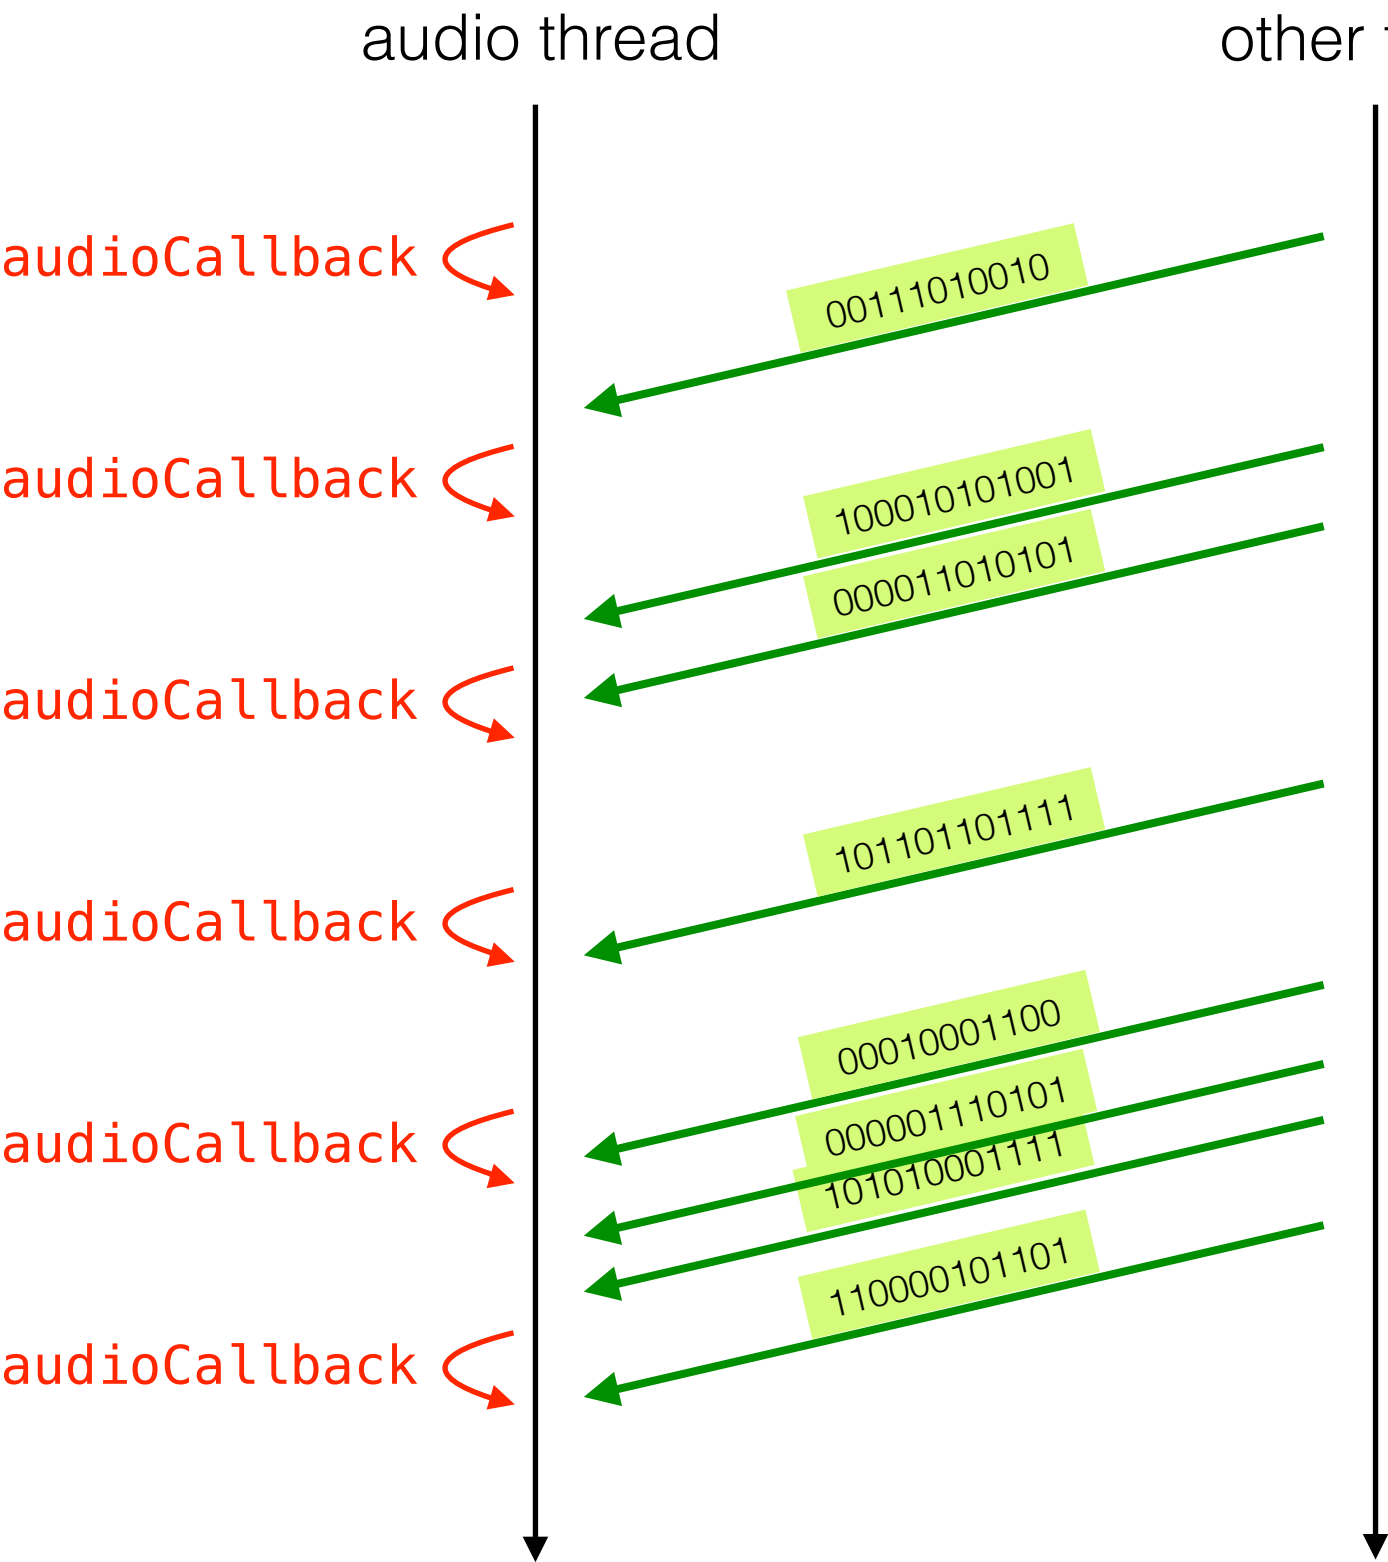
        std::lock_guard<std::mutex> lock (m);
        pool.emplace_back (object);
    }

private:
    void timerCallback() override {
        std::lock_guard<std::mutex> lock (m);
        pool.erase(
            std::remove_if (
                pool.begin(), pool.end(),
                [] (auto& object) { return object.use_count() <= 1; } ),
            pool.end());
    }

    std::vector<std::shared_ptr<void>> pool;
    std::mutex m;
};

```

exchange data between threads

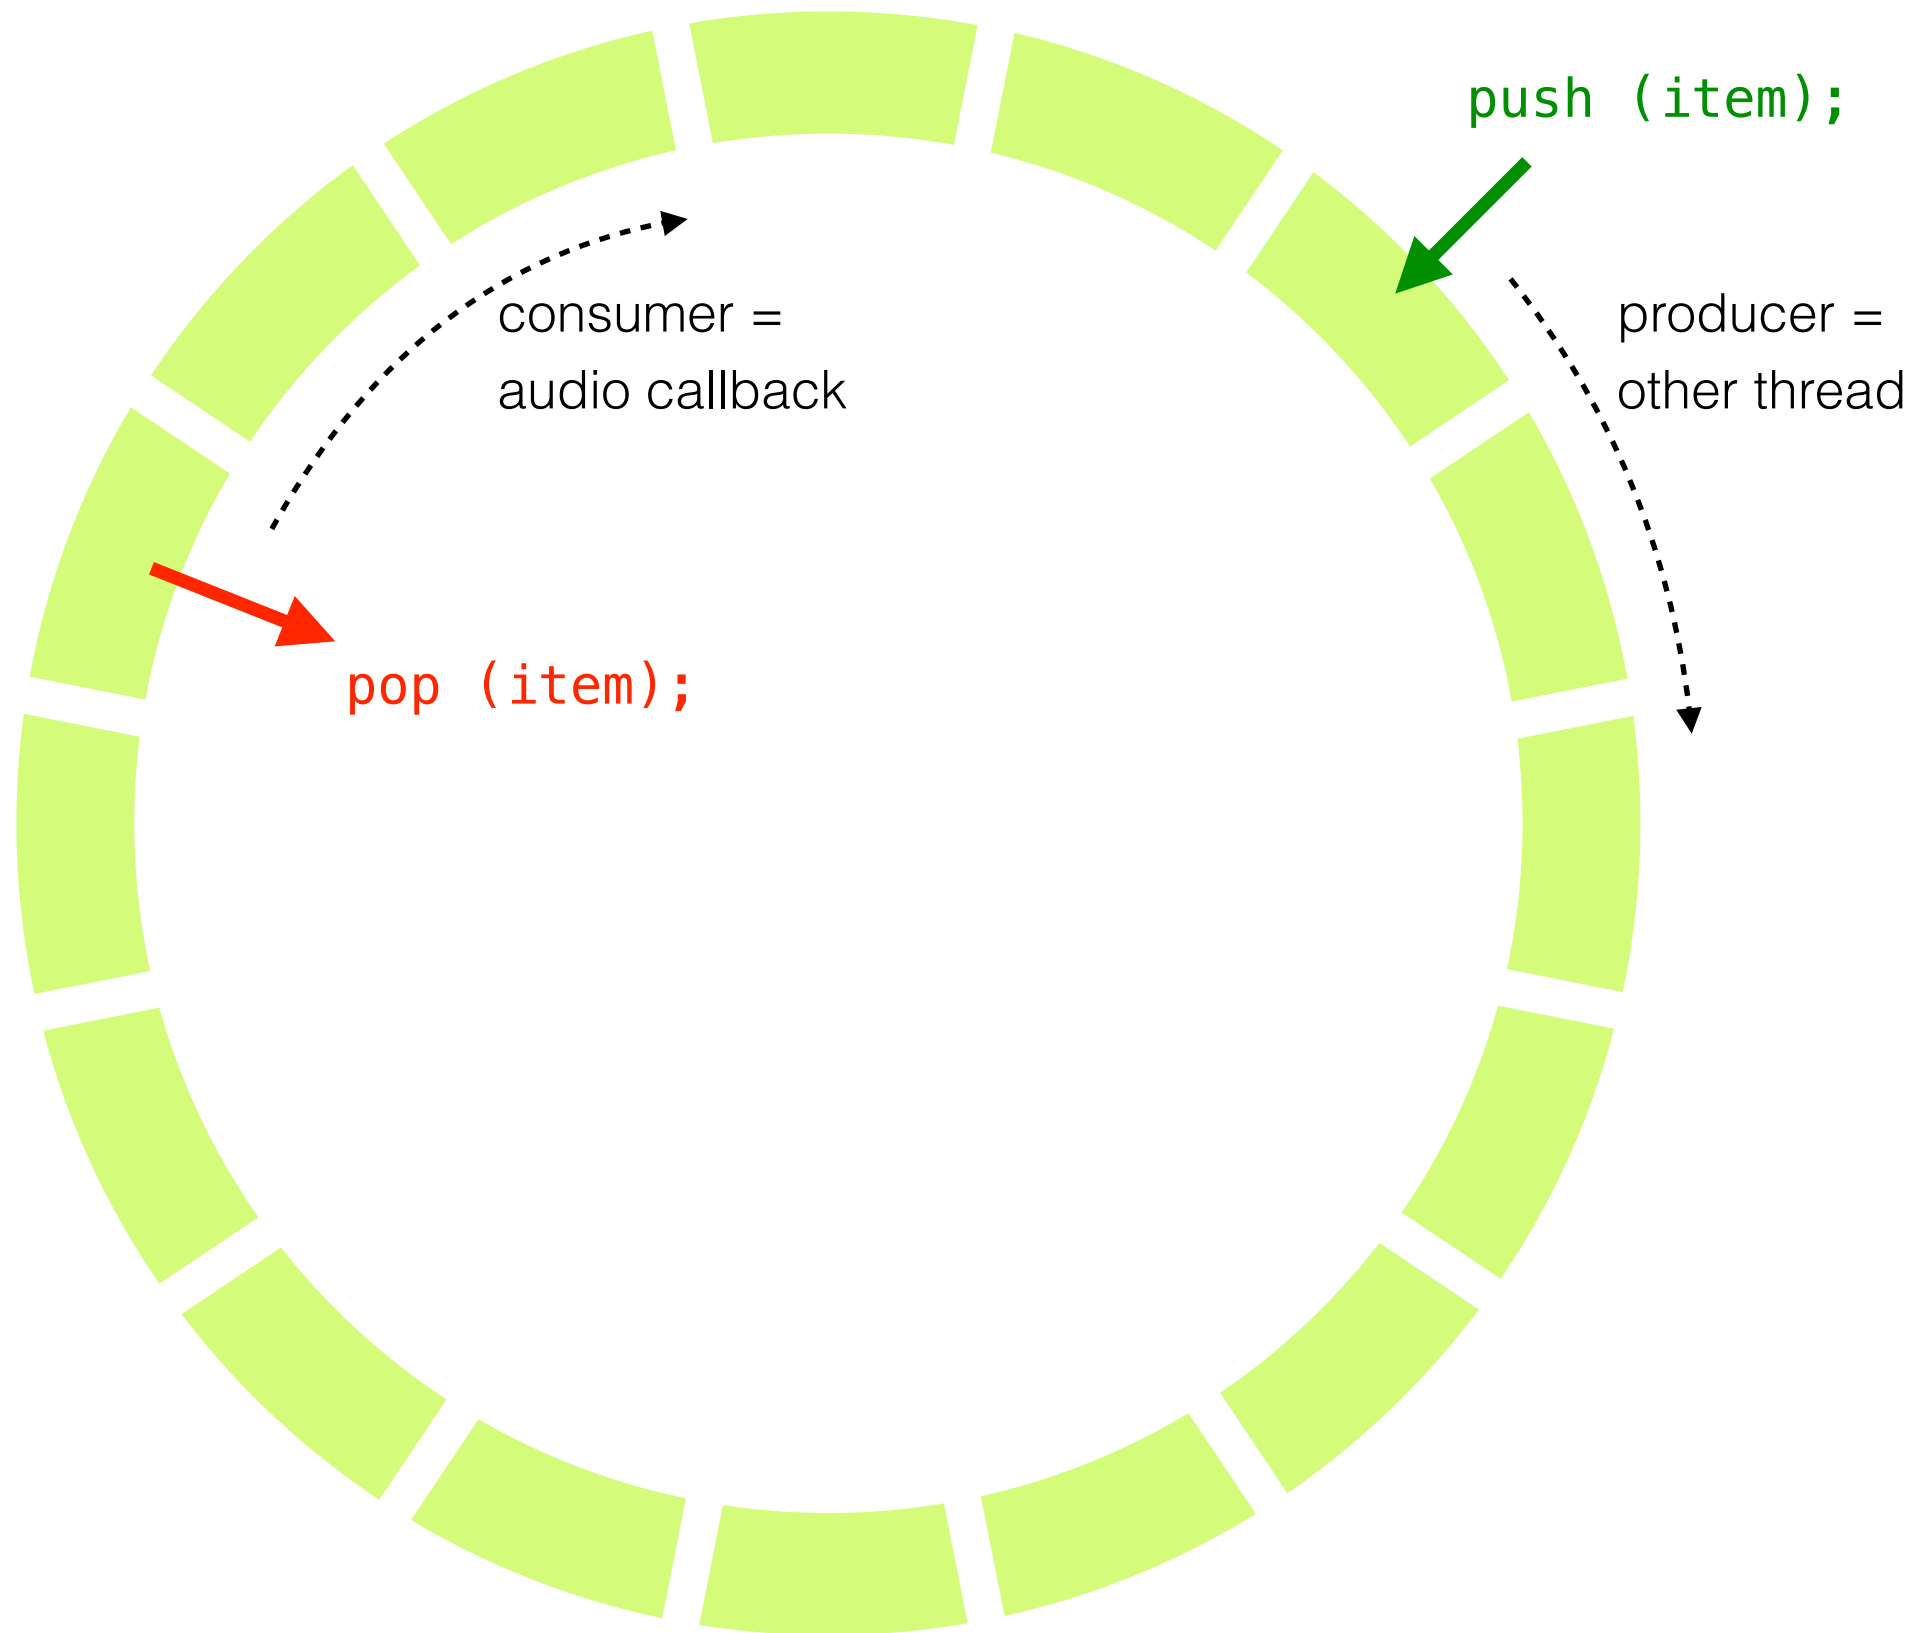


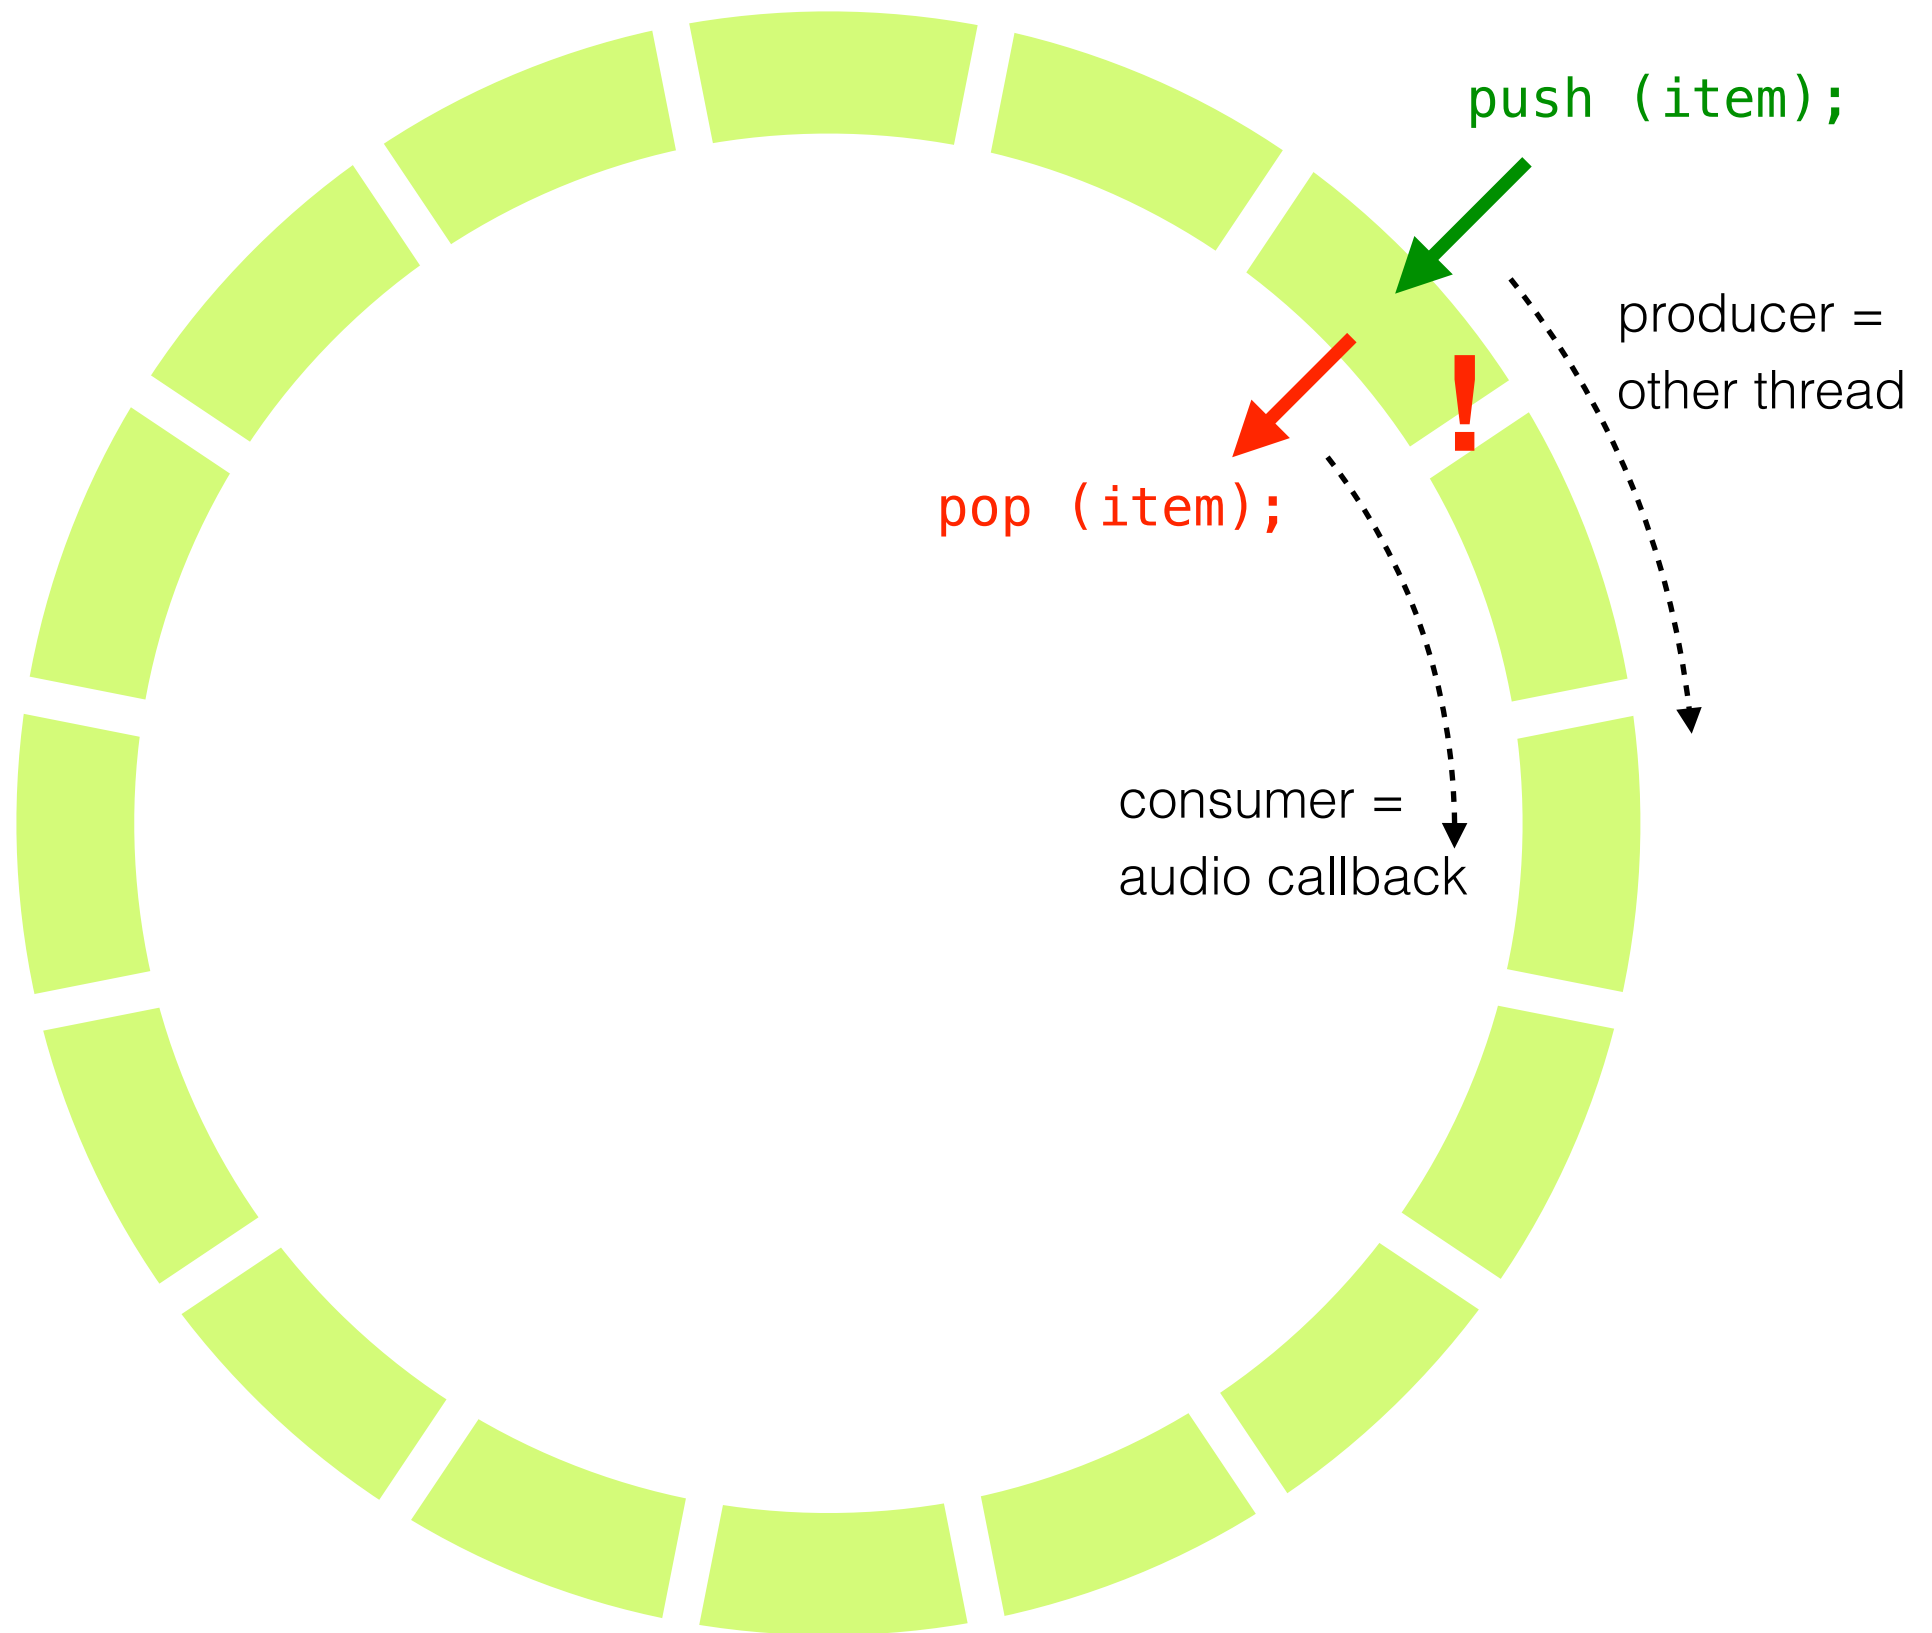
exchange data between threads

= **lock-free queue (fifo)**

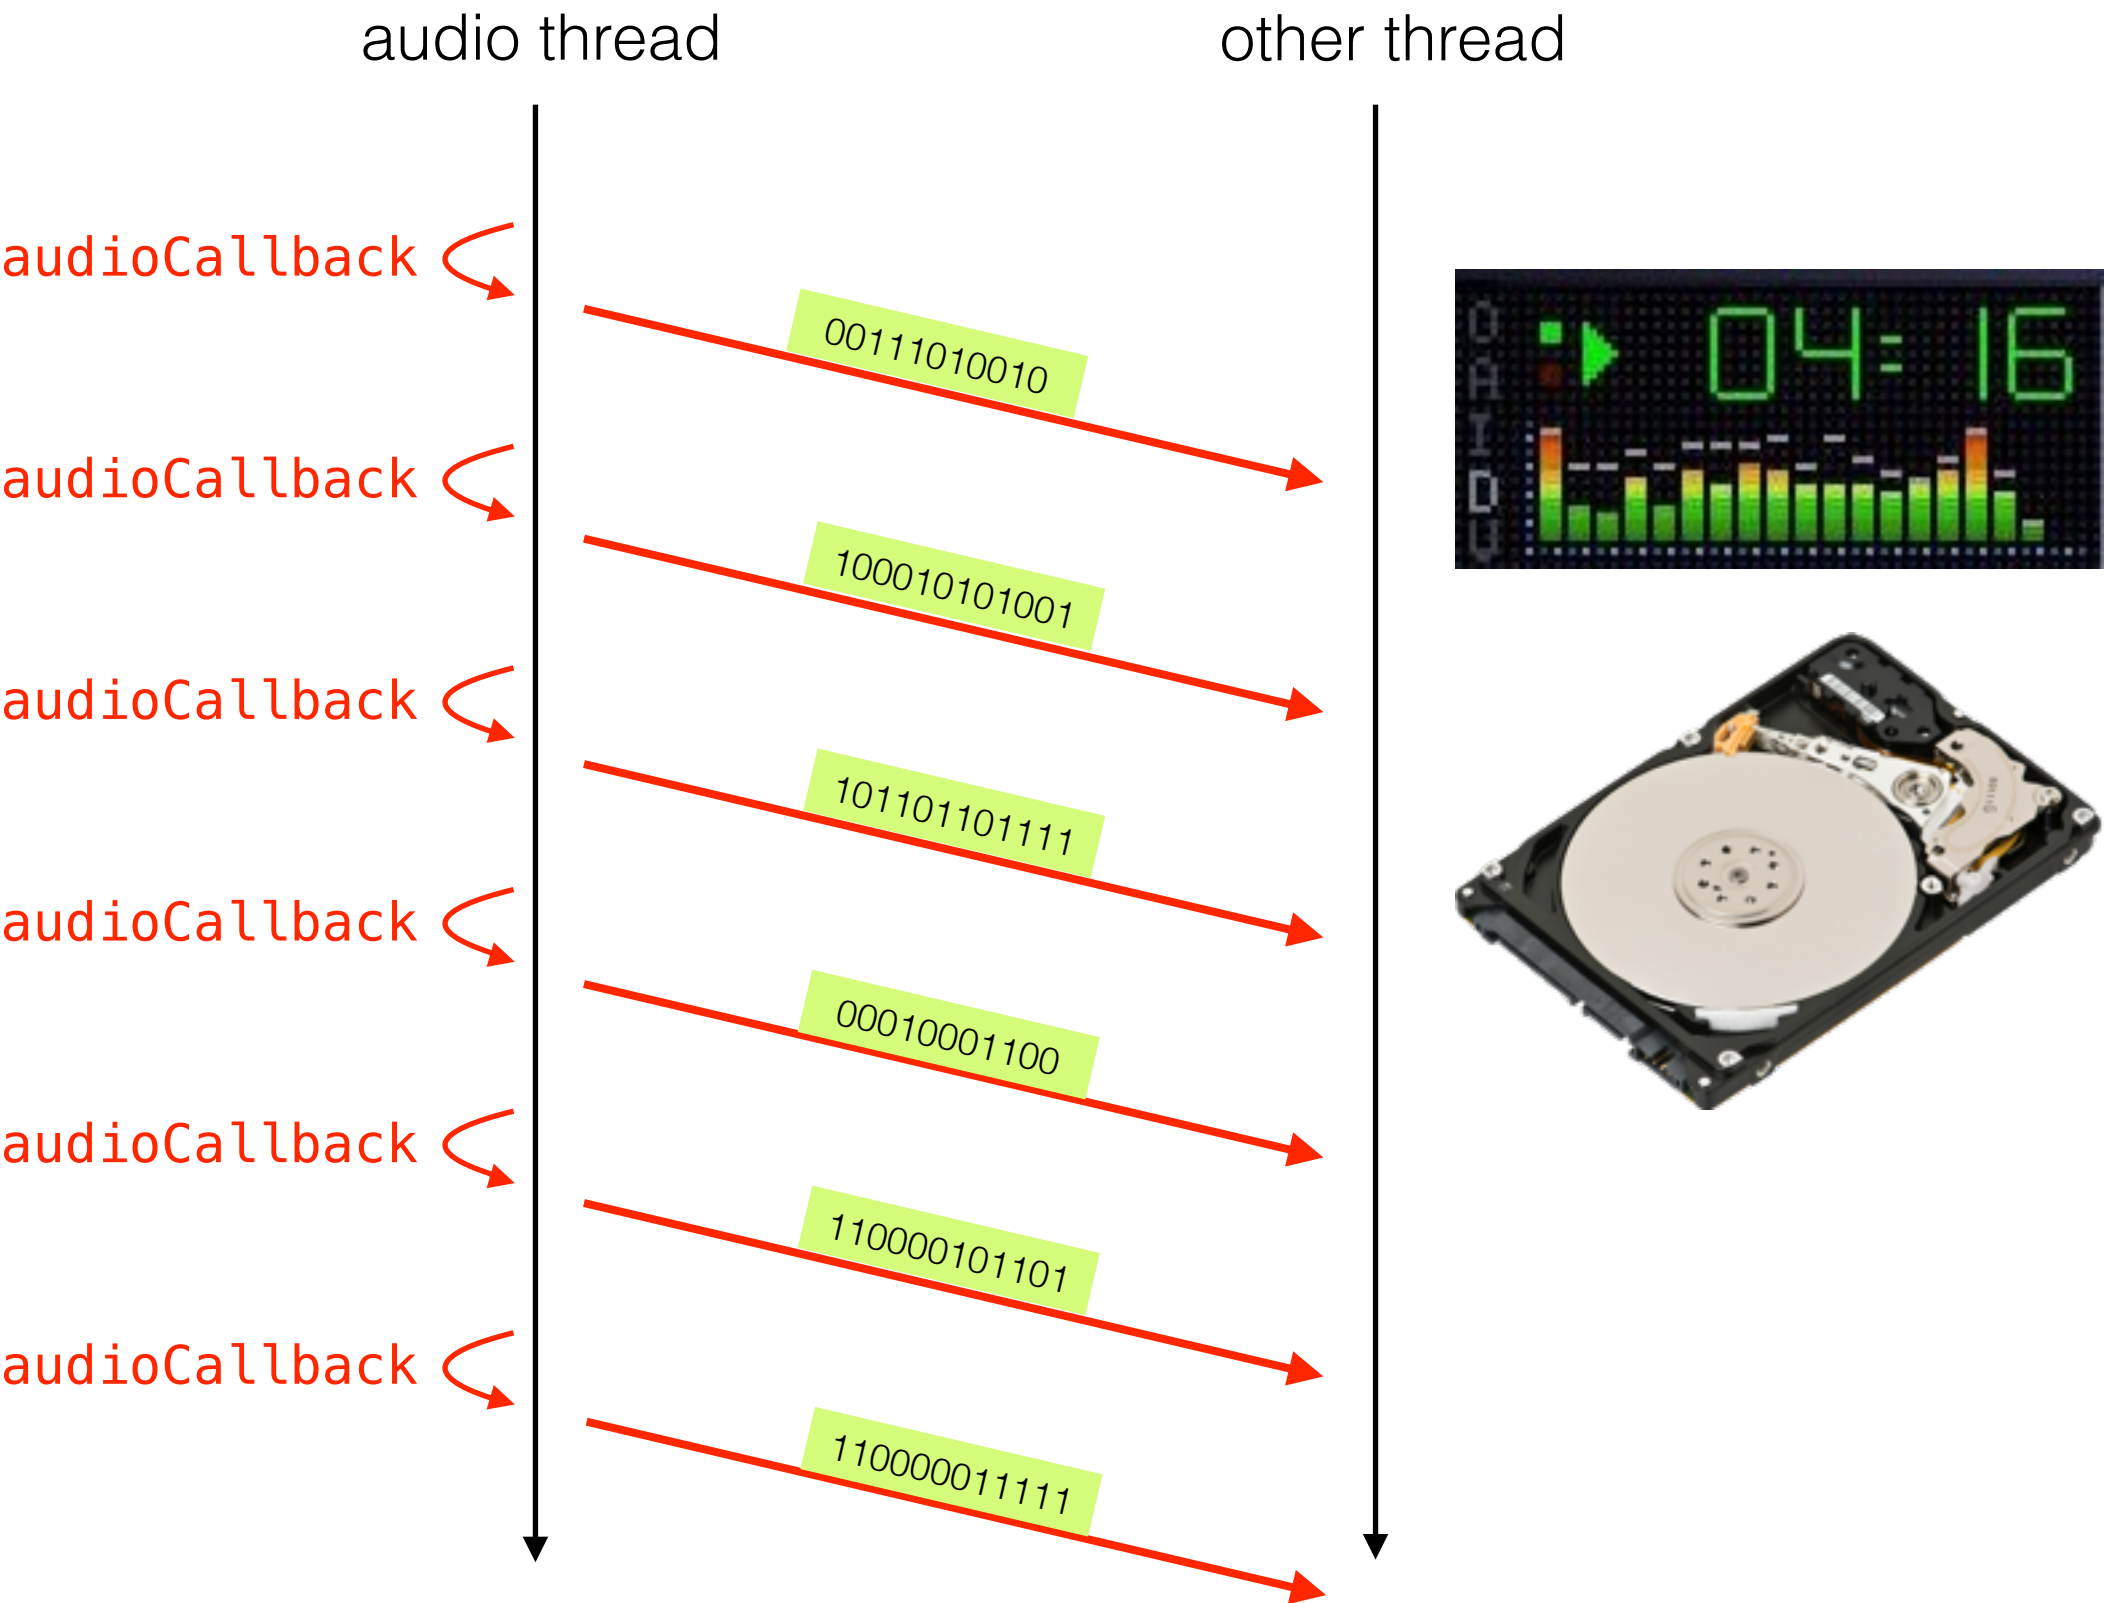
...with fixed maximum size = **lock-free ring buffer**.

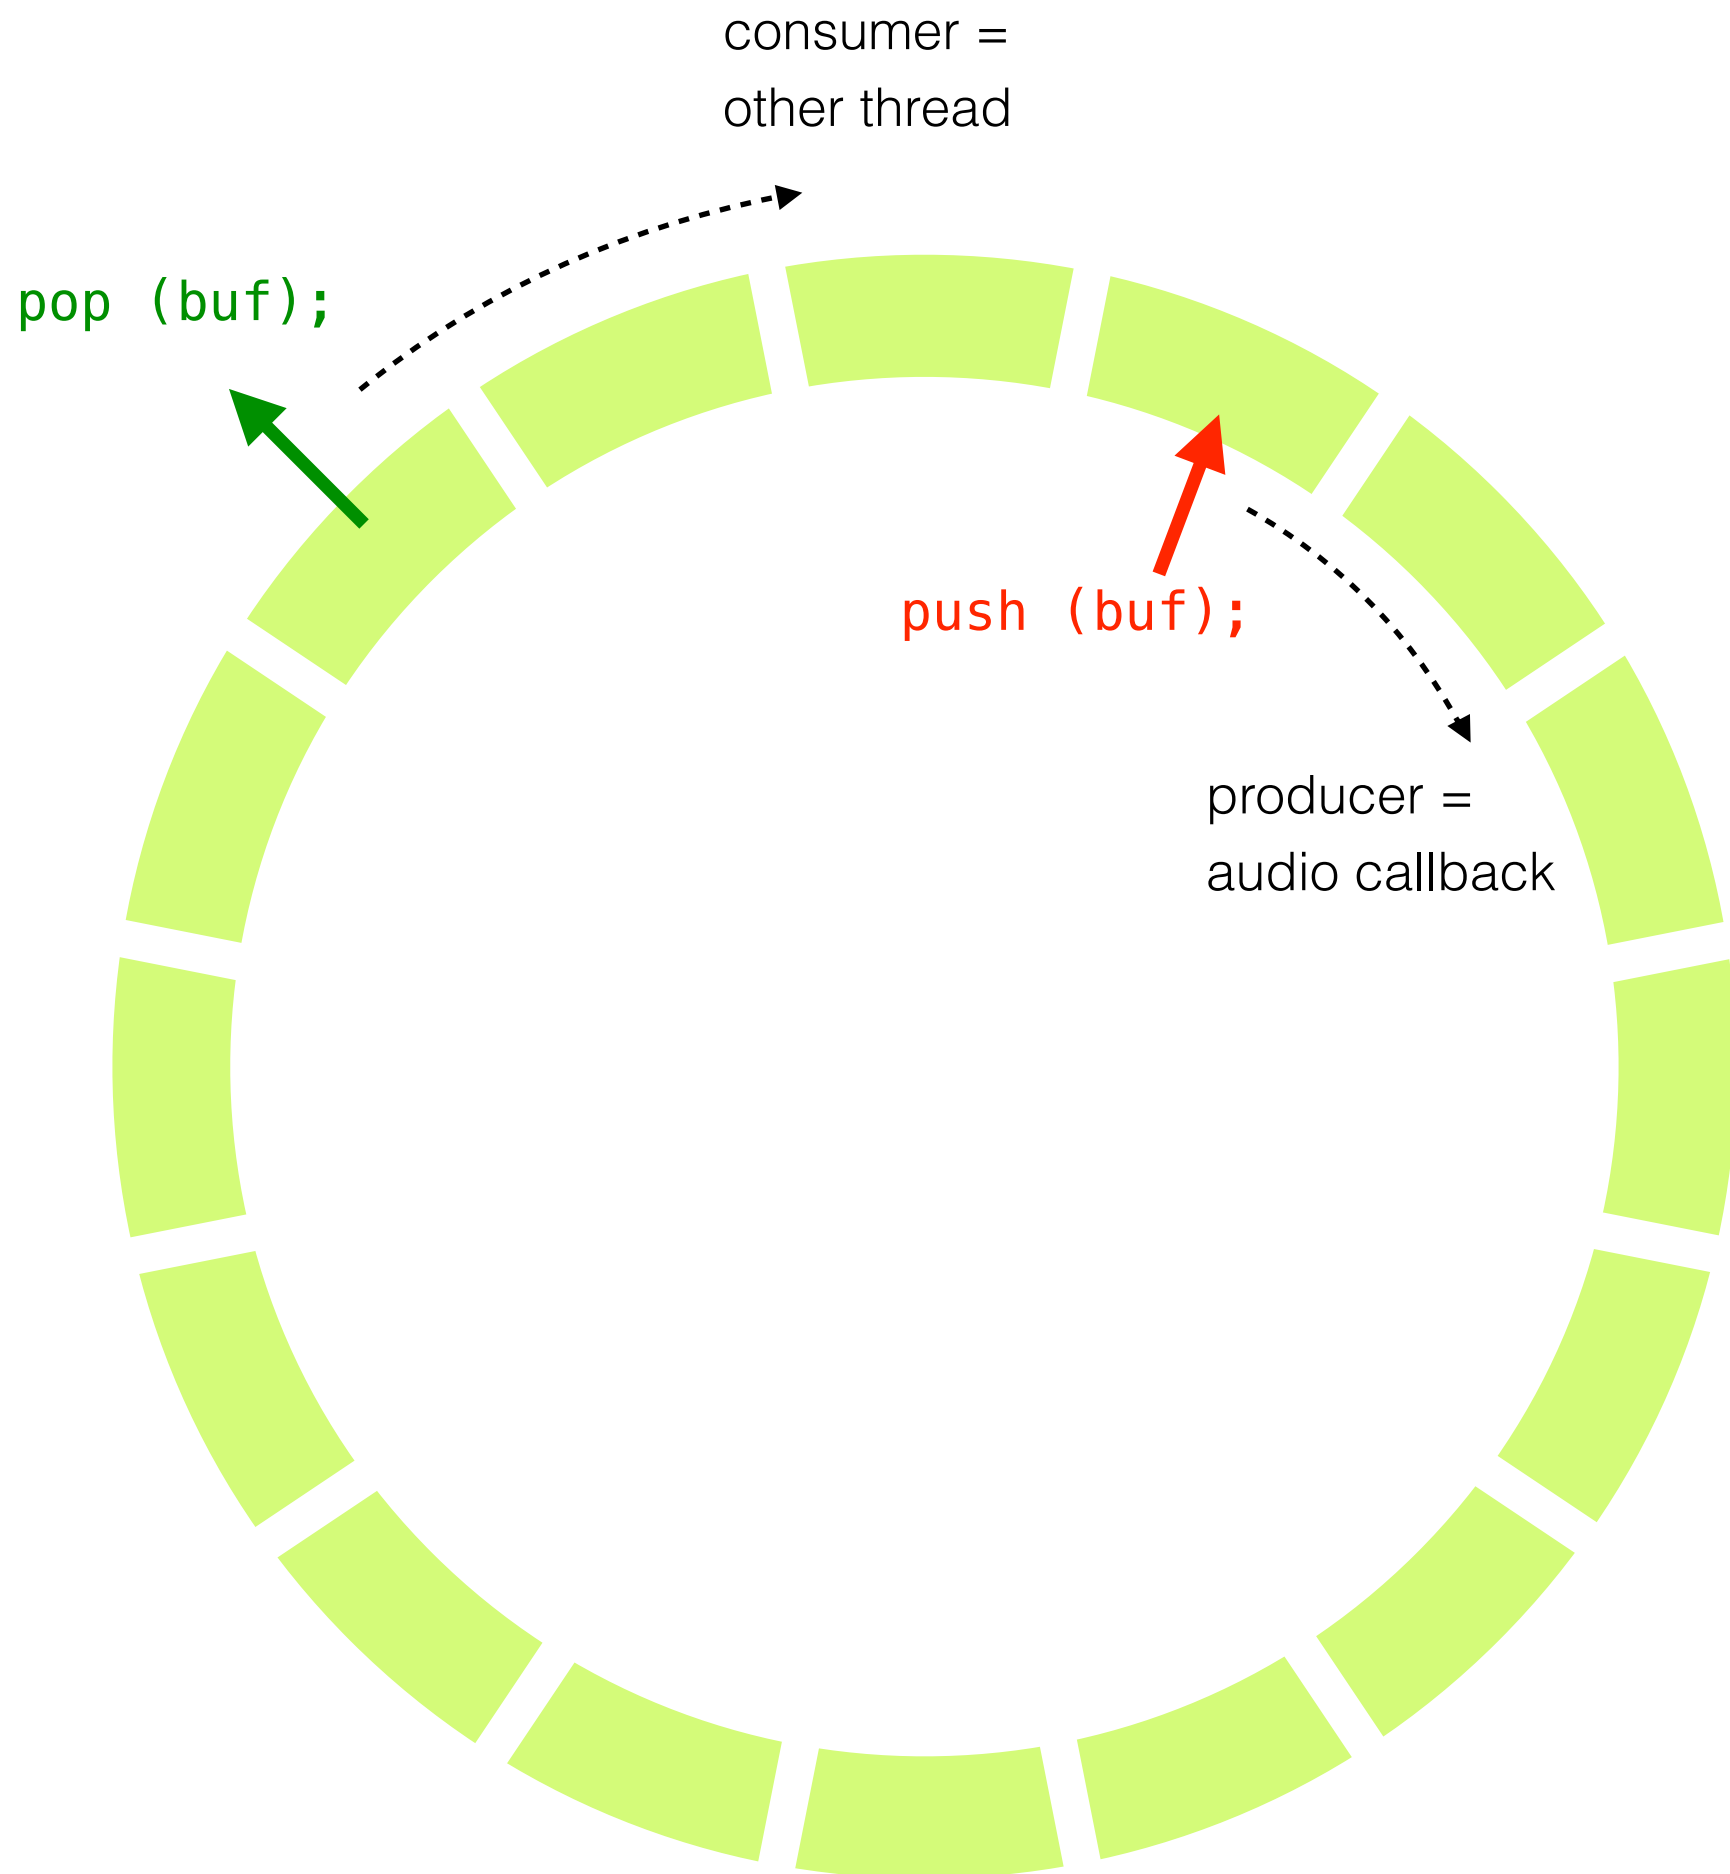
...if possible, single producer / single consumer.





exchange data between threads



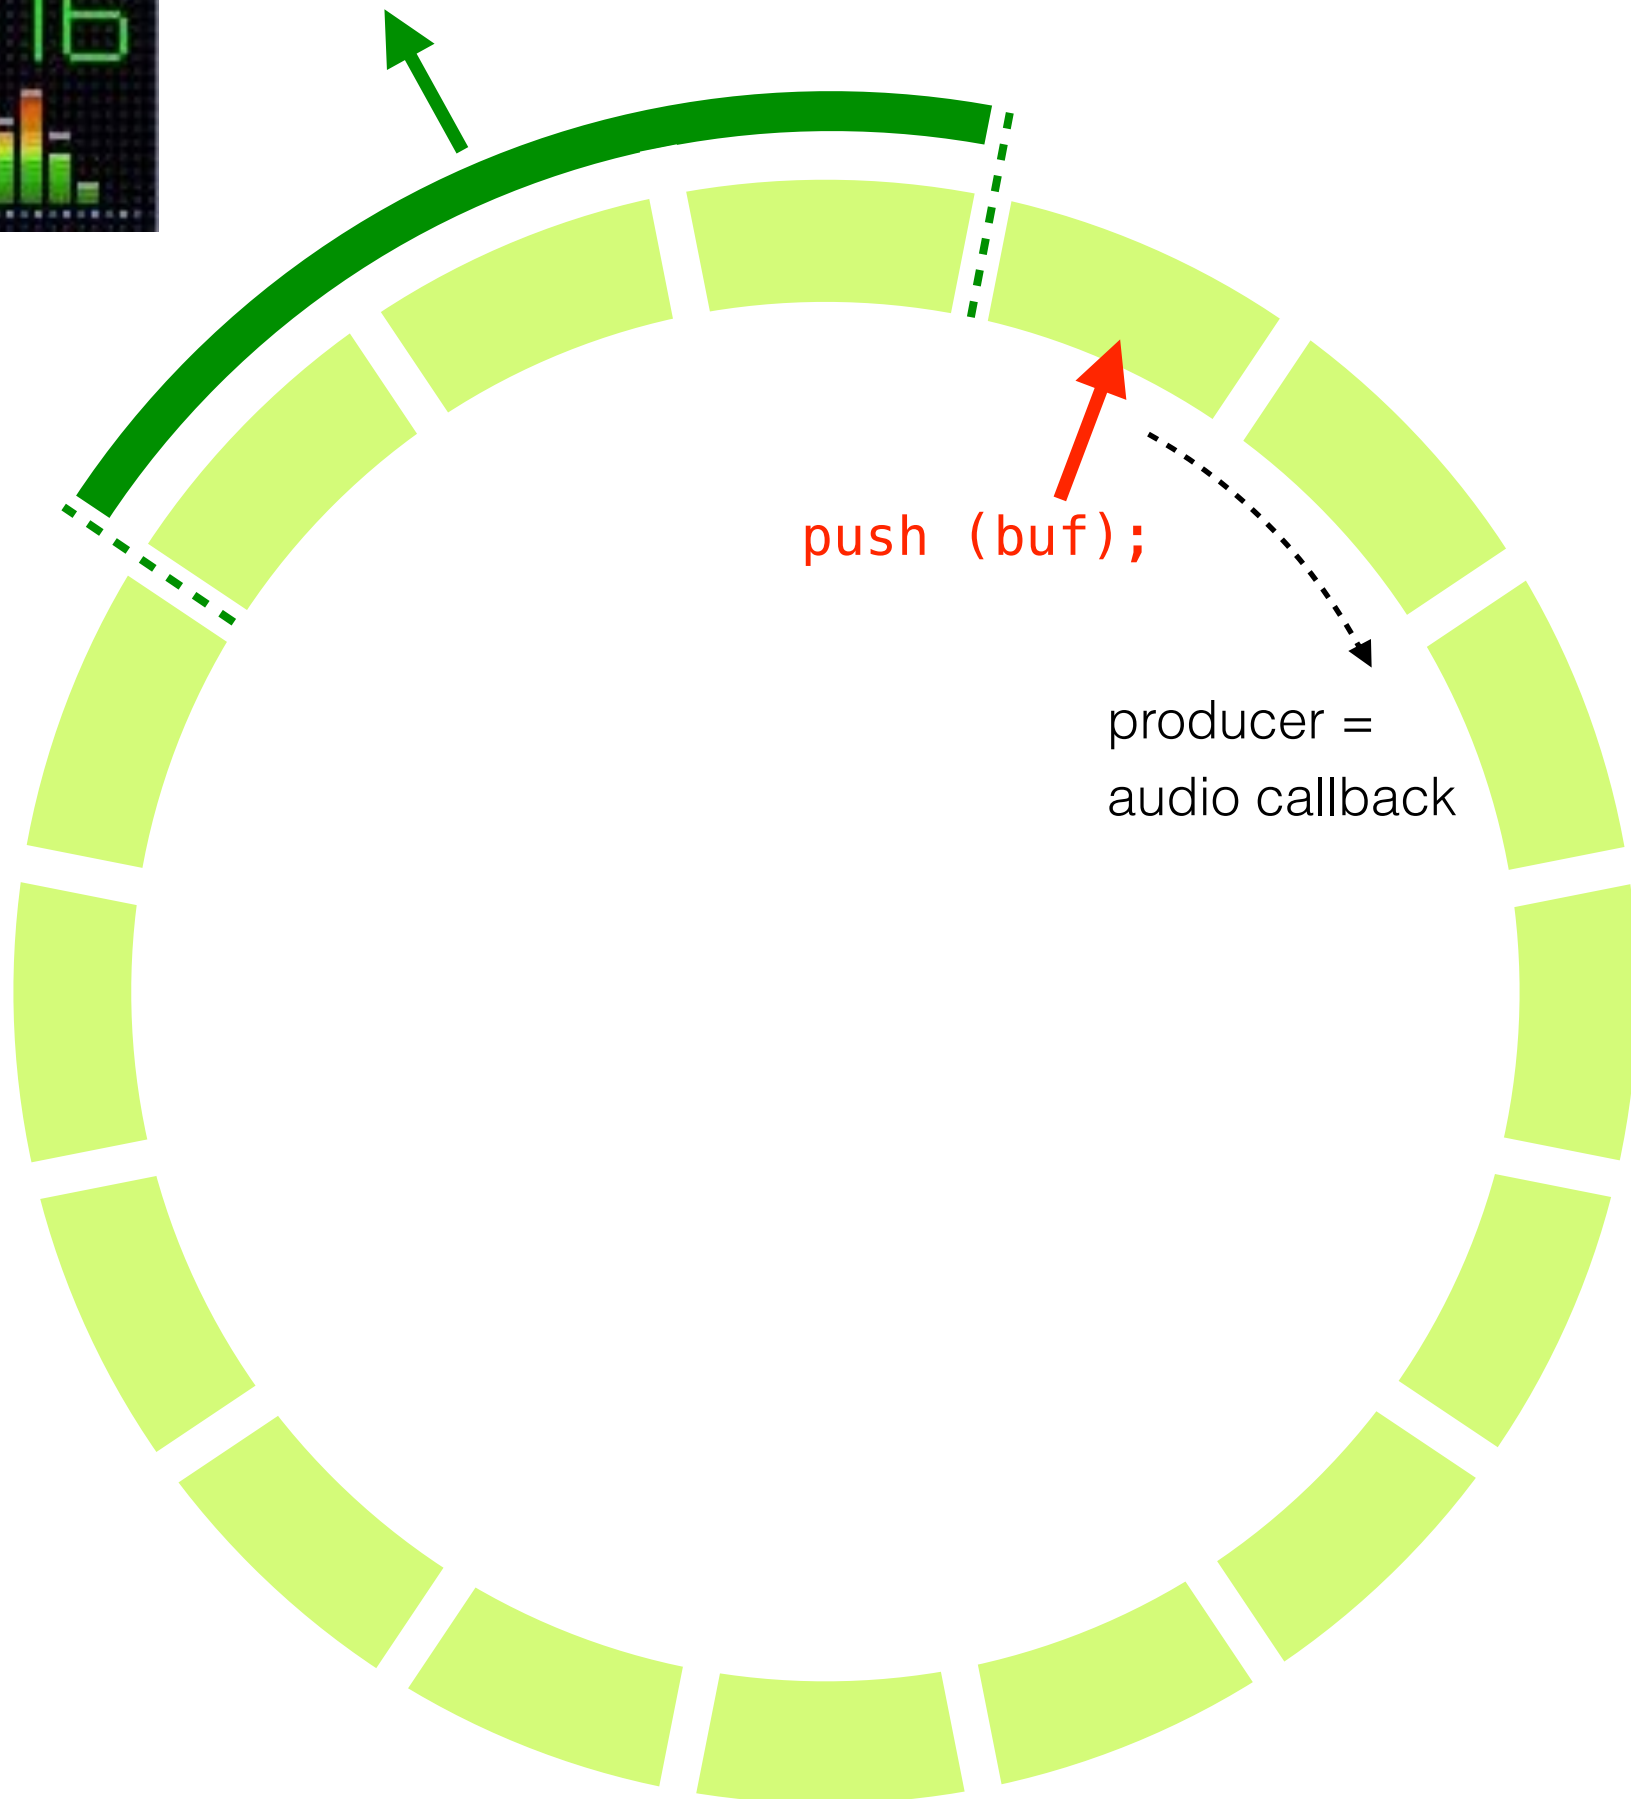




`renderVisualFrame();`

`push (buf);`

producer =
audio callback



A deep space photograph showing a vast field of stars against a black background. The stars vary in brightness and color, with many appearing as small, distant points of light. The text "to be continued..." is centered in the image in a yellow, monospace-style font.

to be continued...