


Executors for C++ - A Long Story

Struggling for a Base Concurrency Building Block

Detlef Vollmann


vollmann engineering gmbh, Luzern, Switzerland

Executors for C++, CppCon, September 2015



Part 1


The mission



Motivation: async


```
std::async([](){ std::cout << "Hello "; });  
std::async([](){ std::cout << "World!\n"; });
```

- No concurrency
- No real control over execution agent
 - `launch::async` and `launch::deferred` insufficient



Async Problem


- A future destructor normally doesn't block
- But blocks if the future is created by async
- Lifetime control is important in C++
- For async there's no other mechanism to control its lifetime
- Problem is different behaviour of ~future()



Motivation: Pipelines


```
pipeline::plan restaurant(  
    orders  
    | pipeline::parallel(chef, 3)  
    | pipeline::parallel(waiter, 4)  
    | end);  
  
thread_pool pool;  
  
pipeline::execution work(restaurant.run(&pool));
```

- Executors as building blocks for higher level abstractions




Executor Requirements

- Run tasks
- Control some lifetime aspects



Part 2


All beginning is ... easy



Original Executor Interface


```
class executor{
public:
    virtual ~executor();
    virtual void add(function<void()> closure) = 0;
    virtual size_t
        uninitiated_task_count() const = 0;
};
```

- (Not really the original interface.)




Default Executor

```
shared_ptr<executor> default_executor();
void set_default_executor(
    shared_ptr<executor> executor);
```



Concrete Executors

- `thread_pool`
- `serial_executor`
- `loop_executor`
- `inline_executor`
- `thread_executor`




`async`

```
async(launch::executor,
      [](){ std::cout << "Hello!\n"; });
```

- Uses `default_executor`
 - we need just a little bit more to shutdown the `default_executor`

```
async([](){ std::cout << "Hello!\n"; });
```


- Could probably also use `default_executor`
 - without breaking any existing code
 - but still blocks on future destructor



`async`

```
thread_pool myPool;
async(myPool,
      [](){ std::cout << "Hello!\n"; });
```


- General way to launch a task on an executor



Motivation: Pipelines


```
pipeline::plan restaurant(  
  orders  
  | pipeline::parallel(chef, 3)  
  | pipeline::parallel(waiter, 4)  
  | end);  
  
thread_pool pool;  
  
pipeline::execution work(restaurant.run(&pool));
```

- This can easily be implemented based on the initial proposal




Mission Accomplished

- async problem solved
 - Just some more detail work
- Accepted February 2014 by Concurrency SG into Concurrency TS



Part 3


The real discussion begins



Abstract Base Class


```
virtual void add(function<void()> closure) = 0;
```

- No template concept
- Not part of the type
 - Not really important for functions
 - Important for structures
- Can cross binary interfaces
- Sometimes simply too costly



Part 4

More requirements




.then

- Proposed continuation .then also allows for an executor:

```
auto f = std::async([](){  
    std::cout << "Hello "; });  
f.then(myPool,  
    [](){ std::cout << "World\n"; });
```

- Without executor, how does .then know on which executor to run best?




Data Concentrator

```
RTExecutor rtExec0(0);
RTExecutor rtExec80(80);


ConcentratorT dataConcentrator{
    wrap(rtExec80, ReadDev(1, in1))
    , wrap(rtExec0, ReadDev(2, in2))
    , wrap(rtExec0, StoreData(out))
};
dataConcentrator.run();
```

- Concentrator like pipeline
 - Two producers, one consumer
 - One producer has higher priority




ASIO

- ASynchronous Input/Output
- Wants to run continuation on thread where OS I/O returns
- Wants to run concurrently or co-operative
- Wants to avoid overhead of futures
- Wants to run on user-defined executors
 - with support for system specific asynchronous events
 - signals/interrupts, timers, mailboxes, ...
- Grown out of lot of experience
 - ASIO specific




Part 5

New proposals



Executors and Async Ops

- ASIO based N4046 by Chris Kohlhoff
- executor and execution_context
 - executor is a light-weight handle
 - execution_context actually holds the threads and tasks
 - execution_context can be used to wait on everything to shut down.
- Proposed concrete executors:
 - system_executor (like thread_executor)
 - strand (like serial_executor)
 - thread_pool
 - loop_executor



Customization Points


- Continuation token
 - direct continuation on same thread
 - synchronization mechanism
 - concurrency mechanism
- Execution interface
 - dispatch()
 - post()
 - defer()
- get_associated_executor()
 - generally required to use
 - allows for arbitrary info from task to executor



Executors and Schedulers, R5


```
class executor{
public:
    template<class Func> void spawn(Func&& func);
};
```

- As template based concept
 - with an interface for type erasing abstract base class
- task_wrapper to associate an executor with a task




Executor Traits

- N4406 "Integrating Executors with Parallel Algorithm Execution"
- Required interface as traits
- Executor semantics
 - concurrent
 - parallel
 - weakly parallel
- Future type
- Task starting
- Bulk task starting




Executor Traits

- N4406 is very specific for parallel algorithms
- Not a proposal for a specific executor interfaces
- Traits allow for implementation that's not provided by the executor
 - bulk interface
 - future based interface




Part 6

Status quo




Proposal Status

- Original (modified) Google proposal accepted into Concurrency TS February 2014 (Issaquah)
- ASIO based proposal presented June 2014 in Rapperswil, tentatively accepted as new base:
 - remove N3785 from TS: SF-F-N-A-SA 6-7-5-2-0
 - More work on N4046 for TS: 10-8-0-0-0
 - Apply N4046 to TS without significant changes: 4-2-3-5-2
- R4 of the Google proposal was presented at SG1 meeting September 2014 in Redmond
 - (Re-)Start with Chris Mysisen's proposal? SF-F-N-A-SA 9-5-4-0-2
- Traits proposal discussed May 2015 in Lenexa, no vote



Part 7

Rethinking



Layers

**User
programs**

Application

**Library
components**

Containers `async` ASIO FlowGraph
 `.then`
Parallel pipeline Event
Algorithms Loop

**Building
blocks**

allocator executor




More Info = Better ...

- There's a lot of information about a task that may be useful for an executor implementation
 - relationship to spawning task
 - long/short running
 - blocking/non-blocking
 - repetitions
 - priority
 - information return
 - ...
- All very specific to some executors/domains
- Possibly nothing of them needs to be directly in the executor interface
- But there must exist mechanisms for information transfer
 - only some of them need to be known by intermediate mechanisms




Part 8

Still something else




Coroutines

- Coroutines are an important part of asynchronicity
- ASIO works together with coroutines
- The current async/resume approach doesn't seem to mix well with executors
- This needs to change




Part 9

Where to go?




Summary

- Executors are an important building block for concurrency
- They should be low level
- The final interface is still very controversial
- Missing experience




Still Open

- What is a task?
- What information does it carry?
- How?
- What is the right interface for the middle layer?
- What do we need to standardize?



Outlook

- ASIO based proposal very complex
- Provides three customization points
 - blocking hints/requirements, continuation token, associated executor
- Chris Mysens's proposal provides only one of them
 - associated executor
- Coroutines currently not co-operative



Questions

- ???

