


<functional>: What's New, And Proper Usage

Stephan T. Lavavej ("[Steh-fin Lah-wah-wade](#)")

Senior Developer - Visual C++ Libraries

stl@microsoft.com

Getting Started



- Please hold your questions until the end
 - Write down the slide numbers
- Everything here is Standard
 - Unless otherwise specified
- Almost everything here is available in VS 2015
 - Rewrote <functional> for Standard conformance
 - Minor issues will be listed at the end

Lambdas (C++11)

C++14: init-captures, generic lambdas

Example

```
vector<string> v{ "hydrogen", "helium",  
    "lithium", "beryllium", "boron", "carbon",  
    "nitrogen", "oxygen", "fluorine", "neon" };
```

```
stable_sort(v.begin(), v.end(),  
    [](const auto& l, const auto& r) {  
        return l.size() < r.size(); });
```

```
for (const auto& e : v) {  
    cout << e << endl;  
}
```

Lambdas Aren't Magical

- A lambda expression:
 - Defines a class
 - With `operator()(Args)` and maybe data members
 - Constructs an object
- Lambda syntax is convenient
 - Handwritten function objects are verbose
- Remember:
 - Lambdas are a Core Language feature
 - `std::function` is a Standard Library feature
 - Lambdas aren't `std::functions`

Lambdas Aren't Magical, Really

- Stateless lambdas: convertible to function pointers
 - Handwritten function objects can do that too
 - It's just `operator FunctionPointer()`
- Remember:
 - A lambda defines a class and constructs an object
 - Lambdas aren't functions
 - Lambdas aren't function pointers
- Never refer to lambdas as "anonymous functions"
 - Or I will make this face >:-[/]

invoke() (C++17)

Terminology: Function Objects

- Usable like functions: `func(args)`
 - Function pointers
 - Even in C, `fp(args)` means `(*fp)(args)`, thanks to DMR
 - Classes with `operator()(Args)`
 - Including lambdas
 - Classes with conversions to function pointers
 - Obscure!
- References to functions are similarly usable
 - Technically, reference types aren't object types

Terminology: Callable Objects

- Callable in a generalized sense
 - Function objects: `func(args)`
 - Pointers to member functions (PMFs): `(obj.*pmf)(args)`
 - Pointers to member data (PMDs): `obj.*pmd`
- The Core Language wants different syntax (hiss!)
 - `pmf(obj, args)` could be permitted, but isn't
- The Standard Library wants uniform syntax (purr!)
 - *INVOKE()* was imaginary in TR1 and C++11/14
 - `invoke()` is available in C++17

invoke(callable, args...)

- Function objects
 - `invoke(func, args...)` is `func(args...)`
- PMFs
 - `invoke(pmf, obj, rest...)` is `(obj.*pmf)(rest...)`
 - `invoke(pmf, ptr, rest...)` is `((*ptr).*pmf)(rest...)`
- PMDs
 - `invoke(pmd, obj)` is `obj.*pmd`
 - `invoke(pmd, ptr)` is `(*ptr).*pmd`
- Base PMFs/PMDs can be invoked on Derived things
- `invoke()` handles both raw pointers and smart pointers

Example

```
template <typename Range, typename Callable>
    void transform_print(const Range& r, Callable c) {
        for (const auto& e : r) {
            cout << invoke(c, e) << endl;
        }
    }

vector<pair<int, int>> v{{4, 40}, {5, 50}, {6, 60}};
transform_print(v,
    [](const auto& p) { return p.first * p.first; });
transform_print(v,
    &pair<int, int>::second);
```

Things That Use `invoke()`

- `<functional>`
 - `std::function`
 - `reference_wrapper`
 - `bind()`
 - `mem_fn()`
- `<type_traits>`
 - `result_of`
- `<future>`
 - `packaged_task`
 - `async()`
- `<mutex>`
 - `call_once()`
- `<thread>`
 - `std::thread`

Recommendations

- In non-generic code, `invoke()` isn't very useful
 - Unless you really hate PMF syntax
- In generic code, `invoke()` can simplify things
 - Take/store arbitrary callable objects and arguments
 - Give them to `invoke()`, let it decide what to do
- Don't special-case PMFs/PMDs
 - Inspecting PMF types is a headache
 - Implementing `invoke()` behavior is extremely difficult

result_of (C++11)

C++14: SFINAE, result_of_t

result_of: Type Trait For invoke()

- C++11 `<type_traits>` (was TR1 `<functional>`)
- `result_of_t<Callable>` is incorrect
- `result_of_t<Callable(Args...)>` is:
 - `decltype(invoker(declval<Callable>(), declval<Args>()...))`
- `declval()` is declared but never defined:
 - `template <typename T> add_rvalue_reference_t<T>
declval() noexcept;`
- In C++14, the `::type` SFINAEs away

Example

```
template <typename T, typename Callable>
    auto transform_sort(const vector<T>& v, Callable c) {
        vector<decay_t<result_of_t<Callable&(const T&)>>> ret;
        for (const T& t : v) { ret.push_back(invoke(c, t)); }
        sort(ret.begin(), ret.end());
        return ret;
    }
const vector<string> v{ "hydrogen", "helium", "etc." };
auto lambda = [](const string& s) { return s.size(); };
for (const auto& e : transform_sort(v, lambda)) {
    cout << e << endl;
}
```


result_of Is Tricky

- It answers "what's the type of this invocation?"
 - But it uses **different syntax** than the real invocation
- cv-qualifiers and value categories can matter
 - `declval<Callable/Args>()` must match real callable/args
 - Especially when C++11 ref-qualifiers are involved
- The real invocation might do extra work
 - `bind()` extensively manipulates its arguments
 - `async()` decays its arguments
- `result_of` is TR1-era tech, predating `decltype`

Recommendations

- Avoid using `result_of`
 - If you must use it, be careful
 - Audit existing usage for bugs, I bet you'll find some
- Use `decltype(STUFF)`, `decltype(auto)`, or `auto`
 - Which generic programmers already need to understand
- In general, avoid computing the same thing through different mechanisms
 - If repetition is necessary, prefer exactly repeating text
 - `decltype(STUFF) matching return STUFF;` is easy to see

mem_fn() (C++11)

Example

```
struct Element {  
    bool is_metallic() const;  
};
```

```
size_t count_metals(const vector<Element>& v) {  
    return count_if(v.begin(), v.end(),  
        mem_fn(&Element::is_metallic));  
}
```

mem_fn() Isn't Fun, Really

- Good: Usually terse
- Bad: Resistant to optimization
- Ugly: Won't compile in certain situations
 - Overloaded member functions (need `static_cast`)
 - Templated member functions (use `static_cast`)
 - Avoid explicit template arguments, Don't Help The Compiler
 - Default arguments (no workaround)
- Unnecessary with anything powered by `invoke()`

Recommendations

- Avoid using `mem_fn()`
 - Algorithm inner loops often affect performance
 - As code evolves, `mem_fn()` is fragile
 - Auditing existing usage is low priority, though
- Use lambdas, especially generic lambdas
 - They're slightly more verbose
 - But they optimize away
 - And they always compile, like other member function calls

Transparent Operator Functors (C++14)

Example

```
vector<int> ints{ 6, 3, 10, 5, 16, 8, 4, 2, 1 };
```

```
vector<string> strs{  
    "O'Neill", "Carter", "Jackson", "Teal'c" };
```

```
sort(ints.begin(), ints.end(), greater<>());
```

```
sort(strs.begin(), strs.end(), greater<>());
```

```
for (const auto& e : ints) { cout << e << endl; }
```

```
for (const auto& e : strs) { cout << e << endl; }
```


Would You Like To Know More?

- Transparent Operator Functors
 - "Don't Help The Compiler"
 - GoingNative 2013, slides 37-48
 - Gained constexpr before C++14 shipped
- Heterogeneous Associative Lookup
 - "STL Features And Implementation Techniques"
 - CppCon 2014, slides 33-40

Recommendations

- Use `greater<>` etc. by default
 - Except when you need implicit conversions (very rare)
- Advantages:
 - Avoids truncation/signedness bugs
 - Avoids unnecessary temporaries
 - Avoids unnecessary copies
 - Less verbose
- Unlike `mem_fn()`, library machinery is OK here
 - Operators are known in advance, perfectly special-cased

bind() (C++11)

Example

```
const vector<int> v{ 1, 4, 9, 16, 25, 36, 49,  
    64, 81, 100, 121, 144 };
```

```
cout << count_if(v.begin(), v.end(),  
    bind(less<>(), _1, 50)) << endl; // 7
```

```
cout << count_if(v.begin(), v.end(),  
    bind(less<>(), 50, _1)) << endl; // 5
```

How `bind()` Works

- `auto b = bind(callable, bound_args...);`
 - Later: `b(unbound_args...);`
- `callable` and `bound_args` are copied or moved
 - Then they're passed as **lvalues** to `invoke()`
 - So `b` can be called repeatedly
 - With `b`'s constness, via `const-overloaded operator()`
- Some bound arguments are special:
 - Placeholders: `_1` perfectly forwards first unbound arg
 - `reference_wrapper<T>`: unwrapped via `get()` to `T&`
 - Nested `bind()`: called with perfectly fwded unbound args
- Unused unbound arguments are ignored

bind() Problems

- Same performance/compiler issues as `mem_fn()`
 - With function pointers in addition to PMFs/PMDs
- Misuse emits ultra-disgusting compiler errors
- Syntax isn't normal C++, especially nested `bind()`
- No short-circuiting for `logical_and`/`logical_or`
- Surprising behavior: bound args passed as lvalues
 - Affects `unique_ptr`, etc.
- Surprising behavior: immediate vs. delayed calls
- Placeholders and nested `bind()` can move twice

Recommendations

- Avoid using `bind()`
- Use lambdas, especially generic lambdas
- `bind()`: good idea in 2005, bad idea in 2015
 - In C++, we usually prefer Library solutions to Core
 - But the Library is terrible at building up function objects
 - Lambdas were added to the Core Language for a reason
 - STL maintainers rarely recommend avoiding the STL
 - `bind()`'s terseness just isn't worth the price

reference_wrapper (C++11)

C++17: trivially copyable

Class Definition

```
template <typename T> class reference_wrapper {  
public:  
    typedef T type;  
    reference_wrapper(T&) noexcept;  
    reference_wrapper(T&&) = delete;  
    operator T&() const noexcept;  
    T& get() const noexcept;  
    template <typename... Args>  
        result_of_t<T&(Args&&...)>  
        operator()(Args&&...) const;  
};
```

Example

```
vector<int> v(8);  
const auto b = v.begin(); const auto e = v.end();  
typedef uniform_int_distribution<int> Dist;  
auto d20 = [urng = mt19937(1729), dist = Dist(  
    1, 20)]() mutable { return dist(urng); };  
generate(b, e, d20); // 2 11 8 18 10 11 16 2  
generate(b, e, d20); // 2 11 8 18 10 11 16 2  
generate(b, e, ref(d20)); // 2 11 8 18 10 11 16 2  
generate(b, e, ref(d20)); // 18 4 12 10 8 13 14 10
```

Recommendations

- Almost all algorithms take function objects by value
 - And are allowed to copy them
- Use `ref()` to pass function objects by reference
 - Avoid explicit template arguments, Don't Help The Compiler
- `reference_wrapper` is useful elsewhere
 - Like `std::thread`'s constructor
 - But first, learn why it uses `DECAY_COPY()`
- Be aware of the 3 functions that unwrap
 - `make_pair()`, `make_tuple()`, and `bind()`
 - `make_tuple(x, ref(y), cref(z))` is `tuple<X, Y&, const Z&>`

Removed Old `<functional>`
Stuff (C++17)

Erased... From Existence

- Deprecated in C++11, removed in C++17
- unary_function/binary_function
 - Provided result_type, etc.
- ptr_fun()
 - Wrapped function pointers with result_type, etc.
- mem_fun()/mem_fun_ref()
 - Strictly superseded by mem_fn()
- bind1st()/bind2nd()
 - Strictly superseded by bind()

Recommendations

- Never use the old `<functional>` stuff
- Remove any existing usage
- C++98/03 algorithms/containers never needed `unary_function/binary_function/ptr_fun()`
- VS 2015: `/D_HAS_AUTO_PTR_ETC=0`
 - Also controls `auto_ptr` and `random_shuffle()`

`std::function` (C++11)

C++14: SFINAE

C++17: Converts non-void to void

Example

```
int sum_squares(int x, int y) {  
    return x * x + y * y; }  
vector<function<int (int, int)>> v;  
v.emplace_back(plus<>());  
v.emplace_back(multiplies<>());  
v.emplace_back(&sum_squares);  
for (int i = 10; i <= 1000; i *= 10) {  
    v.emplace_back([i](int x, int y) {  
        return i * x + y; }); }  
for (const auto& f : v) { cout << f(4, 5) << endl; }
```


How `std::function` Works

- `function<Ret (Args)>` is a wrapper
 - Stores a callable object of arbitrary type
 - Templated on call signature, not callable object type
 - Type erasure, powered by virtual functions (or equivalent)
- Useful when code can't be templated
 - Separately compiled code
 - Virtual functions
 - Container elements

CopyConstructible Required

- The STL usually delays requirements
 - `list<T>` doesn't require `T` to have `operator<()`
 - `list<T>::sort()` requires `T` to have `operator<()`
- `std::function` is special, due to type erasure
 - `function(F)` requires `F` to be CopyConstructible
 - Even though it stores `move(f)`
 - Even if you never copy `std::function`
- Can't store movable-only function objects
 - Design limitation; alternatives are being investigated

Small Functor Optimization

- Arbitrary callable objects can be arbitrarily large
 - Eventually, dynamic memory allocation is necessary
- Small callable objects can be stored locally
 - Guaranteed for function pointers and `reference_wrapper`
 - Otherwise, "small" is the implementer's decision
- `is_nothrow_move_constructible` needed for SFO
 - Because `function::swap()` must be `noexcept`

Magic Numbers

Toolset	VS 2015 x86	VS 2015 x64	libc++ 3.7.0 x64	libstdc++ 5.2.0 x64
std::function	40	64	48	32
SFO Max	32	48	24	16
std::string	24-28	32-40	24	32

Ambiguous C++11, Valid C++14

```
void meow(const function<int (int)>& f) {  
    cout << f(3) << endl; }  
void meow(const function<int (int, int)>& g) {  
    cout << g(4, 5) << endl; }  
meow([](int n) { return n * 11; });  
meow([](int x, int y) { return x * 10 + y; });
```

- `function<Ret (Args)>(F)` is now constrained
 - `F` must be Callable as `Ret (Args)`

Invalid C++14, Valid C++17

```
int x = 1729;  
function<void (int&)> f(  
    [](int& r) { return ++r; });  
f(x);  
cout << x << endl;
```

- Non-void can't be implicitly converted to void
 - According to Core - but the Library makes its own rules!
- Also applies to `packaged_task`, obscure `bind<R>()`

Unresolved Library Issues

- `std::function`'s `operator()` is `const`
 - Can store function objects with non-`const` `operator()`
 - Violates the STL's `const`/multithreading conventions

- This compiles, but really shouldn't:

```
string meow(int x) { return to_string(x); }
```

```
function<const string& (int)> hiss(&meow);
```

- VS warns: returning address of local variable or temporary
- I think I can fix this in the Standard, but it's not trivial

Recommendations

- `std::function` is awesome!
- But use it **only when necessary**
 - When possible, use templates and `auto`
- `std::function` inherently has nonzero costs
 - Time: Type erasure prevents inlining
 - Space: SFO buffer and type erasure consume bytes
 - Codegen: Emits code whether you use it or not
- Avoid unnecessary copies, moves, temporaries

More Info

More Info

- Almost everything here is available in VS 2015
 - Not yet implemented: C++14 `result_of/function` SFINAE
 - Bugs: `async()` uses `bind()`, `packaged_task` uses `function`
 - My `mem_fn()` was sneaky in RTM, fixed in Update 1
- C++17 Working Paper
 - <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2015/n4527.pdf>

Questions?

`stl@microsoft.com`