

Large Scale C++ With Modules: What You Should Know

Gabriel Dos Reis





Where We Are (Today)

USE-DATE.CXX

```
#include <iostream>
#include "Calendar/date.h"

int main() {
    using namespace Chrono;
    Date date { 22, Month::Sep, 2015 };
    std::cout << "Today is " << date << std::endl;
}</pre>
```

CALENDAR/DATE.H

```
#ifndef CHRONO DATE INCLUDED
#define CHRONO DATE INCLUDED
#include <iosfwd>
#include <string>
#include "Calendar/Month.h"
namespace Chrono {
     struct Date
           Date(int, Month, int);
           int day() const { return d; }
           Month month() const { return m; }
           Int year() const { return y; }
     private:
           int d;
           Month m;
           int y;
     std::ostream& operator<<(std::ostream&, const Date&);</pre>
     Std::string to string(const Date&);
#endif // CHRONO DATE INCLUDED
```

A Better Place to Be

USE-DATE.CXX

```
import std.io;
import calendar.date;

int main() {
    using namespace Chrono;
    Date date { 22, Month::Sep, 2015 };
    std::cout << "Today is " << date << std::endl;
}</pre>
```

CALENDAR/DATE.CXX

```
import std.io;
import std.string;
import calendar.month;
module calendar.date;
namespace Chrono {
     export
     struct Date {
           Date(int, Month, int);
           int day() const { return d; }
           Month month() const { return m; }
           Int year() const { return y; }
      private:
           int d;
           Month m;
           int y;
      export
      std::ostream& operator<<(std::ostream&, const Date&);</pre>
      export
      std::string to string(const Date&);
```

A Better Place to Be

USE-DATE.CXX

```
import std.io;
import calendar.date;

int main() {
    using namespace Chrono;
    Date date { 18, Month::Sep, 2015 };
    std::cout << "Today is " << date << std::endl;
}</pre>
```

CALENDAR/DATE.CXX

```
import std.io;
import std.string;
 import calendar.month;
module calendar.date;
namespace Chrono {
     struct Date
           Date(int, Month, int);
           int day() const { return d; }
           Month month() const { return m; }
           Int year() const { return y; }
     private:
           int d;
           Month m;
           int y;
     export
     std::ostream& operator<<(std::ostream&, const Date&);</pre>
     export
     Std::string to string(const Date&);
```

What Is The Big Deal?

> Well,

```
#include <iostream>
#include "Calendar/date.h"

int main() {
    using namespace Chrono;
    Date date { 18, Month::Sep, 2015 };
    std::cout << "Today is " << date << std::endl;
}</pre>
```

- is 176 bytes of user-authored text file (what user sees)

What Is The Big Deal?

> Well,

```
#include <iostream>
#include "Calendar/date.h"

int main() {
    using namespace Chrono;
    Date date { 18, Month::Sep, 2015 };
    std::cout << "Today is " << date << std::endl;
}</pre>
```

- is 176 bytes of user-authored text file (what user sees)
- expands to (what compiler sees)
 - > 412,326 bytes with GCC 5.2.0 or 234,276% compression expansion
 - > 1,203,953 bytes with Clang 3.6.1– or 684,064% inflation
 - > 1,083,255 bytes with VC++ Dev14 or 615,485% inflation

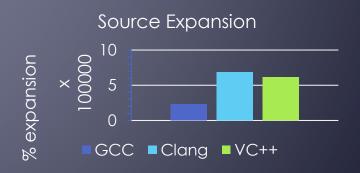
What Is The Big Deal?

> Well,

```
#include <iostream>
#include "Calendar/date.h"

int main() {
    using namespace Chrono;
    Date date { 18, Month::Sep, 2015 };
    std::cout << "Today is " << date << std::endl;
}</pre>
```

- is 176 bytes of user-authored text file (what user sees)
- expands to (what compiler sees)
 - > 412,326 bytes with GCC 5.2.0
 - > 1,203,953 bytes with Clang 3.6.1
 - > 1,083,255 bytes with VC++ Dev14







- > Preprocessor directive #include is textual copy and paste
 - Compiler works hard to process the same entity multiple times
 - Linker works hard to throw all of them away, except one
 - Miserable build throughput

- > Preprocessor directive #include is textual copy and paste
 - Compiler works hard to process the same entity multiple times
 - Linker works hard to throw all of them away, except one
 - Miserable build throughput
- Copy: No consistency guarantee
 - Hard to track bugs (famous "ODR" violation)
 - No component boundaries, brittle enforcement

- > Preprocessor directive #include is textual copy and paste
 - Compiler works hard to process the same entity multiple times
 - Linker works hard to throw all of them away, except one
 - Miserable build throughput
- Copy: No consistency guarantee
 - Hard to track bugs (famous "ODR" violation)
 - No component boundaries, brittle enforcement
- C Preprocessor technology: Impossible to correctly parse/analyze a component
 - Need to know all macros

Pressing Challenges for Modern C++

> Source Code Organization at Large

- Scaling beyond billions of lines of code
- Producing, composing, consuming components with well-defined semantics boundaries

> Paucity of Semantics-Aware Developer Tools

- Serious impediment to programmer productivity
- Great disadvantage vis-à-vis contemporary languages (C#, Java, Ada, etc.)
 - > Reason not to adopt C++
 - Reason to migrate away from C++

> Build time Scalability of Idiomatic C++

- Distributed build, cloud build, etc.
 - > Use **semantics** difference to accelerate build

Aims

- > Give C++ a module system, improving
 - 1. Componentization
 - 2. Isolation (from macros)
 - 3. Build throughput
 - 4. Support for modern semantics-aware developer tools
- > Deliver now, use for decades to come
 - Target: C++17 (yes, it can be done)

> Non Goals:

- Improve or remove the preprocessor

When Can I Use It?

When Can I Use It?

- > MS VC 2015 Update 1 Timeframe
 - Experimental implementation of the module proposal
 - Feedback based on use
 - Evidence of feasibility (for C++17)



- > MS VC 2015 Update 1 Timeframe
 - Experimental implementation of the module proposal
 - Feedback based on use
 - Evidence of feasibility (for C++17)
- Clang has an implementation based on "module maps"

What Hand Are We Dealt?



- > Program = Collection of independently translated units
 - Each TU processed in isolation, without knowledge of peer TUs
- > TUs communicate by brandishing declarations for external names
 - No explicit dependency on TU or component provider
 - No good scalable way to check/verify consistency
- Linker resolves uses of external names to whichever definitions happen to match (somehow)
 - Type-safe linkage problems
 - Opportunities for One Definition Rule (ODR) violation

Basic Linking Model

1.cc (producer of quant)

```
int quant(int x, int y) {
    return x*x + y*y;
}
```

2.cc (consumer of quant)

```
extern int quant(int, int);
int main() {
   return quant(3, 4);
}
```

3.cc (YAPoQ)

```
#include <stdlib.h>
int quant(int x, int y) {
    return abs(x) + abs(y);
}
```

- Valid programs: (a) 1.cc and 2.cc; (b) 2.cc and 3.cc
- Useful, effective, but low-level and brittle
 - Leak implementation details to language specification

The Root Cause: One Definition Rule (ODR)

- > What is ODR anyway?
 - Bjarne Stroustrup:
 - > "I asked Dennis when I started in 1979."
 - [DMR]:
 - "as if there was exactly one section of source text"
- > In the C++ standards
 - Several pages of opaque text about token-for-token comparison, name lookup, overload resolution, template instantiation contexts, etc.
 - Bjarne Stroustrup:
 - "Every single word about "token comparison" is there to workaround absence of a real module system"



Modules: 33,000ft view

Consumption

> As simple as

What is a Module?

Collection of related translation units, with a well-defined set of entry points

 Module interface: set of declarations available to any consumer of a module

> Module unit: TU element of a module

Module name: symbolic reference for a module

Module interface Module Unit (implementation)

My.Module

Production

As simple as

```
import std.io;
import std.string;
import calendar.month;
module calendar.date;
namespace Chrono {
     export struct Date {
           Date(int, Month, int);
           int day() const { return d; }
           Month month() const { return m; }
           Int year() const { return y; }
     private:
           int d;
           Month m;
           int y;
     };
     export std::ostream& operator<<(std::ostream& os, const Date& d)</pre>
         // ...
      export std::string to_string(const Date& d)
         // ...
```

Production

As simple as

```
import std.io;
import std.string;
import calendar.month;
module calendar.date;
export namespace Chrono {
     struct Date {
           Date(int, Month, int);
           int day() const { return d; }
           Month month() const { return m; }
           Int year() const { return y; }
     private:
           int d;
           Month m;
           int y;
     };
     std::ostream& operator<<(std::ostream& os, const Date& d)</pre>
         // ...
     std::string to_string(const Date& d)
         // ...
```

Production

As simple as

```
#include <iostream>
import std.string;
import calendar.month;
module calendar.date;
namespace Chrono {
     export struct Date {
           Date(int, Month, int);
           int day() const { return d; }
           Month month() const { return m; }
           Int year() const { return y; }
     private:
           int d;
           Month m;
           int y;
     };
     export std::ostream& operator<<(std::ostream& os, const Date& d)</pre>
         // ...
     export std::string to_string(const Date& d)
         // ...
```

The Pedestrian's View

- Modules are isolated from macros
 - Interface is "compiled" set of exported entities
 - Not affected by macros defined in the importing TU
 - Conversely, macros defined in a module do not leak out
- A unique place where exported entities are declared
 - A module can be just one TU, or several Tus with a distinguished TU for exports
- Every entity is defined at exactly one place, and processed only once
 - Owned by the defining module
 - Except full semantics analysis of templates
 - Exception is made for "global module" (for seamless integration)
- No new name lookup rules
 - We have too many already, and nobody knows how many
- Modules do not replace header files
 - Macro heavy interfaces are likely to continue using header files, with fairly modularized subcomponents
- > Build time is faster (goal)

The Rules of Engagement

What To Expect

- > Module owns entities in its purview
 - ODR: every entity is defined exactly once
- > Order of consecutive import declarations is irrelevant
- > Modules are isolated from macros
- > Import declarations only makes name available
 - You don't pay for what you don't use

- > Module metadata suitable for use by packaging systems
- > Modules provide ownership

Anatomy of a Module Unit

Module purview

```
#include <iostream>
import Enum.Utils;
                                          // for bits::rep().
module Calendar.Month;
namespace Chrono {
     export enum class Month { Jan = 1, Feb, Mar, Apr, May, Jun, /*... */ };
     constexpr const char* month_name_table[] = {
          "January", "February", /* ... */
    };
     export std::ostream& operator<<(std::ostream& os, Month m)</pre>
         assert(m >= Month::Jan and m <= Month::Dec);</pre>
         return os << month name table[bits::rep(m) - 1];</pre>
```

Anatomy of a Module Unit

Module purview

```
#include <iostream>
import Enum.Utils;
                                          // for bits::rep().
module Calendar.Month;
namespace Chrono {
     export enum class Month { Jan = 1, Feb, Mar, Apr, May, Jun, /*... */ };
     constexpr const char* month_name_table[] = {
          "January", "February", /* ... */
    };
     export std::ostream& operator<<(std::ostream& os, Month m)</pre>
         assert(m >= Month::Jan and m <= Month::Dec);</pre>
         return os << month name table[bits::rep(m) - 1];</pre>
```

Namespace partition

A Word on Ownership

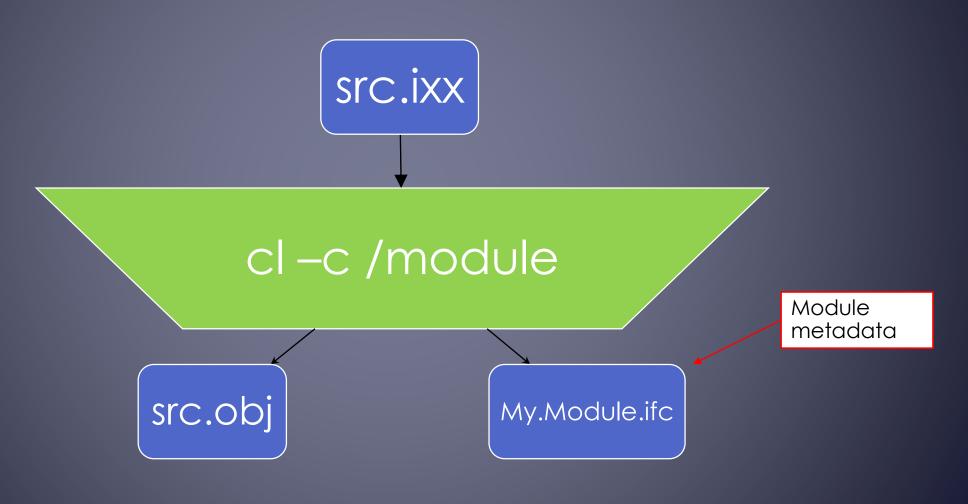
Impoverished Linking Abstractions

- > Strings and bytes
 - Name "mangling" or name "decoration"
 - Unfortunate leakage to language specification
- > Standard "linkage" far behind the practice and needs of our time
- > Examples:
 - GCC and Clang support linkage "visibility"
 - > default
 - > hidden
 - > internal
 - > protected
 - VC++ supports:
 - > dllimport
 - > dllexport

Tool Support

VC++ ongoing implementation

Production: Compiling a module interface



Production: Compiling a module interface

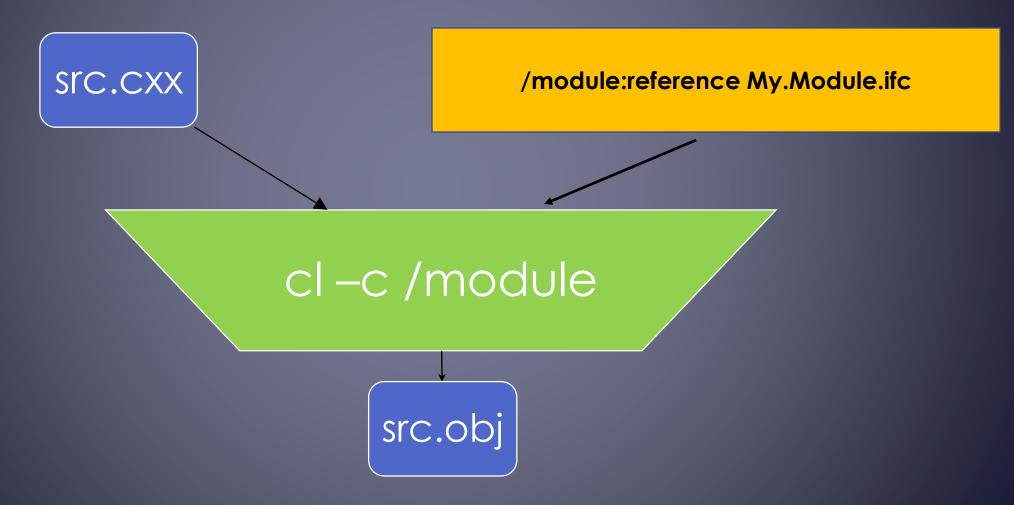
src.cxx

cl-c/module /module:interface

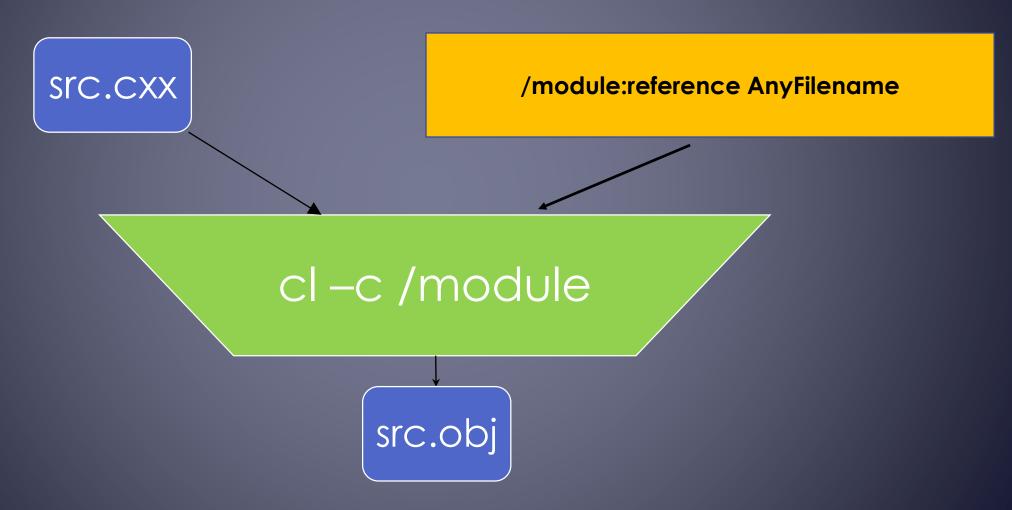
src.obj

My.Module.ifc

Consumption



Consumption



Compiler Options

> /module

- > Turn on module support
- > New keywords: module, import, export

> /module:interface

Force the compiler to interpret source code as module interface definition

> /module:reference <filename>

Look for a compiled module interface (IFC) in the file designated by the path

> /module:search < directory>

> Search directory for referenced files

Transition paths

- > Can one consume a header file as a module?
 - Only if the header file "behaves" well
 - /module:export vec.cxx /module:name std.vector
- > Ablity to selectively export "well behaved" macros
 - /module:exportMacro <macroName>

Compiled Module Interface (IFC)

> Binary Format Designed to represent C++

- Intended to be open (ideally used by all C++ implementations)
- Recognizable by C++ programmers and implementers alike
- An open-source reference implementation (e.g. on GitHub)
- Compact, efficient, complete

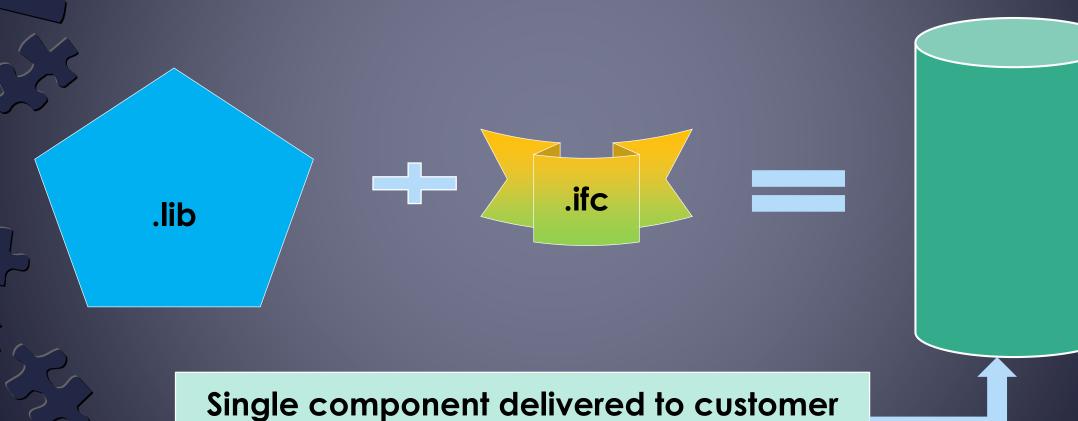
> Structure:

- Set of homogeneous tables representing all relevant entities
- "pointers" are represented by typed indices (all 32-bit wide)
- A "header" describing table locations, size, etc.
- Principle: every index is well typed.
- Deterministically produced by input module source file

> Tooling

- Inspecting; embedding into static LIB or DLL; etc.
- IDE integration

Enhancing Libraries with Module Interfaces



No header file!

PCON 2015



User Feedback

Syntax

"This is simple and elegant. Please do not make it ugly because it has to be C++"

- OK. I will try.

Standardization

- "Can I get it in C++17?"
 - We are trying.
- > "Pretty please, give me modules now"
 - We are trying
- "What about the IFC format"
 - After modules.
- > "Really??? Are you kidding?"
 - No, but I can use some help

Controlling Visibility

- "What about inaccessible members? Are they exported too?"
 - Yes, but I hope we find a good solution
- > "Come on!"



- > See talks
 - by Steve and Ayman: "What is New in Visual C++ 2015 and Future Directions", Thursday
 - Neil MacIntosh on Static Analysis and Safe Buffer Types
 - gdr on "Contracts", tomorrow



Thanks!

Questions?