

Transactional Memory in Practice

Brett Hall
Principal Software Engineer



“May You Live in
Interesting Times”

2

- often said to be an ancient chinese curse

“May You Live in Interesting Times”

- actually not an ancient Chinese curse

3

1936: British diplomat hears it, writes it in memoir

"Better to be a dog in a peaceful time, than to be a man in a chaotic period."

“May You Live in Interesting Times”

- actually not an ancient Chinese curse

The multi-core era is our
“interesting times”

4

- Multi-core era is an “interesting time”,
- forcing us to figure out concurrency
- when dealing with concurrency the Best course would be to share nothing
 - hard to avoid sharing state in most instances though
 - need to synchronize memory access...

Current Synchronization Constructs

- atomics
- mutexes

Atomics

- Too fiddly for programming in the large
- Too fiddly for programming in the small as well except for limited circumstances...

Atomics

- Too fiddly for programming in the large
- Too fiddly for programming in the small as well except for limited circumstances...by people who know what they're doing...

Atomics

- Too fiddly for programming in the large
- Too fiddly for programming in the small as well except for limited circumstances...by people who know what they're doing...who still have a hard time getting it right

Mutexes

- Not too fiddly for programming in the small
- Can work for programming in the large...but don't compose

9

- Work OK for small systems where you can get them to play nicely with each other and avoid deadlock
- But deadlock destroys the ability to compose solutions
- you can try to come up with all sort of ways to avoid it
 - but eventually there's a decent chance that you'll run into something like...

```
void some_function ()
{
    auto lock = std::lock_guard<std::mutex>(some_mutex);
    some_other_function ();
}
```

what does some_other_function do?
need to know to avoid deadlock

```
void some_function ()  
{  
    auto lock = std::lock_guard<std::mutex>(some_mutex);  
    some_library_function_that_you_dont_have_src_for ();  
}
```

Have to hope that the documentation tells you whether a lock is taken or not, and hope that the documentation is correct

```
template <typename F>
void some_function (F no_way_to_know_what_this_func_does)
{
    auto lock = std::lock_guard<std::mutex>(some_mutex);
    no_way_to_know_what_this_func_does ();
}
```

All bets are off
have to hope that caller respects any requests you make

```

void some_function ()
{
    auto lock = std::lock_guard<std::mutex>(some_mutex);
    some_other_function ();
}

void some_function ()
{
    auto lock = std::lock_guard<std::mutex>(some_mutex);
    some_library_function_that_you_dont_have_src_for ();
}

template <typename F>
void some_function (F no_way_to_know_what_this_func_does)
{
    auto lock = std::lock_guard<std::mutex>(some_mutex);
    no_way_to_know_what_this_func_does ();
}

```

The “asking nicely” pattern doesn’t scale

13

- All cases require non-local reasoning about lock order
- All rely on the “asking nicely” pattern, and that doesn’t scale

We could use some better synchronization options

- Actors (aka CSP)
- Transactional Memory

14

Actors = share very little
actors require more code restructuring
will say what TM is in a minute
transactional memory is more flexible
actors can be implemented on top of TM, but going the other way is difficult

What is Transactional Memory?

15

- Start here since interview candidates give blank stares when asked about TM

Transactions

```
transaction
{
    const auto a = var_a;
    const auto b = var_b;
    const auto var_c = a + b;
}
```

16

- Somehow we start a transaction
 - I'll get into more detail in a bit
- transactions behave "atomically"...

“Atomically”, part one

```
//Thread 1
transaction
{
    const auto a = var_a;
    const auto b = var_b;
    const auto var_c = a + b;
}

//Thread 2
transaction
{
    var_a = 5;
}
```

We see a consistent view of memory in our transaction

- First part of “atomically”
- We need to see a consistent view of memory in our transaction
- transaction behaves as though there's no other threads running around messing with stuff
 - could implement with one lock, but that would be a fancy way of writing single threaded code

“Atomically”, part two

```
//Thread 1
transaction
{
    var_a = 10;
    //Do some other stuff
    //...
    //...
    //...
    //...
    //...
    var_b = 5;
}

//Thread 2

transaction
{
    assert (var_a != 10);
    assert (var_b != 5);
}

transaction
{
    assert (var_a == 10);
    assert (var_b == 5);
}
```

Our writes become visible to other threads all at once

18

- writes aren't visible until thread 1 commits
- they all become visible at once
 - if you see one change you can see them all
 - assuming no other thread makes another change, then you might miss a change
- If this is implemented right there's no way to get deadlock

Why did we use
transactional memory?



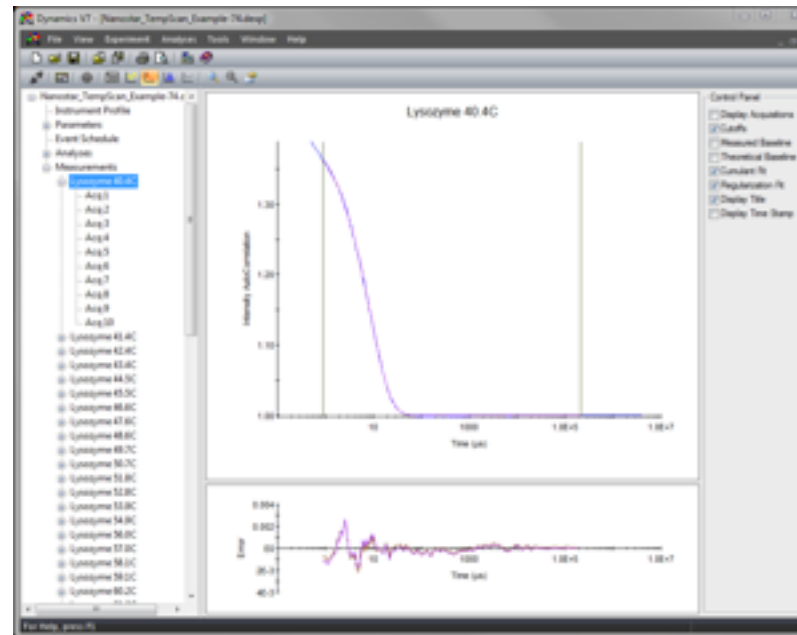
20

- Go back 9 years to when I started at Wyatt
- when I inherited Dynamics it was in this state (took a year or so to realize this)
 - the architecture was never meant the amount of data it was now being tasked with
 - thread safety was through the "do it in the GUI thread" pattern
 - program would stutter and lock up for minutes at a time as calculations were done
 - impossible to use multiple cores for calculations
 - attempting to fix this lead to ...



21

- fixing the threading issues led to a cascading series of changes that was going to require large parts of the program to be rewritten
- competitive pressures for features caused us to shelve fixing this for a few years
- In the interim was playing with Haskell STM
- wrote my own STM system in C++ for kicks
- it worked and TM seemed like a good fit for parts of Dynamics
 - originally only going to use it in a limited scope so that we could easily replace it if there were problems
 - ended up all over the place, more on that later
- The higher ups were more comfortable with using my TM system than with Haskell



22

- Before we get too much farther, should tell you a bit about Dynamics
- 800K ~ 900K lines of code
 - not huge, but not trivial either
- files are broken down into measurements that are further broken down into acquisitions
 - calculations on acquisitions are independent
 - measurements depend on their acquisitions, but not on other measurements
 - lots of easy parallelism to exploit

Our System

Starting a transaction

```
template <typename Func>
auto atomically (Func f, const atomically_args& args = atomically_args {})
    -> decltype (f (std::declval<transaction>{}));

void run_transaction ()
{
    const auto result = atomically ([](transaction& t) {return 0;});
}
```

24

- atomically creates the transaction object and passes it to the function that the caller passes in
 - transaction is not copyable and only atomically can create one
- atomically handles validating and committing or rolling back
- more on the "args" in a bit
- transaction has some methods, but we don't need to worry about that for now
- If the transaction function throws an exception then the transaction is rolled back
- kind of a boring transaction...

Transactional Variables

```
template <typename Type>
class variable
{
public:
    const Type& get (transaction& t) const;
    void set (const Type& value, transaction& t);
};

variable<int> var1;
variable<std::string> var2;

void run_transaction_with_variables ()
{
    const auto result = atomically ([&](transaction& t)
    {
        var2.set ("transaction run", t);
        return var.get (t);
    });
}
```

25

- variables make transactions more interesting
- this is the skeletal version, also have some convenience functions for reading, setting outside a transaction
 - the convenience functions run a transaction that just does the get/set
- system is explicit since only variable's are transacted
 - anything else accessed in a transaction is not transacted
 - This makes our system explicit

Optimistic Concurrency

```
variable<int> var_a (10);
variable<int> var_b (5);

//Thread 1
void transaction1 ()
{
    atomically ([[]](transaction& t)
    {
        auto a = var_a.get (t);
        auto b = var_b.get (t);
        //Do stuff with a and b
        //...
        //...
        //...
        //conflict on var_a
        //=> roll back and try again
    });
}

//Thread 2
void transaction2 ()
{
    atomically ([[]](transaction& t)
    {
        var_a.set (20, t);
    });
}
```

26

- we journal the reads and writes
- at end of transaction we validate the reads
 - if a variable was changed by a transaction on another thread we roll back and start over
 - if no changes then we publish our writes to the rest of the system
- we don't validate writes
 - if there was no read of a variable then its old value was of no concern and if it changed it doesn't matter

Nested Transactions

```
variable<int> var;

int not_a_child (transaction& t)
{
    return var.get (t);
}

int child ()
{
    return atomically ([[] (transaction& t)
                      {
                          return var.get (t);
                      }]);
}

void parent ()
{
    atomically ([[] (transaction& t)
                {
                    auto not_child_result = not_a_child (t);
                    auto child_result = child ();
                }]);
}
```

27

- have to allow nested transactions else we lose composability
 - library only system in C++ has no choice, Haskell doesn't allow nesting, a compiler based system could disallow it
- have to call atomically when in a transaction to get a child transaction
 - most of the time you want to call a function and pass it the transaction object instead (less overhead)
- when child commits its read and writes are merged into the parent
- more details to come when we get to some examples

SG5 Technical Specification

See N3919 and N4514

28

- ISO C++ committee study group for transactional memory
- N4514 voted out of committee as a TS at last meeting
- N3919 is more readable and gives motivation for what their doing

Starting a transaction

```
atomic_noexcept
{
    //throwing from here result in undefined behavior
}

atomic_cancel
{
    //throwing from here causes the transaction to rollback
}

atomic_commit
{
    //throwing from here causes the transaction to commit
}
```

29

- cancel requires the exception to be “transaction safe”
 - limited list of exceptions fits this description
 - supposed to prevent the observation of rolled back actions
- No requirements on implementations except that atomic blocks behave “atomically”
 - don't have to use optimistic concurrency

Transactional Variables

```
atomic_cancel  
{  
    ++i;  
    x = y + 5;  
}
```

30

- everything done in a transaction is transacted
- “implicit system” (the Wyatt system is “explicit”)
- might add “escape actions” later that allow one to mark part of the transaction as not being needed to be transacted
- have to be careful about not accessing the same data in and out of a transaction at the same time
 - leads to data races

Transaction Safety

```
void implicitly_safe ()
{
    //does only "safe" stuff
}

void explicitly_safe () transaction_safe
{
    //does only "safe" stuff, otherwise won't compile
}

void unsafe ()
{
    //does stuff that isn't safe
}

atomic_cancel
{
    implicitly_safe (); // this is OK
    explicitly_safe (); // this is OK
    unsafe (); //won't compile
}
```

31

- only functions "safe" can be called within a transaction
- "safe" functions only do "safe" stuff (see N4515)
 - generally I/O and volatile access is not allowed
 - things get complicated for virtual functions (transaction_safe_dynamic)
- can be explicitly marked safe, safety of implementation enforced by compiler
- calling other functions in a transaction that aren't "transaction_safe" won't compile
- only "safe" stuff can be done in a transaction
- wyatt system has a run-time check for functions that are explicitly marked "unsafe"
 - exception is thrown if unsafe function is called within a transaction
 - best we can do in a library

Nested Transactions

```
void child ()  
{  
    atomic_cancel  
    {  
        //...  
    }  
}  
  
atomic_cancel  
{  
    child ();  
}
```

32

- nesting transactions is allowed
- works similarly to the Wyatt system

The Canonical Synchronization Example

The Canonical Synchronization Example

Bank Accounts

Simple Transfer

```
//Wyatt
variable<int> account1 (10);
variable<int> account2 (0);

void transfer (int x)
{
    atomically ([](transaction& t)
    {
        account1.set (account1.get (t) - x, t);
        account2.set (account2.get (t) + x, t);
    });
}

//SGS TS
int account1 = 10;
int account2 = 0;

void transfer (int x)
{
    atomic_cancel
    {
        account1 -= x;
        account2 += x;
    }
}
```

35

- Note: no chance of deadlock if another thread does the same transfer in the opposite order
- No other thread can observe either operation without observing the other
 - can't see debit of account1 without seeing credit of account 2
- No race: If either account balance changes while we're in our transaction then we start over and try again until we can do it without any conflicts

Transfer with notification

```
void send_email ()
{
    //send and email about a transfer
}

void transfer_with_notification (int x)
{
    atomically ([](transaction& t)
    {
        account1.set (account1.get (t) - x, t);
        account2.set (account2.get (t) + x, t);
        t.after (send_email); //Sends after commit
    });
}
```

36

- Going to be sticking with the Wyatt system from now on, SG5 doesn't have some of these features yet
- "After" queues an action to take when the TOP-LEVEL commit finishes
- multiple actions can be queued
- needed to interface with non-transactional code
 - can't just wait for the transaction commit because we might be in a nested transaction
- SG5 is considering something similar

Waiting for a balance

```
void transfer_when_greater (int limit, int x)
{
    atomically ([](transaction& t)
    {
        if (account1.get (t) <= limit)
        {
            retry (t);
        }

        account1.set (account1.get (t) - x, t);
        account2.set (account2.get (t) + x, t);
    });
}
```

37

- can't use condition variables since we don't have mutexes
- retry sleeps until a variable that's been read changes
 - then transaction starts over
- SG5 also considering something along these lines
- there are some issues to look out for, more in a bit

Pitfalls

Pitfalls

Unfortunately it's not all unicorns and rainbows

39

- never reached "Rainbows and Unicorns" status though
 - see urban dictionary (this one isn't good)

Side-effects

```
void transfer_with_notification_broken (int x)
{
    atomically ([](transaction& t)
    {
        account1.set (account1.get (t) - x, t);
        account2.set (account2.get (t) + x, t);
        send_email (); //Sends before commit, maybe more than once
    });
}
```

40

- Need to avoid side-effects in transactions
 - SendEmail will be called more than once for the same transfer if we have a conflict
 - could be called ANY number of times
- Most of the time it's obvious
 - every once in a while it's harder to track down, but not any worse than a difficult to find deadlock
- in SG5 system this couldn't happen since SendEmail wouldn't be transaction safe and the compiler would catch it
- in our case we have to use something else...

NO_ATOMIC

```
void send_email (NO_ATOMIC)
{
    //send and email about a transfer
}

void transfer_with_notification_broken_but_not_as_bad (int x)
{
    atomically ([](transaction& t)
    {
        account1.set (account1.get (t) - x, t);
        account2.set (account2.get (t) + x, t);
        send_email (); //crashes here, normally caught in testing
    });
}
```

- declares an arg of type "no_atomic"
 - "no_atomic" constructor throws if it is called within a transaction
- Get a run-time error when SendEmail is called in transaction
- caught by testers most of the time

Starvation

```
variable=double> value;

//thread 1
void long_transaction ()
{
    atomically ([6]([transaction& t])
    {
        auto val = value.get (t);

        //do a bunch of stuff
        //...
        //yep, still doing stuff
        //...
        //...
        //need to do enough stuff so
        //that the transaction on
        //thread 2 can commit and cause
        //us to have a conflict every time
    });
}

//thread 2

void short_transaction_that_causes_conflicts ()
{
    while (true)
    {
        atomically ([6]([transaction& t])
        {
            auto old_value = value.get (t);
            value.set (old_value + 5.0, t);
        });
    }
}
```

42

- short transaction continually causes conflicts in a longer transaction
- until a couple of months ago we had never seen this happen
- One of the optional arguments to `atomically` allows you to specify how many conflicts are acceptable
 - either throws an exception when the limit is hit or “runs locked”
 - “run locked” = no other thread can commit until our transaction does
 - controlled by another argument
 - had put this in preemptively in some spots, but had never seen the code fire
 - of course has a bug when we finally needed it (weird edge case we didn’t have a unit test for)
 - heavy-handed response, but works for us given how rare this is
 - other ways to handle, but much more complicated to implement

Retry Deadlock

```
variable<bool> value1 {false};
variable<bool> value2 {false};

//thread 1
void set_1_and_wait_for_2 ()
{
    atomically ([](transaction& t)
    {
        value1.set (true, t);
        if (!value2.get (t))
        {
            retry (t);
        }
    });
}

//thread2
void set_2_and_wait_for_1 ()
{
    atomically ([](transaction& t)
    {
        value2.set (true, t);
        if (!value1.get (t))
        {
            retry (t);
        }
    });
}
```

43

- retry is extremely useful, but has some issues
- two threads, each is waiting on the other to do something that will be done once the other thread commits, but neither can commit because they're waiting on the other
- never saw this in really code, but is theoretically possible
- really we need to split up the transaction...

Split the transactions

```
variable<bool> value1 {false};
variable<bool> value2 {false};

//thread 1
void set_1_and_wait_for_2 ()
{
    atomically ([](transaction& t)
    {
        value1.set (true, t);
    });
    atomically ([](transaction& t)
    {
        if (!value2.get (t))
        {
            retry (t);
        }
    });
}

//thread2
void set_2_and_wait_for_1 ()
{
    atomically ([](transaction& t)
    {
        value2.set (true, t);
    });
    atomically ([](transaction& t)
    {
        if (!value1.get (t))
        {
            retry (t);
        }
    });
}
```

Nested, split transactions

```
variable<bool> value1 {false};
variable<bool> value2 {false};

//thread 1
void set_1_and_wait_for_2 ()
{
    atomically ([](transaction& t)
    {
        value1.set (true, t);
    });
    atomically ([](transaction& t)
    {
        if (!value2.get (t))
        {
            retry (t);
        }
    });
}

atomically ([](transaction& t)
{
    set_1_and_wait_for_2 ();
});

//thread2
void set_2_and_wait_for_1 ()
{
    atomically ([](transaction& t)
    {
        value2.set (true, t);
    });
    atomically ([](transaction& t)
    {
        if (!value1.get (t))
        {
            retry (t);
        }
    });
}

atomically ([](transaction& t)
{
    set_2_and_wait_for_1 ();
});
```

45

- broken again
- modified from example given by Hans Boehm to the SG5 mailing list
 - SG5 is wrestling with this now in discussions of retry in TS
- Fixed by applying NO_ATOMIC to the set_.... functions
 - relying on commit happening can be considered a side-effect
- Can be more subtle, the second transaction could be buried in library code

Inconsistent Values

```
variable<size_t> value1 {10};
variable<size_t> value2 {2};

void thread1 ()
{
    atomically ([&](transaction& t)
    {
        auto a = value1.get (t); //a = 10
        //do some stuff
        //...
        //...
        //...
        //...
        //...
        //until the other thread has committed
        auto b = value2.get (t); //b = 11
        auto values = std::vector<int>(a - b); //oops
    });
}

void thread2 ()
{
    atomically ([&](transaction& t)
    {
        value1.set (15, t);
        value2.set (11, t);
    });
}
```

46

- quirk of our system
- value1 > value2 is a class invariant that ensures we only allocate a sane amount of memory
 - in this case we will probably try to allocate a 4gb vector
- normally our transaction would have a conflict and we'd restart and see the consistent values
 - memory allocation is a side-effect, but most of the time it's idempotent (especially if RAII is used) so it doesn't matter
- to solve this case use...

Validate

```
variable<size_t> value1 {10};
variable<size_t> value2 {2};

void thread1 ()
{
    atomically ([](transaction& t)
    {
        auto a = value1.get (t); //a = 10
        //do some stuff
        //...
        //...
        //...
        //...
        //...
        //until the other thread has committed
        auto b = value2.get (t); //b = 11
        validate (t); //restart here until no conflicts
        auto values = std::vector<int>(a - b);
    });
}

void thread2 ()
{
    atomically ([](transaction& t)
    {
        value1.set (15, t);
        value2.set (11, t);
    });
}
```

47

- validate() runs transaction validation and restarts transaction if there's a conflict
 - can't get to the memory allocation if things are inconsistent
- has only happened a few times
- Looking into ways to avoid this problem without killing performance

Fine vs. Coarse Grained

```
struct fine_grained
{
    variable<int> m_member1;
    variable<int> m_member2;
    variable<int> m_member3;
    variable<int> m_member4;
};

struct coarse_grained
{
    int m_member1;
    int m_member2;
    int m_member3;
    int m_member4;
};
variable<std::shared_ptr<const coarse_grained>> m_coarse;
```

48

- anything shared between threads needs to be transacted
- decision is what level to transact at
 - fine grained: individual members are transactional variables
 - good for often updated structure
 - coarse grained: individual members not transactional, but whole structure is always stored immutably in a transactional variable
 - good for read lots, write rarely
 - can become point of contention if too much writing going on

Weak Atomicity

```
void exploiting_weakness ()
{
    exhaust_port* port = atomically ([&transaction& t)
    {
        return fly_down_the_trench (t);
    });
    shoot_proton_torpedo (port);
}
```

49

- our system is weakly atomic
- the exhaust port is stored in a transactional variable
 - our system can't stop you from mutating the object outside of a transaction
 - defeats the whole point of having the system
- to guard against this...

Require Immutable or “Internally Transacted”

```
//Immutable
variable<std::shared_ptr<const exhaust_port>> port;

//Internally Transacted
struct exhaust_port
{
    //all mutable state is stored in transactional variables...
    variable<ray_shielding> m_shielding;

    //...or accessed through transactional methods
    void shoot_with_proton_torpedo (transaction& t);
};
```

50

- ... as a policy we required stuff stored in transactional variables to be either immutable or “internally transacted”
 - “internally transacted” = all mutable state is in transactional variables
- enforced through code review

Invasive



51

- original plan was to only use TM in data store and maybe experiment parameters
 - it sort of crept like the kudzu above to take over a lot more
 - this was due to TM usefulness and complication of interfacing transactional and non-transactional code
- Used TM for everything except low-level instrument communication (data reception)
- GUI isn't transactional (MFC), but calls transactional code

No one's heard of it



52

- When I ask about TM in job interviews usually just get blank stares in response
 - Hasn't really mattered, people come up to speed quick
- But there's also not a lot of places to turn for "best practices"
 - lots of academic papers on implementing TM systems
 - but you're mostly on your own when it comes to figuring out how to apply the system
- that's it for pitfalls...

Performance?



53

- Much like this car: not the fastest, but fits our purposes well

Performance

- Dynamics is not a “low latency” application (for the most part)
 - Not a game or trading bot
 - We’re good at “hiding” our latencies
 - Calculations are almost ideally set-up for transactional memory
 - Some instrument communication is low latency, but doesn’t use TM
- With these caveats we get plenty of performance

54

- Dynamics isn't a low latency app
- Calcs taking even 10 seconds extra (out of minutes) is acceptable
- push as much as possible onto background threads (TM facilitates this)
- prioritize calcs based on what the user is looking at
 - sometimes the user is looking at something that requires everything to be calculated, they just have to wait but the GUI stays responsive
- Data reception has latency requirements
 - have to keep up with data stream or get disconnected
 - don't use TM here (lock free queues instead)
- Performance is good, even with our bone-headed implementation
 - as more cores become common place we will need to upgrade out implementation so it scales better

Calculation Structure

```
calculate ()  
{  
    params = getParams ();  
    data = getData ();  
    auto res = CalculateResult (params, data);  
    StoreResult ();  
}
```

55

- Calculations were the main thing we wanted to push into background threads
- Calculations are structured such that only a few transactional reads need to be done at the start and then some transactional writes at the end
- the transactional overhead negligible compared to what Calculate result is doing

Performance

- More overhead
- Bigger non-determinism price to pay than with mutexes
- Can have less contention

56

- TM May Still be using mutexes under the covers
 - otherwise using lock-free stuff
 - don't know how to do retry without a mutex and condition variables
 - either way still have to pay sync price
 - Also have to pay for journaling and validation (and maybe non-locality if system is indirect)
- Don't hold locks for long, so less contention possibility (especially for read only stuff), but...
- When there is contention you might have to pay higher price
 - repeated work when there's conflicts
 - an example...

Performance

```
std::mutex m;  
int a;  
int b;  
  
void use_a_and_b ()  
{  
    auto lock = std::lock_guard<std::mutex>(m);  
    //do a bunch of read-only stuff with a and b  
}  
  
variable<int> a;  
variable<int> b;  
  
void use_a_and_b ()  
{  
    atomically {[](transaction& t)  
    {  
        //do a bunch of read-only stuff with a and b  
    }  
}
```

57

- Only one thread can call use_a_and_b at a time when using mutexes
- Any number of threads can use the transactional version
 - transactional version might have more overhead, but the contention win might weigh in its favor as long as writes are RARE
- the mutex version can probably be restructured to have less contention (or use a shared_mutex)
 - that has to be done by hand
 - in transactional case its automatic

Hardware Transactional Memory

- Exists today: Intel TSX
 - Transacts cache-lines
 - Implicit
 - Bounded

58

- haven't actually used it, just read about it
- was in high-end Xeon Haswell chips
 - silicon had bug
 - disabled until later broad well chips, still high end only
 - in skylake, might be more evenly distributed
- cache-lines transacted => have to watch out for cache ping-ponging
- implicit like TS
 - all memory access is transacted while in a transaction
- Bounded
 - limited buffers for reads and write
 - have to fall back on software TM if cache spills
 - probably an expensive context switch

How'd it work out?

- TL;DR: Great!
 - Just keep in mind what our app is like
- New people get up to speed quickly
- More fun than mutexes
 - Optimizes for programmer time

59

- Had two new people join the team since we started using it
 - one junior, one senior
 - Neither had TM experience
- learning curve is short
 - Usually stop making obvious mistakes in less than a month
- mutexes set the bar pretty low
 - TM has lower cognitive load for the most part

Open Source?

60

- Always get asked about this
- last year the answer was “no” when I gave a lightning talk,
- this year the answer is...

Open Source?

Soon!

61

- Got the go ahead from the higher ups at Wyatt to do it a few weeks ago
- Haven't had time to clean it up
 - some embarrassing code that we've been living with that needs to be fixed
 - open sourcing it is already improving our code quality
- Other goodies along with TM system
 - DeferredResult, Channel, etc.
- watch my blog for details...

Resources

- My blog: <https://backwardsincompatibilities.wordpress.com/>
- Wyatt: <http://www.wyatt.com>
- SG5 papers:
 - N3919: motivation and details
 - N4514: standard-ese
 - N4438: my paper about our experience at Wyatt
- GCC experimental TM system

62

- blog url is in my conference profile
- No reference implementation for SG5 stuff yet, but plans for GCC and Clang implementations
- for gcc just google “gcc transactional memory”
 - has been in there since 4.7
- Intel had an experimental compiler
 - it's listed as “retired” on their website