

# The Unexceptional Exceptions

Fedor G Pikus, Chief Scientist  
Mentor Graphics, Design2Silicon Division

CPPCon, September 2015

# Outline

---

- The motivation
- Introduction to exceptions
- Other ways to handle errors
- Handle errors not exceptions
- Simple guidelines for error handling
  - with and without exceptions
- Using exceptions in legacy code
- Additional reading

# The Motivation

---

- The main goal of this class is not to teach how to use exceptions
  - it's not that hard
- The main goal is to teach to not be afraid to use exceptions
  - also not to be awed by them

# Exceptions

---

- Exceptions are a language facility for handling abnormal situations at run time
  - Usually these situations are errors
- Exceptions interrupt normal control flow
  - They are used when normal control flow becomes impossible or undesired
- Exceptions create alternative control flow
  - Exception handling code is executed only when exception occurs

# Basic Exceptions

---

```
int main() {  
    try {  
        Widget* w = create_widget();  
        if(w=NULL) throw 1;  
    }  
    catch (int e) {  
        cout<<"Exception " << e << endl;  
    }  
}
```

# Exceptional Tools

---

- Exception is generated using “throw”
  - throw *object*;
  - Any data type can be thrown
  - Standard provides exception base class `std::exception`
  - Exception must have copy constructor
  - Exception may contain data
    - throw `myStatus("Oops, bad file");`
    - You can throw a temporary object

# Exceptional Circumstances

---

- What happens when exception is thrown?
- Execution of current scope ends
  - Code between throw and closing “}” is skipped
  - Control jumps to the closing “}”
  - Local variables are destructed
  - Stack frame is unwound (exiting the scope)
- The process is repeated until exception is caught
- If exception is not caught in main() the program terminates
  - Compiler can determine that the exception is not caught, and terminate the program immediately (without unwinding stack)

# Basic Exceptions Example

```
void do_stuff() {  
    Widget w;  
    while (1) {  
        Gadget g;  
        throw 1;  
        transmogrify(g,w); // Not called  
    }                       // g destroyed (may be)  
}                           // w destroyed (may be)  
  
int main() {  
    do_stuff();  
    more_stuff();          // Not called  
}                          // Not caught
```



# Basic Exceptions Example

```
void do_stuff() {
    Widget w;
    while (1) {
        Gadget g;
        throw 1;
        transmogrify(g,w); // Not called
    }                      // g destroyed
                          // w destroyed
}

int main() {
    try {
        do_stuff();
    } catch { int i } { cout << "Caught " << i << endl; }
    more_stuff();          // Called
}
```

# Catching Exceptions

- Exception is handled with try-catch blocks
  - `try {do_stuff();} catch(type e) {handle();}`
  - You can catch
    - by reference: `catch(int& e) {...}`
    - by value: `catch(int e) {...}`
      - copy will be made, usually undesirable (except for basic types)
    - by pointer: `catch(int* e) {...}`
      - avoid if possible (who should delete the pointer?)
  - `catch(...)` catches any type

# Catching Exceptions

---

- You can catch different types separately

```
try { do_stuff(); }
```

```
catch(int e) {  
    cout << "Error number " << e; }
```

```
catch(std::exception& e) {  
    cout << "Exception " << e.what(); }
```

```
catch(...) {  
    cout << "Something bad happened"; }
```


# Throwing Polymorphically

---

- Can derived class be thrown?
  - Yes, anything can be thrown

# Throwing Polymorphically

```
class D : public B {};  
  
void f(B& throw_if_trouble) {  
    if (trouble()) throw throw_if_trouble;  
}  
  
int main() {  
    D status;  
    try { f(status); }  
    catch(D& e) { cout << "It's a D!"; }  
    catch(...) {  
        cout << "It's weird...";  
    }  
}
```



# Throwing Polymorphically

---

- Can derived class be thrown?
  - Yes, anything can be thrown
- Is throw itself polymorphic?
  - No, throw is not a virtual function
  - Type of the object to throw is determined at compile time
- Can exceptions be thrown polymorphically?
  - Yes, using virtual functions

# Throwing Polymorphically

---

```
class B {  
    virtual void raise() { throw *this; }  
};  
  
class D : public B {  
    virtual void raise() { throw *this; }  
};  
  
void f(B& throw_if_trouble) {  
    if (trouble()) throw_if_trouble.raise();  
}
```

# Re-throwing Exceptions

- Can exception handler itself throw?
  - Yes, but what?
  - The “current exception”

```
try { ... }  
catch(...) {  
    cout << “Oops!”;  
    throw;  
}
```

- The exception continues to propagate



# Re-throwing Exceptions

- Re-throw is often used to handle exceptions “in stages”
  - Exception can be modified along the way

```
void transmogrify() {  
    try { ... }  
    catch(status& s) {  
        s.trace("Transmogrification failed");  
        throw;  
    }  
}
```

# Re-throwing Exceptions

- Exception dispatcher uses re-throw

```
void dispatcher() {  
    try { throw; }  
    catch(exception1& e) {  
        ... handle exception1 ... }  
    catch(exception2& e) {  
        ... handle exception2 ... }  
}  
  
try { something_dangerous(); }  
catch(...) { dispatcher(); }
```

- Allows code reuse in exception handling

# Alternatives to Exceptions

---

- First, what's wrong with exceptions?

# Exceptional Problems

---

- Exceptions introduce separate execution flow
- Exception-handling code runs while the normal execution flow is suspended
  - The program is in who-knows-what state, interrupted at any random time, but exception-handling channel has to run anyway
  - Exceptions (usually) don't happen when things are going well, so the program is likely in a bad state already
- The goal of exception handling is to restore the program to a well-defined state
  - How well defined?
  - What can we expect from the program that throws?
- We need some guarantees, otherwise anything is possible

# Alternatives to Exceptions

---

- First, what's wrong with exceptions?
- Exception-safe code is hard to write

# Alternatives to Exceptions

- Main alternative to exceptions is return codes (or error codes)

```
int f() {  
    Widget* w = create_widget();  
    if (w == NULL) return -1;  
}
```

- Return codes use the same control flow for normal and exceptional situations

# The Main Difference

---

- Return codes can be ignored
  - Return codes often are ignored
  - `printf("This never fails\n");`
    - `printf()` does return an error code!
  - If an error code is ignored, program continues
- Exceptions continue to propagate until they are caught
  - Propagate all the way to `main()` if not caught
  - If an exception is ignored, program terminates

# Exceptional Flexibility

---

- Return codes are constrained by the type of return value
  - Return values may be used in normal flow
  - “Impossible” return value may not exist
- Exceptions can be of any type
  - Several unrelated types can be thrown
  - New types can be added any time
  - Exceptions are more flexible
- But there is a price for the flexibility



# Exceptions and Interface

---

- Return values are part of the interface
  - At least the type is, the values must still be documented separately
- Exceptions are mostly invisible in the interface
  - what can this function throw?  
`Widget* create_widget();`
    - anything at all...
- The separate execution channel – here it is

# Exceptions and Interface

- But there are `throw()` declarations! And `noexcept()`!
  - Yes, but they quite don't do what the name implies

`Widget* create_widget() throw();`

- There is no compile-time check that an exception is not thrown
- If such exception is thrown, the program will terminate (run-time check)

- No-throw guarantee can be useful, `noexcept` more so

`void test_widget() throw(); // C++03`

`void test_widget() noexcept(true); // C++11`

- Still not enforced at compile time – the program may throw and terminate at run time
- no-throw functions may throw and catch internally!

# Point of No Return

---

- Constructors have no return values!
  - Exceptions can be thrown in constructors
  - If exception is thrown, the object was not constructed successfully
- Destructors have no return values...
  - Destructors should not throw either
  - Destructors can be called while exception is propagating and stack is unwound
    - Two exceptions cannot propagate at the same time
  - In C++11, destructors are `noexcept()` by default

# Exceptions vs Error Codes

---

## Exceptions

- Cannot be ignored
- Own execution flow
- Propagates until handled
- Implicit interface
- Flexible
- Can be used in constructors

## Return codes

- Can be ignored
- Same execution flow
- Must be propagated manually
- Explicit interface
- Constrained
- Cannot be used in constructors

# Exception Safety is Hard

---

```
Widget* make_widget(Log* l) {  
    l->open();  
    Widget* w = new Widget();  
    l->record(w);  
    l->close();  
    return w;  
}
```

- What could possibly go wrong?
  - Everything...

# Exception Safety is Hard

---

```
Widget* make_widget(Log* l) {  
    l->open();
```

- `open()` can throw

- function exits, no harm

```
Widget* w = new Widget();
```

```
l->record(w);
```

```
l->close();
```

```
return w;
```

```
}
```

# Exception Safety is Hard

---

```
Widget* make_widget(Log* l) {  
    l->open();  
    Widget* w = new Widget();
```

```
    ■ operator new can throw  
      — function exits, log file remains open  
    l->record(w);  
    l->close();  
    return w;  
}
```

# Exception Safety is Hard

```
Widget* make_widget(Log* l) {  
    l->open();  
    Widget* w = new Widget();  
    l->record(w);
```

■ `record()` can throw

— log file remains open, Widget is leaked

— What if Widget contains a lock?

```
    l->close();  
    return w;  
}
```



# No Exceptions - No Problems

---

- Rewrite everything to use error codes...

```
Widget* make_widget(Log* l) {  
    l->open();  
    Widget* w = new Widget();  
    l->record(w);  
    l->close();  
    return w;  
}
```

- ... and things become much simpler

# See No Evil

---

```
Widget* make_widget(Log* l) {  
    l->open();
```

- `open()` can return error code  
— la-la-la, I can't hear you...

```
Widget* w = new Widget();  
l->record(w);  
l->close();  
return w;  
}
```

# Hear No Evil

---

```
Widget* make_widget(Log* l) {  
    l->open();  
    Widget* w = new Widget();  
    ■ memory allocation can fail, return NULL  
      — la-la-la, I can't hear you...  
    l->record(w);  
    l->close();  
    return w;  
}
```

# Say No Evil

---

```
Widget* make_widget(Log* l) {  
    l->open();  
    Widget* w = new Widget();  
    l->record(w);  
    ■ record(w) can fail, return error code  
      — la-la-la, I can't hear you...  
    l->close();  
    return w;  
}
```

# No Exceptions - Same Problems

- Rewrite everything to use error codes...

```
Widget* make_widget(Log* l) {  
    if (l->open()<0) return NULL;  
    Widget* w = new Widget();  
    if (!w) return NULL;  
    if (l->record(w)<0) return NULL;  
    if (l->close()<0) return NULL;  
    return w;  
}
```

- ... and things become just as bad, only uglier

# Handle Errors, not Exceptions

---

- The fundamental problem is handling errors
  - not handling exceptions
- After an error is detected, the program must remain in a well-defined state
- As a minimum, no resources must be leaked
  - basic error safety guarantee
- Ideally, operation which caused error is undone
  - strong error safety guarantee
- It's hard because something went wrong ...

# Exception Safety Guarantees

---

- No-throw guarantee
  - “We don’t use exceptions here”
  - Neither do our libraries? Neither does C++ runtime...?
    - Do you allocate memory? Use dynamic casts on references?
  - What happens if something goes wrong? See the alternatives...
- No guarantees
  - “We have no idea what we are doing”
  - What happens if something goes wrong?
- If an exception is thrown, the state of the program is undefined

# Exception Safety Guarantees

---

## ■ Basic exception safety guarantee

- If the exception is thrown, the program remains in a defined state
- Defined state: some actions are possible, with known results
- On an object that throws, some member functions may be called (possibly only the destructor)
- Defined state is not always a very useful state
- Termination is a well-defined state (see “not very useful” above)

## ■ Strong exception safety guarantee

- Any action either succeeds or has no effect (“commit or rollback”)
- This is very hard to do, may be expensive, sometimes impossible
- More useful in a narrow context, like a library: any member function either succeeds or leaves the object unchanged
- Still very expensive – STL does not provide such guarantee



# Practical exception safety guarantee

- Some operations may provide strong guarantee
  - STL `push_back()` – either succeeds, or throws and the container remains unchanged
- Most operations provide basic guarantee
  - Basic invariants are maintained (at least destructors can be called on all objects)
  - No resources are permanently leaked (memory, locks, threads,...)
    - “permanently” – you may need to destroy some objects first
  - Total cleanup on throw is not required, but there must be a way to free resources (usually an object owns them and remains valid and accessible to the caller)
- Some operations provide no-throw guarantee
  - `std::swap` and user-defined swap – very important!
  - All destructors

# Handle Errors, not Exceptions

---

- Exception safety guarantees are just a particular case of error handling guarantees
- If the program attempts something and there is an error:
  - Errors cannot happen (similar to “we don’t use exceptions” but not quite as easy)
  - Anything is possible (still “we have no idea what we’re doing”)
  - The program remains in a defined state and leaks no resources (basic guarantee again)
  - Every operation will either succeed or be undone completely (strong guarantee)
- Whether errors are handled by exceptions or not, some guarantees of a defined state must be given

# Guidelines

---

- Know and specify error handling guarantees

# Exceptions are Just a Tool

```
Widget* make_widget(Log* l) {  
    l->open();  
    Widget* w = NULL;  
    try { w = new Widget(); }  
        catch{ l->close();  
            throw;  
        }  
    try { l->record(w); }  
        catch{ delete w;  
            l->close();  
            throw;  
        }  
}
```

- This looks a lot like error codes...

# Nothing Exceptional

```
Widget* make_widget(Log* l) {  
    if (l->open() < 0) return NULL;  
    Widget* w = NULL;  
    w = new Widget();  
    if (!w) {  
        l->close();  
        return NULL;  
    }  
    if (l->record(w) < 0) {  
        delete w; l->close();  
        return NULL;  
    }  
}
```

- This looks a lot like try-catch blocks ...

# Errors Not Exceptions

- Think about handling errors, ignore the details

```
Widget* make_widget(Log* l) {  
    l->open();  
    if (error) return;  
    Widget* w = new Widget();  
    if (error) { l->close(); return; }  
    l->record(w);  
    if (error) { delete w; l->close(); return; }  
}
```

# Errors Not Exceptions

- In a complex operation, all steps which succeeded before an error must be undone  
—or at least left in a well-defined state

```
Widget* make_widget(Log* l) {  
    l->open();  
    on_error: { l->close(); }           // From this point on  
    Widget* w = new Widget();  
    on_error: { delete w; }           // From this point on  
    l->record(w);  
}
```

# Errors Not Exceptions

- We need an action which is performed automatically when we leave the scope
  - for any reason (throw, return, break)

```
Widget* make_widget(Log* l) {  
    l->open();  
    on_error: { l->close(); }  
    Widget* w = new Widget();  
    on_error: { delete w; }  
    ...  
}
```

- Local objects are destructed when leaving scope
- What about `l->close()` when no errors happened?



# Guidelines

---

- Know and specify error handling guarantees
- Handle normal cleanup in normal channel

# Exceptionally Simple

```
struct LogGuard {  
    Log* l;  
    LogGuard(Log* l) :l(l) { l->open(); }  
    ~LogGuard() { l->close(); }  
}
```

```
template<class T> struct ScopedPtr {  
    T* p_;  
    ScopedPtr(T* p): p_(p) {}  
    ~ScopedPtr () { delete p_;}  
}
```

- C++11 provides `unique_ptr<T>` - use it
  - I want to show what's happening inside

# Exceptionally Simple

---

```
Widget* make_widget(Log* l) {  
    LogGuard LG(l);  
    // If anything happens log will close  
    ScopedPtr <Widget> w = new Widget();  
    // If anything happens w will be  
    // deleted  
    l->record(w);  
    return w;  
}  
// Log is closed automatically – yay!  
// w is deleted automatically - bummer
```

# Exceptionally Resourceful

---

- We have converted the error handling problem to a resource management problem
  - Standard resource management technique in C++ is RAII – resource acquisition is initialization
- Resources which are acquired in a scope must be released upon leaving the scope
  - unless they are still owned by someone else
- C++ has many ways to manage resources and their ownership

# Exceptionally Simple

```
Widget* make_widget(Log* l) {  
    LogGuard LG(l);  
    // If anything happens log will close  
    ScopedPtr <Widget> w = new Widget();  
    // If anything happens w will be  
    // deleted  
    l->record(w);  
    // w should not be deleted now  
    // we need a way to release ownership  
    return w; // w is ScopedPtr not Widget*  
} // Log is closed automatically – yay!
```

# Manual Release

```
template<class T> struct ScopedPtr {  
    T* p_;  
    ScopedPtr (T* p): p_(p) {}  
    ~ ScopedPtr () {  
        delete p_;  
    }  
    T* release() {  
        T* tmp = p_;  
        p_ = NULL;  
        return tmp;  
    }  
}
```

# Exceptionally Simple

---

```
Widget* make_widget(Log* l) {  
    LogGuard LG(l);  
    // If anything happens log will close  
    ScopedPtr <Widget> w = new Widget();  
    // If anything happens w will be  
    // deleted  
    l->record(w);  
    return w.release(); // No errors after this!  
} // Log is closed automatically
```

# Guidelines

---

- Know and specify error handling guarantees
- Handle normal cleanup in normal channel
- Use RAII to automate resource cleanup



# Exceptional Automation

- Using resource management consistently makes programming even easier
  - reference counting or ownership transfer
- ```
unique_ptr<Widget> make_widget(Log* l) {  
    LogGuard LG(l);  
    unique_ptr<Widget> w = new Widget();  
    l->record(w);  
    return w;                // pointer moved out of w  
}
```
- Before C++11 you could use `autoptr` (and it actually did what you wanted)
- Where are the try-catch blocks?

# Do or Don't Do, There Is No Try

---

- We are handling errors not exceptions
  - exceptions are just a tool, we could use error codes
- Error handling is mostly about making sure that resources are not leaked and program invariants are maintained
  - these actions must be taken even without errors
  - this must be done in normal control flow
  - not error-only control flow like catch blocks or error tests
- But at some point errors must be handled too

# Reasons to Try (and Catch)

---

- Stop exceptions from propagating
  - must be done somewhere (or the program terminates)
- Process the exception, propagate the error
  - throw different exception or return an error code
    - useful when adding exceptions to “legacy” code
- Modify exception and re-throw it
  - special case, rarely used
- Clean up the try block after the error
  - NOT a reason to catch
  - use automatic scope-specific cleanup instead

# Reasons to Try (and Catch)

---

- Actions which must be done in error channel only should be done in catch blocks
  - Stop exceptions from propagating
  - Process the exception, propagate the error
  - Modify exception and re-throw it
- Actions which must be done in error and normal channels should not be done in try blocks
  - Clean up the try block after the error

# Guidelines

---

- Know and specify error handling guarantees
- Handle normal cleanup in normal channel
- Use RAII to automate resource cleanup
- Use catch blocks only when necessary

# Unnecessary Error Handling

---

- Manual error cleanup
- Not taking advantage of language features available for error handling
- Bad organization of error objects

# Unnecessary Error Handling

---

- Manual error cleanup
  - deleting raw pointers, closing files, etc
  - Explicitly reclaiming non-memory resources
- Not taking advantage of language features available for error handling
- Bad organization of error objects

# Resizing a Buffer

```
template <class T> Vector<T>::resize() {  
    T* new_buf=allocate(2*size);  
    try {  
        uninitialized_array_copy(buf,buf+size,new_buf);  
    } catch(...) {  
        deallocate(new_buf);  
        throw;  
    }  
    for(int i = size-1; i >= 0; i--) (buf+i)->~T();  
    deallocate(buf);  
    buf=new_buf; size *= 2;  
}
```

- If resizing fails, the whole operation is undone



# Copying Data

```
template <class T> void uninitialized_array_copy(
    T* begin, T* end, T* dest) {
    T* start = dest;
    try {
        while(begin != end)
            new(dest++) T(*begin++);
    } catch(...) {
        while(start != dest)
            start++->~T( );
        throw;
    }
}
```

- If copying fails, previous copies are destroyed

# Resizing a Buffer, no Catch

```
template<class T> struct Buffer {
    T* p;
    Buffer(size_t s): p(allocate(s)) {}
    ~Buffer() { deallocate(p); }
    void swap(T* q) noexcept(true)
        { std::swap(q,p); } // swap() does not throw!
}

template <class T> Vector<T>::resize() {
    Buffer<T> new_buf(2*size);
    uninitialized_array_copy(buf,buf+size,new_buf.p);
    for(int i = size-1; i >= 0; i--) (buf+i)->~T();
    new_buf.swap(buf);
    size *= 2;
}
```

# Copying Data, No Catch

```
template<class T> struct CopyGuard {
    T* b; T*& e; bool release;
    CopyGuard(T* b, T*& e): b(b),e(e),release(false){}
    ~CopyGuard() {
        if (!release) while(b!=e) b++->~T();
    }
    void success() { release = false; }
}

template <class T> void uninitialized_array_copy(
    T* begin, T* end, T* dest) {
    CopyGuard<T> CG(dest, dest);
    while(begin != end) new(dest++) T(*begin++);
    CG.success();
}
```

# Unnecessary Error Handling

---

- Manual error cleanup
- Not taking advantage of language features available for error handling
  - Return codes require more explicit error handling than exceptions
    - when exceptions are used correctly
- Bad organization of error objects

# Exceptions vs Error Codes Again

---

- Using return codes:

```
Widget* make_widget() {  
    if (create_widget()==NULL) return NULL;  
}
```

```
Widget* process_widget() {  
    if (make_widget()==NULL) return NULL;  
}
```

- Errors must be propagated manually

# Exceptions vs Error Codes Again

- Using exceptions poorly:

```
Widget* make_widget() {  
    try { create_widget(); }  
    catch(...) { ...handle error... throw; }  
}
```

```
Widget* process_widget() {  
    try { make_widget(); }  
    catch(...) { ...handle error... throw; }  
}
```

# Exceptions vs Error Codes Again

- Using exceptions well:

```
Widget* make_widget() {  
    ScopeGuard SG(...);  
    create_widget();  
}  
  
Widget* process_widget() {  
    ScopeGuard SG(...);  
    make_widget();  
}
```

- Exceptions continue to propagate until caught

# Guidelines

---

- Know and specify error handling guarantees
- Handle normal cleanup in normal channel
- Use RAII to automate resource cleanup
- Use catch blocks only when necessary
- Prefer exceptions to error codes



# Unnecessary Error Handling

---

- Manual error cleanup
- Not taking advantage of language features available for error handling
- Bad organization of error objects
  - error objects could be exceptions or codes
  - different error objects in different subsystems
  - inconsistent error objects
  - error objects must be designed for the whole system
    - may be unavoidable if the system is composed from subsystems with different error handling

# Error Codes and System Design

- Subsystem 1: database

```
Database<Record>::store( Record r ) {  
    if ( file.fail() ) throw DB_IO_problem();  
}
```

- Subsystem 2: music jukebox

```
Jukebox::add_song( SongRecord s ) {  
    try { database.store(s); }  
    catch(DB_IO_problem e)  
        { throw Jukebox_IO_problem(); }  
}
```

- Error codes can suffer from the same problem

# Guidelines

---

- Know and specify error handling guarantees
- Handle normal cleanup in normal channel
- Use RAII to automate resource cleanup
- Use catch blocks only when necessary
- Prefer exceptions to error codes
- Consider error objects part of system design

# Error Codes and System Re-Design

- Subsystem 1: database (new)

```
Database<Record>::store( Record r ) {  
    if ( file.fail() ) throw DB_IO_problem();  
}
```

- Subsystem 2: music jukebox (old)

```
int Jukebox::add_song( SongRecord s ) {  
    try { database.store(s); }  
    catch(DB_IO_problem e)  
        { return Jukebox_IO_problem; }  
}
```

- Errors are repackaged at system boundaries

# Guidelines

---

- Know and specify error handling guarantees
- Handle normal cleanup in normal channel
- Use RAII to automate resource cleanup
- Use catch blocks only when necessary
- Prefer exceptions to error codes
- Consider error objects part of system design
- **Handle/convert errors at subsystem boundaries**

# Becoming Exceptional

---

- Exceptions make error handling easier
- Most of the time we work on existing projects not start new projects
- Many existing projects do not use exceptions
  - most of those that do try too hard and catch too much
- Rewriting the whole system to use exceptions may be impossible
- How to start using exceptions?

# Becoming Exceptional

---

- How to start using exceptions?

# Becoming Exceptional

- How to start using exceptions?
  - That's not the goal, the goal is to handle errors

```
char* read_record(const char* filename){  
    int fd = open(filename, O_RDONLY);  
    char buffer[1024];  
    int n = read(fd, buffer, 1024);  
    if (n != 1024) return NULL;  
    ...  
    close(fd);  
}
```

- Using RAII to close the file is more important than a particular error handling technique



# Becoming Exceptional

- How to start using exceptions?
  - That's not the goal, the goal is to handle errors **well**

```
char* read_record(const char* filename){  
    int fd = open(filename, O_RDONLY);  
    char buffer[1024];  
    int n = read(fd, buffer, 1024);  
    if(n != 1024) return NULL;  
    ...  
    close(fd); // has error code too
```

- But don't neglect the advantages of exceptions over error codes either

# Becoming Exceptional

---

- How to start using exceptions for handling errors?
- Maintain defined state, avoid leaking resources
  - use RAII for automatic resource cleanup
  - this technique does not depend on using exceptions
- Use exceptions to handle errors in new subsystems
- Convert errors at subsystem boundaries
  - catch exceptions before returning control to the old system, return error codes as that system expects

# Legacy and 3<sup>rd</sup> party code

---

- Code from different sources may use different error handling approaches
- Convert error handling on the interfaces if necessary
- Know and maintain error handling guarantees
- Challenges:
  - Error handling is not always visible in the interfaces
  - Error handling is usually the worst documented part
- There is no substitute and no shortcuts for knowing what “defined state” is and what’s guaranteed

# Legacy and 3<sup>rd</sup> party client code

- “Integrating 3<sup>rd</sup> party code” usually means using someone else’s libraries to build applications
- What if the situation is reversed: you are building a library to serve a variety of clients?
- Convert errors at subsystem boundaries
  - Need something to convert
- Consider designing a library with multiple error-handling mechanisms

```
void Transmogrify(Doohickey& d, int* error_code=NULL);
```

or

```
void Transmogrify(Doohickey& d);
```

```
void Transmogrify(Doohickey& d, int& error_code) noexcept;
```

# Guidelines

---

- Know and specify error handling guarantees
- Handle normal cleanup in normal channel
- Use RAII to automate resource cleanup
- Use catch blocks only when necessary
- Prefer exceptions to error codes
- Consider error objects part of system design
- Handle/convert errors at subsystem boundaries
- Convert legacy code to exceptions gradually

# Guidelines

---

- Know and specify error handling guarantees
- Handle normal cleanup in normal channel
- Use RAII to automate resource cleanup
- Use catch blocks only when necessary
- Prefer exceptions to error codes
- Consider error objects part of system design
- Handle/convert errors at subsystem boundaries
- Convert legacy code to exceptions gradually

# Additional Reading

---

- Herb Sutter, "Exceptional C++", Addison-Wesley 1999
- Herb Sutter, "More Exceptional C++", Addison-Wesley 2002
- Bjarne Stroustrup, "The C++ Programming Language", Addison-Wesley 2004, Appendix E (Standard-Library Exception Safety), <http://www.research.att.com/~bs/3rd.html>
- Bjarne Stroustrup, "Programming with Exceptions", 2001, <http://www.informit.com/articles/article.aspx?p=21084>
- Herb Sutter, "Exception Safety and Exception Specifications: Are They Worth It?", <http://www.gotw.ca/gotw/082.htm>
- Mark Radford, "Designing C++ Interfaces - Exception Safety", 2001, <http://accu.org/index.php/journals/444>
- Herb Sutter, "Unmanaged Pointers in C++: Parameter Evaluation, auto\_ptr, and Exception Safety", DDJ 12.01.2002, <http://www.ddj.com/cpp/184403851>

# Additional Reading

---

- Herb Sutter, "Exception-Safe Function Calls", <http://www.gotw.ca/gotw/056.htm>
- Andrei Alexandrescu and Petru Marginean, "Generic: Change the Way You Write Exception-Safe Code - Forever", DDJ 12.01.2000
- David Abrahams, "Lessons Learned from Specifying Exception-Safety for the C++ Standard Library", [http://www.boost.org/community/exception\\_safety.html](http://www.boost.org/community/exception_safety.html)
- Stephen Dewhurst, "C++ Gotchas", Addison-Wesley 2002, "Gotcha 64: Throwing String Literals"
- Stephen Dewhurst, "Flexible Memory Management in C++", Software Development West 2008
- Marshall Cline, "Exceptions and error handling", <http://www.parashift.com/c++-faq-lite/exceptions.html> (Exceptions FAQ)



# Additional Reading

---

- Jon Kalb, <http://exceptionsafecode.com/>
- David Abrahams, "Exception Safety in Generic Components", [http://www.boost.org/community/exception\\_safety.html](http://www.boost.org/community/exception_safety.html)
- Scott Meyers, "Declare functions as noexcept if possible", <http://scottmeyers.blogspot.com/2014/03/declare-functions-noexcept-whenever.html>
- Andrzej Krzemiński, "noexcept – what for?", <https://akrzemi1.wordpress.com/2014/04/24/noexcept-what-for/>
- TS Filesystem library, <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2014/n4099.html> - supports both exceptions and error codes

# Guidelines

---

- Know and specify error handling guarantees
- Handle normal cleanup in normal channel
- Use RAII to automate resource cleanup
- Use catch blocks only when necessary
- Prefer exceptions to error codes
- Consider error objects part of system design
- Handle/convert errors at subsystem boundaries
- Convert legacy code to exceptions gradually
- Bonus guideline

The background features a dark blue gradient with numerous thin, bright green and blue lines radiating from the top and bottom edges towards the center, creating a starburst or light-ray effect.

# QUESTIONS?



[www.mentor.com](http://www.mentor.com)

# Bonus Nightmare Material

---

# Bonus Nightmare Material

---

- C++ now supports threads
  - In practice, C++ supported threads for a long time
- Most Unix/GCC systems use POSIX Threads (PThreads)
  - `pthread_create()`, `pthread_join()`, `pthread_lock()`, ...
- What about `pthread_cancel()`?
  - If a thread is terminated by `pthread_cancel()`, are destructors called?

# Bonus Nightmare Material

---

- Glibc implements `pthread_cancel()` as an exception
- Destructors are called as the stack is unwound – yay!
- You can catch it – yay! yay..?

```
try {  
    wait_on_socket();  
} catch (...) {  
    cout << "We just killed pthread_cancel()" << endl;  
}
```

- In practice, “cancel” exception is special and the program will abort

# Bonus Nightmare Material

- Glibc implements `pthread_cancel()` as an exception
- Destructors are called as the stack is unwound – yay!
- You can really catch it – yay!

```
try {  
    wait_on_socket();  
} catch (abi::__forced_unwind&) {  
    throw; // Sorry, didn't mean to mess with pthread_cancel  
} catch (...) {  
    cout << "Something else is wrong" << endl;  
}
```

- You probably want to `#ifdef` that for Glibc systems



# Bonus Nightmare Material

---

- Glibc implements `pthread_cancel()` as an exception
- Destructors are called as the stack is unwound – yay!
- You can really catch it – yay!
- Any cancellation point can throw now – oh #^&\$\*!
  - That's most system calls...
  - Here goes the `noexcept()` guarantee
  - Aren't destructors `noexcept()`? Yes, they are...
  - Destructors can close files (and often do)
  - `close()` is a cancellation point...
- Bonus guideline – avoid `pthread_cancel()`